

Introduction to Naming

COS 316: Principles of Computer System Design

Amit Levy & Wyatt Lloyd



- Today: Naming schemes in systems
- Next, naming in:
 - UNIX file system
 - Git
 - Naming in Networking

Assignment 1 is out Due 9/16. *Hint: you really want to attend precept!*

Why do systems need to name things?

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names
 - File names instead of subdivisions of magnetic disks

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names
 - File names instead of subdivisions of magnetic disks
- Abstracts differences in implementation: expose different resources using the same names

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names
 - File names instead of subdivisions of magnetic disks
- Abstracts differences in implementation: expose different resources using the same names
 - File names for data on disks, solid-state drives, or network storage

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names
 - File names instead of subdivisions of magnetic disks
- Abstracts differences in implementation: expose different resources using the same names
 - File names for data on disks, solid-state drives, or network storage
- Isolate applications: the same name refers to different resources for different applications

Why do systems need to name things?

- Abstract complexity: expose underlying resources using simpler names
 - File names instead of subdivisions of magnetic disks
- Abstracts differences in implementation: expose different resources using the same names
 - File names for data on disks, solid-state drives, or network storage
- Isolate applications: the same name refers to different resources for different applications
 - `chroot`: give each application its own root of the file system

So many examples of names!

<code>http://www.princeton.edu/news</code>	URL
<code>www.princeton.edu</code>	hostname
<code>140.180.223.42</code>	IP address
<code>main</code>	procedure name
<code>pc</code>	ARM register name
<code>aalevy@cs.princeton.edu</code>	email
<code>~/Documents/lecture.pdf</code>	file name

Why are names important?

- The primary way applications/users interact with a system

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems
- Choice of naming scheme can dictate system trade-offs:

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems
- Choice of naming scheme can dictate system trade-offs:
 - Performance

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems
- Choice of naming scheme can dictate system trade-offs:
 - Performance
 - Security

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems
- Choice of naming scheme can dictate system trade-offs:
 - Performance
 - Security
 - Portability

Why are names important?

- The primary way applications/users interact with a system
- Often the first step in designing a system
 - A central design decision
 - Can make implementation easier or obvious
- A good place to start understanding systems
- Choice of naming scheme can dictate system trade-offs:
 - Performance
 - Security
 - Portability
 - Generality

Naming and system trade-offs

Let's compare alternate naming systems from the physical world:

Building numbers in a city block

Naming and system trade-offs

Let's compare alternate naming systems from the physical world:

Building numbers in a city block

New York

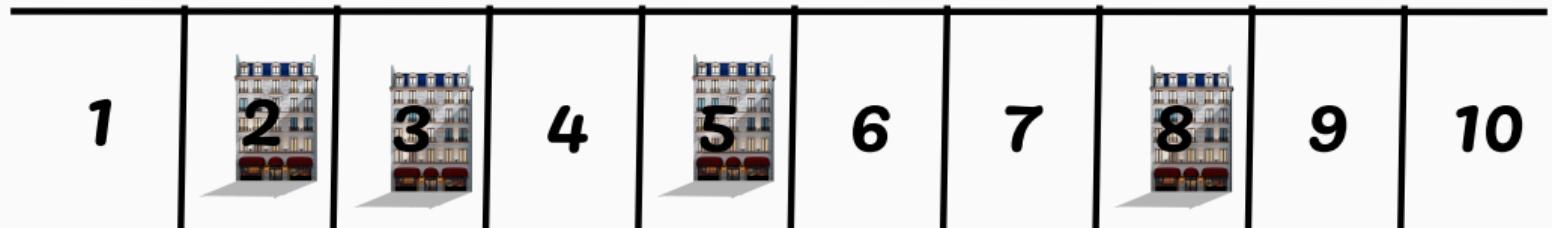
-VS-

Tokyo



Figure 1: How do we name each of these buildings?

New York



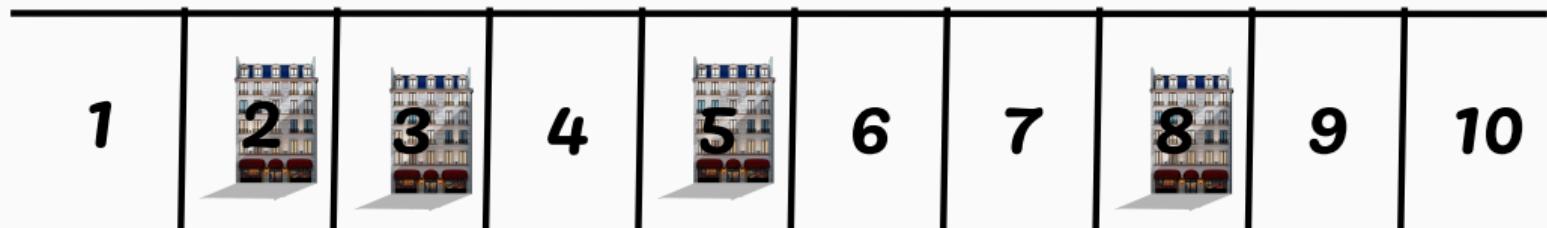
How do buildings get their name?



How do buildings get their name?

1. Subdivide the block into N plots numbered 1-N
2. Each building fits in exactly one of the plots
3. Name buildings with the number of their plot

New York

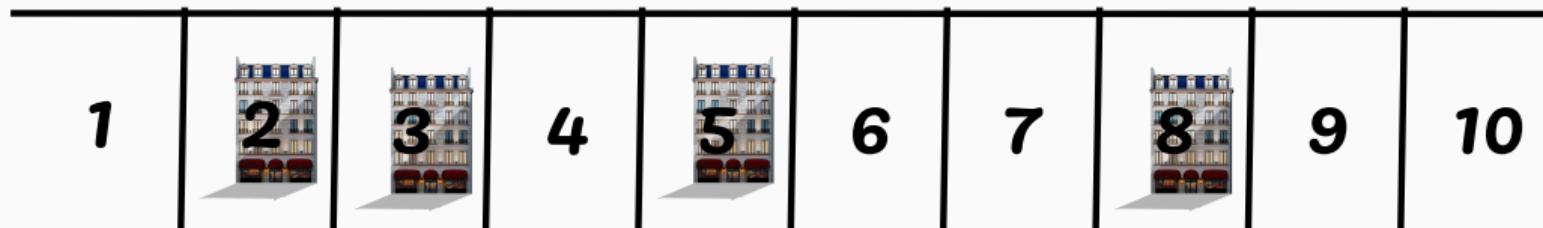


How do buildings get their name?

1. Subdivide the block into N plots numbered 1-N
2. Each building fits in exactly one of the plots
3. Name buildings with the number of their plot

So: 344 10th Ave is one plot south of 346 10th Ave

New York



How do buildings get their name?

1. Subdivide the block into N plots numbered 1-N
2. Each building fits in exactly one of the plots
3. Name buildings with the number of their plot

So: 344 10th Ave is one plot south of 346 10th Ave

But... what happens when we build the $N+1$ th building?



1987



1993



2012



1964

How do buildings get their name?



1987



1993



2012



1964

How do buildings get their name?

Buildings numbered from oldest to newest.

Whenever a new building is built, give it the next highest number on the block



1987



1993



2012



1964

How do buildings get their name?

Buildings numbered from oldest to newest.

Whenever a new building is built, give it the next highest number on the block

Finding a building is hard! Need to check buildings until you find the right number.



1987



1993



2012



1964

How do buildings get their name?

Buildings numbered from oldest to newest.

Whenever a new building is built, give it the next highest number on the block

Finding a building is hard! Need to check buildings until you find the right number.

But... virtually infinite possibility for expansion!

Different naming and trade-offs

	New York Addressing	Tokyo Addressing
Easy allocation	no :(YES
Easy lookup	YES	no :(

Naming Scheme Framework

- **Values:** What is it that we're naming?
 - Disk sectors?
 - Network nodes?
 - Users?
- **Names:** What's the format of a name?
 - Alphanumeric strings up to 32 characters
 - Non-zero integers
 - 128-bit numbers
- **Allocation mechanism:** How does the system create new names and values?
- **Lookup mechanism:** How does the system map from names to values?

Applying Framework to Building Numbers

	New York Addressing	Tokyo Addressing
Values	Physical buildings	Physical buildings
Names	Numbers 1-N	Numbers ≥ 1
Allocation	Pre-defined plot numbers	New building gets next highest number
Lookup	Search ordered set	Visit every building and check

Virtual Memory

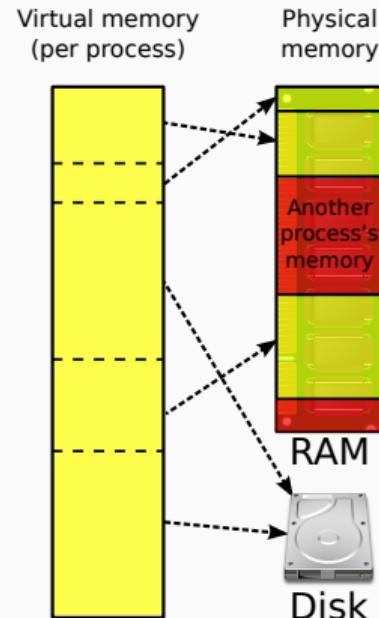


Figure 2: Virtual Memory

What does virtual memory give us?

- Isolation: `0xdeadbeef` maps to different values for different applications
- Abstraction over storage media: virtual addresses can name memory, storage, remote memory, HW register, ...
- Flexibility in provisioning
 - Resource efficiency: physical memory can be less sparse than virtual memory
 - Overcommit physical memory (e.g. swap to disk when necessary)

Virtual Memory as a Naming Scheme

- Values?
- Names?
- Allocation?
- Lookup?

Virtual Memory: Values

- *Physical memory address (i.e. 32-bit or 64-bit address up to size of RAM)*

Virtual Memory: Values

- *Physical memory address (i.e. 32-bit or 64-bit address up to size of RAM)*
- On-disk file and offset in file

Virtual Memory: Values

- *Physical memory address (i.e. 32-bit or 64-bit address up to size of RAM)*
- On-disk file and offset in file
- Some hardware registers (e.g. a network card configuration registers)

Virtual Memory: Values

- *Physical memory address (i.e. 32-bit or 64-bit address up to size of RAM)*
- On-disk file and offset in file
- Some hardware registers (e.g. a network card configuration registers)
- Remote memory

Virtual Memory: Names

Pointer-sized (e.g. 32-bit or 64-bit) addresses and process identifiers

$(PID, \text{virtual_address})$

e.g.

$(3487, 0xdeadbeef)$

Virtual Memory: Allocation

```
void *mmap(void *addr, size_t length) (simplified)
```

- Application chooses an unused name: an address not yet allocated for it
- Kernel keeps a list of unused physical 4KB memory pages
- Kernel allocates “value” by removing a physical page from the list
- Kernel adds new mapping between virtual address and physical to the application’s “page table”
 - in-memory data structure understood by the virtual memory hardware that maps virtual addresses to physical addresses

Virtual Memory: Lookup

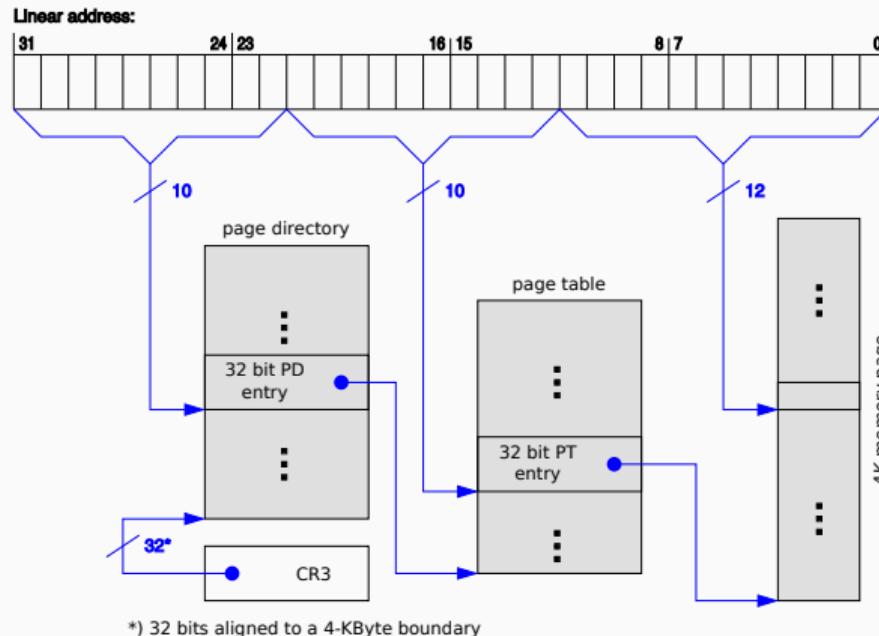


Figure 3: Two-level page table structure in x86

Virtual Memory: Lookup

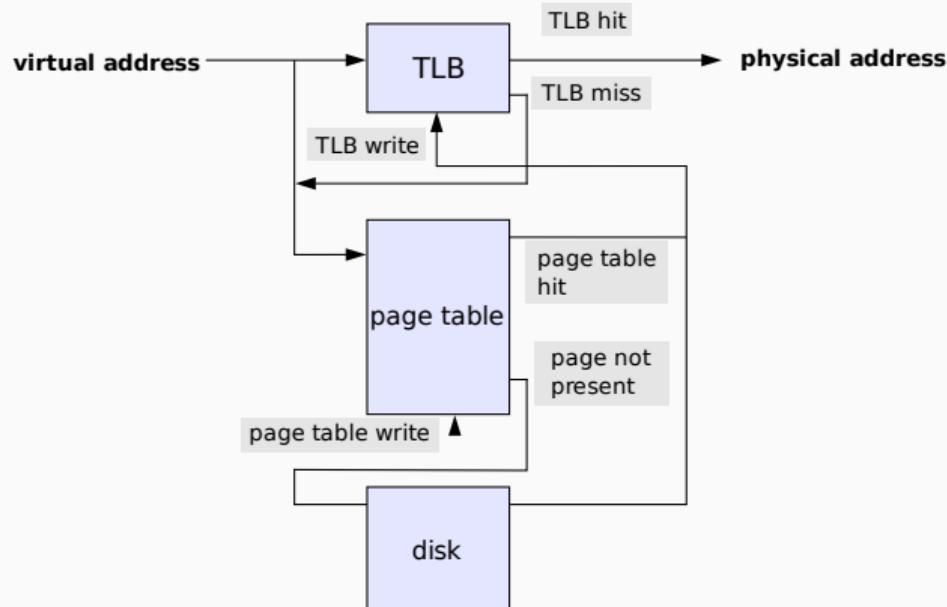


Figure 4: Virtual-to-physical translation

Virtual Memory: Lookup

For this to work, the OS needs to do some housework when context switching:

- Set CR3 register to point to process's page table
- Invalidate the TLB
 - Mark entire TLB as invalid—simple but can cause unnecessary slow down
 - Associate process IDs with each TLB entry

Virtual Memory: Trade-offs with Alternative Naming Schemes

- Single shared address space (identity mapping)
 - Better performance, since no translation hardware/software needed
 - But can't isolate using names
- Swap out all memory for one process at a time (original UNIX)
 - Simple and efficient to implement in hardware
 - Can't run applications in parallel
 - Expensive to switch between applications
- Segmentation: virtual address are low-order bits of physical address + segment register
 - Relatively simple hardware (just concatenate segment register and virtual address)
 - Isolates concurrent applications using names
 - Much coarser grain: all virtual memory must be contiguous in RAM
 - Can't share memory between applications

Summary

- Systems use names to abstract complexity, abstract differences, isolate applications
- Naming choices can dictate system trade-offs
 - performance: fast allocation vs. fast lookup
 - isolation
- Framework for naming:
 - Values
 - Names
 - Allocation mechanism
 - Lookup mechanism