

CPU Caches

COS 316: Principles of Computer System Design

Amit Levy & Wyatt Lloyd

Why do we cache?

Use caches to mask performance bottlenecks by replicating data nearby

Design decisions that characterize a cache

- *Look-aside vs. Look-through*
 - determines who is responsible for writing/fetching data from backing store
- *Write-through vs. Write-back*
 - determines whether items changed in the cache are written immediately to the backing store (write-through) or only upon eviction (write-back)
- *Write-allocate vs. Write-no-allocate*
 - determines whether we allocate space for an item when fetching and storing it (write-allocate) or only when fetching (write-no-allocate) it
- *Eviction policy*
 - determines which item(s) to evict when we run out of space in the cache




Figure 1: CPU Connected Directly to Memory




Figure 1: CPU Connected Directly to Memory

Which combination of look-aside vs look-through, write-through vs. write-back, and write-allocate vs. write-no-allocate would you choose?

Locality: When a cache might be useful

- Useful data tends to continue to be useful



Figure 2: Temporal locality

- Useful data tends to be located “near” other useful data




Figure 3: Spatial locality

CPU caches are particularly constrained:

- Size: typically many orders of magnitude smaller than backing store
 - E.g. 64KB L1 Cache vs 64GB main memory. **6 orders of magnitude!**
- Performance: speed, power consumption, physical die space
- General purpose workloads

CPU caches are particularly constrained:

- Size: typically many orders of magnitude smaller than backing store
 - E.g. 64KB L1 Cache vs 64GB main memory. **6 orders of magnitude!**
- Performance: speed, power consumption, physical die space
- General purpose workloads

The trick: exploit physical memory naming scheme

CPU caches exploit both kinds of locality:

- Exploit temporal locality by remembering the contents of recently accessed memory
- Exploit spatial locality by fetching blocks of data around recently accessed memory




Figure 4: CPU Cache's View of Memory Address




Figure 5: CPU Cache's View of Memory Address

- Addresses with the same tag are added to cache together
 - Spatial locality: bytes around previously accessed byte already in the cache
- Size of block offset determines block size:
 - n bits of block offset means blocks are 2^n bytes
 - E.g. 6 offset bits means 64 byte blocks

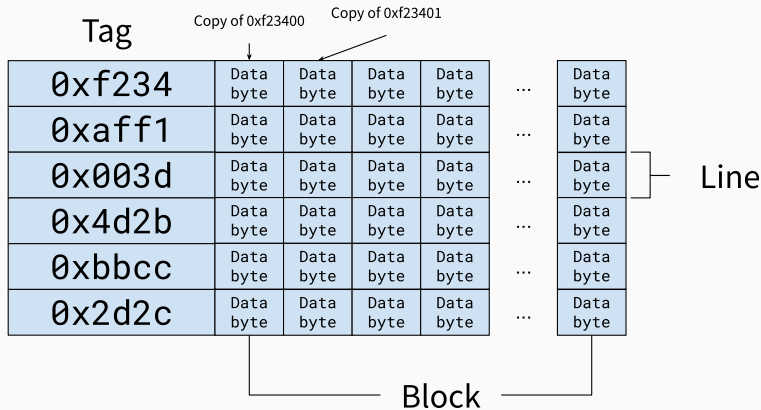



Figure 6: CPU cache stores a block at each Line

Exercise

Starting from 3-line cache that uses 4-bits for the offset, which of the following accesses, performed in order, are hits or misses?

1. 0xff1200df
2. 0xff1200d3
3. 0x01cd3310
4. 0x01cd3310
5. 0xff1200df



Cache Read Algorithm

1. Look at memory address on processor
2. Search cache tags to find a matching block
3. Found in cache?
 - Hit: return data from cache at offset from block
 - Miss:
 - 3.1 Read data block from main memory
 - 3.2 Add data to cache
 - 3.3 Return data from cache at offset from block

Cache Read Algorithm

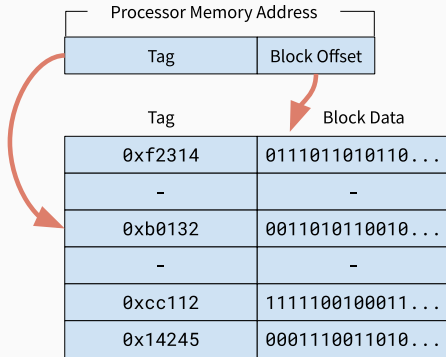
1. Look at memory address on processor
2. Search cache tags to find a matching block
3. Found in cache?
 - Hit: return data from cache at offset from block
 - Miss:
 - 3.1 Read data block from main memory
 - 3.2 Add data to cache
 - 3.3 Return data from cache at offset from block

Which line do we evict for the new block?

Three common placement policies:


- Fully Associative
 - Evict with: LRU, FIFO, NLRU, ...
- Direct Mapped
 - Eviction is trivial
- N-way Associative
 - Combination of both

Fully Associative



Check all lines in the cache for a matching tag


Fully Associative



Check all lines in the cache for a matching tag

What's the disadvantage of fully associative cache?


Direct Mapped



Index size determines number of indices

Check tag at line with matching index: if equal "hit", "miss" otherwise

Direct Mapped

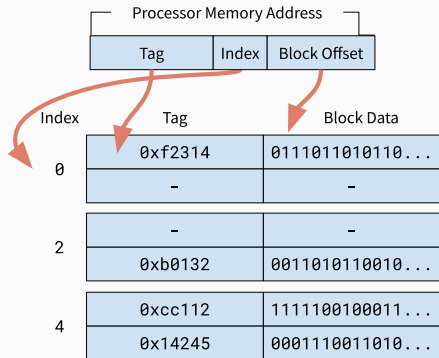


Index size determines number of indices

Check tag at line with matching index: if equal "hit", "miss" otherwise

What's the disadvantage of a direct mapped cache?

N-way Associative




Check *all* tags at line with matching index: if equal "hit", "miss" otherwise

N = number of lines in each set

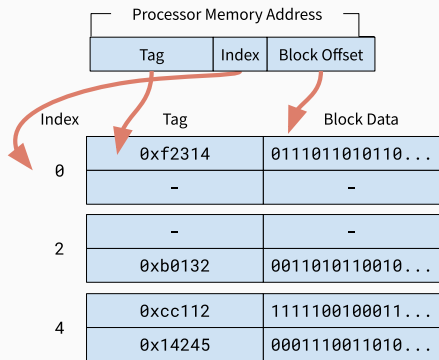
Index size determines number of sets

Exercise: N-way Associative



How many index bits for a 2-way set associative cache with 128 cache lines?

Exercise: N-way Associative



How many index bits for a 2-way set associative cache with 128 cache lines?

128 cache lines, 2 lines per set, how many sets? $128/2 = 64$, how many bits? $\log_2(64) = 6$

- Next time: Web caching with CDNs
- Problem set due tonight

