CDN Caching

COS 316: Principles of Computer System Design

Lecture 9

Amit Levy & Wyatt Lloyd

Anatomy of a Web Page Fetch

- Web page = HTML file + embedded images/objects
- HTML page does not embed objects in it
  - *Q: Why not?*

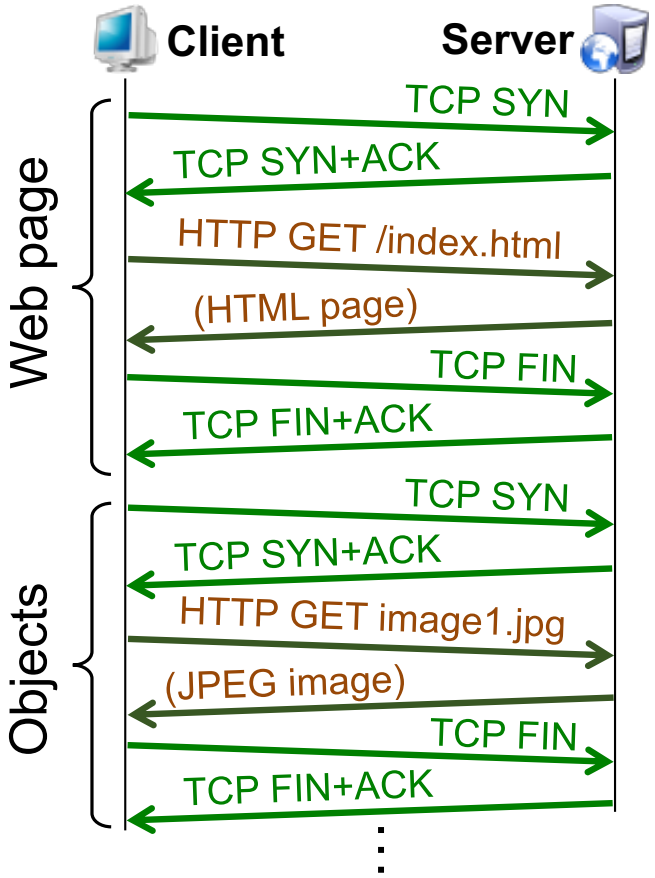

Serve It All From Datacenters?

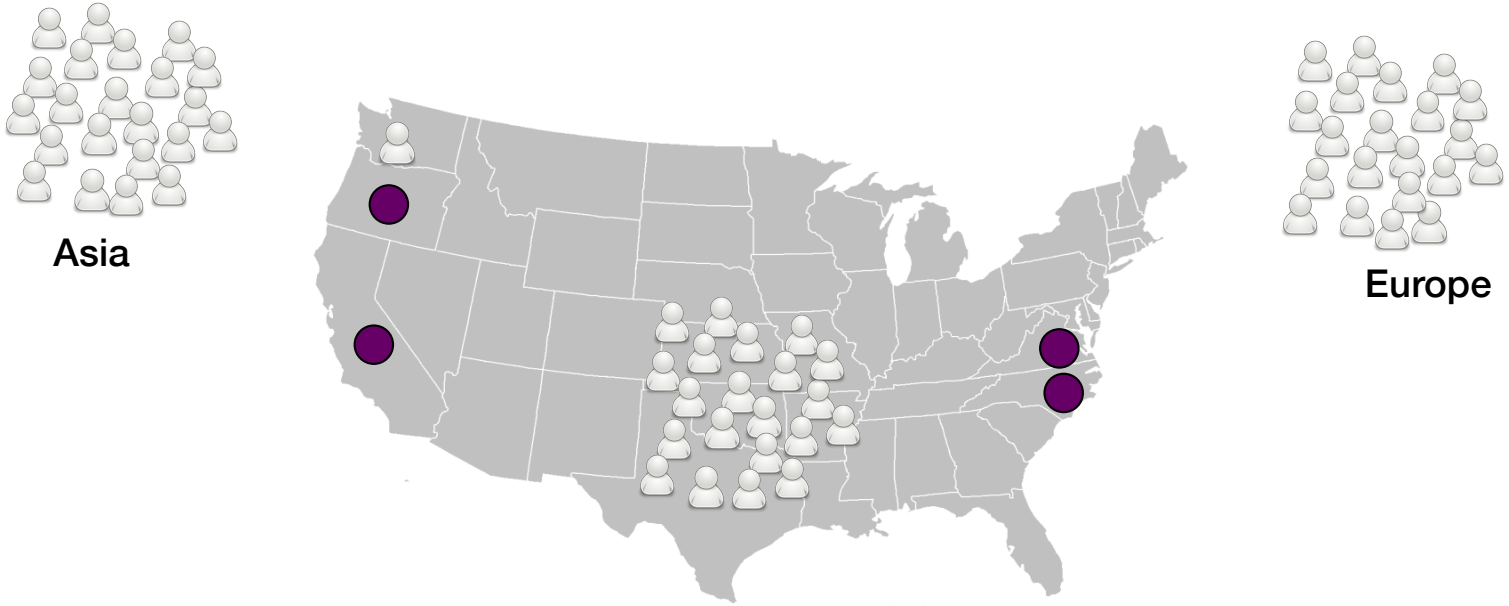

- High latency for many clients ☹
- Use a ton of bandwidth to send the same objects over and over ☹

Serve Objects From CDN

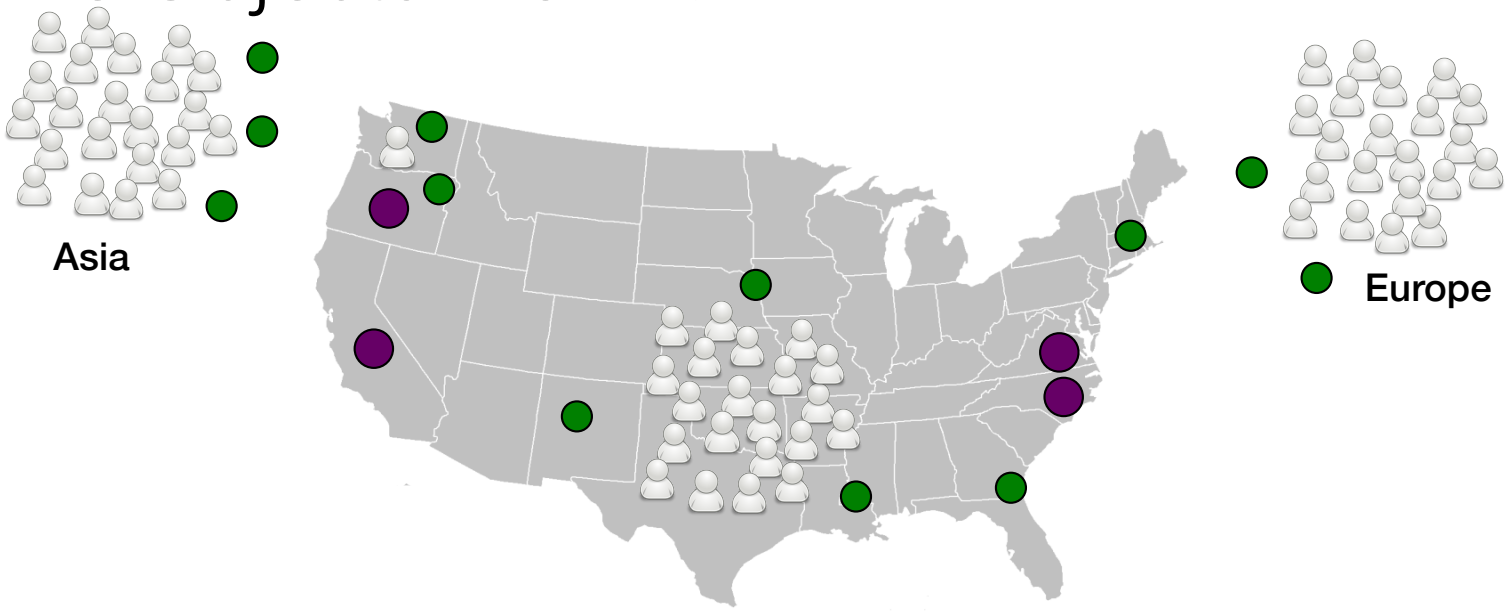

- Lower latency because content is closer ☺
- Less global bandwidth ☺

CDN Locations

- Where can we stick machines?
- Inside our datacenters
- At colocation facilities (other people's DCs)
- At Internet Exchange Points
- Inside others' networks

Not *that* many machines (limited space and/or high cost)

CDN Locations vs Datacenters

- There are many more CDN locations
  - e.g., Akamai has 1000+
- CDNs are closer to people
  - Because there are more of them
  - Placed for locality
- CDN locations are much wimpier
  - Maybe 1 rack (~40 machines)
  - …
  - Maybe 8 racks (~320 machines)

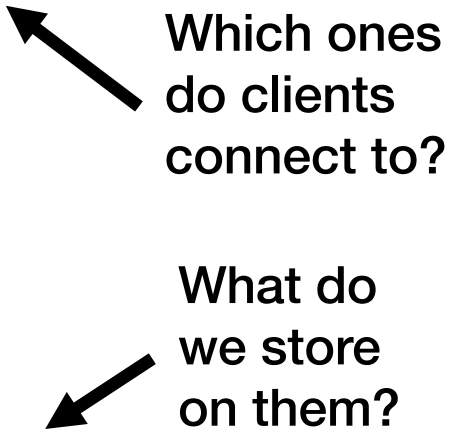

CDN Locations Store What?

- Store everything?
  - Exabytes of data for some services
  - 1 rack (40 machines):
    - 144 GB memory? * 40 -> 5.8TB memory … not even close…
    - 10x4 TB HDD? * 40 -> 1.6 PB … not even close…
- Need to store a subset of objects!
  - *Q: But which objects to store?*

CDN Cache Hit

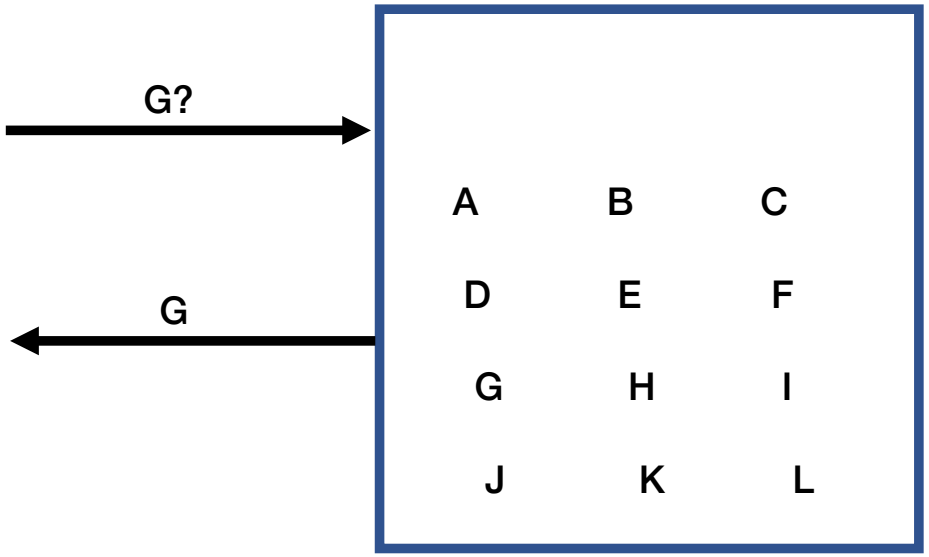

CDN Cache Miss

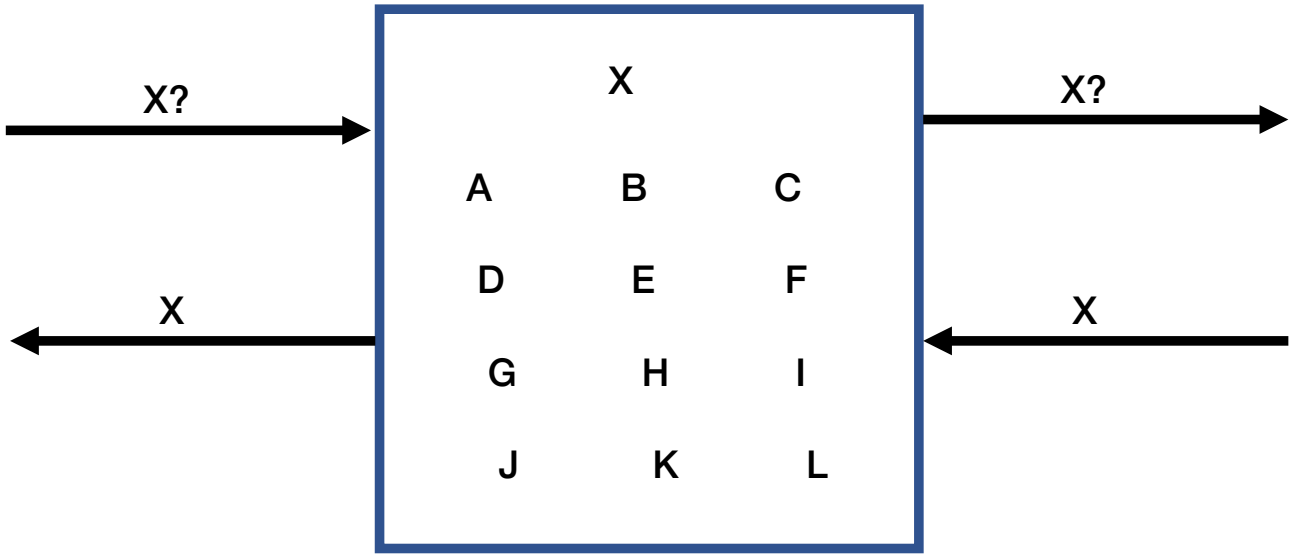

If I want to store X, what do I get rid of to make space?

Cache Algorithms 101

- First In First Out (FIFO)
- Least recently used (LRU)
- Least frequently used (LFU)
- Belady (Offline optimal)

- (Note: all fully associative today)

First-In-First-Out (FIFO)

- Evict objects added to cache longest ago
- Very simple!

- 3 item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g

- Can we do better?

Least Recently Used (LRU)

- Evict object used longest ago
  - “Objects used more recently are more likely to be accessed again”
  - Exploits temporal locality
- Implementation: Update access time for every hit
- 3 item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h

Least Frequently Used (LFU)

- Evict object with fewest hits
  - “Objects used more often are more likely to be accessed again”
  - If tie, use LRU
- Implementation: Update access count for every hit
- 3 item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m

Belady (Offline Optimal Caching)

- What is the best a caching algorithm could do?
- Offline: uses knowledge of the future
  - (Can't use in practice)

- Evict the object with the furthest **next** access time
  - Worst object to keep in the cache

- 3 item cache example:
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m

Effectiveness of Algorithms for CDN Caching?

[Figures from Qi Huang's 2013 SOSP Talk]

Edge Cache with Different Sizes

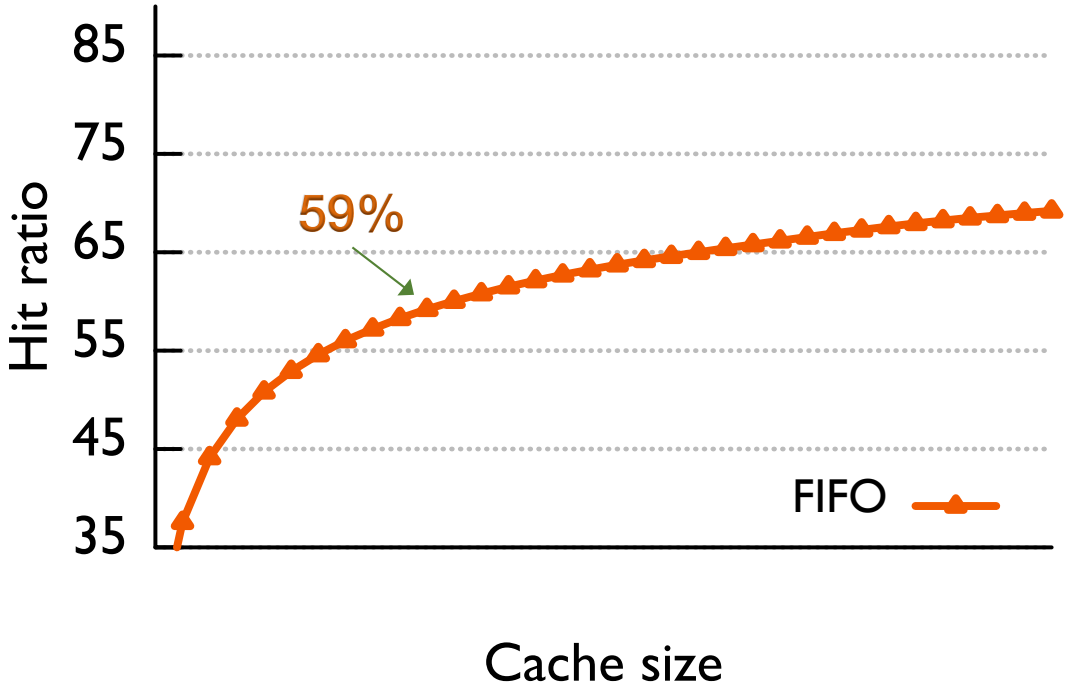

- **Facebook's San Jose CDN edge cache circa 2013**

Edge Cache with Different Sizes

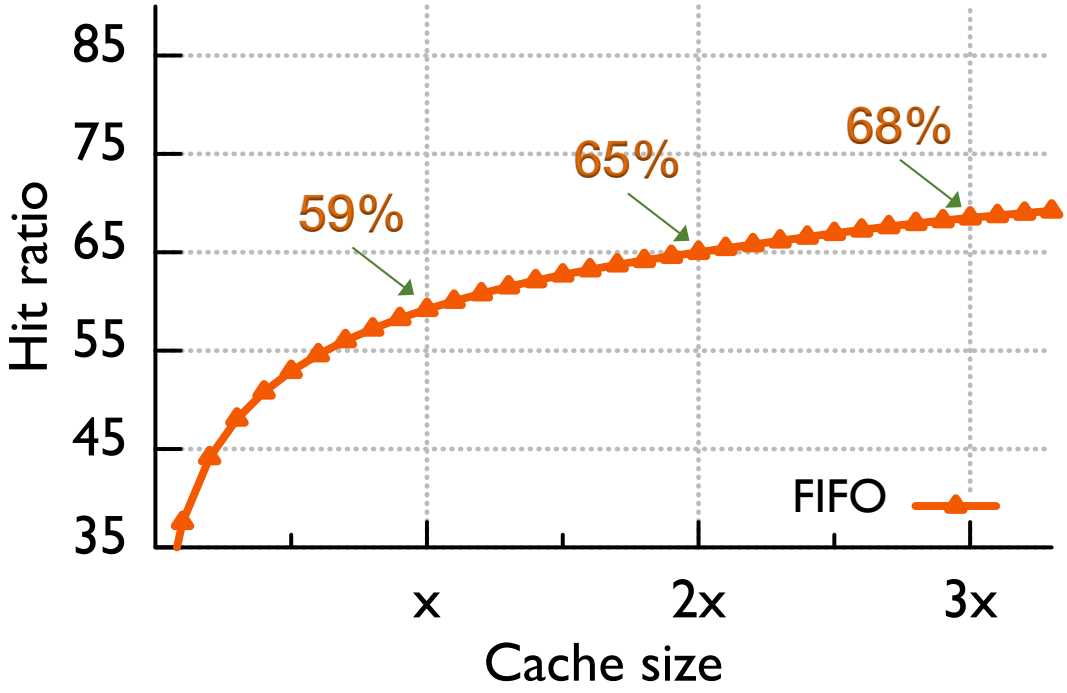

- "x" estimates deployment size (59% hit ratio)

Edge Cache with Different Sizes

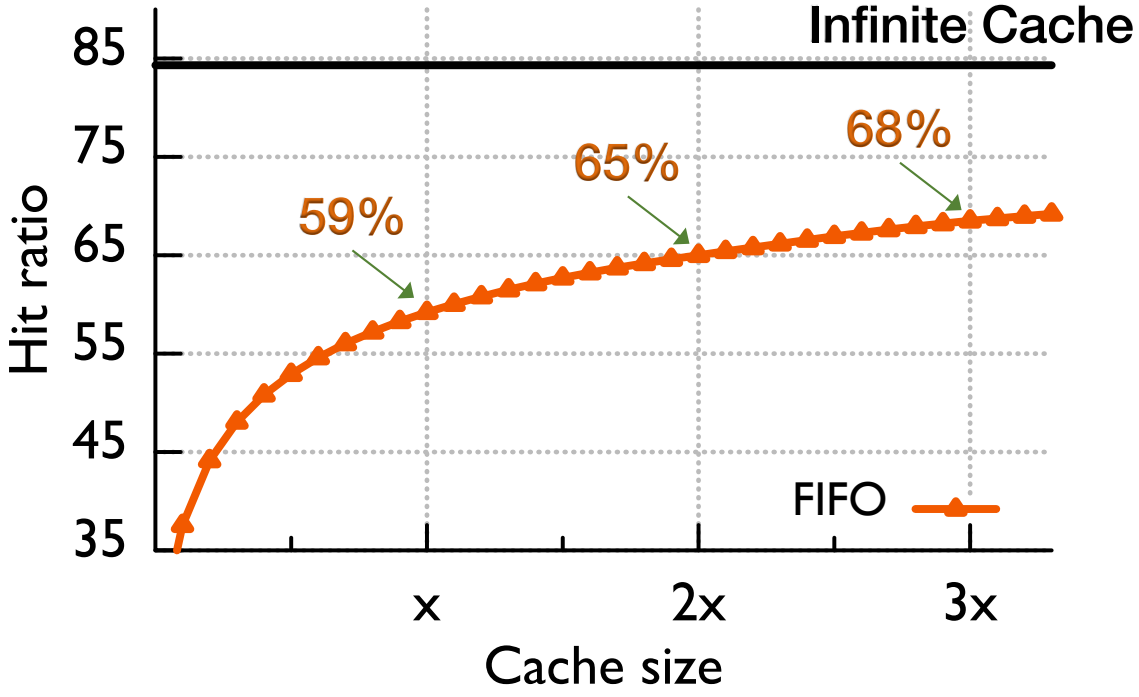

- "Infinite" size ratio needs 45x of capacity

Edge Cache with Different Algos

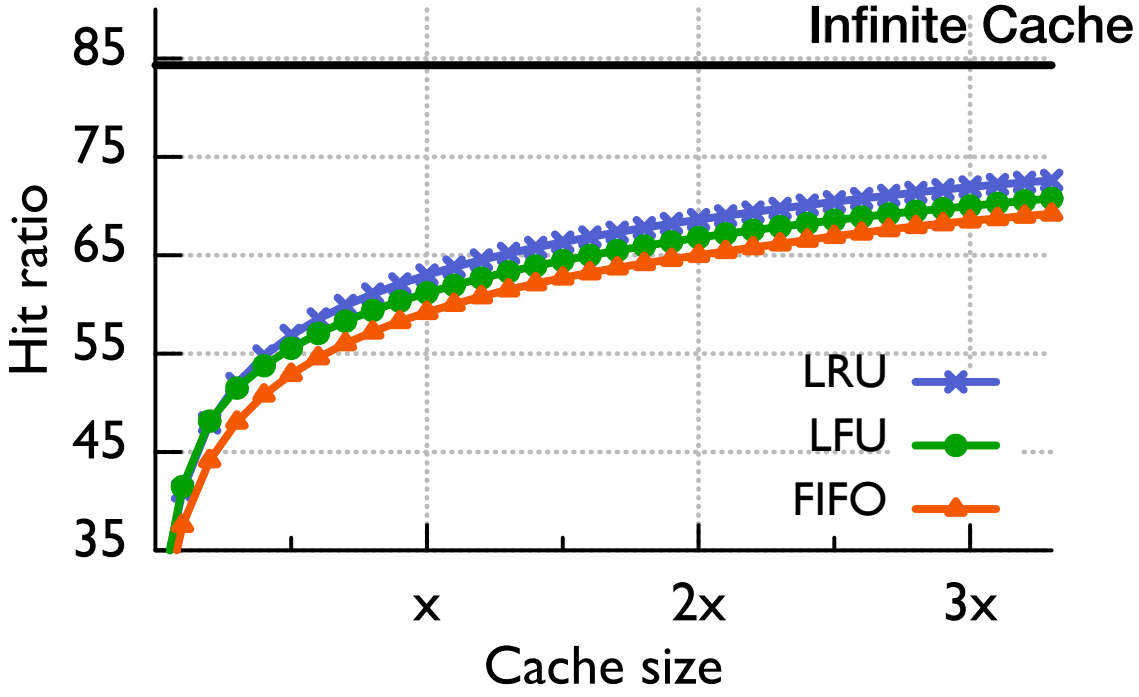

- LRU > LFU > FIFO

Edge Cache with Different Algos

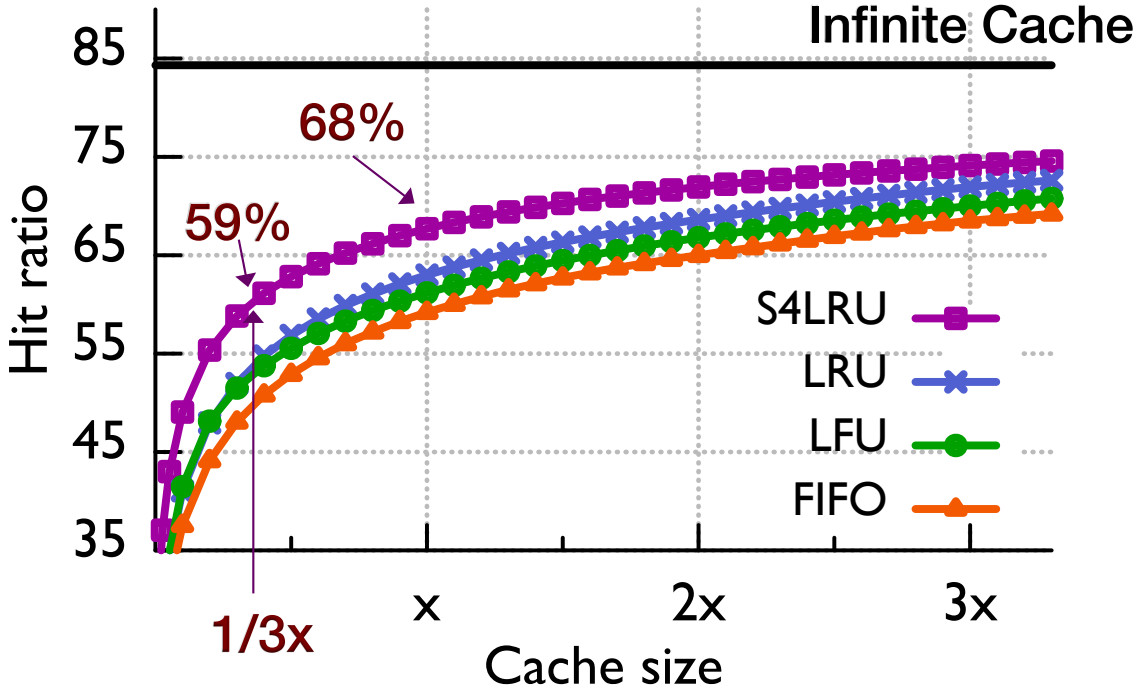

- S4LRU is a more complex algorithm, uses recency and frequency

Edge Cache with Different Algos

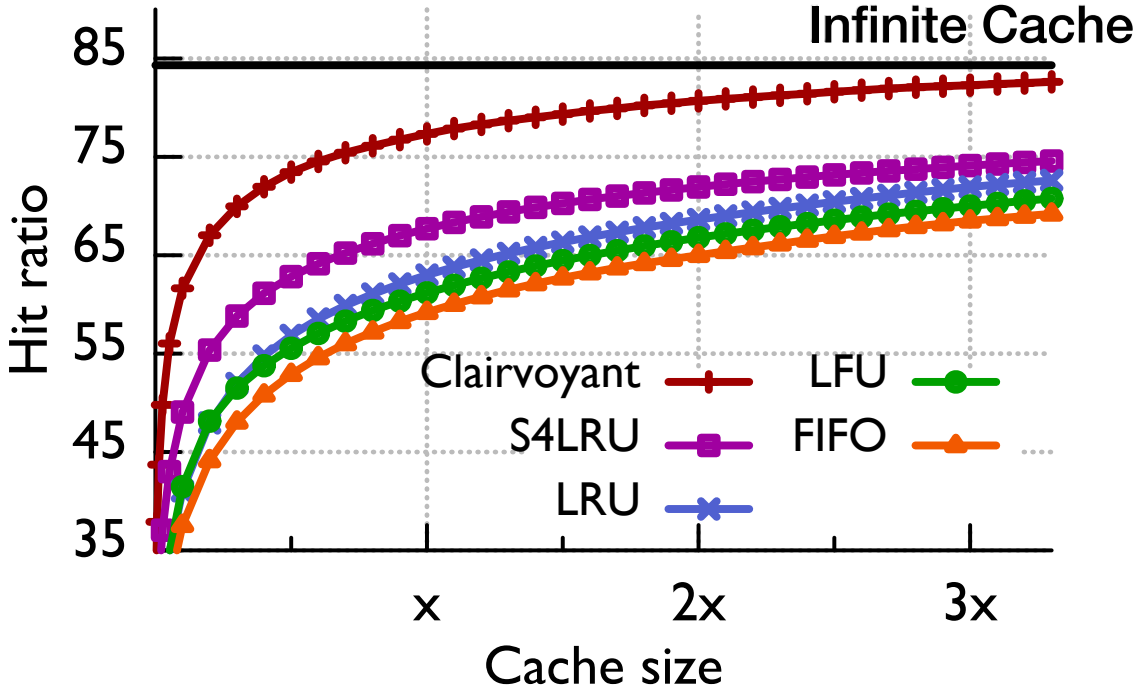

- Clairvoyant (Bélády) shows we can do much better!

A Hierarchy of CDN Caches

[Figures from Qi Huang's 2013 SOSP Talk]

Geo-distributed Edge Cache (FIFO)

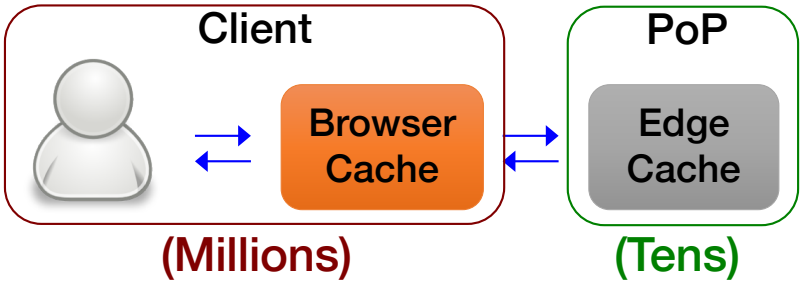

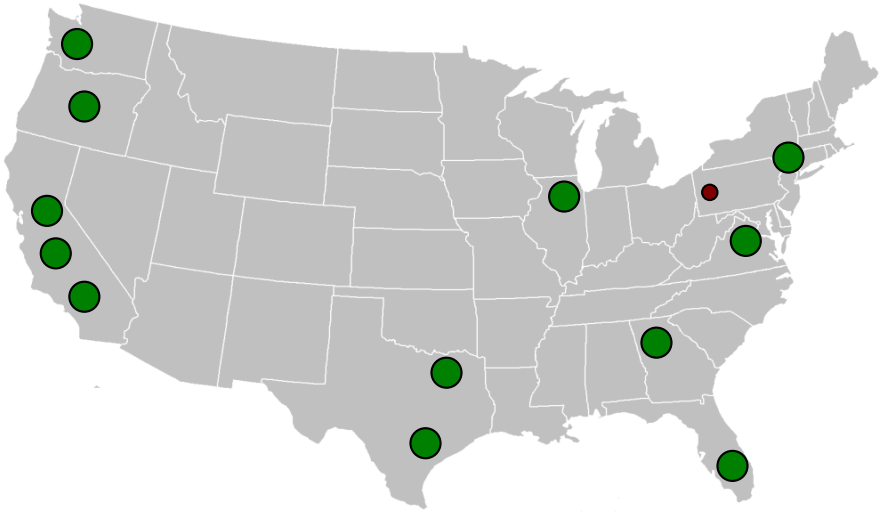

Geo-distributed Edge Cache (FIFO)

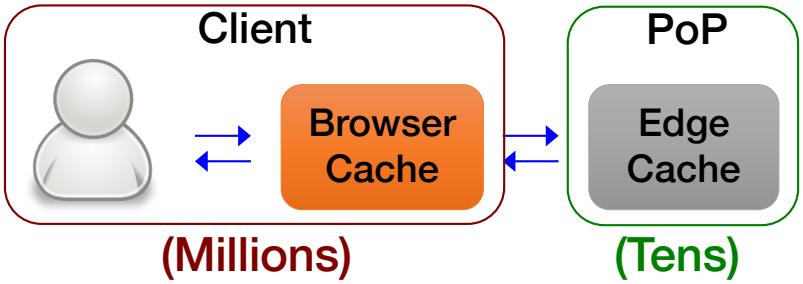

Purpose

1. **Reduce cross-country latency**
2. **Reduce Data Center bandwidth**

Geo-distributed Edge Cache (FIFO)

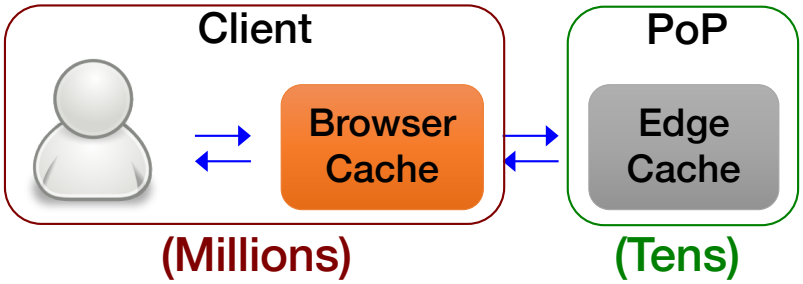

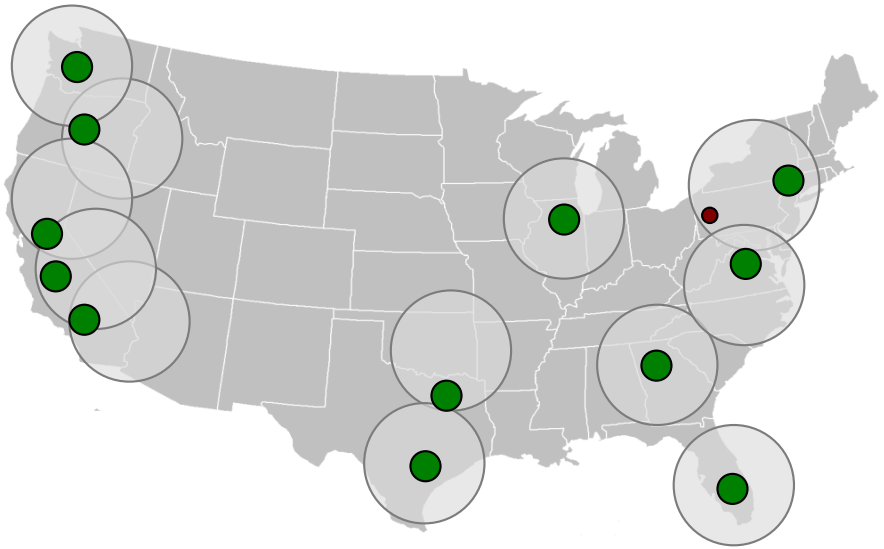

Geo-distributed Edge Cache (FIFO)

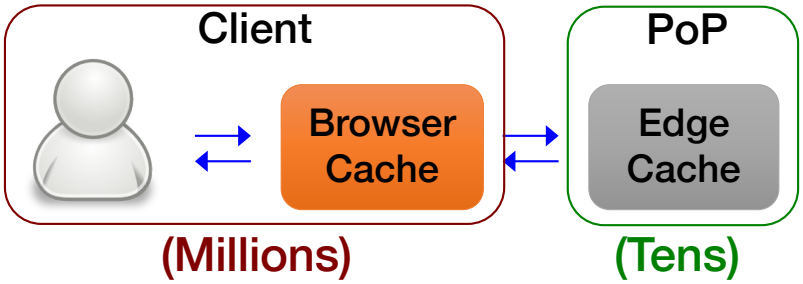

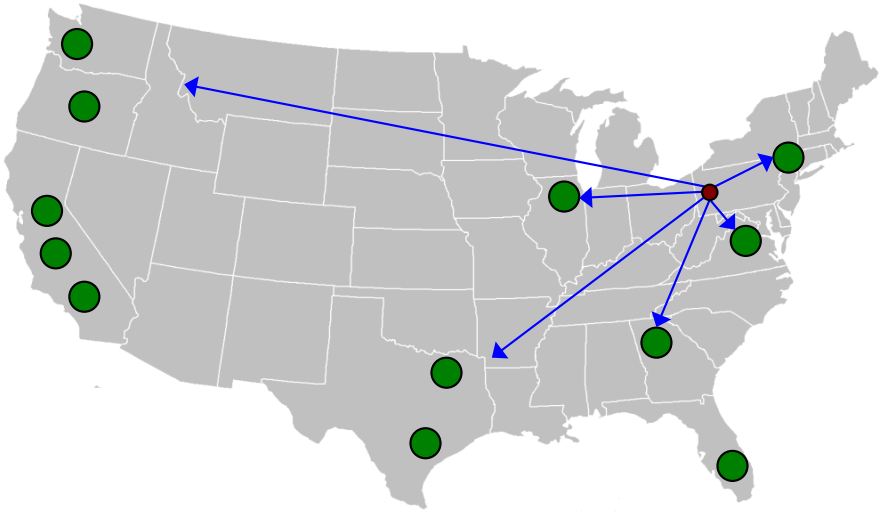

Single Global Origin Cache (FIFO)

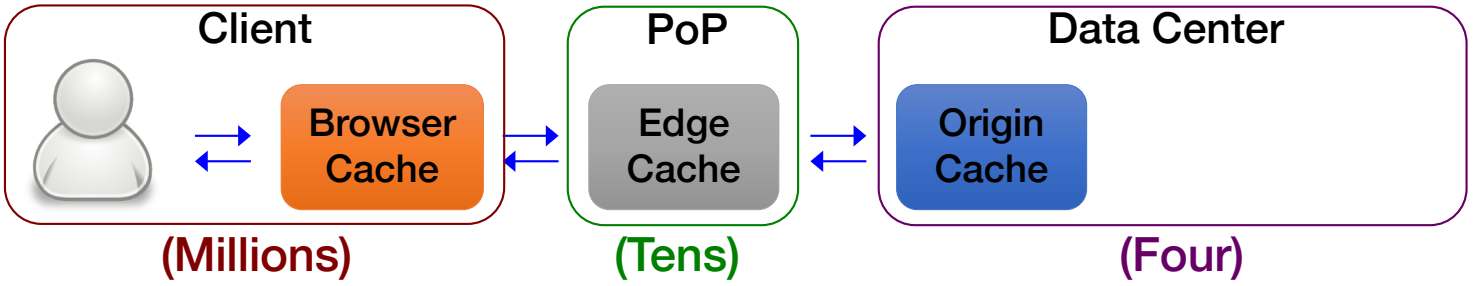

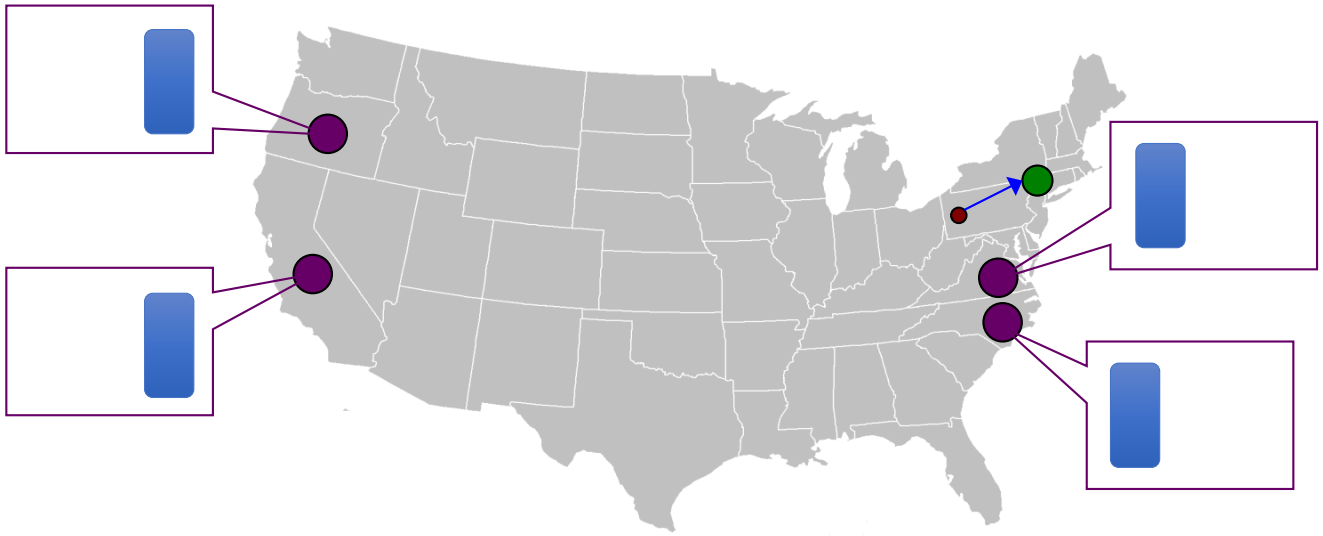

Single Global Origin Cache (FIFO)

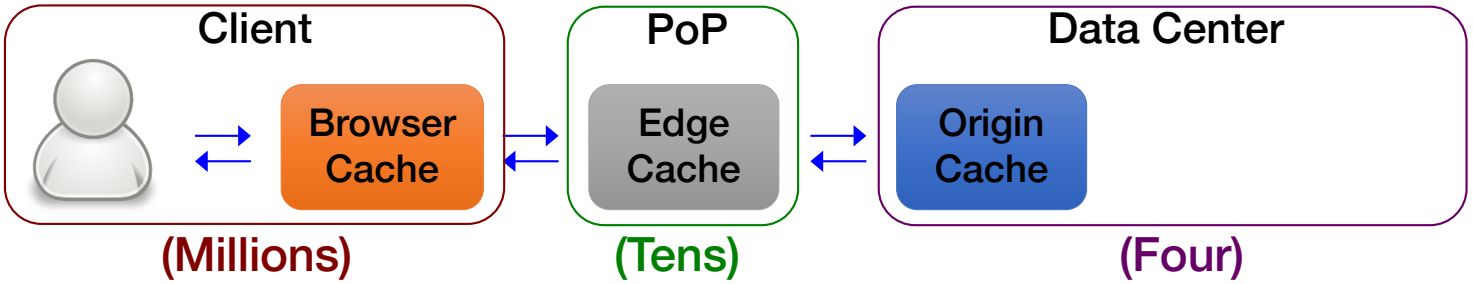

Purpose

1. **Minimize I/O-bound operations**

Single Global Origin Cache (FIFO)

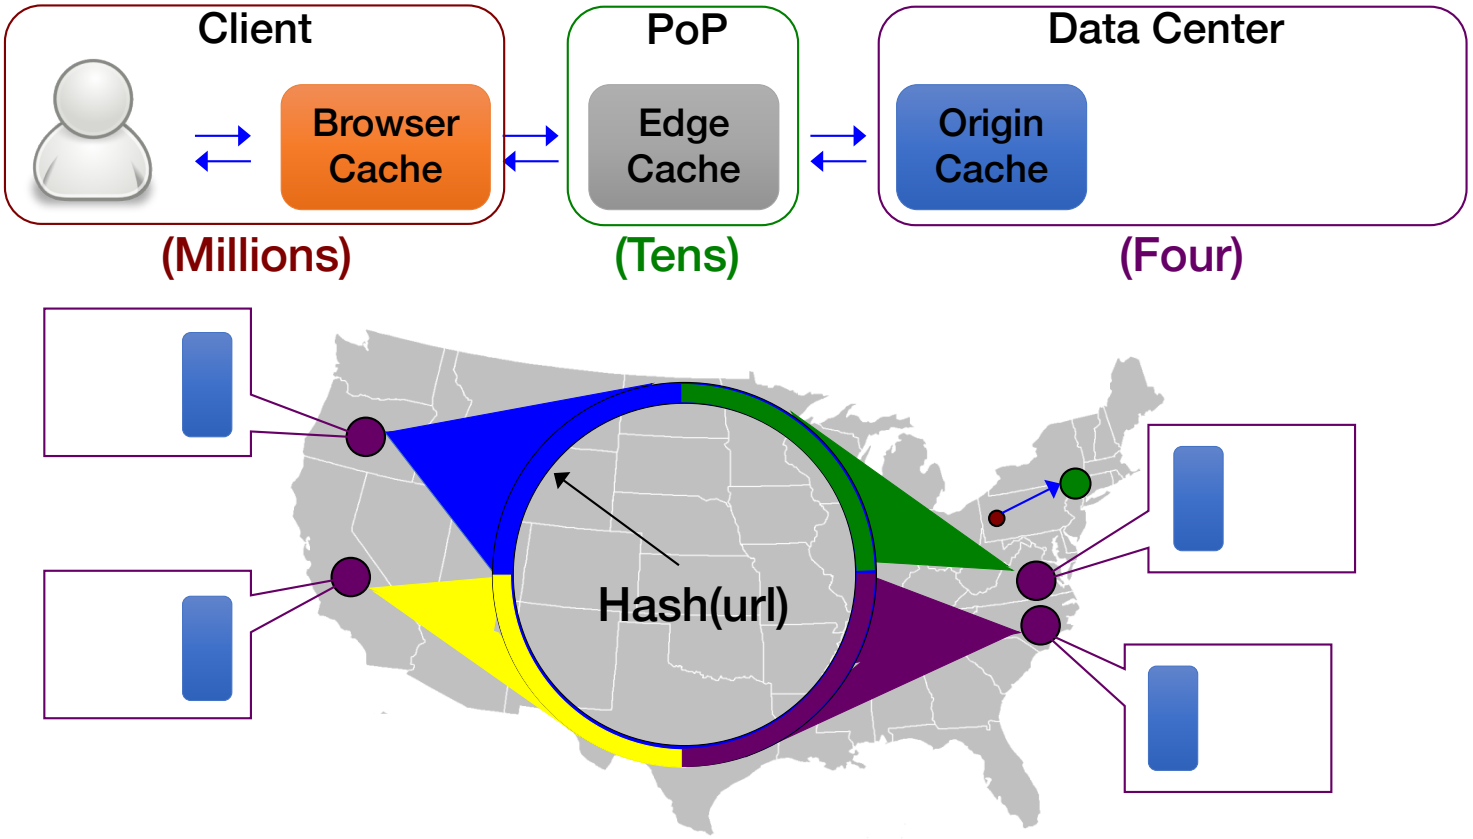

Single Global Origin Cache (FIFO)

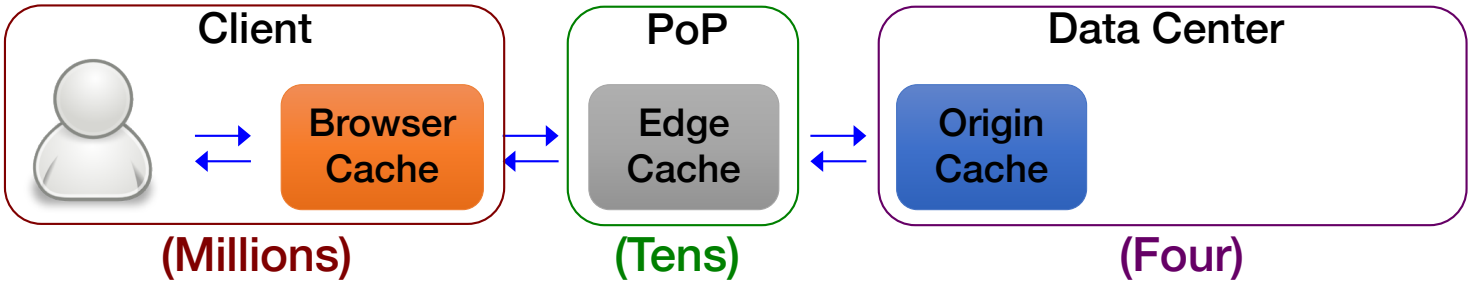

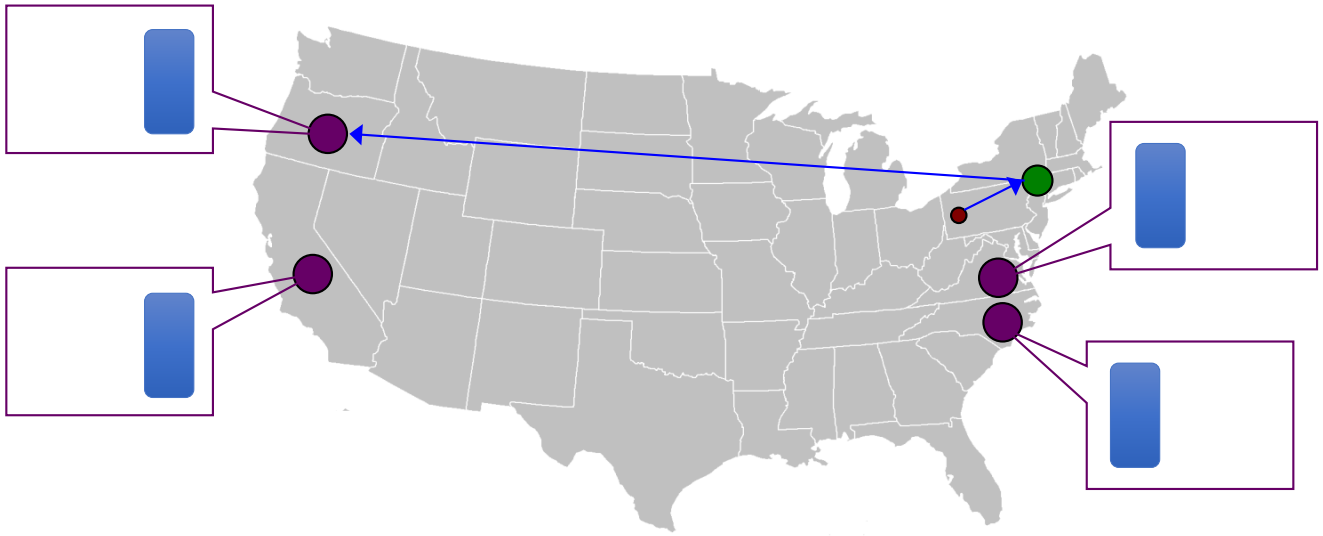

Backend

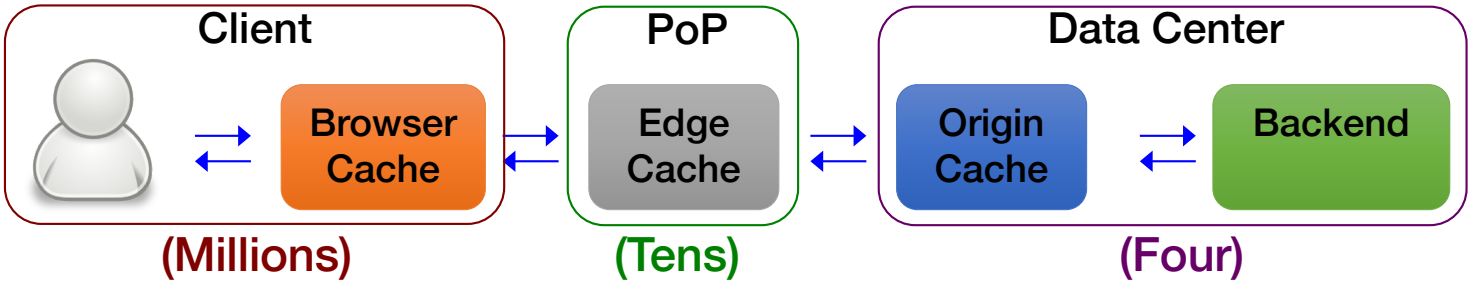
Client
Browser Cache
PoP
Edge Cache
Data Center
Origin Cache
Backend
(Millions)
(Tens)
(Four)

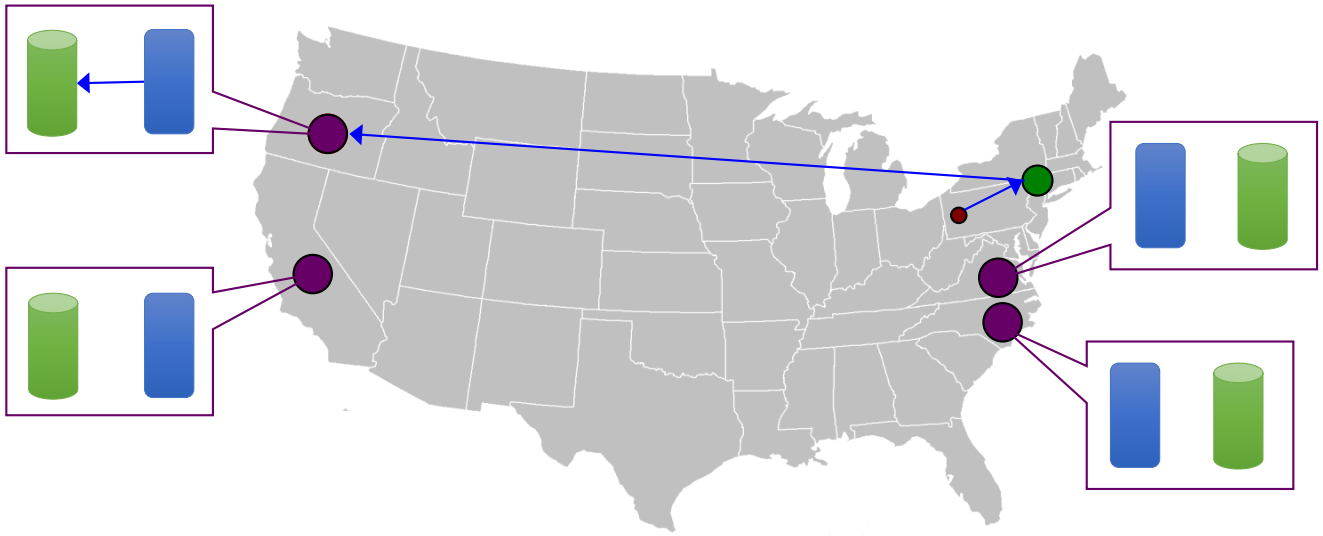

CDN Effectiveness

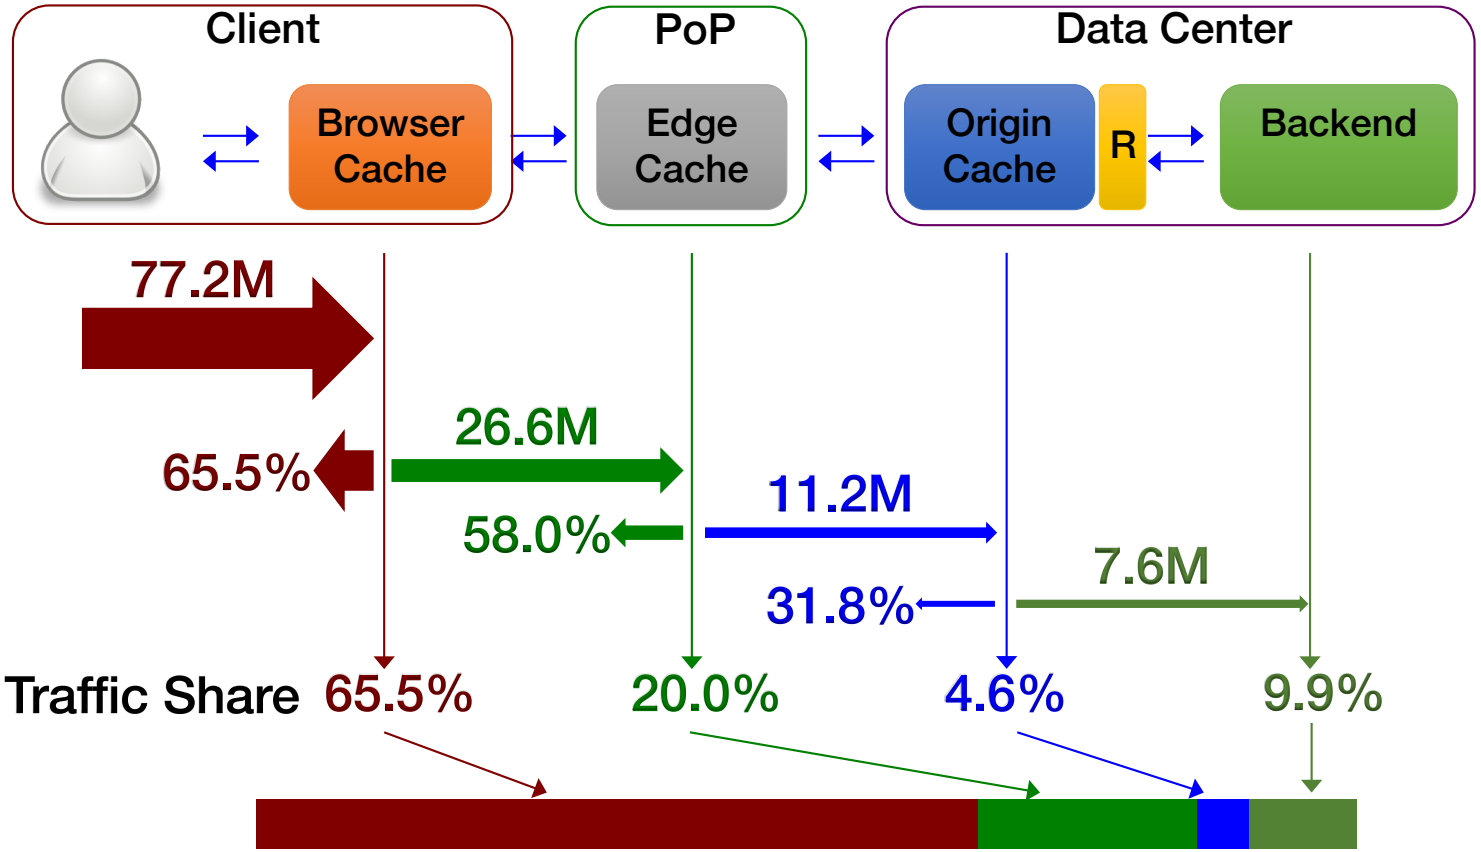

Content Distribution Networks

- Serve “objects” in web pages, and much more, e.g., video segments
- Caching Algorithms: What do we store on them?
  - FIFO, LRU, LFU, Belady, … active area of research
- CDN Hierarchy
  - Smaller / more localized caches → bigger / more general caches