

# Web Caching



COS 316: Principles of Computer System Design

Lecture 9

Amit Levy & Jennifer Rexford

1

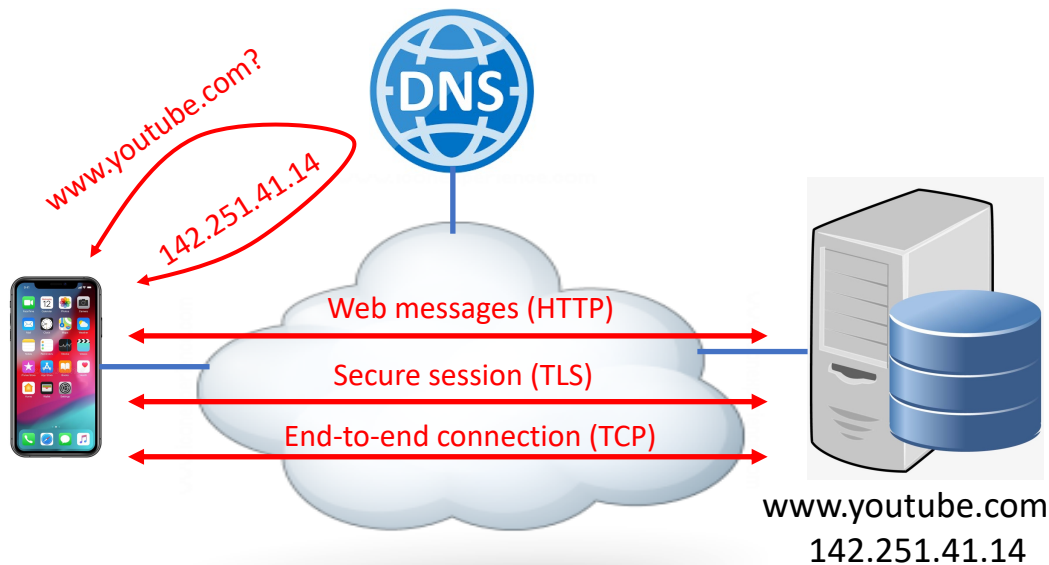
## Downloading a Web Page

User visits <https://www.youtube.com>



2

## Downloading a Web Page (https://www.youtube.com)



3

## Multiple Problems


- User latency
  - Round-trips to query multiple DNS servers
  - Multiple round-trips with the Web server
  - Delivery of a (possibly large) Web item
- Server overhead
  - Handling many requests from many clients
  - Financial costs to deploy enough servers
- Network bandwidth
  - Traffic on many links in multiple networks
  - Financial costs for the affected networks



4

## Caching to the Rescue: Domain Name System


- What to cache?
  - Mapping of popular names to IP addresses
    - E.g., www.youtube.com → 142.251.41.14
  - Mapping of *parts* of names to DNS server IPs
    - E.g., .com top-level domain → 192.26.92.30
- Where to cache?
  - Local DNS server (e.g., for the campus)
  - Client machine (e.g., user's browser)
- How to avoid stale information?
  - Cached entries have a limited "time to live"



5

## Caching to the Rescue: Communication Channel

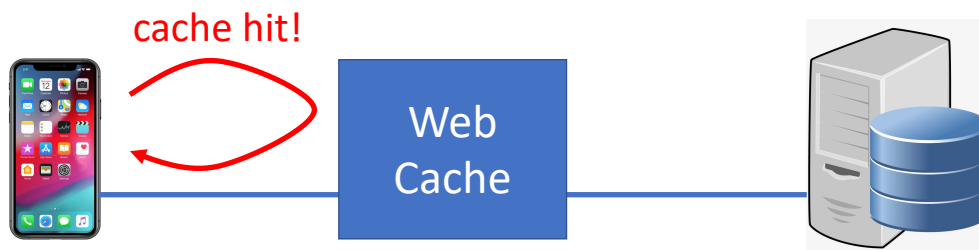
- End-to-end communication
  - TLS: confidentiality, integrity, and authenticity
  - TCP: ordered, reliable delivery of byte stream
- Establishing the channel is expensive
  - Communication delays, creating data structures, and computing keys
- Exploit temporal locality by reusing the channels



6

## Caching to the Rescue: Web Objects

- Cache Web items closer to the client
  - Reduce latency
  - Reduce server overhead
  - Reduce use of network bandwidth



7

## Web Caching: Outline

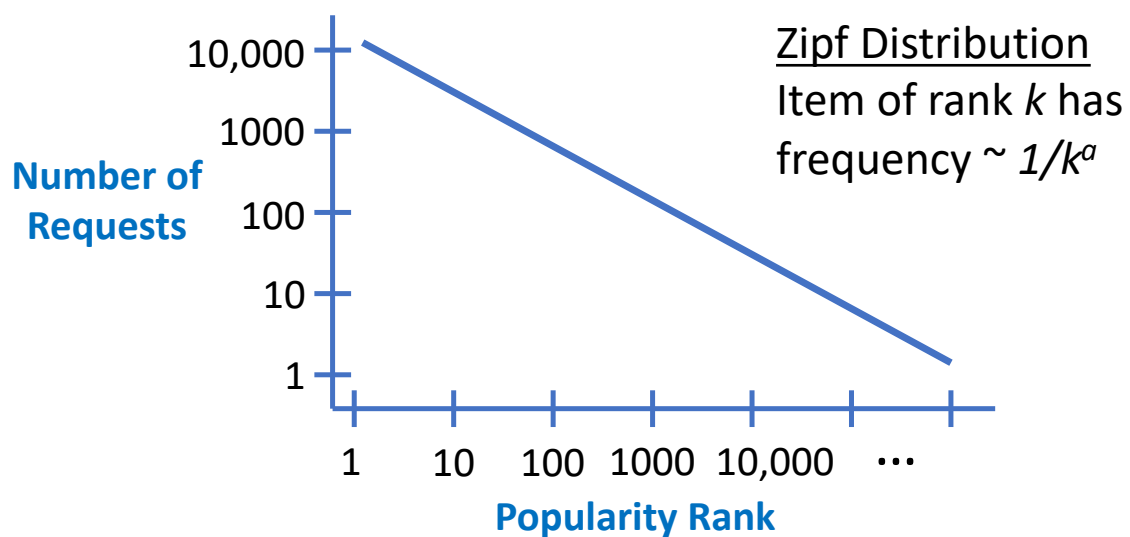
- |   |  |
|---|--|
| <ul style="list-style-type: none"> <li>• Cache replacement               <ul style="list-style-type: none"> <li>• Popularity distributions</li> <li>• Replacement algorithms</li> </ul> </li> <li>• Cache consistency               <ul style="list-style-type: none"> <li>• Dynamic objects</li> <li>• Cache validation</li> </ul> </li> </ul> | <ul style="list-style-type: none"> <li>• Cache placement               <ul style="list-style-type: none"> <li>• Client's web browser</li> <li>• Client's network</li> <li>• Server's network</li> <li>• Third party (CDN)</li> </ul> </li> <li>• Content Distribution Network</li> </ul> |
|---|--|

8

# Cache Replacement

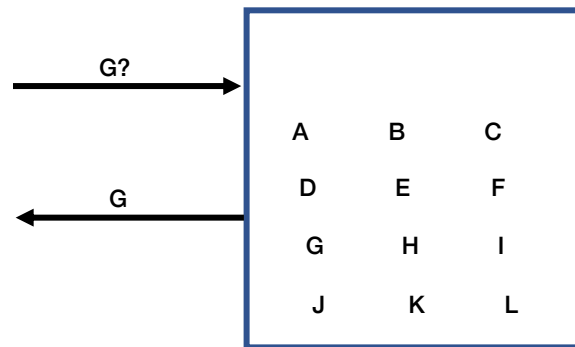
9

Web Caching Should Work Well!



10

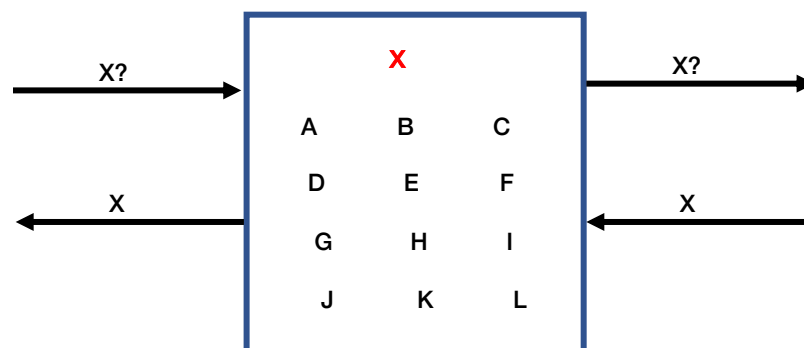
## Web Cache Hit



On cache hit, retrieve the object from the cache!

11

## Web Cache Miss

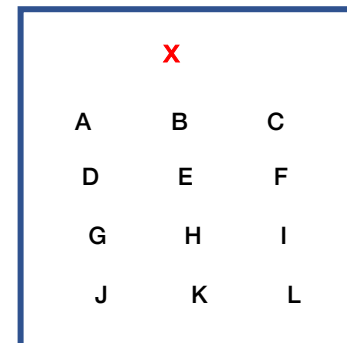


If I want to store X, what do I get rid of to make space?

12

## Cache Replacement Algorithms

- Which object to evict?
  - Least likely to be used again soon
  - Least expensive to fetch again



	X	
A	B	C
D	E	F
G	H	I
J	K	L

- Example algorithms
  - First in first out (FIFO)
  - Least recently used (LRU)
  - Least frequently used (LFU)
- (Note: all fully associative today)

13

## Cache Replacement: First-In-First-Out (FIFO)

- Evict objects added to cache longest ago
- Very simple!
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
- Can we do better?

14

## Least Recently Used (LRU)

- Evict object used longest ago
  - “Objects used more recently are more likely to be accessed again”
  - Exploits temporal locality
- Implementation: Update access time for every hit
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h

15

## Least Frequently Used (LFU)

- Evict object with fewest hits
  - “Objects used more often are more likely to be accessed again”
  - If tie, use LRU
- Implementation: Update access count for every hit
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m

16




## Clairvoyant (Belady): Offline Optimal Caching

- What is the best a caching algorithm could do?
- Offline: uses knowledge of the future
  - (Can't use in practice)
- Evict the object with the furthest **next** access time
  - Worst object to keep in the cache
- Three-item cache example:
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m

17

## Edge Cache with Different Sizes




- Facebook's San Jose CDN edge cache circa 2013

[Figures from Qi Huang's 2013 SOSP Talk]

18

18

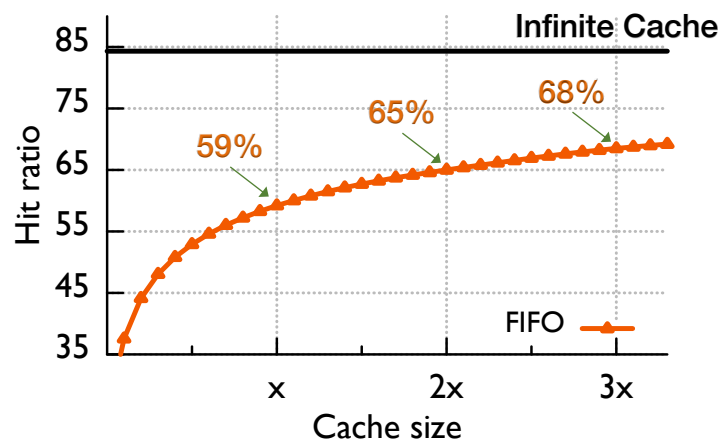
## Edge Cache with Different Sizes



19

19

## Edge Cache with Different Sizes




- “Infinite” size ratio needs 45x of capacity

20

20

## Edge Cache with Different Algos

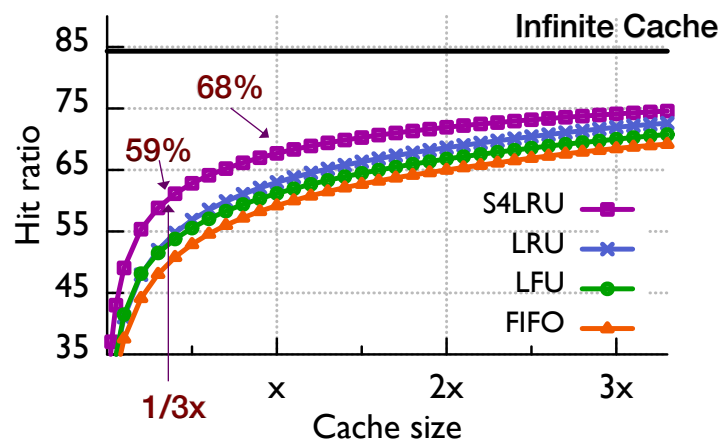


- LRU > LFU > FIFO

21

21

## Edge Cache with Different Algorithms

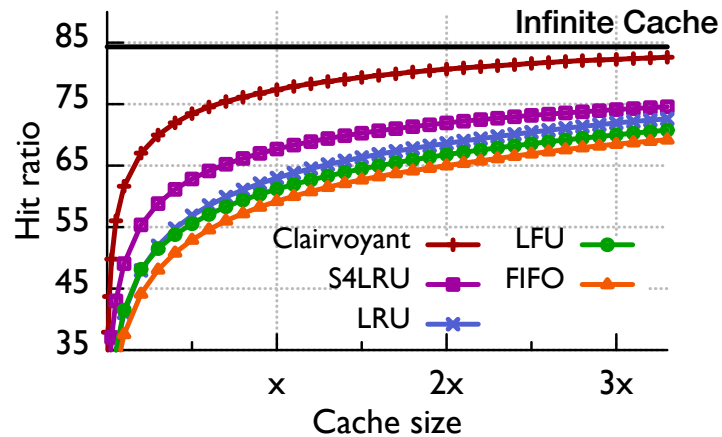


- S4LRU is a more complex algorithm, uses recency and frequency

22

22

## Edge Cache with Different Algos



- Clairvoyant (Bélády) shows we can do much better!

23

23

## Cache Consistency

24

## Some Web Content is Not Cacheable

- Dynamic content
  - E.g., stock prices, scores, web cams
- Content generated by scripts
  - Results depend on the specific parameters
  - E.g., <https://www.google.com/search?q=php+script+url>
- Personalized content
  - E.g., based on cookie sent by the browser
- Encrypted content
  - Cannot decrypt without the appropriate key

Last Updated 9:19pm EST

^IXIS	4,445.89	-19.43	-0.4%
NASDAQ INSURANCE			
^GSPC	1,367.59	+1.85	+0.1%
S&P 500			
AAPL	525.76	+3.35	+0.6%
APPLE INC.			
T	30.36	+0.02	+0.1%
AT&T INC.			
GOLD	114.86	-0.87	-0.8%
RANDGOLD RESOURCE			
V	116.86	-0.68	-0.6%
VISA INC.			
YHOO	14.86	-0.03	-0.2%
YAHOO! INC.			

Powered by EduLifeLine.com



25

## Cache Consistency Challenges



Shared  
Cache



### Web cache needs to know

- Whether to cache an item
- How long to cache an item
- Whether to check an item's freshness
- Whether it is okay to return a stale item
- Whether the item has sensitive data

### Server knows the content

- Whether the item is dynamic
- How often the item changes
- Whether the item has changed
- Whether stale information is useful
- Whether item contains sensitive data

**Scalability challenge: the server cannot remember everyone has cached a item**

26

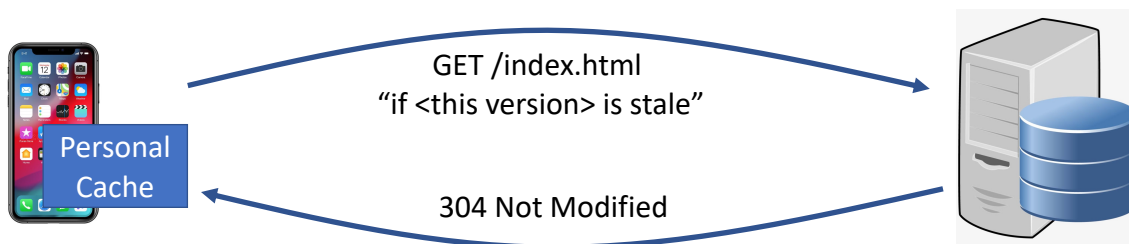
## HTTP Response Messages for Cache Control

- Whether to cache
  - no store: no cache should store it
- Who should cache
  - private: only a private cache (e.g., browser)
  - public: any cache, including shared ones
- How long to cache
  - max-age=N: for N seconds
  - must-revalidate: check with the server (don't return stale item)

**Cache-Control: public, max-age=604800, must-revalidate**

27

## Cache Validation: Client Checks Freshness



### How do they identify the “version”?

- Timestamp
  - When the item was modified by the server
  - E.g., Last-Modified: Wed, 21 Oct 2015 07:28:00 GMT
- Version number
  - Entity tag provided by the server
  - E.g., ETag: "33a64df551425fcc55e4d42a148795d9f25f89d4"

28

# Cache Placement

29

## Client Machine (e.g., Browser)

### Advantages

- Very low latency
- Preserves access bandwidth
- Available when disconnected

### Disadvantages

- Low hit rate due to “cold” misses
- Many cache consistency checks
- Incomplete logs at the server



Personal  
Cache



30

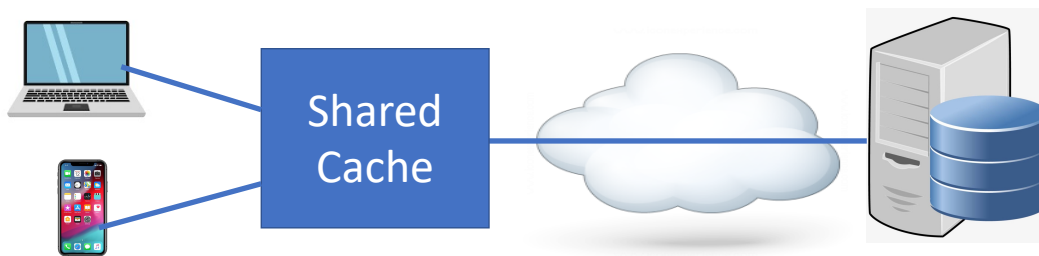
## Client Network (Forward Proxy Cache)

### Advantages

- Low latency
- Preserves enterprise bandwidth
- Hits for locally popular content

### Disadvantages

- Cost to deploy the cache
- Many consistency checks
- Incomplete logs at the server



31

## Server Network (Reverse Proxy Cache)

### Advantages

- High hit rate across global users
- Greater cooperation with server
- Complete request logs for server
- Preserves server bandwidth

### Disadvantages

- Costs to deploy the cache
- Does not reduce latency much
- Consumes wide-area bandwidth




32



## Content Distribution Network (CDN)

- Outsourced caching infrastructure
  - Caching for clients and servers
  - Dedicated equipment and software
  - Trained staff, best practices, etc.
- Coordination with the server
  - Generating non-cacheable content
  - Providing detailed measurement data
- Smart cache placement
  - Many caches: handle large request load
  - Close to many clients: reduce latency




More than 4200 locations in 135 countries

33


## CDN Challenges

- Where to place edge sites?
  - Close to many clients, with reasonable cost
- Where to replicate a server's content?
  - Many edge sites → duplicated data
  - Few edge sites → larger client latency
- How to direct a client to an edge site?
  - Proximity: for low latency
  - Light load: to reduce congestion
- How to manage each cache?
  - Maximize hit rate?
  - Minimize miss penalty?
  - Fairness across origin servers?



34

## CDN Effectiveness



35

## Conclusions

- Downloading a Web page
  - Name resolution, transport connection, secure session, web messages
- Benefits of caching
  - Reduces user latency, server load, and network bandwidth
- Cache replacement
  - Maximize hit rate by trying to predict the future
- Cache consistency
  - Efficient ways to avoid returning unnecessarily stale responses
- Content distribution networks
  - Caching close to clients, while working on behalf of the servers

36