

# Web Caching




COS 316: Principles of Computer System Design

Lecture 9

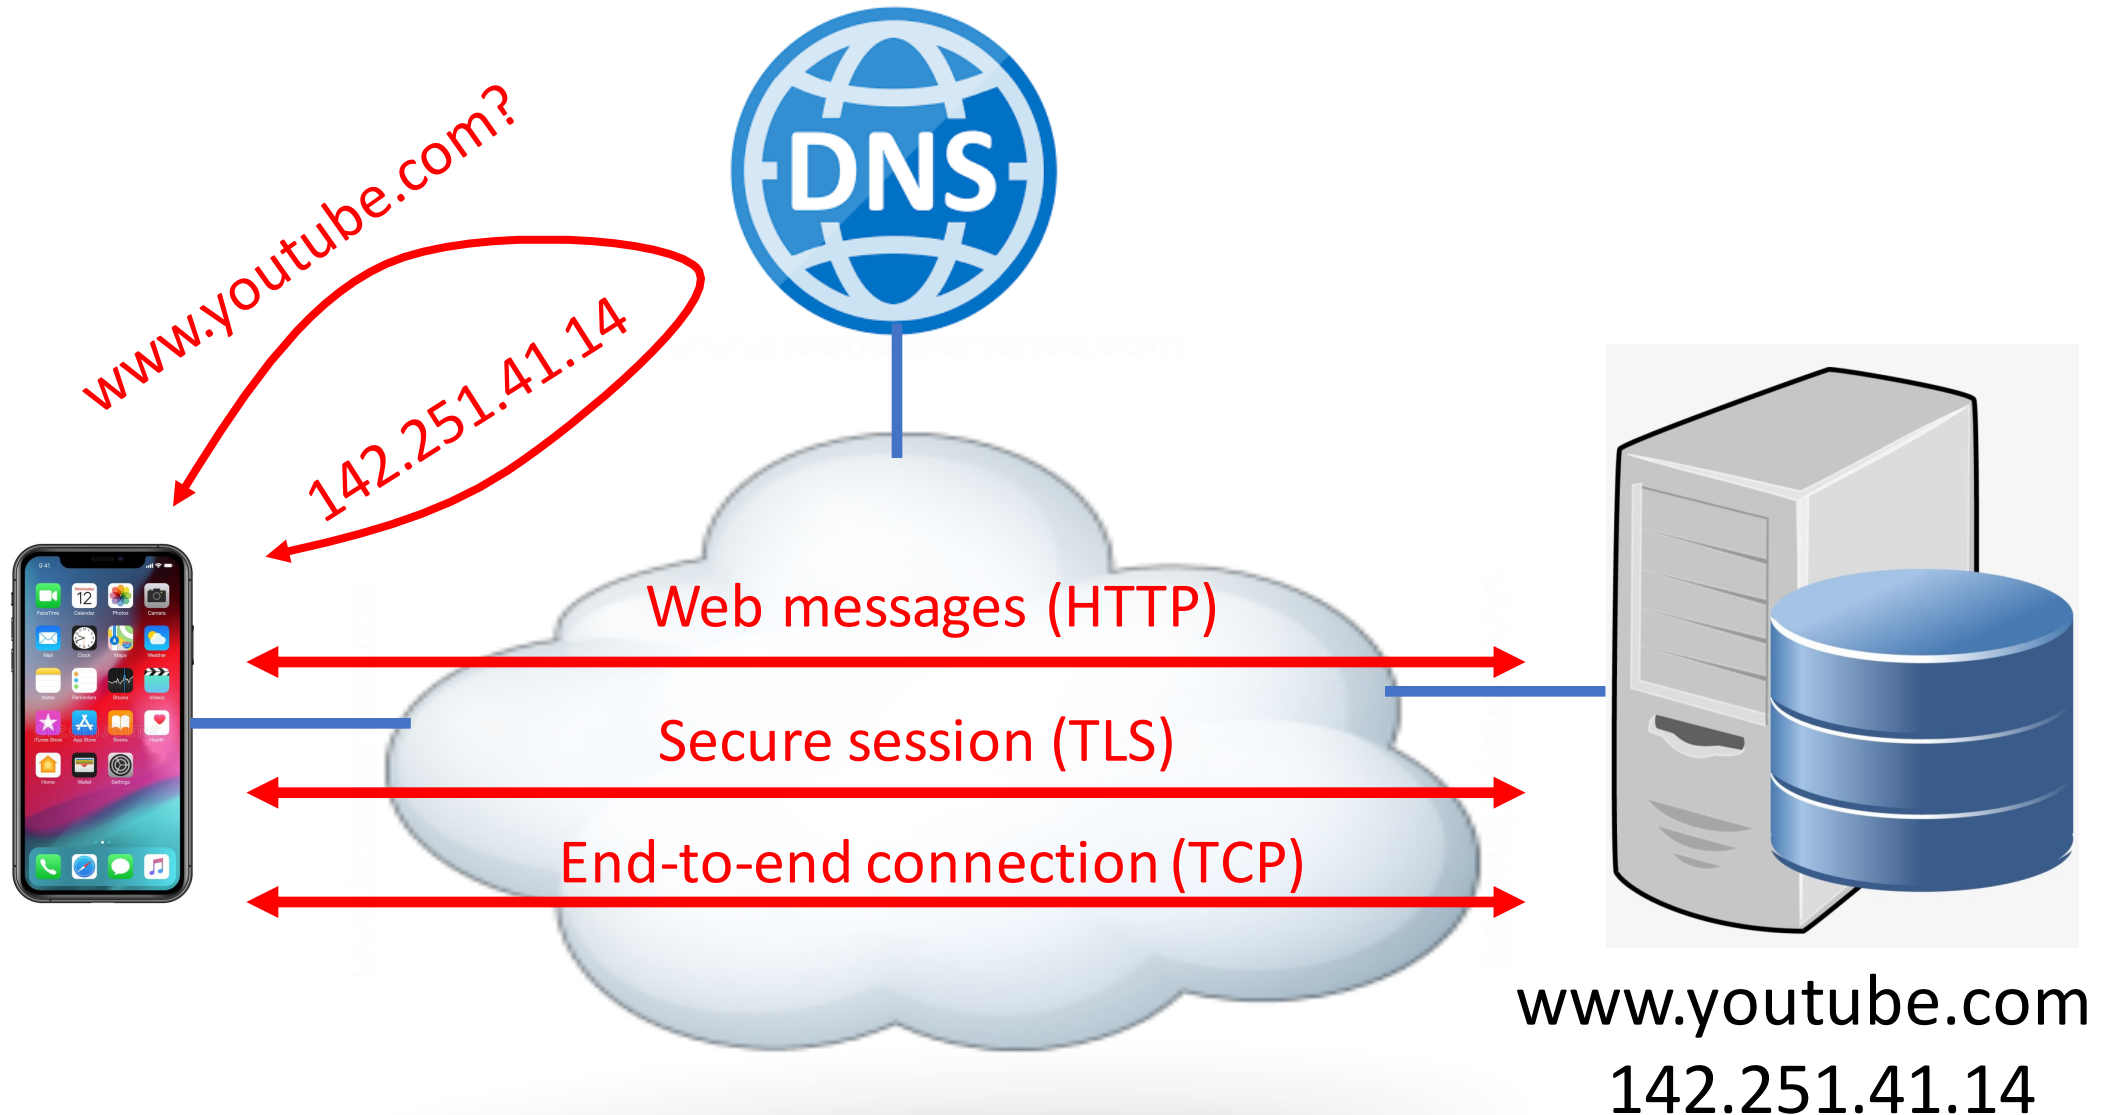
Amit Levy & Ravi Netravali

# Downloading a Web Page

**User visits <https://www.youtube.com>**



# Downloading a Web Page (https://www.youtube.com)



# Multiple Problems

- User latency
  - Round-trips to query multiple DNS servers
  - Multiple round-trips with the Web server
  - Delivery of a (possibly large) Web item
- Server overhead
  - Handling many requests from many clients
  - Financial costs to deploy enough servers
- Network bandwidth
  - Traffic on many links in multiple networks
  - Financial costs for the affected networks




# Caching to the Rescue: Domain Name System

- What to cache?
  - Mapping of popular names to IP addresses
    - E.g., `www.youtube.com` → `142.251.41.14`
  - Mapping of *parts* of names to DNS server IPs
    - E.g., `.com` top-level domain → `192.26.92.30`


# Caching to the Rescue: Domain Name System

- What to cache?
  - Mapping of popular names to IP addresses
    - E.g., `www.youtube.com` → `142.251.41.14`
  - Mapping of *parts* of names to DNS server IPs
    - E.g., `.com` top-level domain → `192.26.92.30`




# Caching to the Rescue: Domain Name System

- What to cache?
  - Mapping of popular names to IP addresses
    - E.g., `www.youtube.com` → `142.251.41.14`
  - Mapping of *parts* of names to DNS server IPs
    - E.g., `.com` top-level domain → `192.26.92.30`
- Where to cache?
  - Local DNS server (e.g., for the campus)
  - Client machine (e.g., user's browser)
- How to avoid stale information?
  - Cached entries have a limited “time to live”



# Caching to the Rescue: Communication Channel


- End-to-end communication
  - TLS: confidentiality, integrity, and authenticity
  - TCP: ordered, reliable delivery of byte stream
- Establishing the channel is expensive
  - Communication delays, creating data structures, and computing keys
- Exploit temporal locality by reusing the channels





# Caching to the Rescue: Web Items

- Cache Web items closer to the client
  - Reduce latency
  - Reduce server overhead
  - Reduce use of network bandwidth




# Web Caching: Outline


- Cache replacement
  - Popularity distributions
  - Replacement algorithms
- Cache consistency
  - Dynamic items
  - Cache validation
- Cache placement
  - Client's web browser
  - Client's network
  - Server's network
  - Third party (CDN)
- Content Distribution Network

# Cache Replacement

# Web Caching Should Work Well!




# Web Cache Hit



On cache hit, retrieve the object from the cache!


# Web Cache Miss



If I want to store X, what do I get rid of to make space?

# Cache Replacement Algorithms

- Which object to evict?
  - Least likely to be used again soon
  - Least expensive to fetch again
- Example algorithms
  - First in first out (FIFO)
  - Least recently used (LRU)
  - Least frequently used (LFU)
- (Note: all fully associative today)



# Cache Replacement: First-In-First-Out (FIFO)

- Evict objects added to cache longest ago
- Very simple!
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
- Can we do better?



# Least Recently Used (LRU)

- Evict object used longest ago
  - “Objects used more recently are more likely to be accessed again”
  - Exploits temporal locality
- Implementation: Update access time for every hit
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h


# Least Frequently Used (LFU)

- Evict object with fewest hits
  - “Objects used more often are more likely to be accessed again”
  - If tie, use LRU
- Implementation: Update access count for every hit
- Three-item cache example:
  - Request stream: a, b, a, c, a, d, a, e, a, f, g
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m


# Clairvoyant (Belady): Offline Optimal Caching

- What is the best a caching algorithm could do?
- Offline: uses knowledge of the future
  - (Can't use in practice)
- Evict the object with the furthest **next** access time
  - Worst object to keep in the cache
- Three-item cache example:
  - Request stream: h, h, h, i, j, k, h
  - Request stream: l, l, m, n, o, m


# Edge Cache with Different Sizes



# Edge Cache with Different Sizes




# Edge Cache with Different Sizes




- “Infinite” size ratio needs 45x of capacity

# Edge Cache with Different Algos



- **LRU** > **LFU** > **FIFO**


# Edge Cache with Different Algorithms



- **S4LRU** is a more complex algorithm, uses recency and frequency



# Edge Cache with Different Algos



- Clairvoyant (Bélády) shows we can do much better!

# Cache Consistency

# Some Web Content is Not Cacheable

- Dynamic content
  - E.g., stock prices, scores, web cams
- Content generated by scripts
  - Results depend on the specific parameters
  - E.g., <https://www.google.com/search?q=php+script+url>
- Personalized content
  - E.g., based on cookie sent by the browser
- Encrypted content
  - Cannot decrypt without the appropriate key

Last Updated 9:19pm EST  

^IXIS	4,445.89	-19.43	-0.4%
NASDAQ INSURANCE			
^GSPC	1,367.59	+1.85	+0.1%
S&P 500			
AAPL	525.76	+3.35	+0.6%
APPLE INC.			
T	30.36	+0.02	+0.1%
AT&T INC.			
GOLD	114.86	-0.87	-0.8%
RANDGOLD RESOURCE			
V	116.86	-0.68	-0.6%
VISA INC.			
YHOO	14.86	-0.03	-0.2%
YAHOO! INC.			

Powered by [EduLifeLine.com](http://EduLifeLine.com)



# Cache Consistency Challenges



## Web cache needs to know

- Whether to cache an item
- How long to cache an item
- Whether to check an item's freshness
- Whether it is okay to return a stale item
- Whether the item has sensitive data

# Cache Consistency Challenges



## Web cache needs to know

- Whether to cache an item
- How long to cache an item
- Whether to check an item's freshness
- Whether it is okay to return a stale item
- Whether the item has sensitive data

## Server knows the content

- Whether the item is dynamic
- How often the item changes
- Whether the item has changed
- Whether stale information is useful
- Whether item contains sensitive data

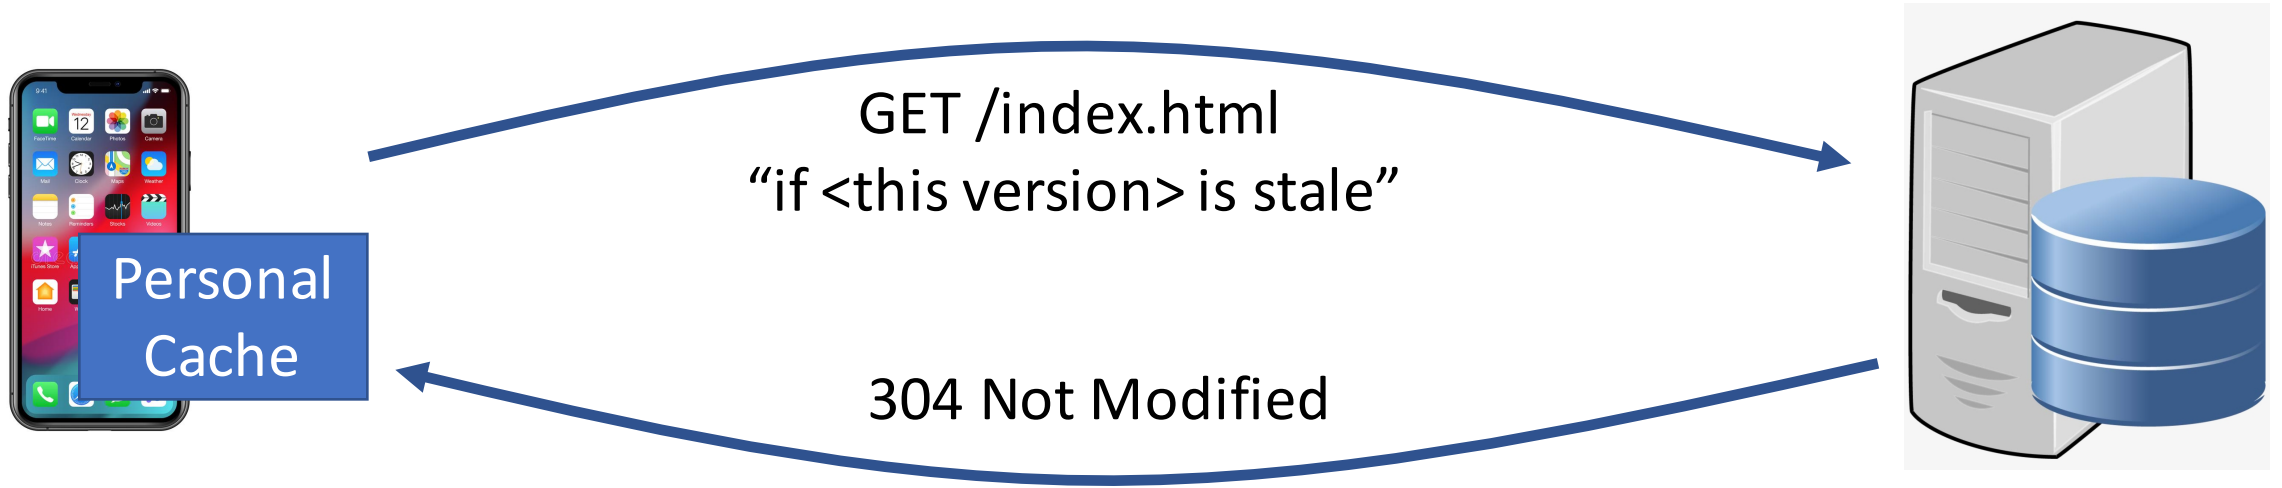
**Scalability challenge: the server cannot remember every client that has cached an item**

# HTTP Response Header for Cache Control

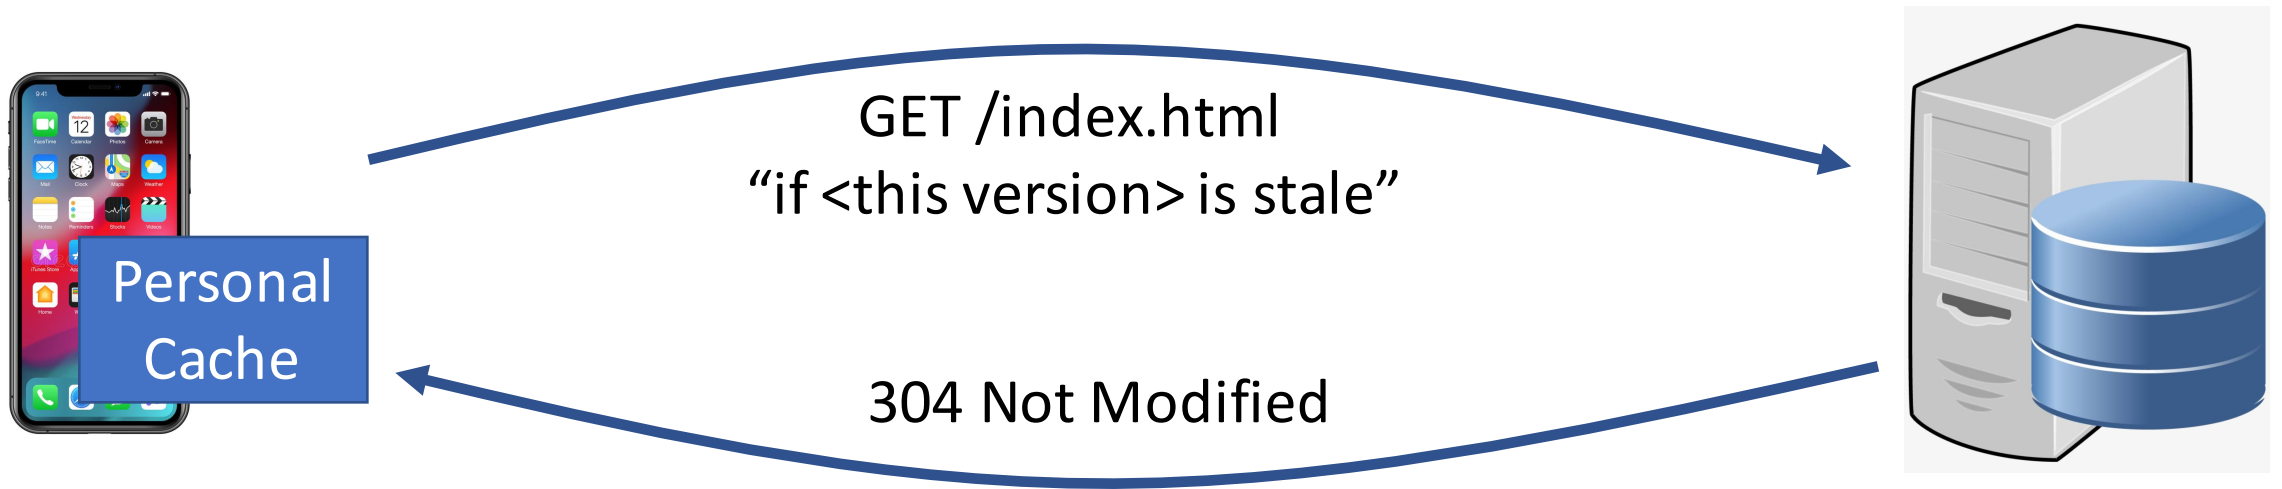
- Whether to cache
  - no store: no cache should store it
- Who should cache
  - private: only a private cache (e.g., browser)
  - public: any cache, including shared ones
- How long to cache
  - max-age=N: for N seconds
  - must-revalidate: check with the server (don't return stale item)

**Cache-Control: public, max-age=86400, must-revalidate**

# Cache Validation: Client Checks Freshness



# Cache Validation: Client Checks Freshness



## How do they identify the “version”?

- Timestamp
  - When the item was modified by the server
  - E.g., Last-Modified: Wed, 21 Oct 2015 07:28:00 GMT
- Version number
  - Entity tag provided by the server
  - E.g., ETag: "33a64df551425fcc55e4d42a148795d9f25f89d4"



# Cache Placement

# Client Machine (e.g., Browser)

## Advantages

- Very low latency
- Preserves access bandwidth
- Available when disconnected

## Disadvantages

- Low hit rate due to “cold” misses
- Many cache consistency checks
- Incomplete logs at the server




# Client Network (Forward Proxy Cache)

## Advantages

- Low latency
- Preserves enterprise bandwidth
- Hits for locally popular content

## Disadvantages

- Cost to deploy the cache
- Many consistency checks
- Incomplete logs at the server



# Server Network (Reverse Proxy Cache)

## Advantages

- High hit rate across global users
- Greater cooperation with server
- Complete request logs for server
- Preserves server bandwidth

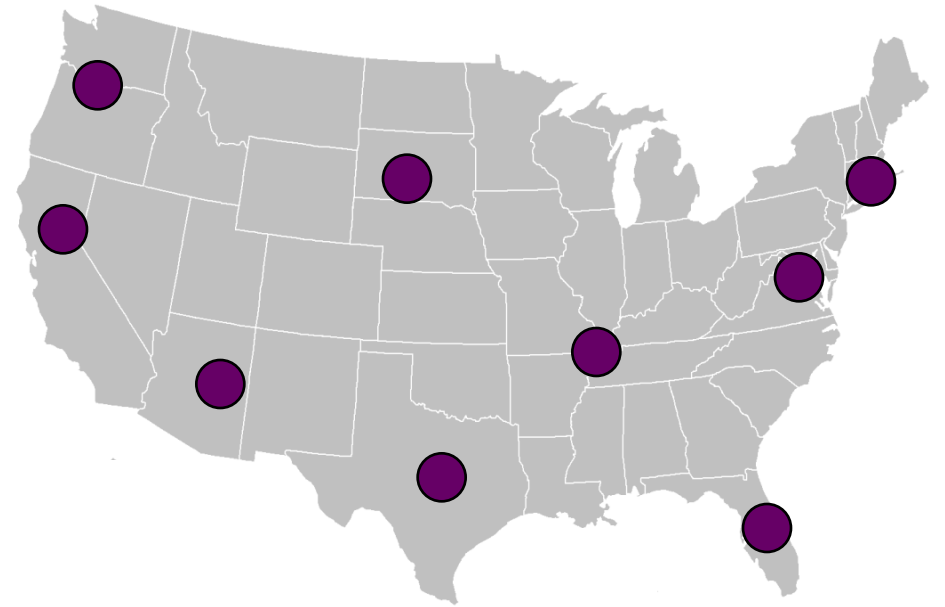
## Disadvantages

- Costs to deploy the cache
- Does not reduce latency much
- Consumes wide-area bandwidth



# Content Distribution Network (CDN)

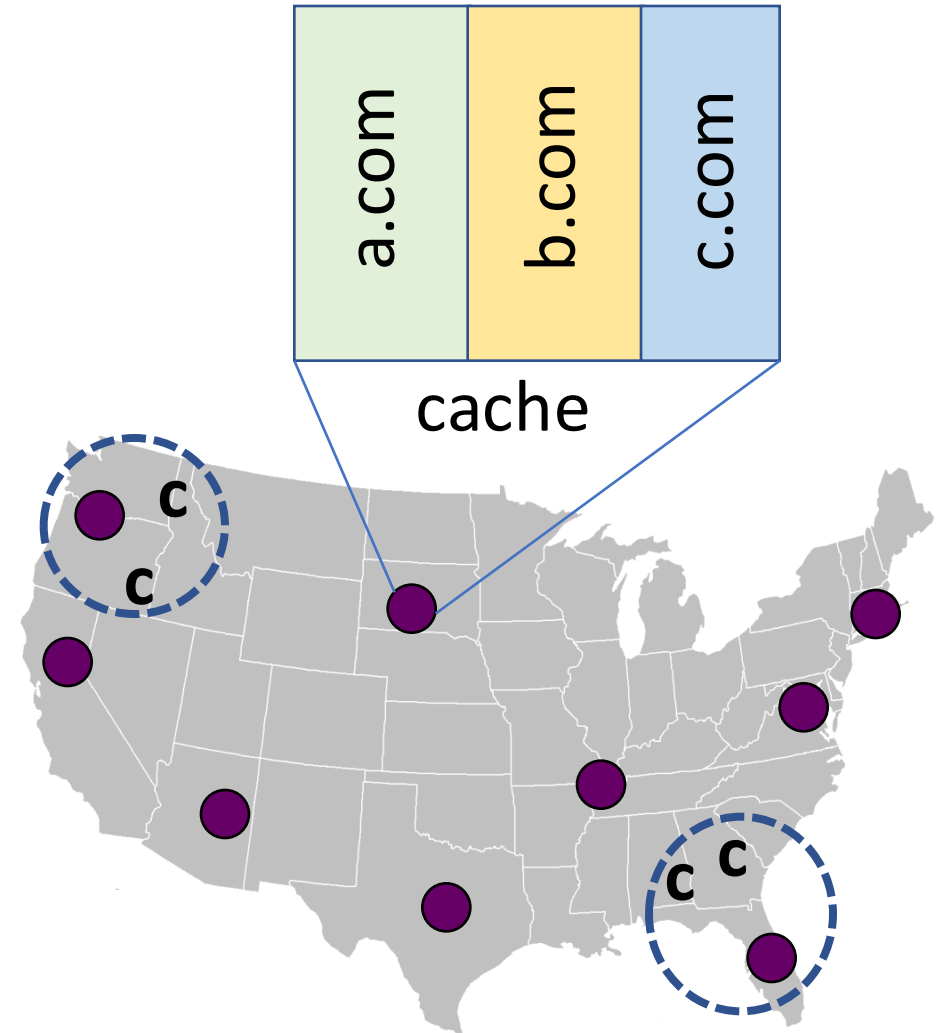
- Outsourced caching infrastructure
  - Caching for clients and servers
  - Dedicated equipment and software
  - Trained staff, best practices, etc.
- Coordination with the server
  - Generating non-cacheable content
  - Providing detailed measurement data
- Smart cache placement
  - Many caches: handle large request load
  - Close to many clients: reduce latency



More than 4200 locations in 135 countries


# CDN Challenges

- Where to place edge sites?
  - Close to many clients, with reasonable cost
- Where to replicate a server's content?
  - Many edge sites → duplicated data
  - Few edge sites → larger client latency
- How to direct a client to an edge site?
  - Proximity: for low latency
  - Light load: to reduce congestion
- How to manage each cache?
  - Maximize hit rate?
  - Minimize miss penalty?
  - Fairness across origin servers?

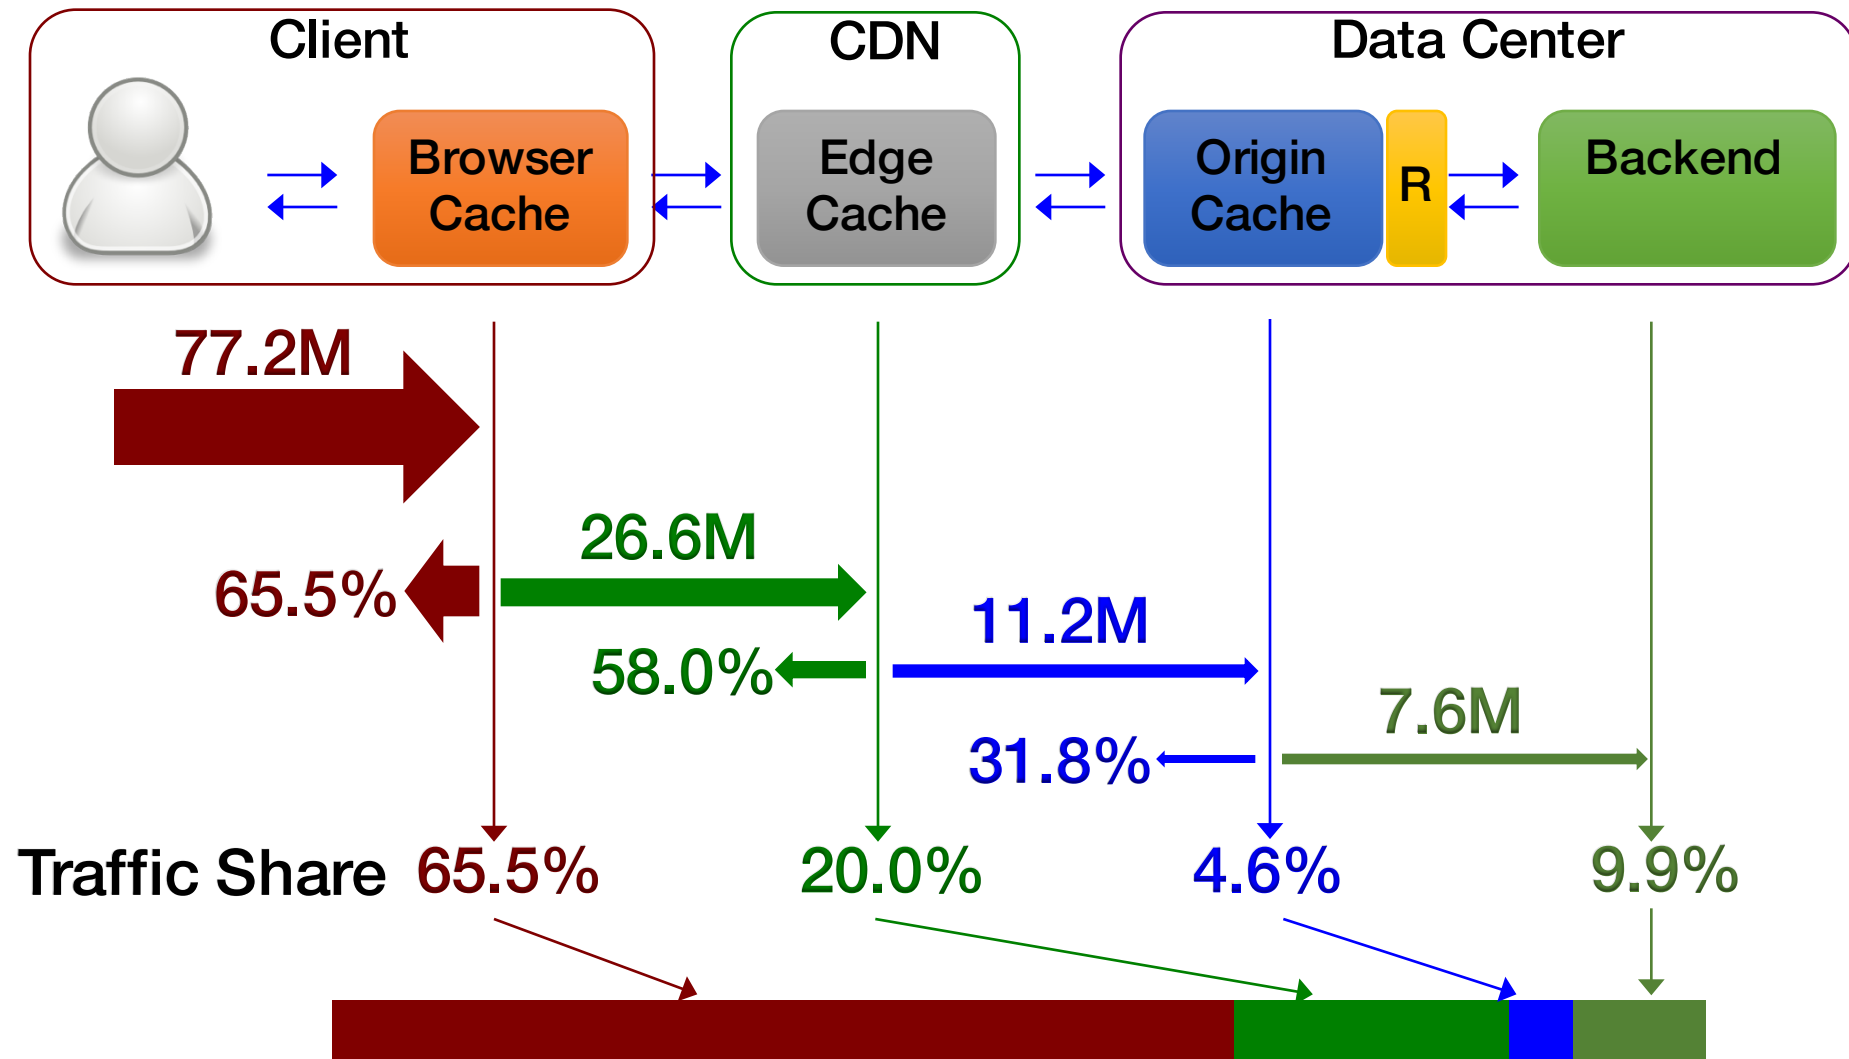


# CDN Challenges

- Where to place edge sites?
  - Close to many clients, with reasonable cost
- Where to replicate a server's content?
  - Many edge sites → duplicated data
  - Few edge sites → larger client latency
- How to direct a client to an edge site?
  - Proximity: for low latency
  - Light load: to reduce congestion
- How to manage each cache?
  - Maximize hit rate?
  - Minimize miss penalty?
  - Fairness across origin servers?



# CDN Effectiveness





# Conclusions

- Downloading a Web page
  - Name resolution, transport connection, secure session, web messages
- Benefits of caching
  - Reduces user latency, server load, and network bandwidth
- Cache replacement
  - Maximize hit rate by trying to predict the future
- Cache consistency
  - Efficient ways to avoid returning unnecessarily stale responses
- Content distribution networks
  - Caching close to clients, while working on behalf of the servers