


# Learning Relaxed Belady for CDN Caching



COS 316: Principles of Computer System Design  
Lecture 10

Amit Levy & Wyatt Lloyd

# Edge Cache with Different Algos



- Clairvoyant (Bélády) shows we can do much better!

# Cutting Edge Research From Princeton!

Learning Relaxed Belady for  
Content Distribution Network Caching.

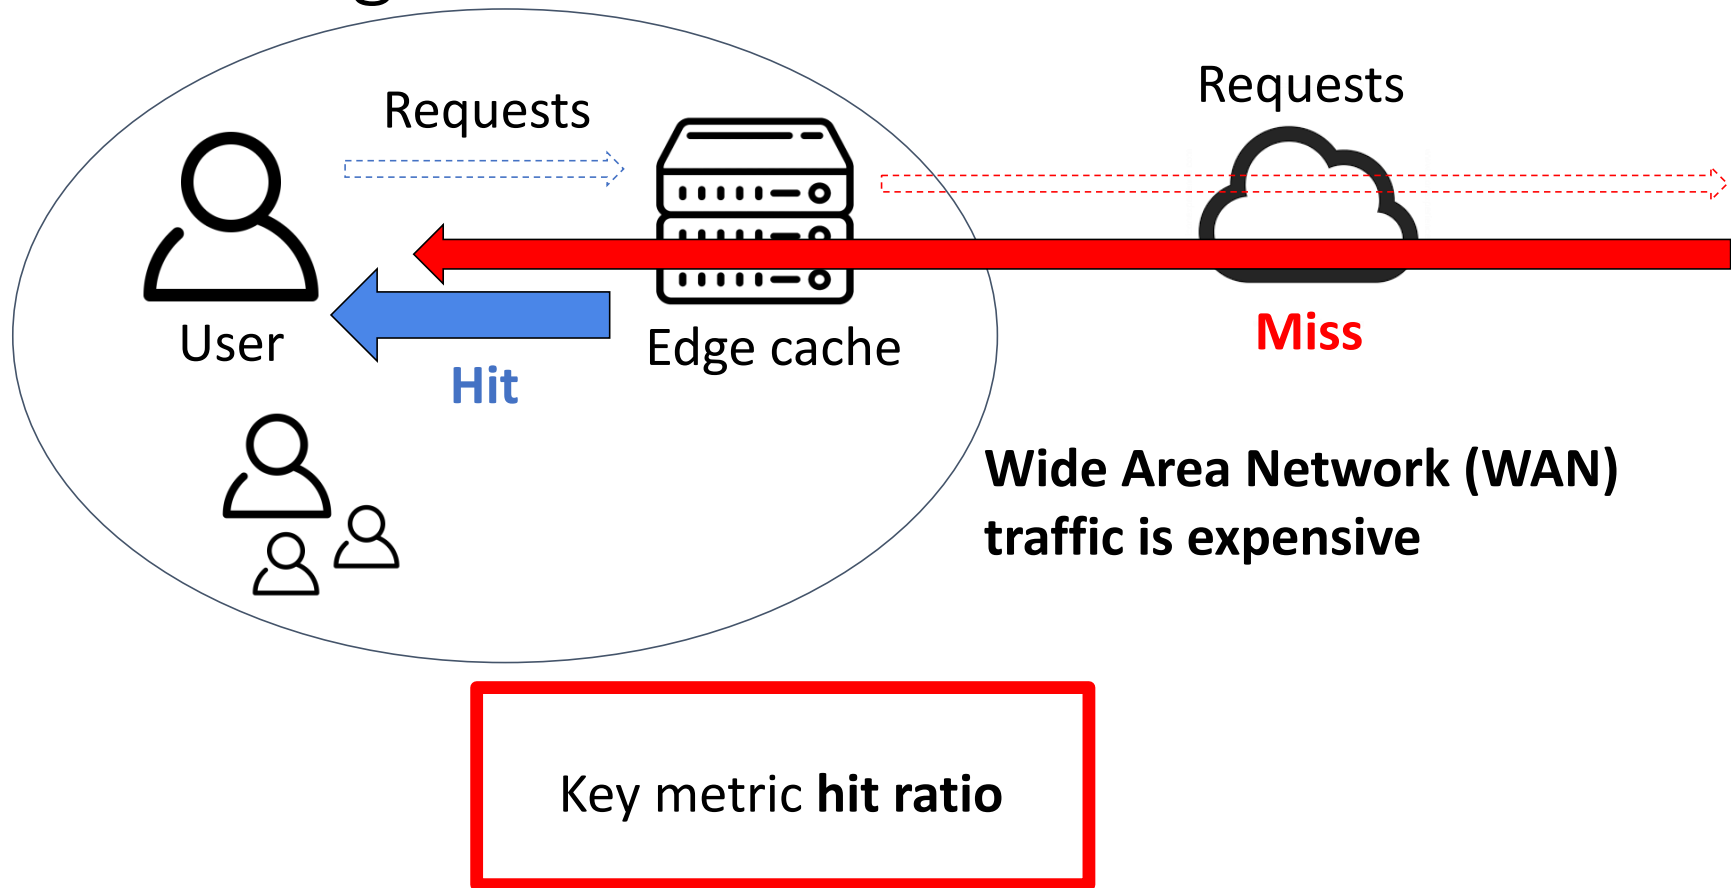
**Zhenyu Song**, Daniel S. Berger,  
Kai Li, and Wyatt Lloyd.

In 17th USENIX Symposium on Networked  
Systems Design and Implementation  
(NSDI 20), February 2020.



Microsoft  
**Research**

# CDN Caching Goal: Minimize WAN Traffic



# Caching Remains Challenging

Heuristic-based algorithms (1965–): **LRU**, **LFU**, GDSF, ARC, ...

- Work well for some workloads, but work poorly for other

ML-based adaptation of heuristics (2017–): UCB, LeCAR, ...

- Also work well for some workloads, but poorly for others

The **Belady** algorithm (1966)

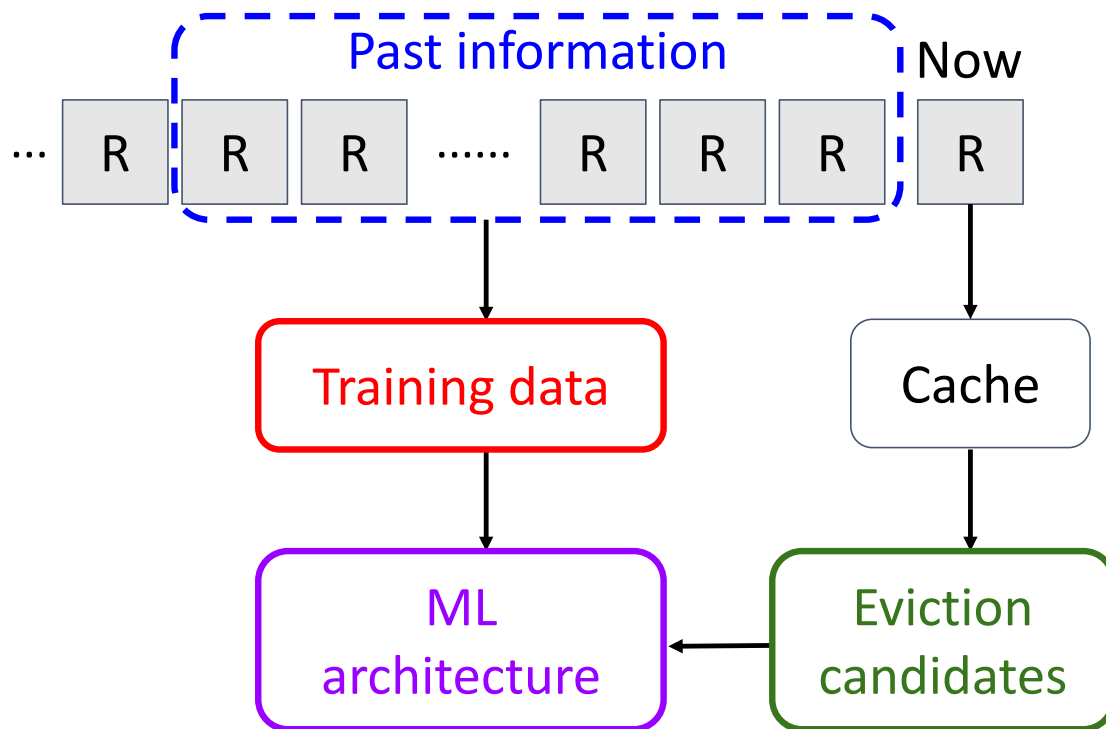
- Offline optimal: requires future knowledge
- Large gap in miss ratio between state-of-the-art and Belady:
- 20–40% on production traces

# Introducing Learning Relaxed Belady (LRB)

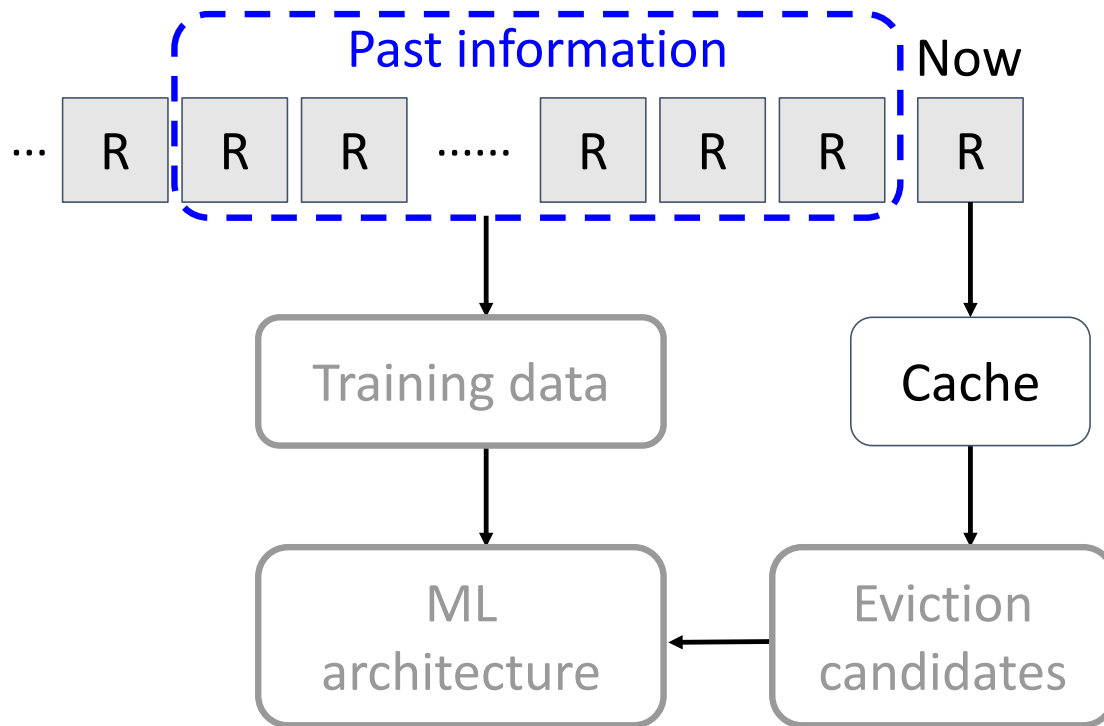
New approach: mimic Belady using machine learning

- Machine-Learning-for-Systems (ML-for-Systems)
  - Enabling technologies
  - When does it make sense?

# General Overview of our Approach



# Challenge 1: Past Information

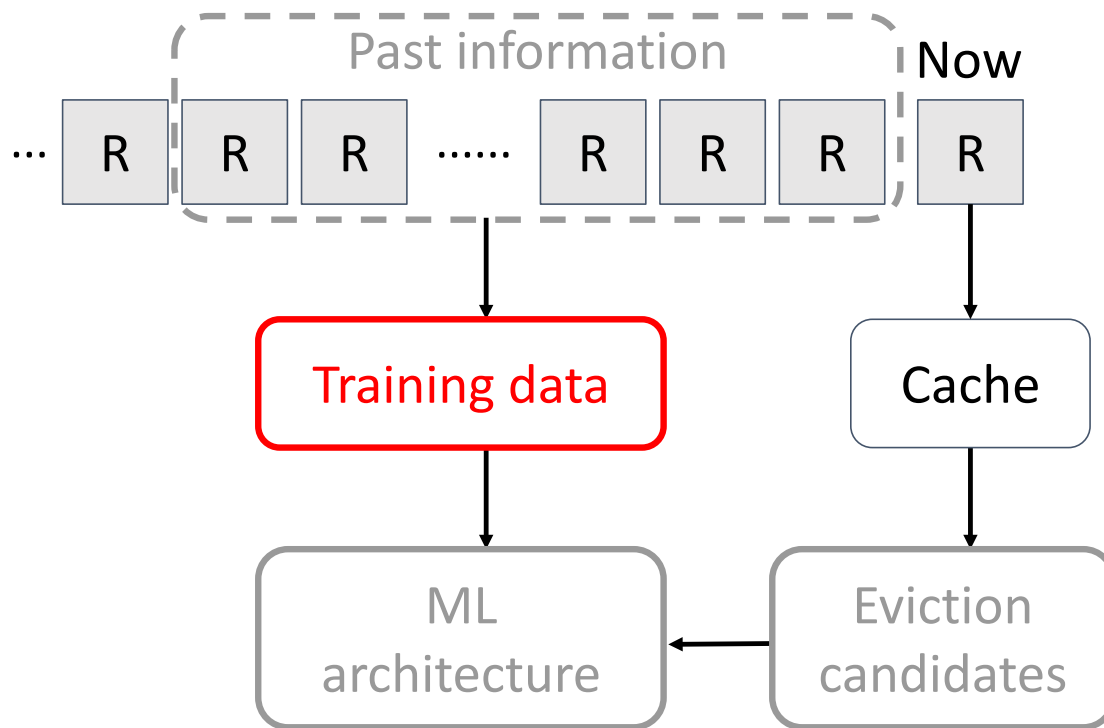


What past information to use?

More data improves training but increases memory overhead



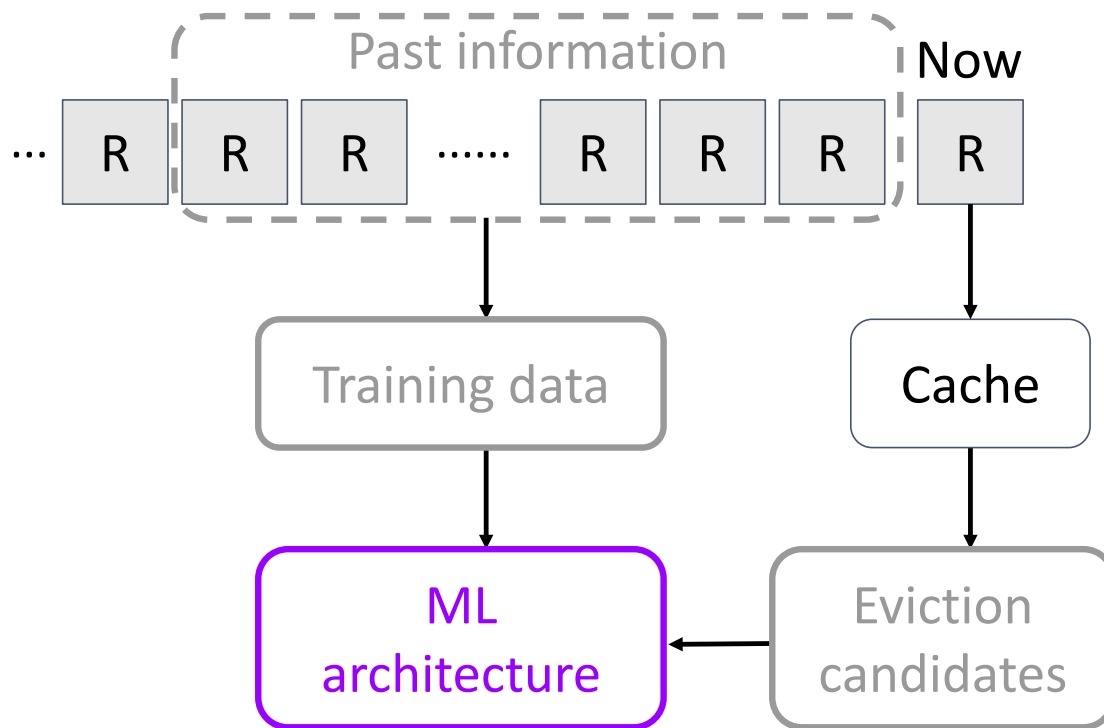
## Challenge 2: Generate Online Training Data



What past information to use?

Generate online training data?

## Challenge 3: ML Architecture



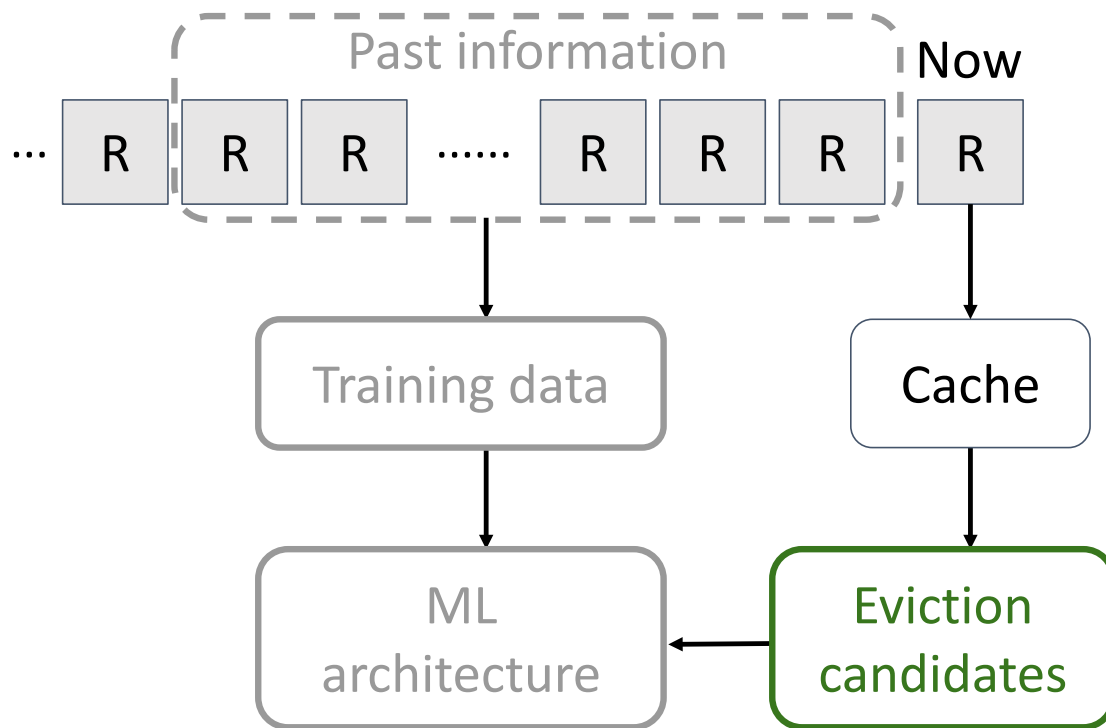
What past information to use?

Generate online training data?

What ML architecture to select?

Large design space:  
features, model, prediction  
target, loss function

## Challenge 4: Eviction Candidates



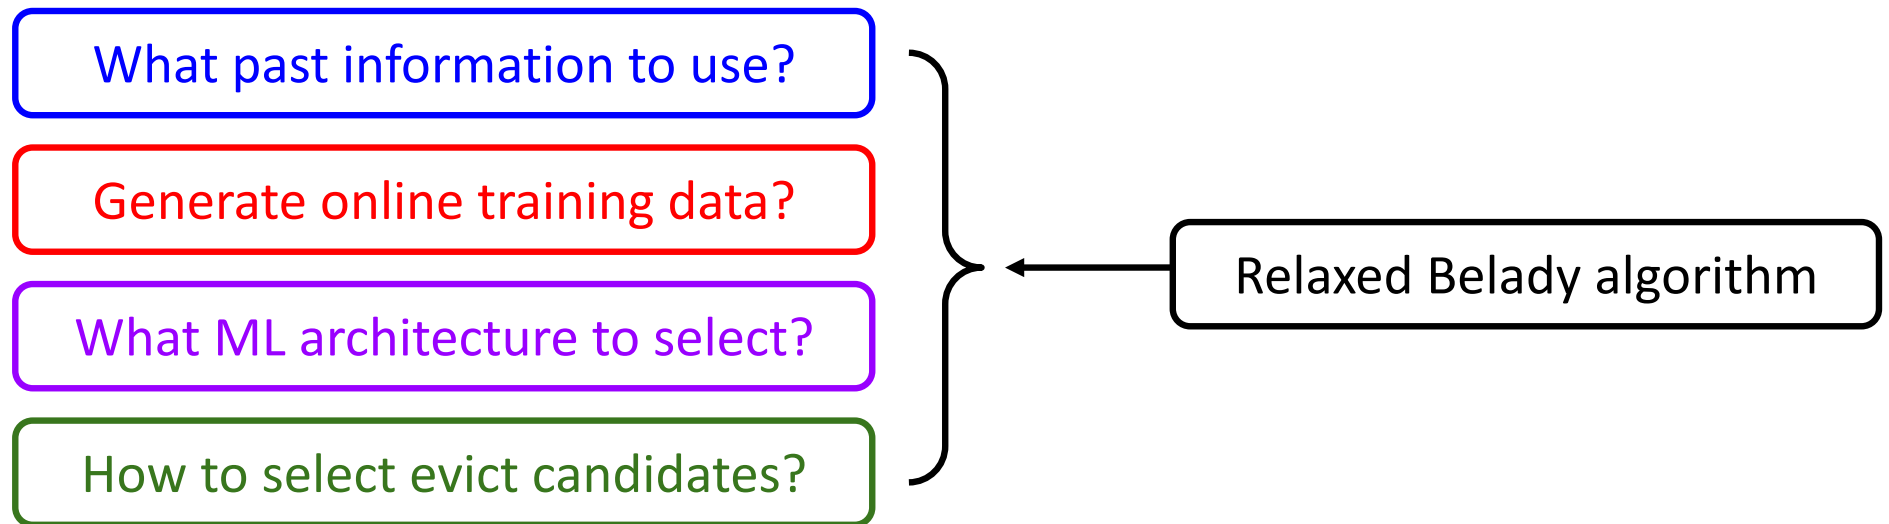
What past information to use?

Generate online training data?

What ML architecture to select?


How to select evict candidates?

## Solution: Relaxed Belady Algorithm



# Challenge: Hard to Mimic Belady Algorithm


Belady: evict object with next access farthest in the future



Mimicking exact Belady is impractical

- Need predictions for all objects → prohibitive computational cost
- Need exact prediction of next access → further prediction are harder

# Introducing the Relaxed Belady Algorithm

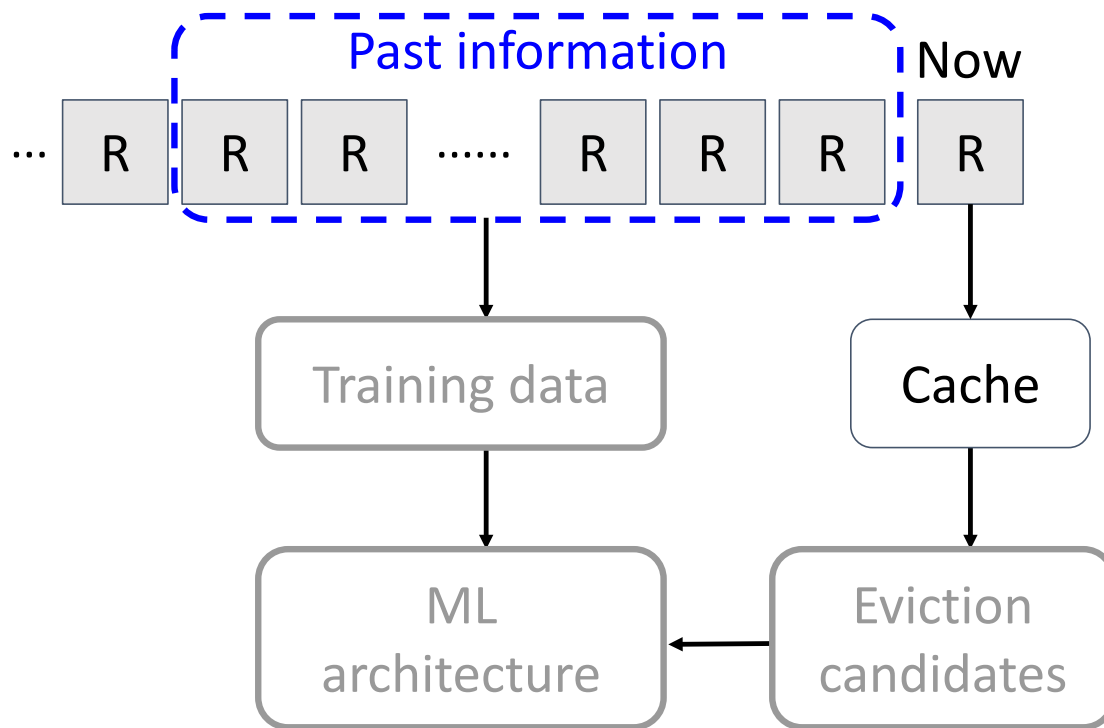


**Observation: many objects are good candidates for eviction**

Relaxed Belady evicts a random object beyond **boundary**

- Do not need predictions for all objects → reasonable computation
- No need to differentiate beyond boundary → simplifies the prediction

# Challenge 1: Past Information




What past information to use?

More data improves training but increases memory overhead

# Track Objects within a Sliding Memory Window

Sliding memory window mimics Belady boundary



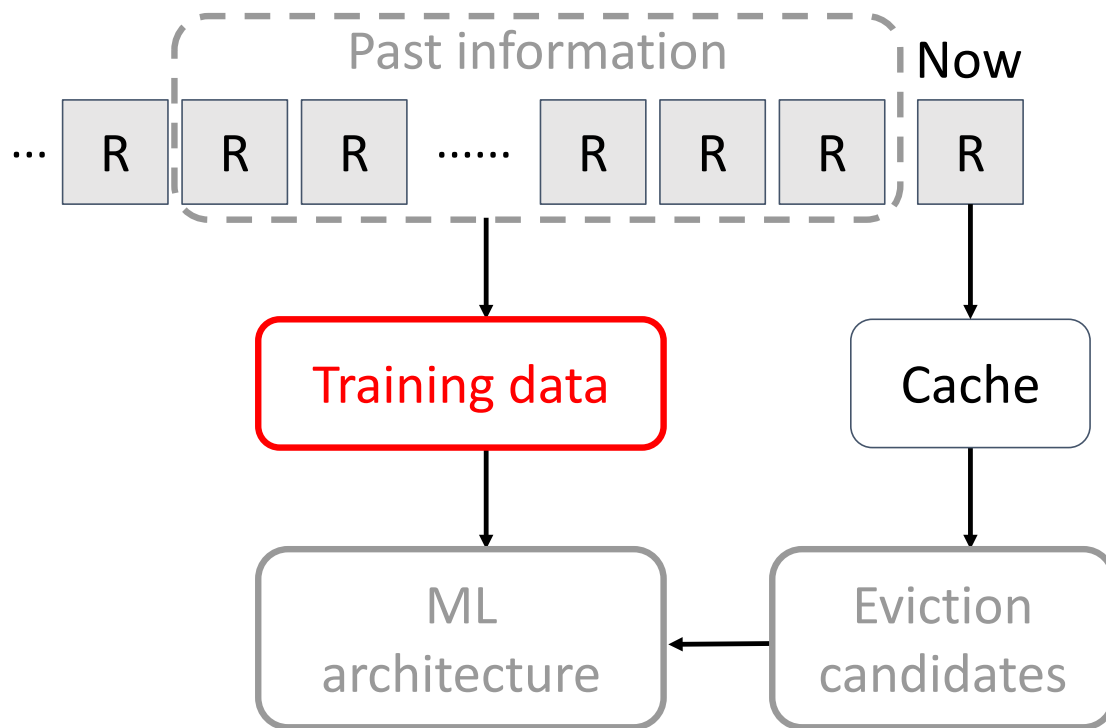
Only track objects within **memory window**

Per object features

Window size is LRB's main hyperparameter



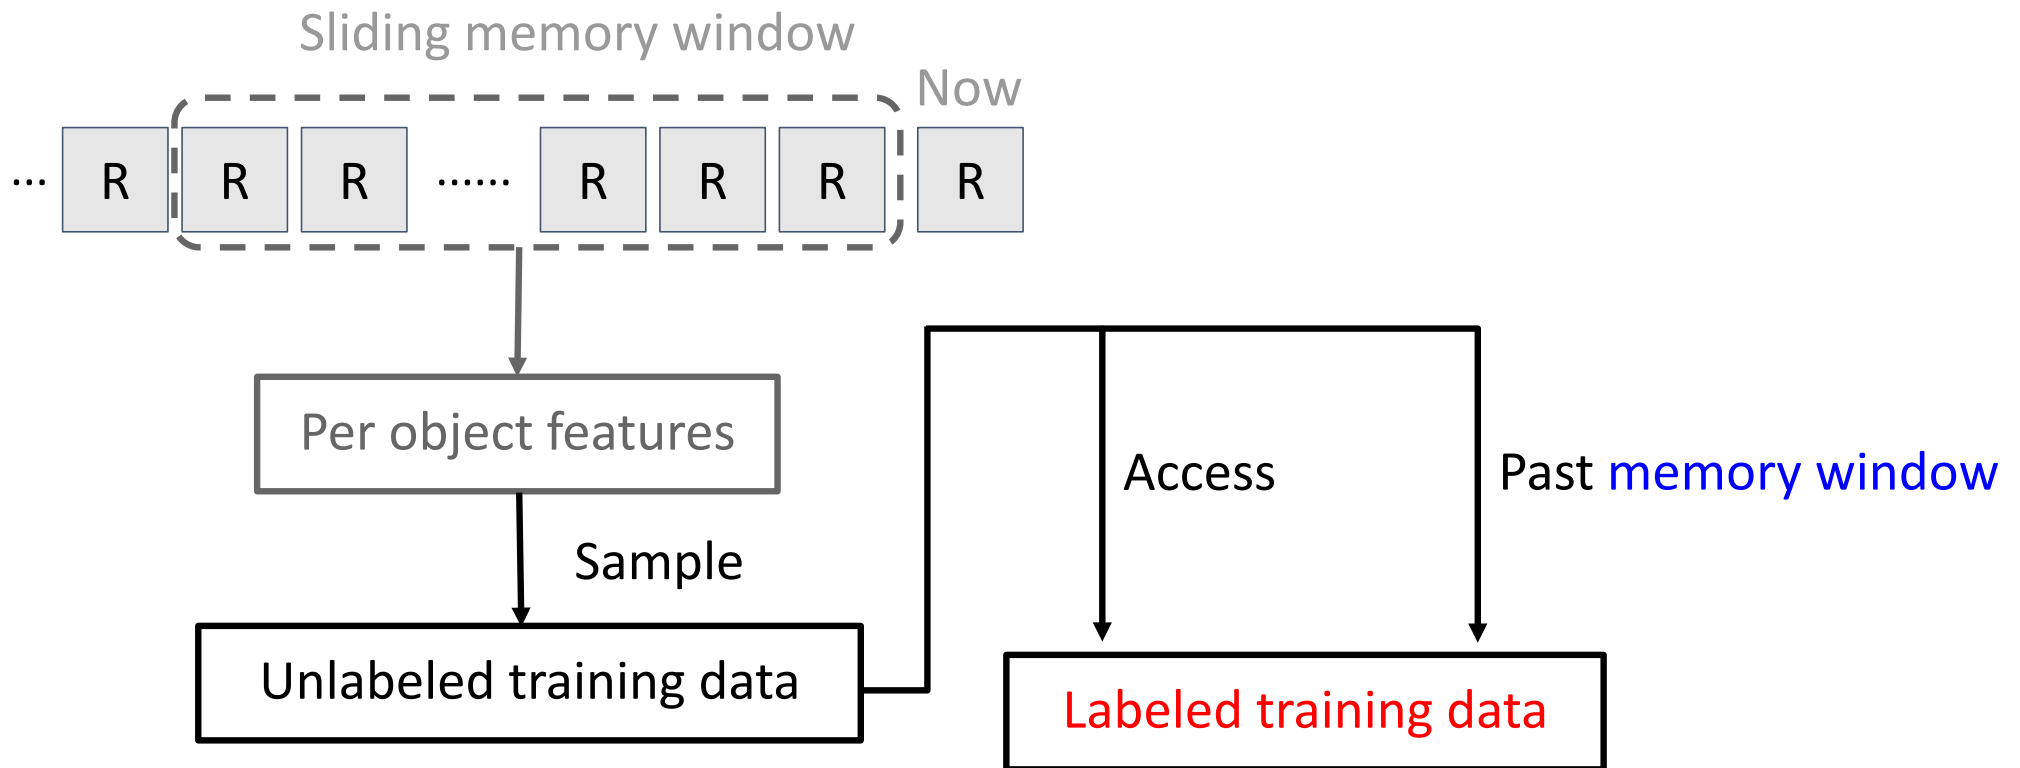
## Challenge 2: Training Data



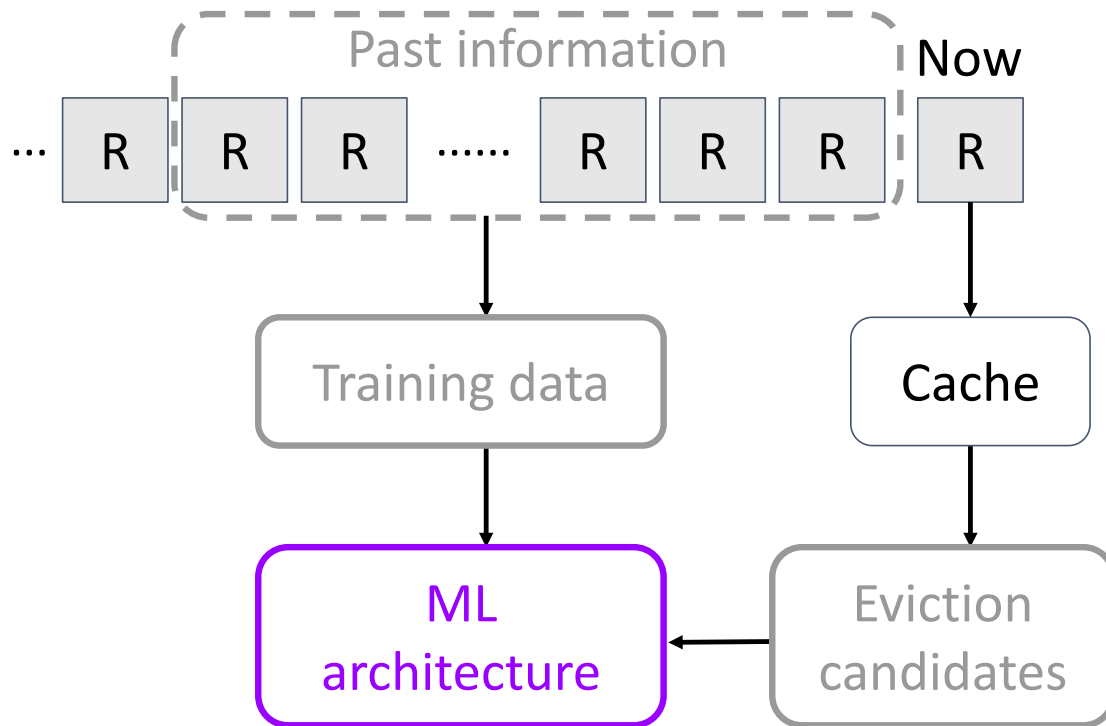
What past information to use?

Generate online training data?

# Sample Training Data & Label on Access or Boundary



## Challenge 3: ML Architecture



What past information to use?

Generate online training data?


What ML architecture to select?

Large potential design space

# Solution 3: Feature & Model Selection

Use good decision ratio to evaluate new designs

Features
Object size
Object type
Inter-request distances (recency)
Exponential decay counters (long-term frequencies)

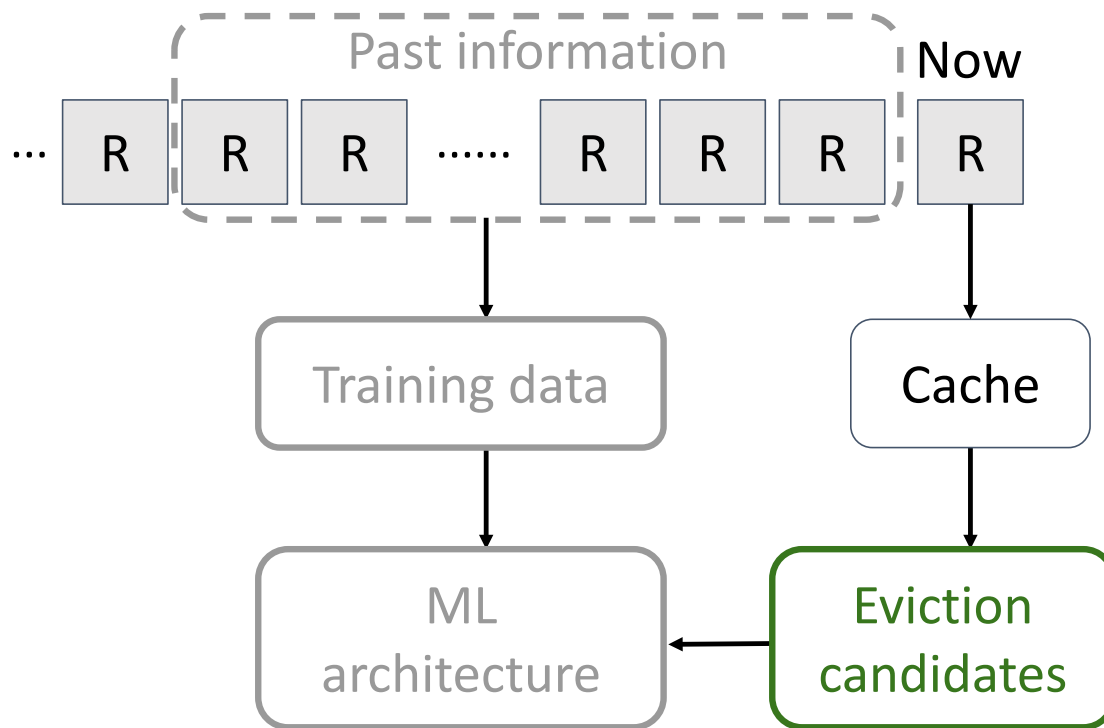


Gradient boosting decision trees

**Lightweight & high good decision ratio**

Training ~300 ms, prediction ~30 us

## Challenge 4: Eviction Candidates



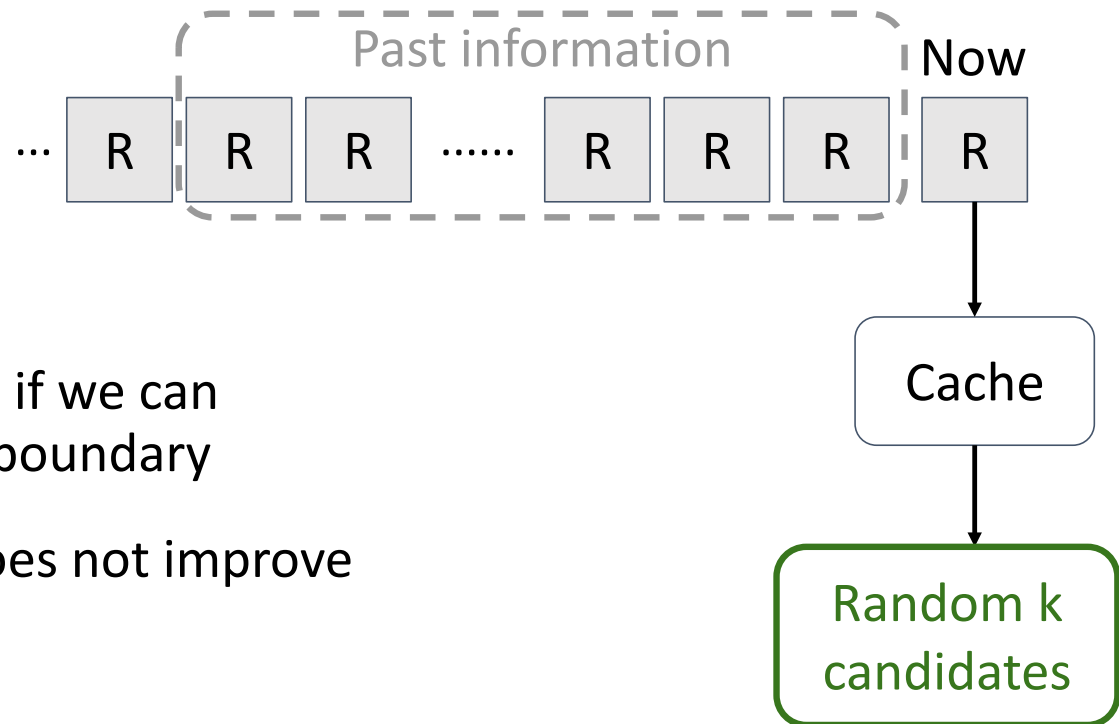
What past information to use?

Generate online training data?

What ML architecture to select?

How to select evict candidates?

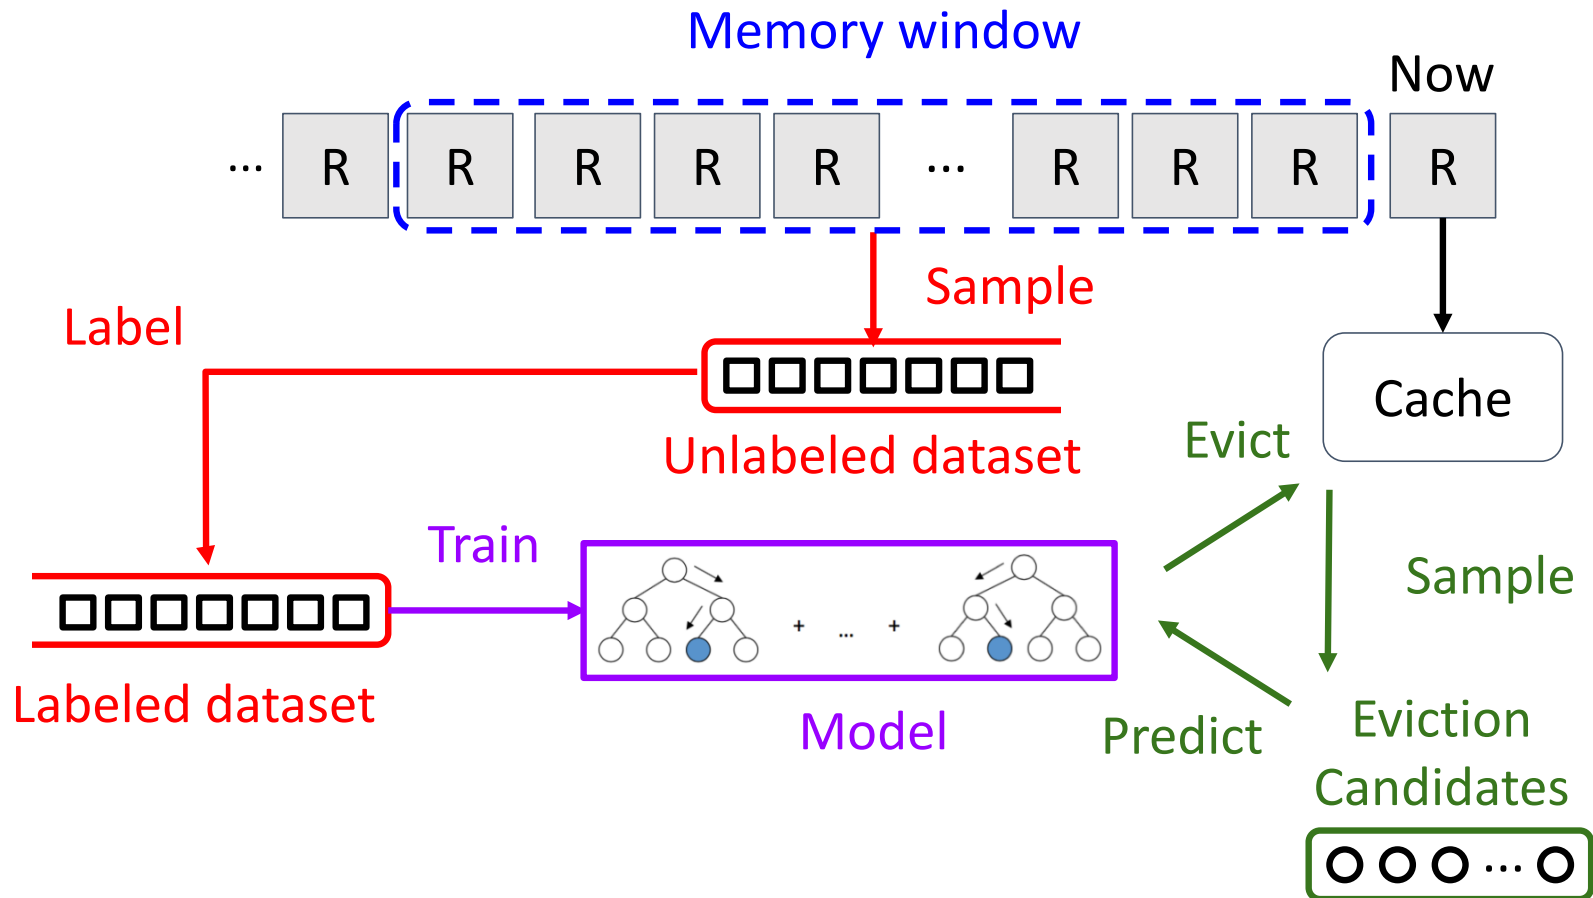
## Solution 4: Random Sampling for Eviction



Can mimic relaxed Belady if we can find 1 object beyond the boundary

k=64 candidates; more does not improve good decision ratio

# Learning Relaxed Belady



# Implementation

- Simulator implementation
  - LRB + 14 other algorithms
- Prototype implementation
  - C++ on top of production system (Apache Traffic Server)
  - Many optimizations




# Evaluation Setup

- Q1: Learning Relaxed Belady (LRB) traffic reduction vs state-of-the-art
- Q2: overhead of LRB vs CDN production system
- Traces: 6 production traces from 3 CDNs
- Hyperparameter ([memory window](#)/model/...) tuned on 20% of trace

# LRB Reduces WAN Traffic

Industry standard

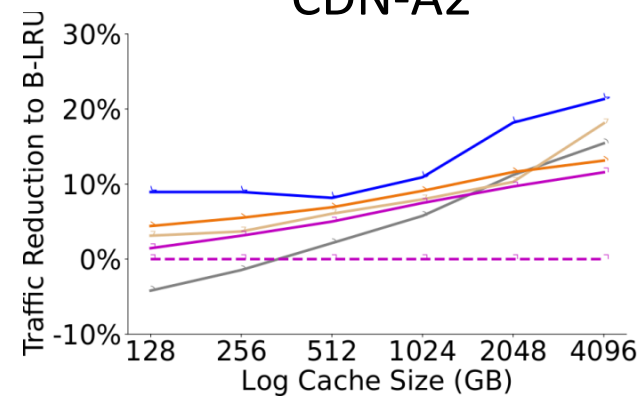
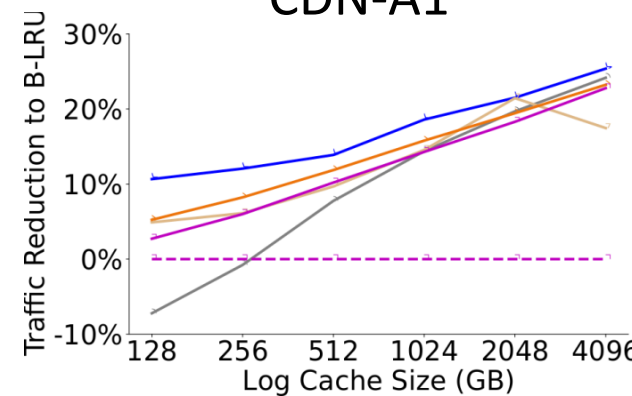




20% traffic reduction over B-LRU  
10% reduction over the best SOA



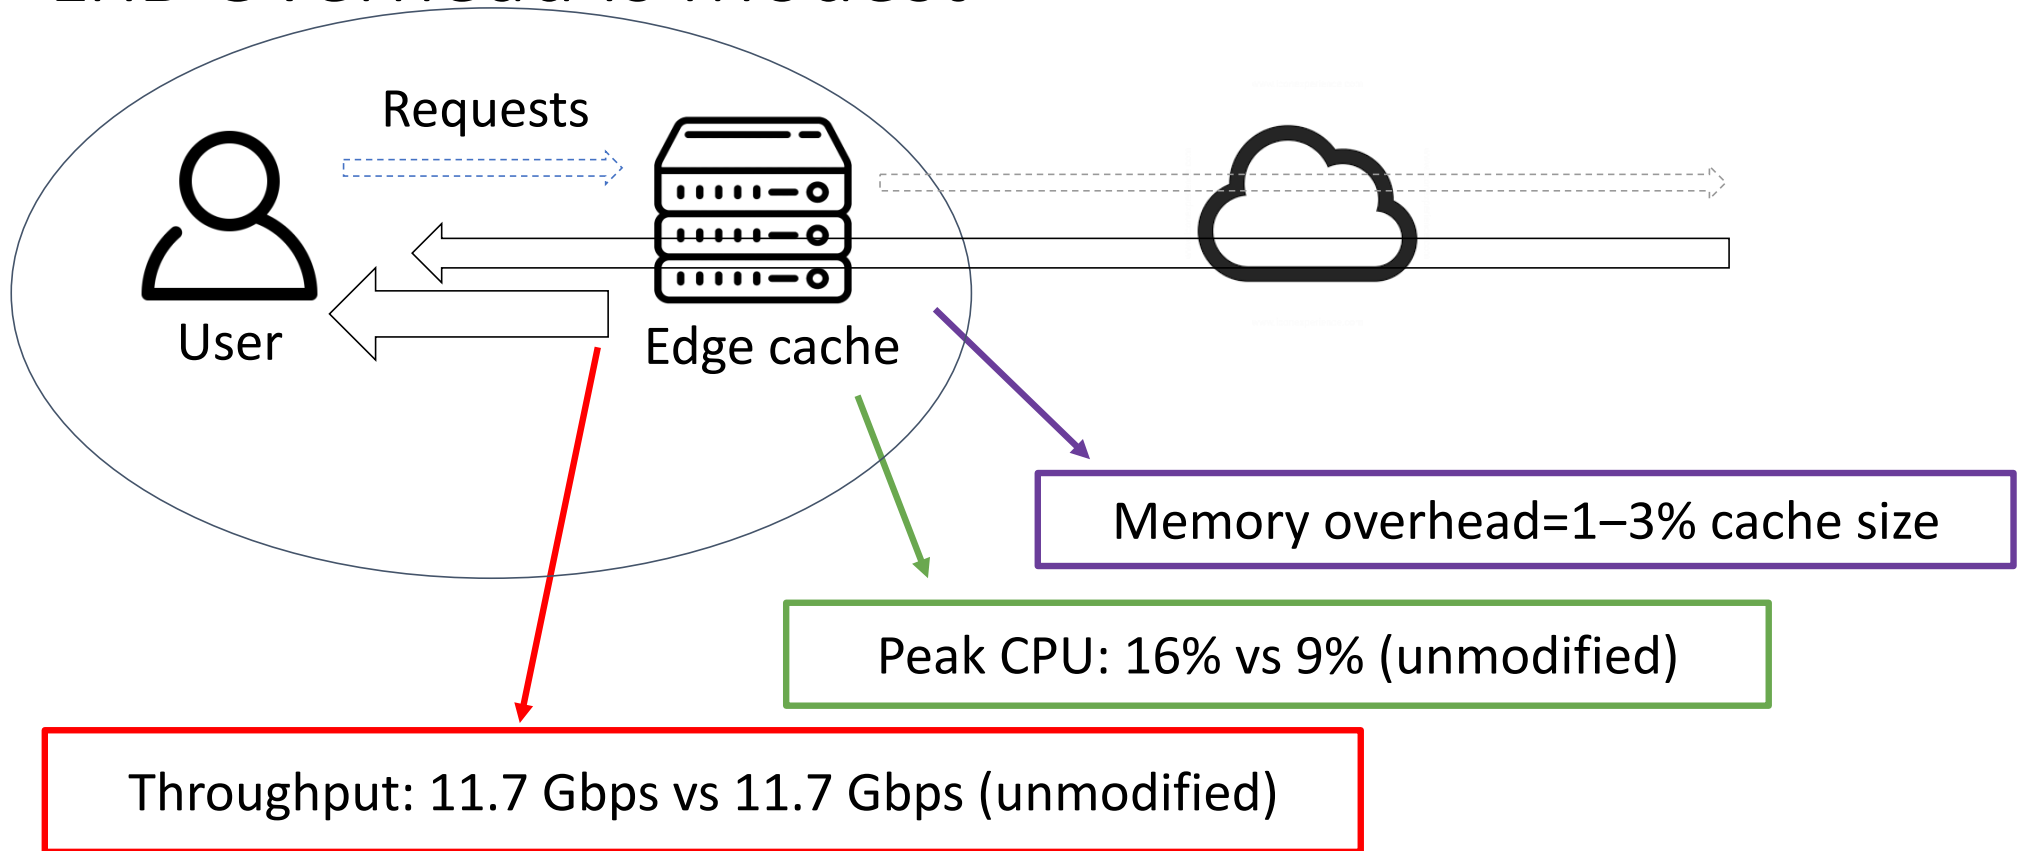
Wikipedia trace

# LRB Consistently Improves on the State of the Art

LRB (Ours)   LFUDA   LRU4   TinyLFU   LeCaR   B-LRU   LRU



# LRB Overhead Is Modest



# Conclusion

- LRB reduces WAN traffic with modest overhead
- ML-for-systems generally promising to replace heuristics
- Key insight: **relaxed Belady**
  - Simplifies machine learning & reduces system overhead



Microsoft  
**Research**

