

NetCache: Balancing Key-Value Stores with Fast In-Network Caching



COS 316: Principles of Computer System Design
Lecture 10

Amit Levy & Jennifer Rexford

1

What is a Key-Value Store?



redis



amazon
DynamoDB

2

Key-Value Store: A Very Simple Database

- Hash table
 - key: an arbitrary string (e.g., "Joe", "jrex.jpg")
 - value: unstructured data (e.g., number, string, complex object)
- Primitive operations
 - put(key, value)
 - value = get(key)
 - delete(key)
- A "dictionary"
 - Array, map...

key	value
"system"	"a set of things working together as parts of a mechanism or an interconnecting network."
"internet"	"the single worldwide computer network that interconnects other computer networks"
"database"	"a comprehensive collection of related data organized for convenient access, generally in a computer."
"cache"	"a temporary storage space or memory that allows fast access to data"

3

Applications of Key-Value Stores


- When to use a key-value store
 - Real-time, random data access
 - Caching for frequently accessed data
 - Simple applications
- Many example applications
 - Online shopping: user id (key) and shopping cart (value)
 - User profile: user id (key) and preferences (value)
 - Movies/music: file name (key) and file contents (value)
 - Blockchain: hash value (key) and block entity (value)



4

High-Speed Key-Value Store


- A need for speed
 - “web pages routinely fetch thousands of key-value pairs” [Facebook]
 - Systems with many users
- A key-value store needs resources
 - Storage: to store many key-value pairs
 - Bandwidth: to handle many operations
- In-memory key-value store
 - Avoid reads and writes to slow disk-based systems
 - But, what if you need more storage and/or bandwidth?



5

Distributed, High-Speed Key-Value Store

- Vertical scaling
 - Add more resources to a server
 - E.g., more processing, memory, or network resources
- Horizontal scaling
 - Add more servers
 - Reduce the responsibilities of each server



6

Horizontal Scaling: Partitioning the Keys


- **Sharding:** breaking data up into partitions
 - Each server handles operations for a portion of the keys
 - Reduces query load and storage space on individual servers

key	value	
"system"	"a set of things working together as parts of a mechanism or an interconnecting network."	Server 0
"internet"	"the single worldwide computer network that interconnects other computer networks"	
"database"	"a comprehensive collection of related data organized for convenient access, generally in a computer."	Server 1
"cache"	" a temporary storage space or memory that allows fast access to data"	

7

Goals in Partitioning the Keys

- Easy to direct requests to the right server
- Even distribution of data storage
- Even distribution of request load



- E.g., range of keys
 - Alphabet: letters A-F, G-L, M-R, S-Z
 - Months: Jan-Mar, Apr-Jun, Jul-Sep, Oct-Dec

8


Better Design: Hash-Based Sharding

- Hashing on the key

- Randomize and distribute requests
- Easy to compute without coordination

- Achieves the goals?

- Easy to direct requests to the right server?
- Even distribution of data storage?
- Even distribution of request load?




9

Distributed, In-Memory Key-Value Stores Need a Cache

10

Challenge: Skewed Workload

low throughput & high tail latency



- Popular items receive far more queries than others
- E.g., “10% of items account for 60-90% of queries in the Memcached deployment at Facebook”

11

Challenge: Dynamic Workloads


low throughput & high tail latency



- The popular items change rapidly
- E.g., popular posts, limited-time offers, trending events


12

Fast, Small Cache for Better Load Balancing



13

Fast, Small Cache for Better Load Balancing




E.g., 100 backends with 100 billions items

[B. Fan et al. **SoCC'11**, X. Li et al. **NSDI'16**]

14

Key-Value Store is *Already* in High-Speed Memory




Cache needs to achieve the aggregate throughput of the storage layer.

15

NetCache


16

Cache Directly in the Network Devices!



17


In-Network Load-Balancing Cache



High-speed switches have hardware designed to process packets of data as they arrive


18

Individual Packets Carry the Messages



19

Cache Hit and Cache Miss



- Cache hit: switch responds directly
- Cache miss: switch forwards to the server, which responds

20


Programmable Packet-Processing Hardware

- Modern switches are increasingly programmable

- Parsing of packet headers
- Match-action processing on header fields
- Registers for storing data across packets




www.p4.org



21

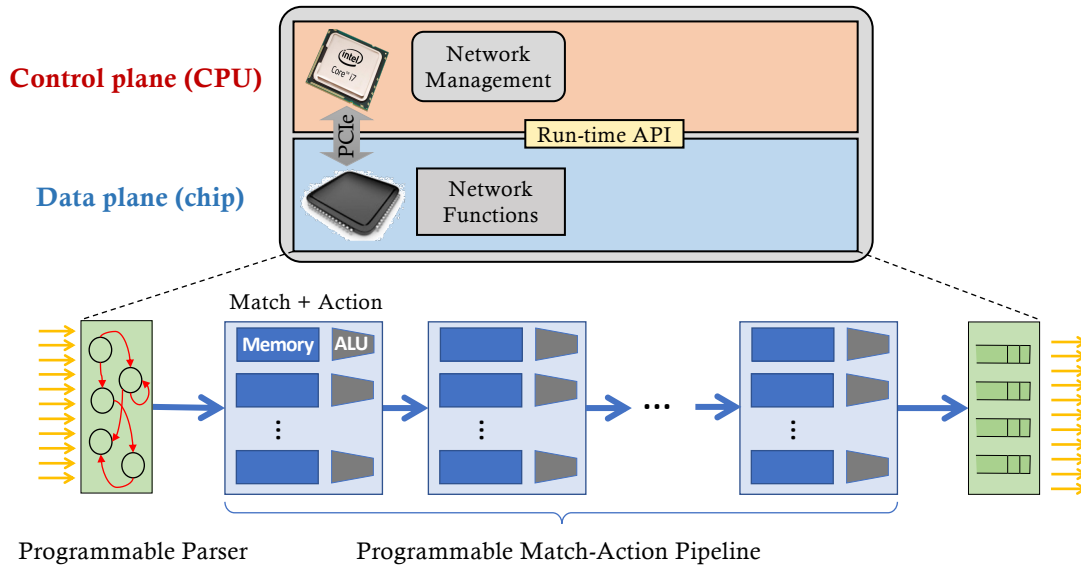
NetCache Uses Programmable Switches

- Programmable parser
 - Parse custom key-value fields in the packet
- Programmable match-action pipeline
 - Read and update key-value data
 - Maintain query statistics for cache updates




22

Control Plane of Programmable Switches



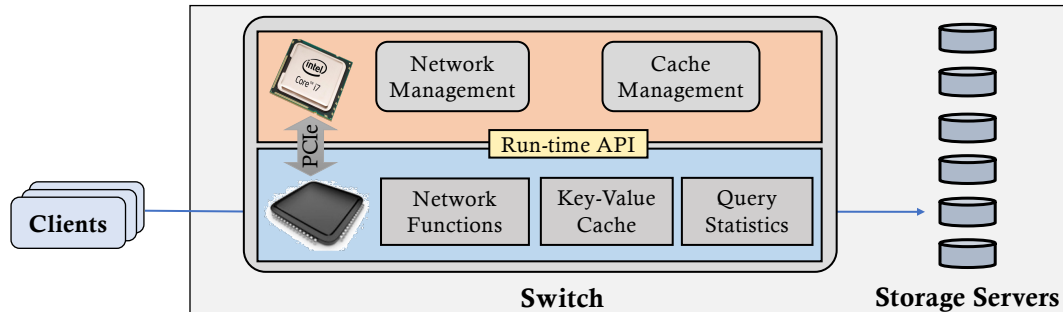
23

Cache Eviction and Replacement



24


NetCache Architecture



- **Switch data plane**
 - Key-value store to serve queries for cached keys
 - Query statistics to enable efficient cache updates
- **Switch control plane**
 - Insert hot items into the cache and evict less popular items
 - Manage memory allocation for on-chip key-value store

25

Data-Plane Query Handling



26

NetCache Query Statistics

27

Query Statistics in the Data Plane


- For cached keys
 - Count the number of accesses to each cached key
 - ... so the control plane can decide which key-value pairs to *evict*
- For uncached keys
 - Identify new “hot” keys
 - ... so the control plane can decide which key-value pairs to *insert*
 - Problem: huge number of uncached keys

28

Compact Data Structure: CountMin Sketch

- Approximate counts without per-key memory

- r rows, each with a hash function $h_i()$
- c columns
- $r \times c$ counters




29

Compact Data Structure: CountMin Sketch

- Approximate answers with less memory

- r rows, each with a hash function $h_i()$
- c columns
- $r \times c$ counters




Estimated count:
 $\min\{77, 80, 79\} = 77$

Report this hot key!

30

NetCache Query Statistics



- Cached key: per-key counter array
- Uncached key
 - Count-Min sketch: identify new hot keys
 - Bloom filter: remove duplicated hot key reports

31

NetCache is a **key-value store** that leverages **in-network data-plane caching** to achieve

billions QPS throughput

&

~10 μ s latency

even under

highly-skewed

&

rapidly-changing

workloads.

32

Conclusion

- Key-value stores
 - Distributed, in-memory hash table
 - Simple, fast, and scalable
- Skewed workloads are still a challenge
 - Some keys are much more popular than others
 - So, some storage servers are overloaded
- In-network caching to the rescue
 - Store a small number of popular keys
 - ... to reduce load (and variability of load) on the servers
- Programmable data-plane hardware
 - An opportunity for new network functionality
 - ... but a challenge due to limited memory and processing