

Introduction to Layering and Network Layering



COS 316: Principles of Computer System Design
Lecture 11

Amit Levy & Ravi Netravali



**“Modularity based on abstraction
is the way things get done”**


2009 Turing Award Lecture

Barbara Liskov

Modularity Through Layering


- Systems on systems on systems though layering
- Each layer hides complexity with abstraction
- Network layers today!

The Problem of Communication



- Re-implement every application for every new transmission medium?
- Change every application on any change to a transmission medium?
- No! But how does the Internet design avoid this?

Solution: Layering

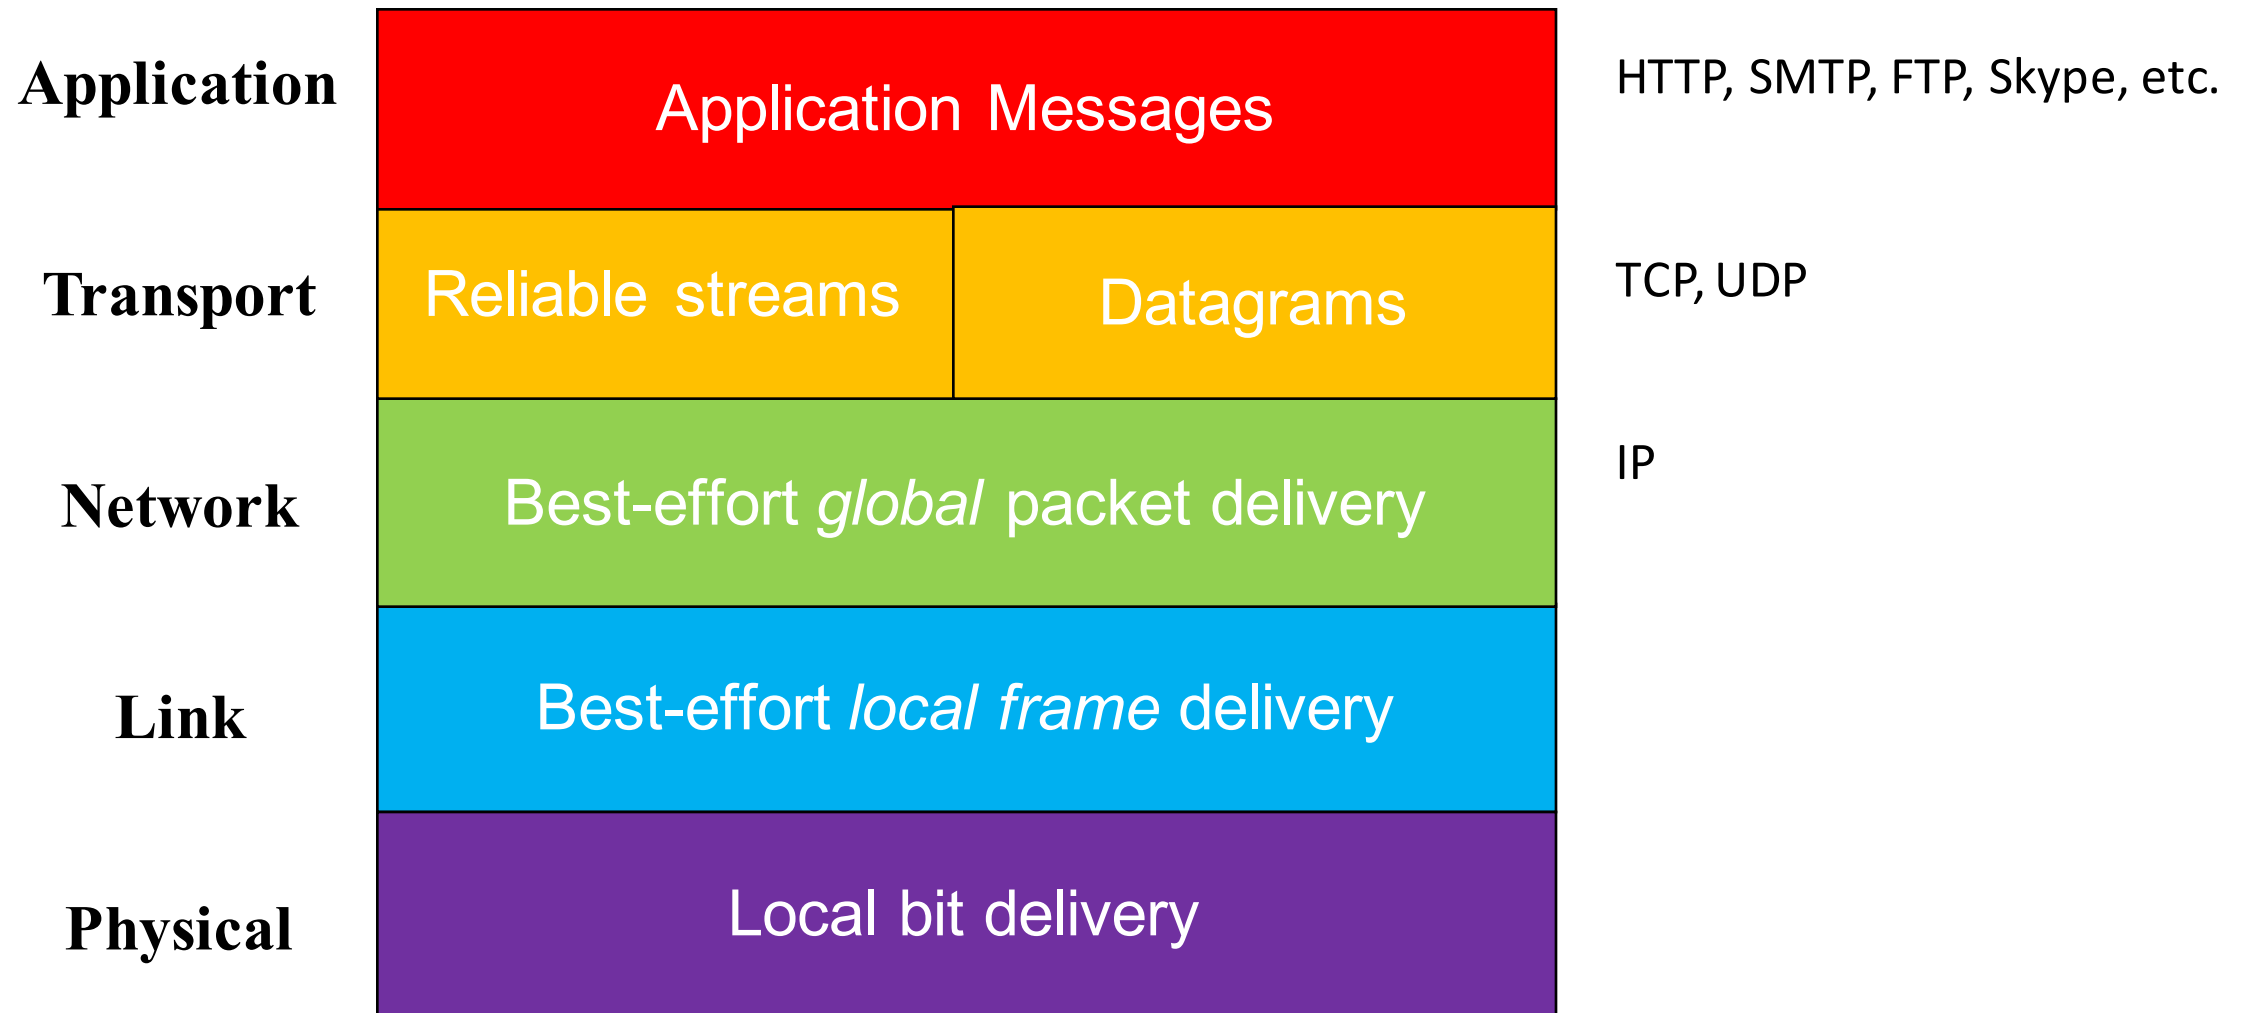


- Intermediate **layers** provide a set of abstractions for applications and media
- New applications or media need only implement for intermediate layer's interface

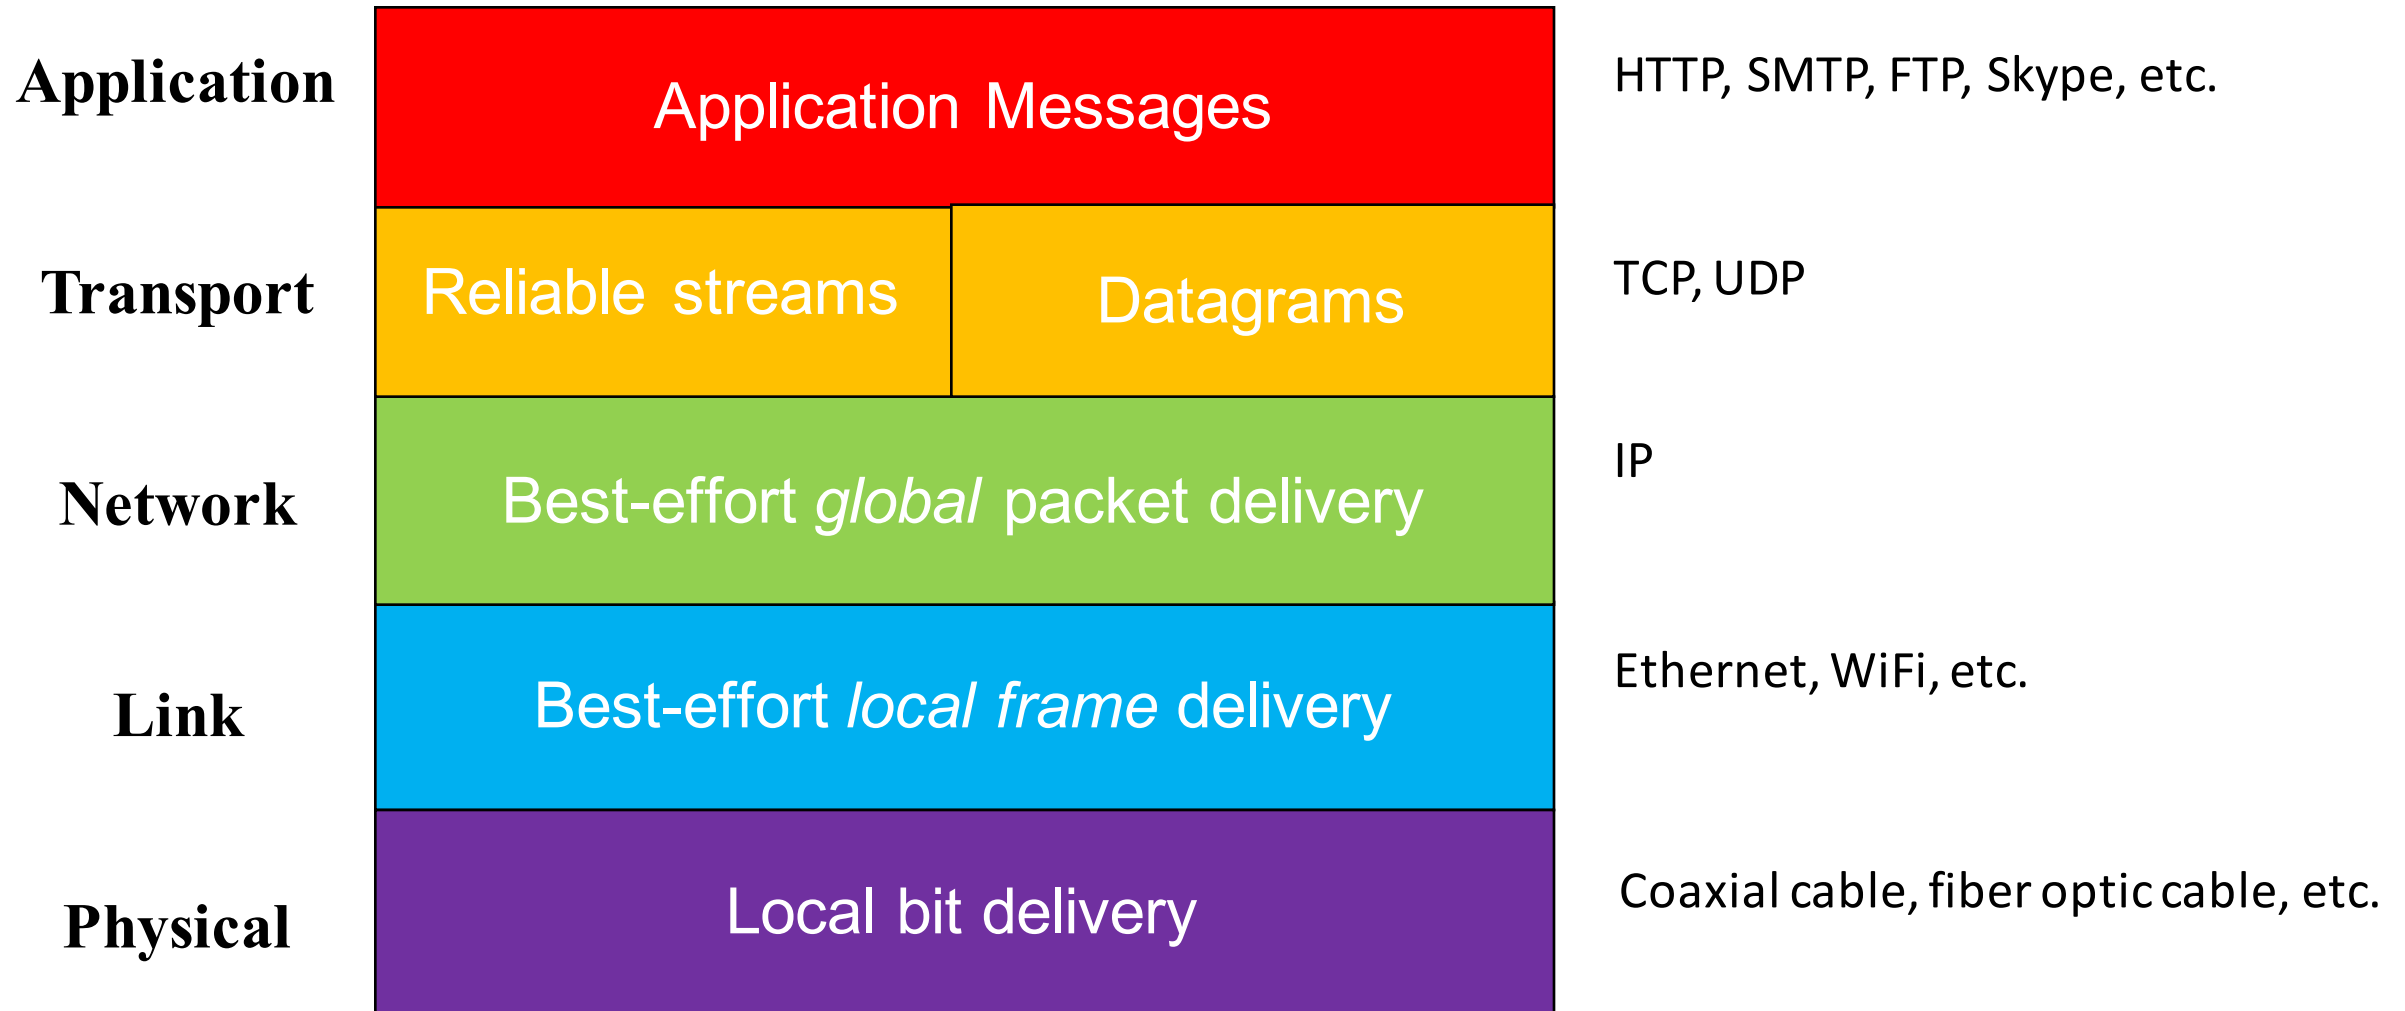
The Art of Layering

- How many layers?
- What goes in each layer?
- What abstraction (interface) does each layer provide?

Internet Protocol Layers




Internet Protocol Layers



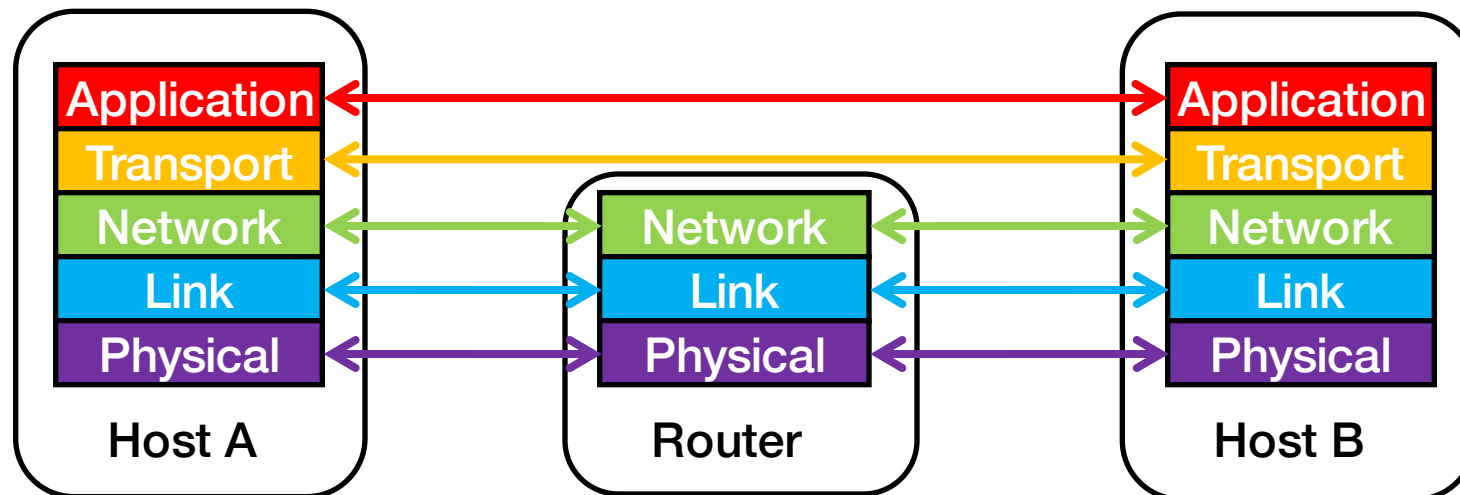
Logical Communication Between Layers

- How to forge agreement on meaning of bits exchanged between two hosts?
- **Protocol:** Rules that govern the format, contents, and meaning of messages
 - Each layer on a host interacts with peer host's corresponding layer via **protocol interface**




Logical Communication Between Layers

- How to forge agreement on meaning of bits exchanged between two hosts?
- **Protocol:** Rules that govern the format, contents, and meaning of messages
 - Each layer on a host interacts with peer host's corresponding layer via **protocol interface**




Physical communication

- Communication goes down to the **physical network**
- Then from **network** peer to peer
- Then up to the **relevant application**




Communication Between Peers

- How do peer protocols coordinate with each other?
- Layer attaches its own **header** (H) to communicate with peer
 - Higher layers' headers, data **encapsulated** inside message
 - Lower layers generally do not inspect higher layers' headers



Communication Between Peers

- How do peer protocols coordinate with each other?
- Layer attaches its own **header** (H) to communicate with peer
 - Higher layers' headers, data **encapsulated** inside message
 - Lower layers generally do not inspect higher layers' headers



IP is the “Narrow Waist” of the Internet

- **The** network-layer protocol
 - Enables portability above and below
- Lots of link layer protocols underneath
- Several transport protocols on top
 - TCP, UDP, QUIC



IP: Best-Effort Global Packet Delivery

- Never having to say you're sorry
 - Don't have to reserve bandwidth and memory
 - Don't have to do error detection and correction
 - Don't have to remember anything across packets
- Easier to survive failures
 - Transient disruptions are okay during failure recovery




IP: Best-Effort Global Packet Delivery

- Never having to say you're sorry
 - Don't have to reserve bandwidth and memory
 - Don't have to do error detection and correction
 - Don't have to remember anything across packets
- Easier to survive failures
 - Transient disruptions are okay during failure recovery
- Can run on nearly any link technology
 - Greater interoperability and evolution
 - RFC 1149: IP Datagrams Over Avian Carriers




IP: Statistical Multiplexing

- Data traffic is bursty
 - Logging in to remote machines
 - Exchanging e-mail messages
- Don't waste bandwidth
 - No traffic exchanged during idle periods
- Better to allow multiplexing
 - Different transfers share access to same links



IP: Statistical Multiplexing

- Data traffic is bursty
 - Logging in to remote machines
 - Exchanging e-mail messages
- Don't waste bandwidth
 - No traffic exchanged during idle periods
- Better to allow multiplexing
 - Different transfers share access to same links



Transport: Application to Application


- Network layer is host-to-host
- Transport layer is port-on-host-to-port-on-host
 - think application to application
 - demultiplexing
 - e.g., port 80 is HTTP, port 443 is HTTPS, port 22 is SSH
- *Why transport and not network layer?*

Transport: Application to Application

- Network doesn't have error detection
- Transport layer does have error detection
- *Why transport and not network layer?*
- *Why not both?*

Transport: Transmission Control Protocol (TCP)

- Ordered, reliable stream of bytes
 - Built on top of best-effort packet delivery at the network layer
- Challenges with IP
 - Lost or delayed packets
 - Corrupted packets
 - Out-of-order packet arrivals
 - Receiver runs out of space
 - Network cannot handle current load

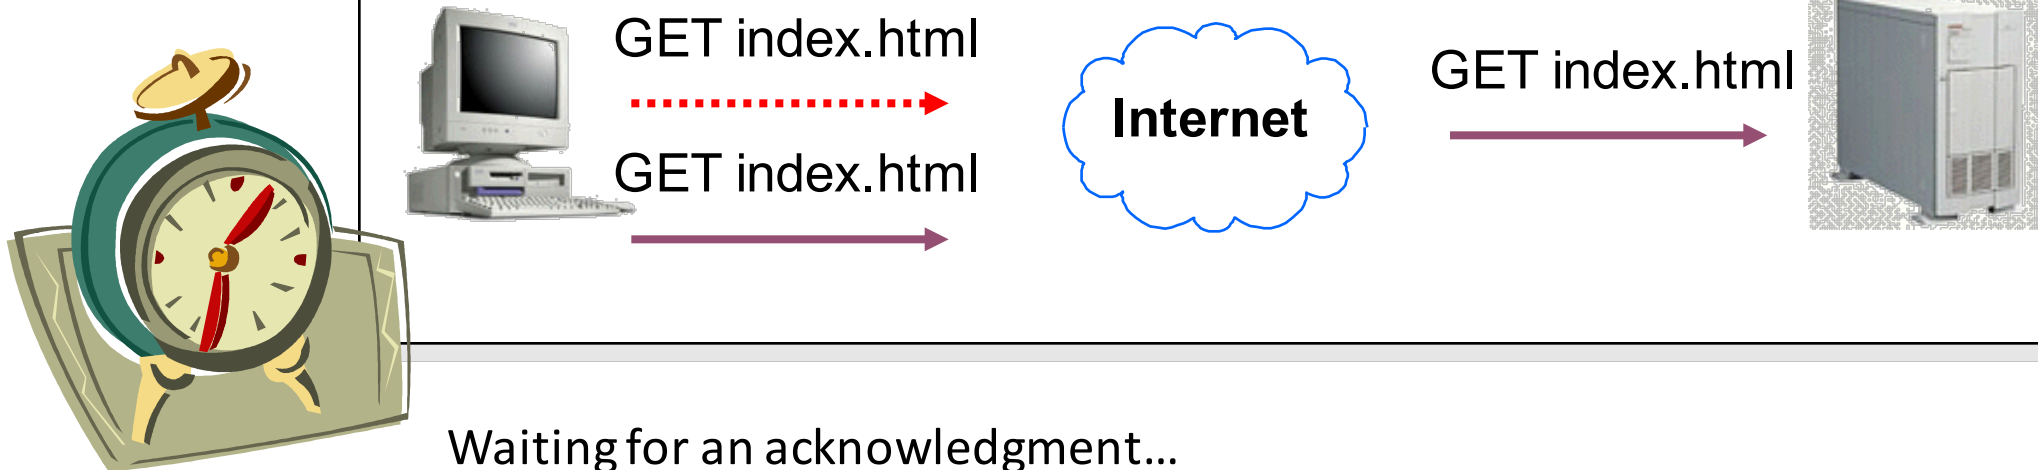


TCP: Lost or Delayed Packets

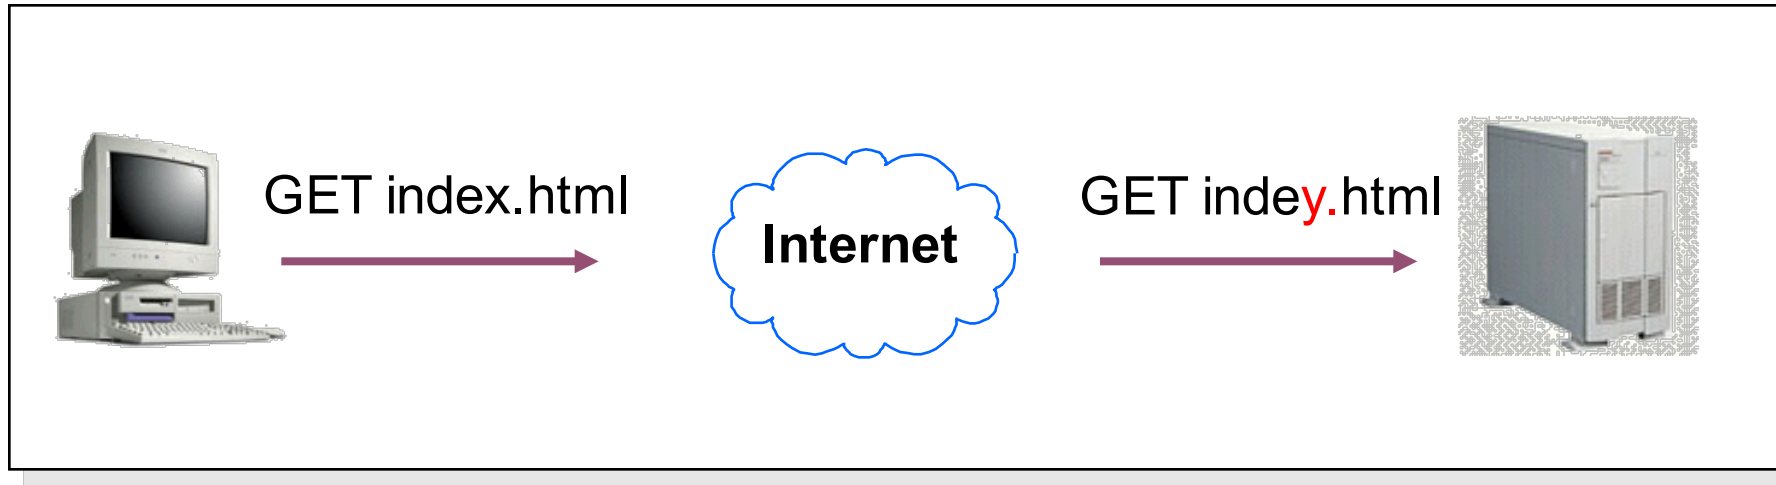
Problem: Lost or Delayed Data



Solution: Timeout and Retransmit




TCP: Corrupted Data




- Sender computes a checksum
134
 - Sender sums up all bytes in the payload + 212
 - And sends the sum to the receiver = 346
 - Receiver checks a checksum
134
 - Receiver sums up all bytes in the payload + 216
 - And compares against the checksum = 350
- Then what?*

TCP: Out-of-Order Packet Arrivals

Problem: Out of Order



Solution: Add Sequence Numbers



TCP: Receiver that Runs Out of Space



?



- Receiver maintains a *window size*
 - Amount of data it can buffer
- Advertises window to the sender
 - Amount sender can send without acknowledgment
- Ensures that sender does not send too much
 - While still sending as much as possible

Flow control!

TCP: Receiver that Runs Out of Space



?



- Receiver maintains a *window size*
 - Amount of data it can buffer
- Advertises window to the sender
 - Amount sender can send without acknowledgment
- Ensures that sender does not send too much
 - While still sending as much as possible

Flow control!

TCP: Network that Cannot Handle the Load

- Problem: Too many packets at once
- Solution: Congestion control
 - Future lecture!



Transport: User Datagram Protocol (UDP)

- Datagram of bytes
 - A message

*UDP does less than TCP,
why do we want UDP too?*

- Challenges with IP
 - Lost or delayed packets X
 - Corrupted packets ✓
 - Out-of-order packet arrivals X
 - Receiver runs out of space X
 - Network cannot handle current load X

Layering & Network Layers Conclusion

- The art of layering
- Network layers
 - Protocol, headers, encapsulation
- IP layer: best-effort global packet delivery between host
- TCP layer: ordered, reliable byte stream between applications