

The Cloud

COS 316: Principles of Computer System Design

Amit Levy & Jennifer Rexford

Layering: When it works we get

- Modularity
 - Application developer and ethernet developer can work independently
- Portability
 - The narrow waist of IP
- Hide complexity with abstraction
 - Simple applications on top of TCP
- Re-use
 - Many transport protocols on top of IP

Layering and Its Discontents

- How layering can *fail* and why abstraction can sometimes be *bad*
 - Poor choices of abstraction
 - Too much abstraction
- When to “pop open the hood”?
- Important, ubiquitous systems sometimes an accident of history

What Problem Does the Cloud Solve?



What Problem Does the Cloud Solve?



=

\$\$\$\$\$\$\$\$\$...

The Cloud Problem

Want to get resources in a datacenter, but:

- Relatively few resources (multiplexing)
- For a short period of time (burst scalability)
- Without having to trust my neighbors (strong isolation)
- For cheap

Why *datacenter* resources?

- Reliable (e.g. redundant power and network)
- Efficient (e.g. close to cheap power, economies of scale)
- Fast (e.g. high speed data links)
- ...

A person is sitting on a rocky cliff, looking out over a vast, endless sea of white clouds. The sun is setting on the horizon, casting a warm, golden glow across the sky and the clouds. The person is wearing a dark jacket and shorts, and is looking towards the right side of the frame. The clouds are thick and billowy, filling the entire lower half of the image. The sky is a mix of blue and orange, with some wispy clouds near the horizon.

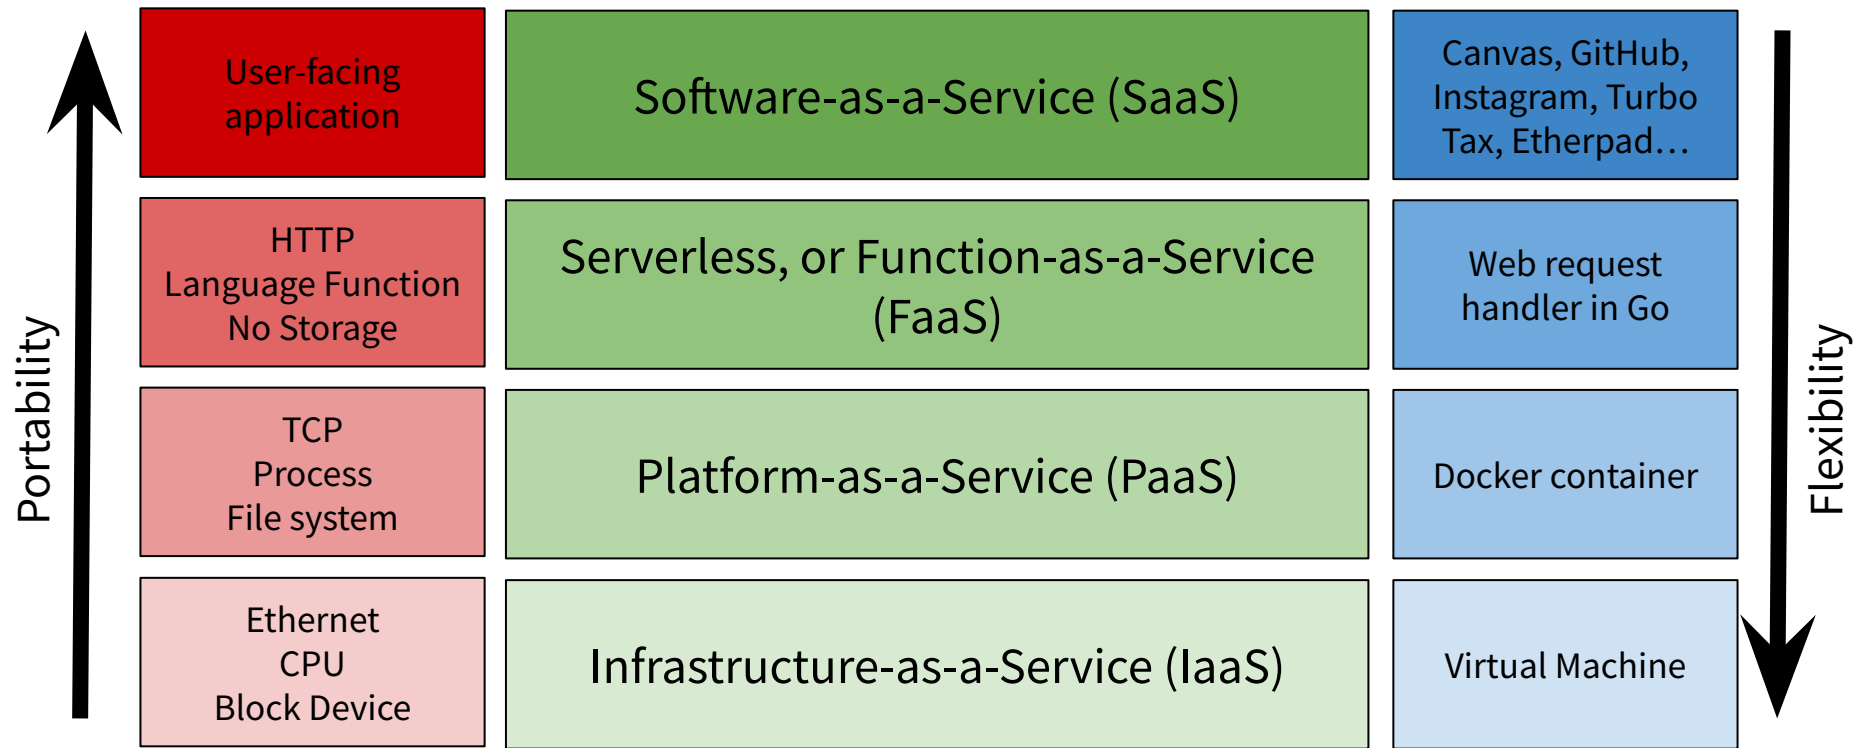
What is The Cloud?

What is The Cloud?


- An abstraction of unbounded **compute** and **storage**
 - We will only focus on compute today
- A set of software *abstraction layers* that enable, variously:
 - Burst scalability
 - Resource multiplexing
 - Strong isolation
 - High utilization
 - Programmer convenience*
- As much an economically-drive organization of resources as an engineering one

*This is almost always better achieved with a library than an enforced layer

OK, But What is The Cloud?



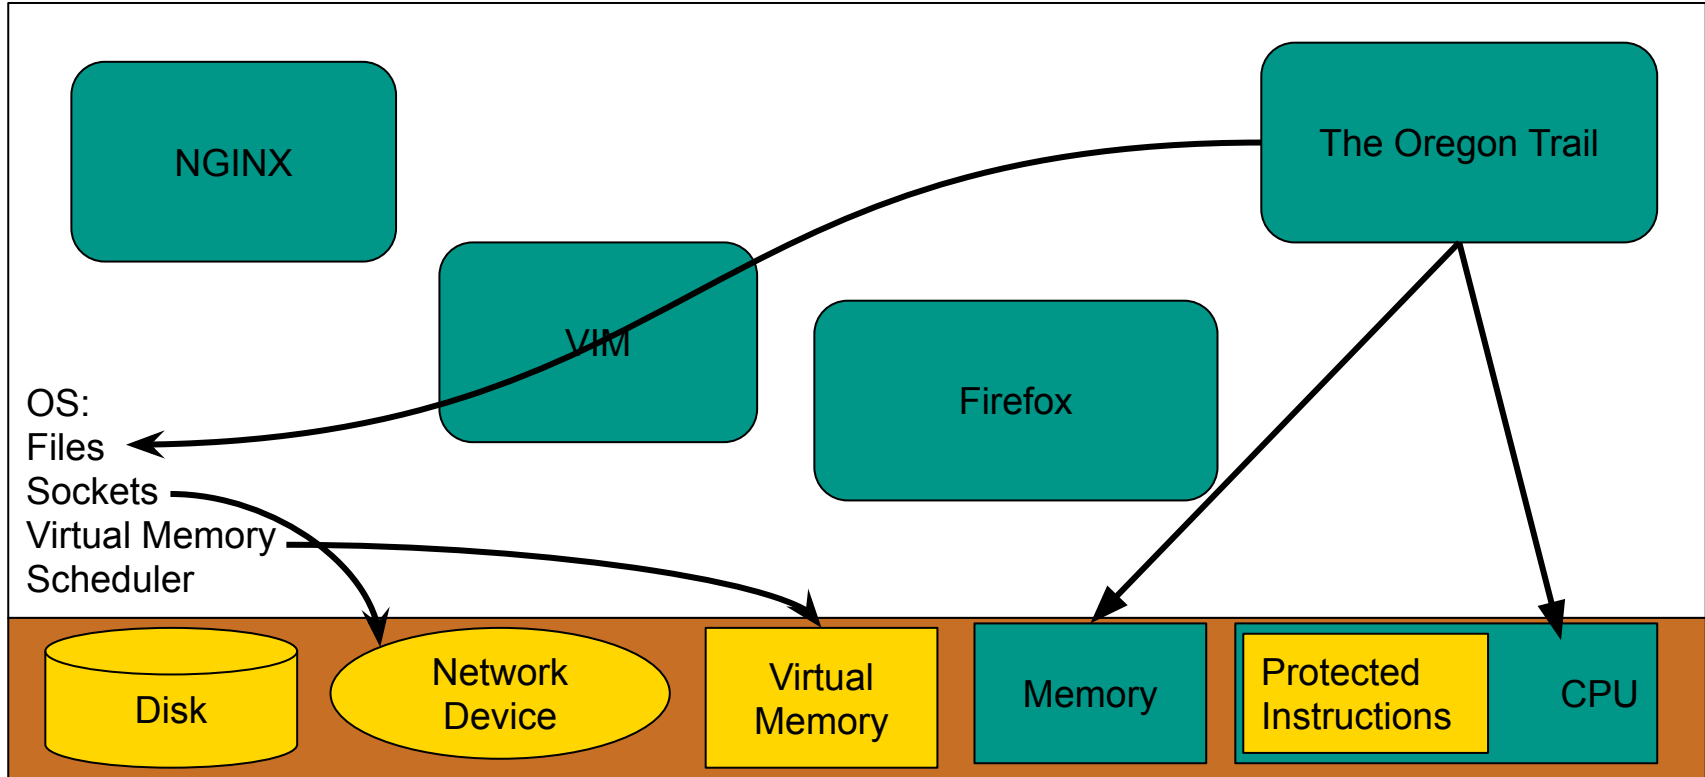
OK, But What is The Cloud?



From Virtual Machines to Cloud Computing



Traditional Operating System



Operating Systems Have *Thick* Abstractions

UNIX File System

- Files divide disk into non-contiguous chunks.
- File system layers need to interpose between applications and disk

Sockets

- Use transport-layer specific abstractions
 - E.g., allocate by tuple of (destination address, destination port, source port)
- Networking subsystem routes each packet, allocates new sockets

Big abstractions & multicore

The year is [some year in the 1990s]...

CPUs have **1 core**

Operating systems are **monolithic** and **centralized**

But...

Single-core CPU acceleration is slowing...

Your next computer will have **many cores**



Scalability: Disco (1997)

“Extensive modifications to the operating system are required to efficiently support scalable machines.”

“[W]e examine the problem of extending modern operating systems to run efficiently on large-scale shared-memory multiprocessors without a large implementation effort. [...] We use virtual machines to run multiple commodity operating systems on a scalable multiprocessor.”

“Disco: Running Commodity Operating Systems on Scalable Multiprocessors”

Edouard Bugnion, Scott Devine, Kinshuk Govil, and Mendel Rosenblum

Ed, Scott & Mendel (along with Diane Green) founded VMWare a year later
in October 1998

Scalability: Disco (1997)

From Edouard Bugnion's job talk at EPFL in 2014:

“First, Disco ran commodity operating systems on scalable MIPS multiprocessors. [...] Second, VMware Workstation is a successful commercial product that allows...”

Work by a few grad students

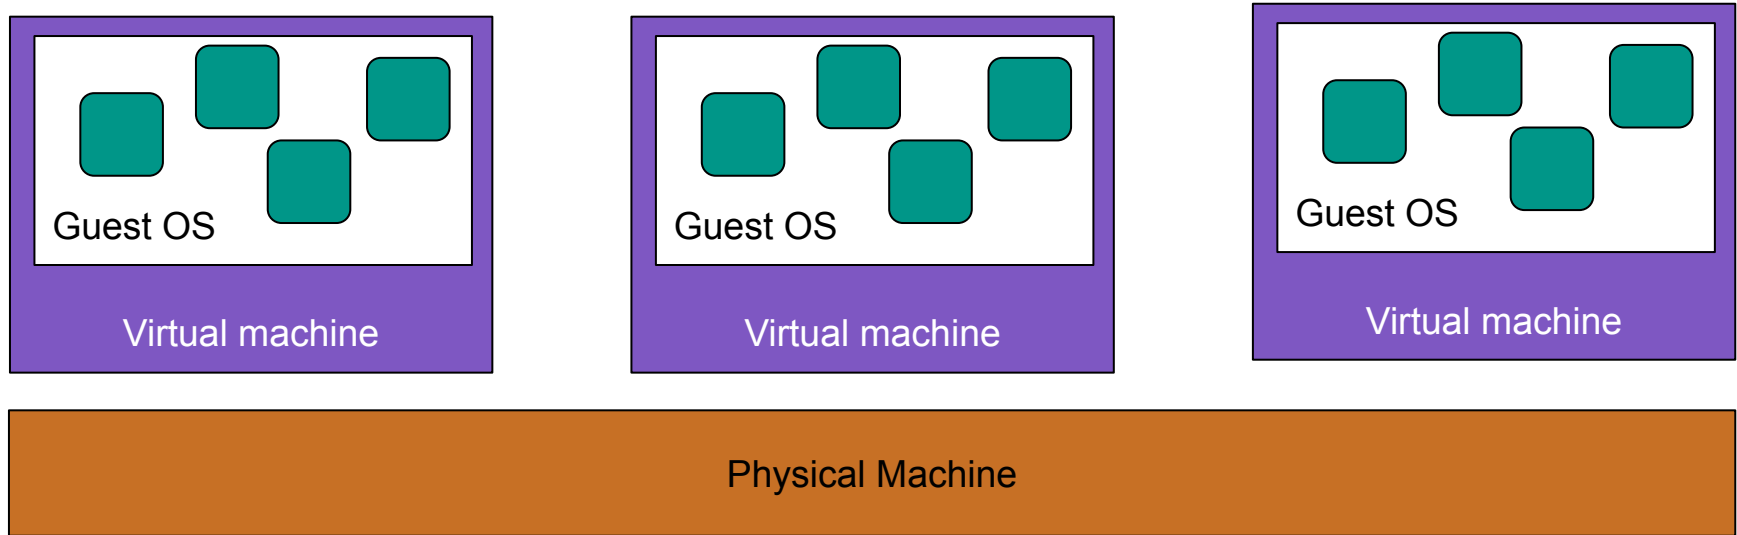
(along with Diane Gre

1

SAME!

Work by 10,000s of employees
at a \$70B Company

Virtual Machines



Virtualization

presents a physical machine as though many guest OSs had exclusive access


Isn't a VMM just an Operating System?

- Yes!
- No: API is hardware resources (disk, network card), not abstract interfaces (file system, socket)

Virtual Machines vs. Processes

	Virtual Machine Interface	Process Interface
<i>Network</i>	Network device (ethernet, WiFi, etc...)	TCP & UDP sockets
<i>Storage</i>	Block device	File System
<i>Compute</i>	CPU	Unprivileged subset of CPU (x86/ARM/RISC-V...)
<i>Memory</i>	Virtual & physical memory addresses	Virtual addresses only

Virtual Machines



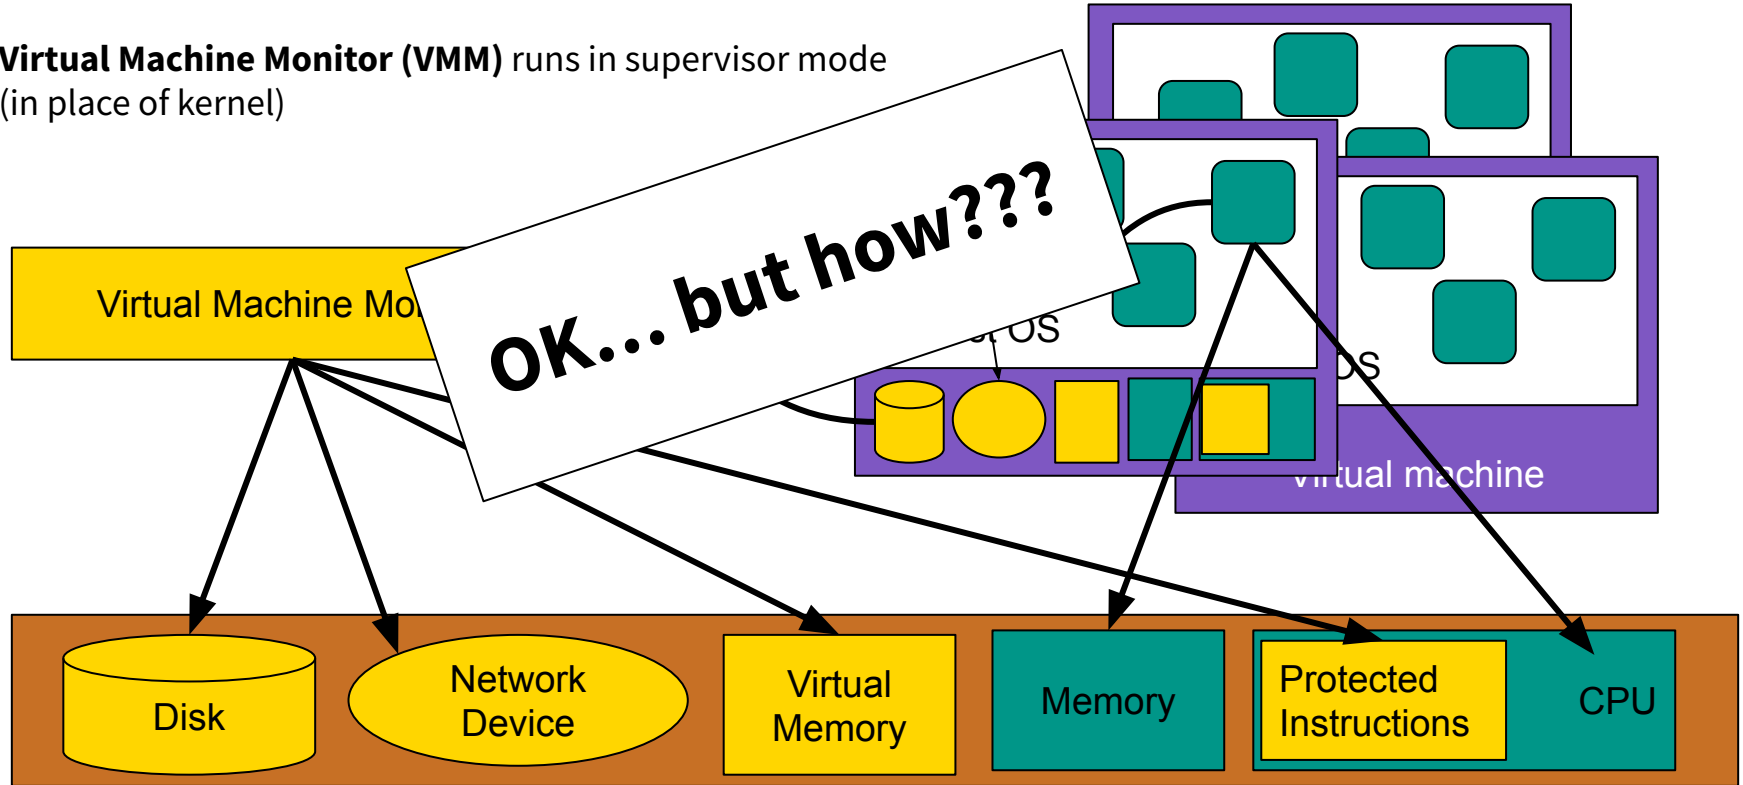
Virtualization

presents a physical machine as though many guest OSs had exclusive access

But How??

Virtual Machines

Virtual Machine Monitor (VMM) runs in supervisor mode
(in place of kernel)



Trap-and-Emulate: System Calls

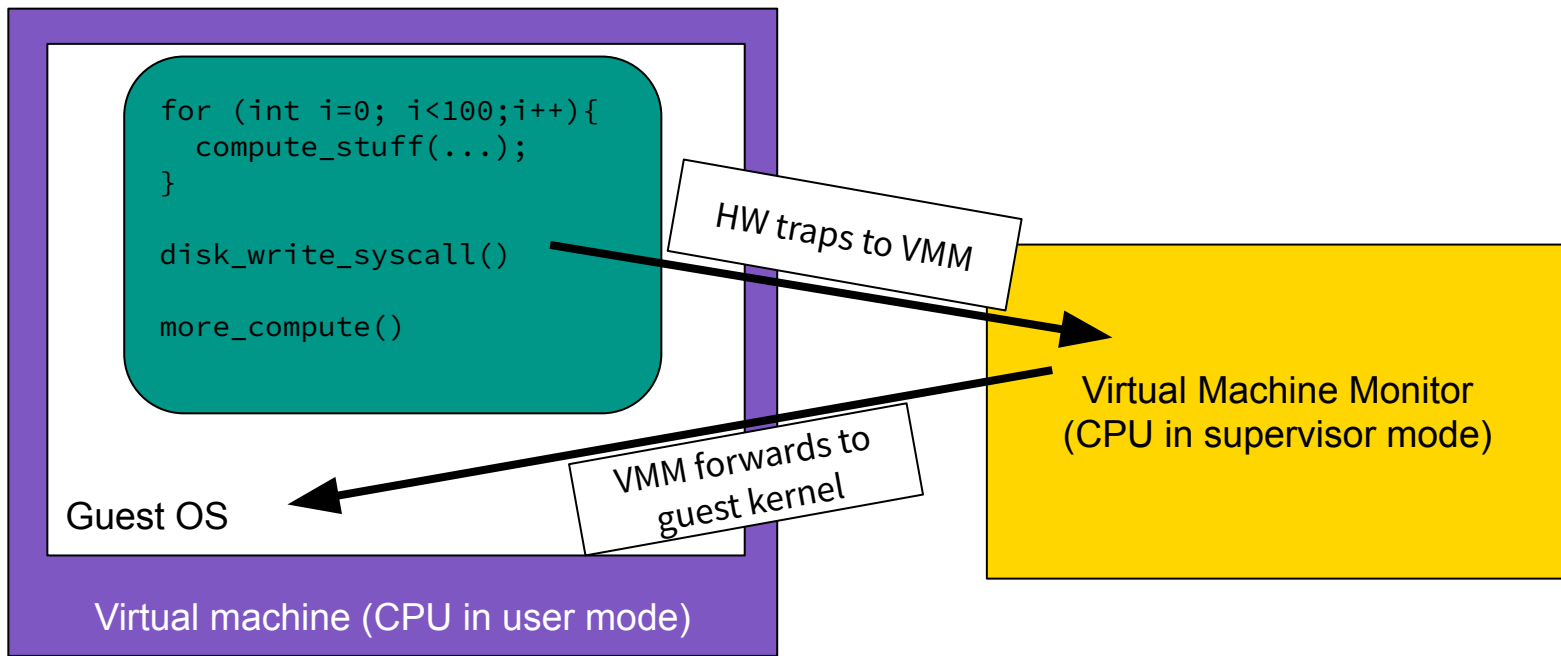
Assumes typical separation of applications and operating systems:

1. Applications interact with hardware by issuing system calls
2. System calls trapped by hardware which switches to OS context
3. OS uses *protected instructions* to access hardware directly

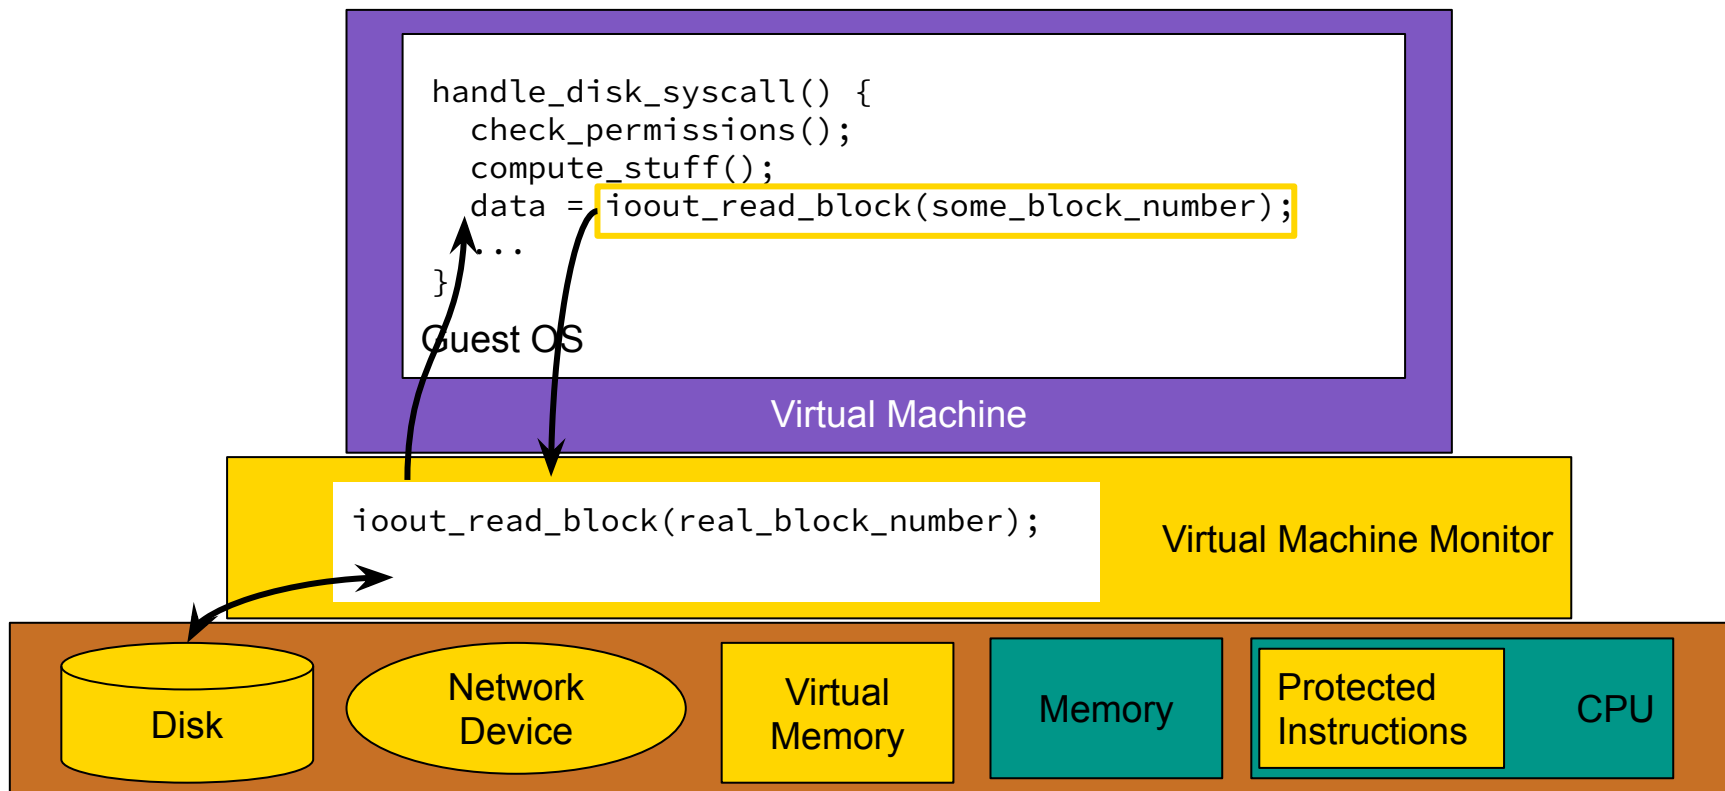
Virtual Machine Monitor “routes” system calls to appropriate guest OS

OS, running in process mode, causes trap through hardware on *protected instruction*

Trap-and-Emulate: System Calls



Trap-and-Emulate: I/O



What if we *can't* trap-and-emulate?

x86 didn't fault when some protected instructions were executed by user programs.

E.g. setting interrupt flags with Pop-Flags (POPFL) instruction silently fails

- Binary translation
 - Before running VM code for the first time, translate it to replace with explicit calls to VMM
 - Disco & early VMWare
- Para-virtualization
 - Modify guest OSs to detect if they are running inside a VM and use different instructions
 - Xen (2003) & most VMMs since
- Modify the hardware
 - Intel VT-X, APICV, VT-d, SR-IOV, GVT-d, and on and on... starting 2005

Flexibility: VMWare Workstation, Parallels

1999-~2006

Use virtual machines alongside physical machines

Run Windows apps on Linux, and Linux apps on Windows

Cloud Computing: Amazon EC2

“Before the advent of Amazon EC2, you had to buy or rent sufficient servers to cover your present needs, and you also had to be able to anticipate [and] forecast [...] for enough hardware to accomodate(*sic*) [...] growth as well as bursts of traffic [...]

With Amazon EC2, you don't need to acquire hardware in advance of your needs. Instead, you simply turn up the dial, spawning more virtual CPUs, as your processing needs grow.”

-Amazon EC2 Announcement, August 25th 2006 *Jeff Barr*

Did VMMs solve single-server scalability?

- Yes! Can harness increasing number of cores to run more virtual machines!
- No: VMMs don't help scale *individual* applications
- No: VMMs don't mediate *sharing* between virtual machines

A decade later:

“An Analysis of Linux Scalability to Many Cores”

Silas Boyd-Wickizer, Austin T. Clements, Yandong Mao, Aleksey Pesterev, M. Frans Kaashoek, Robert Morris, and Nickolai Zeldovich

Careful adaptation of OS kernels broke the scalability barrier

Did VMMs solve The Cloud Problem?

- Yes! Can provision datacenter resources in small chunks, for hours at a time
- No: not small enough chunks! An hour is too long!
 - PaaS, FaaS to the rescue
- Yes: actually VMMs *can* be provisioned in smaller chunks for less time
 - “My VM is Lighter (and Safer) than your Container” Filipe Manco et al, SOSP 2017
 - Firecracker VM, 2018
- Yes/No: Strong isolation between VMs!
 - Well... maybe not... wait till next time

Take Aways

- Traditional *implementations* of OS abstractions poor fit for multicore
- Difficult to re-implement thick abstractions
- VMMs introduced a much simpler machine abstraction
 - Just run multiple OSs to “scale”
 - A decade to take advantage of increased density before OSs caught up
- Would we still build VMWare/EC2 today?
- Up next... what happens when annoying academics get their hands on the cloud?