

Challenges with Layering: a Video Conferencing Case Study



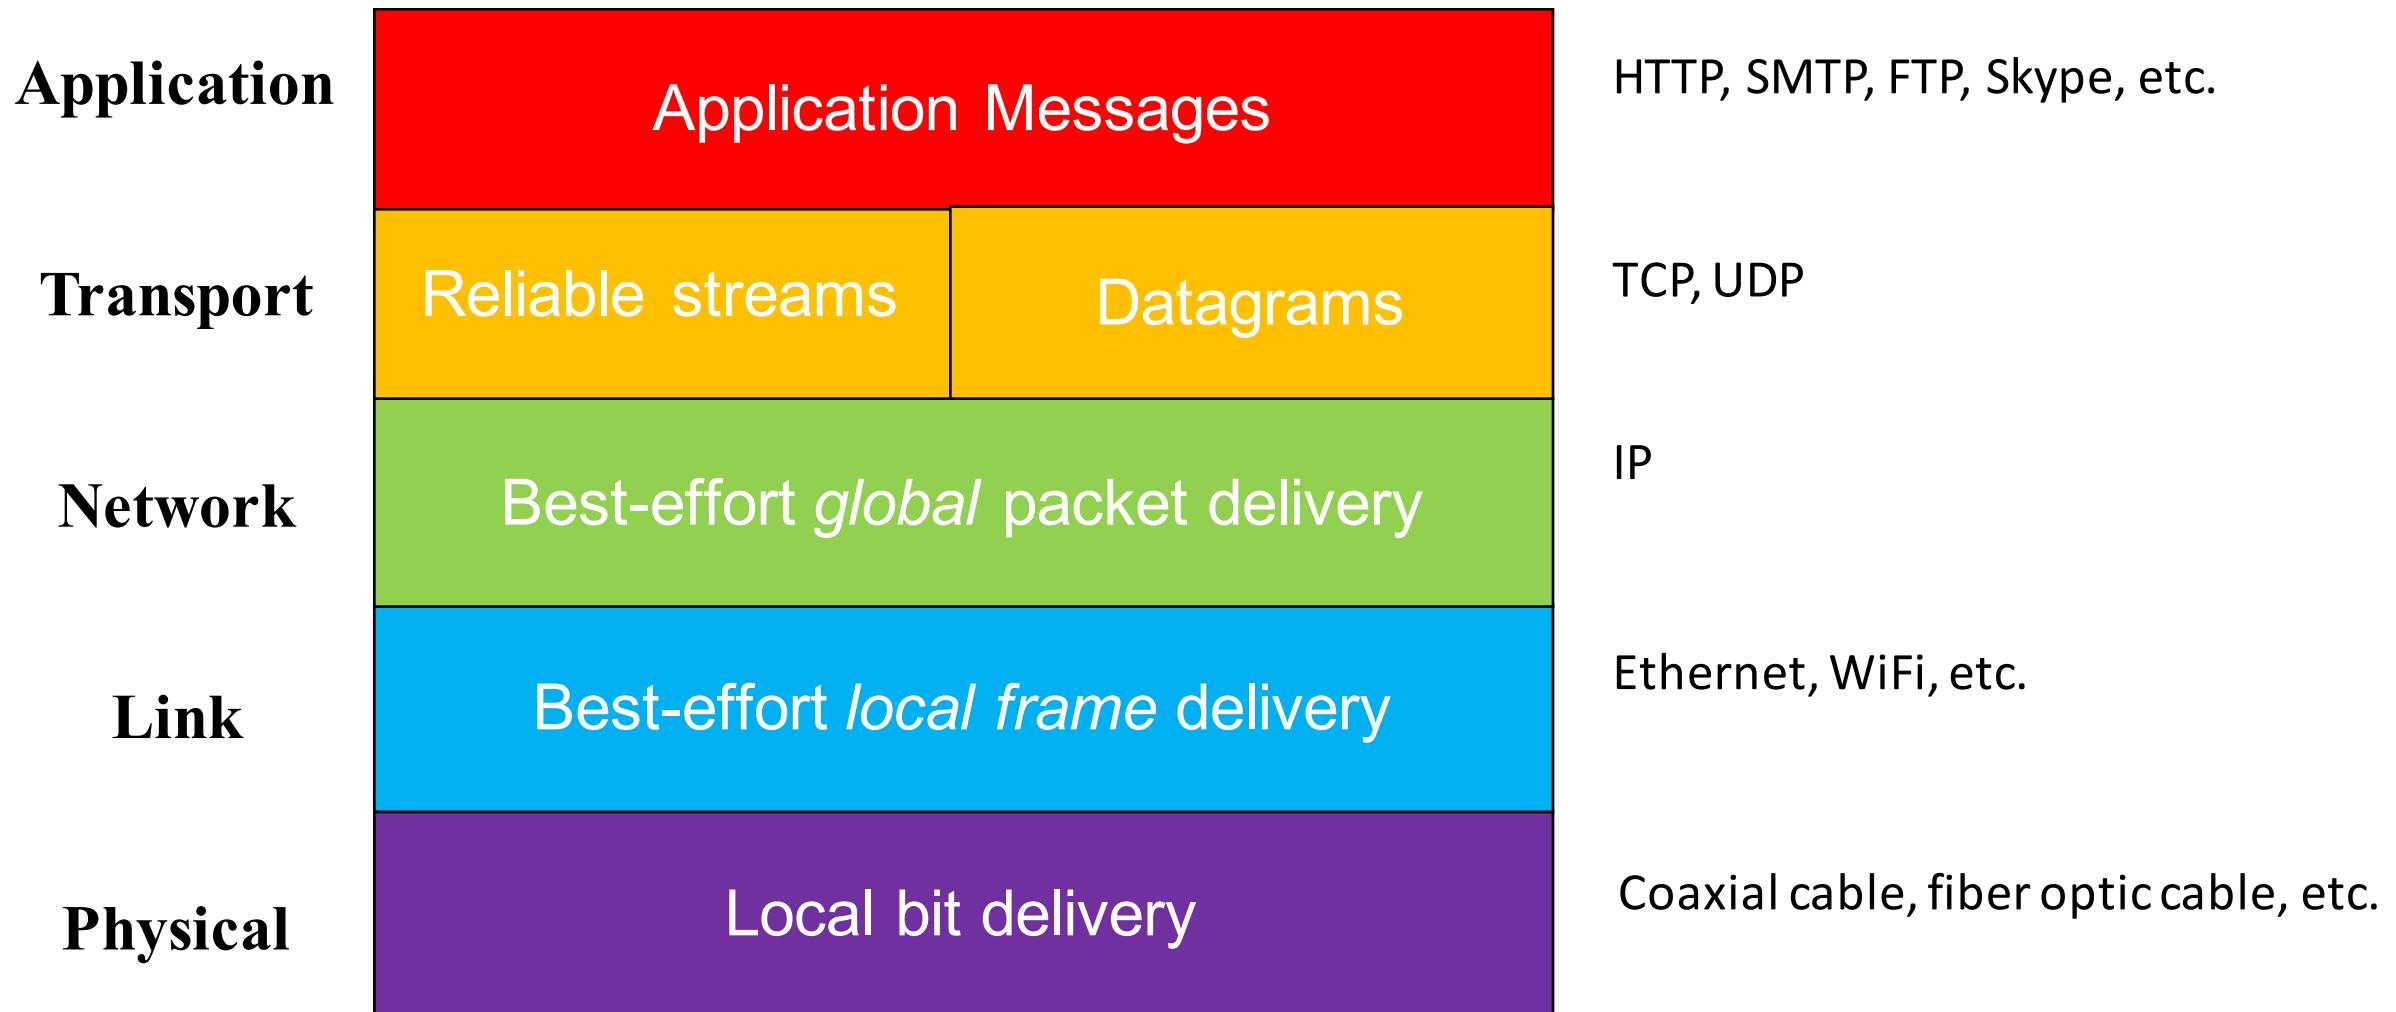
COS 316: Principles of Computer System Design
Lecture 12

Amit Levy & Ravi Netravali

Modularity Through Layering

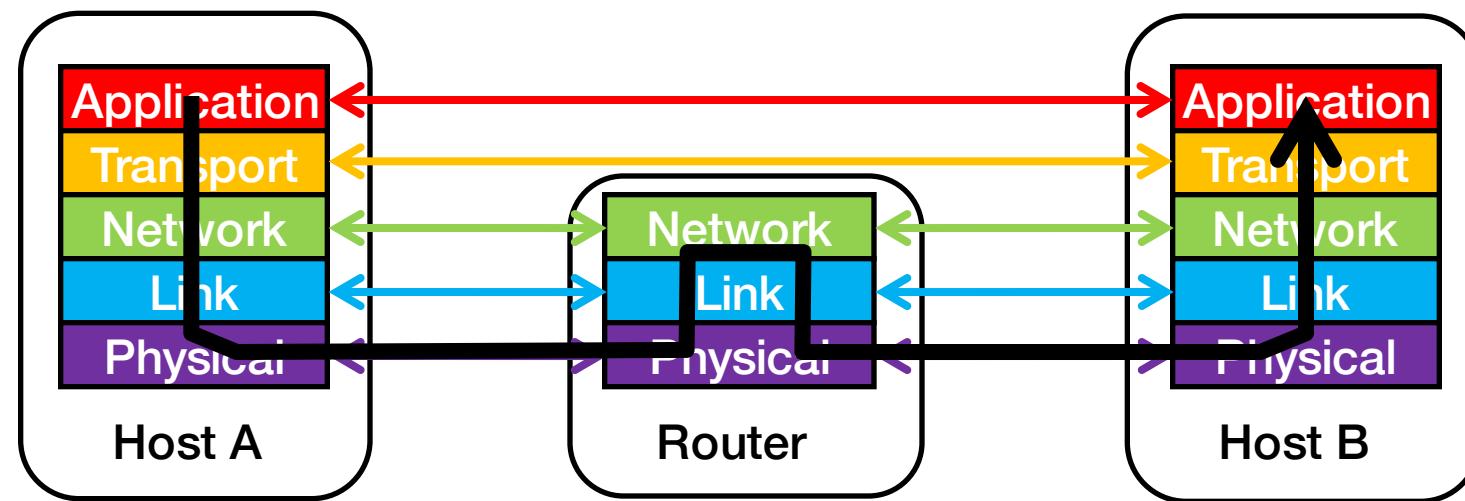
- Systems on systems on systems though layering
- Each layer hides complexity with abstraction
- Pro: simplifies design of each layer (and overall system)

Internet Protocol Layers



Physical communication

- Communication goes down to the **physical network**
- Then from **network** peer to peer
- Then up to the relevant application



Cons of Layering?

- Focus on generality rather than top-to-bottom optimality
- Lack of visibility across layers
- Potential lack of coordination across layers

Salsify: Low-Latency Network Video Through Tighter Integration Between a Video Codec and a Transport Protocol

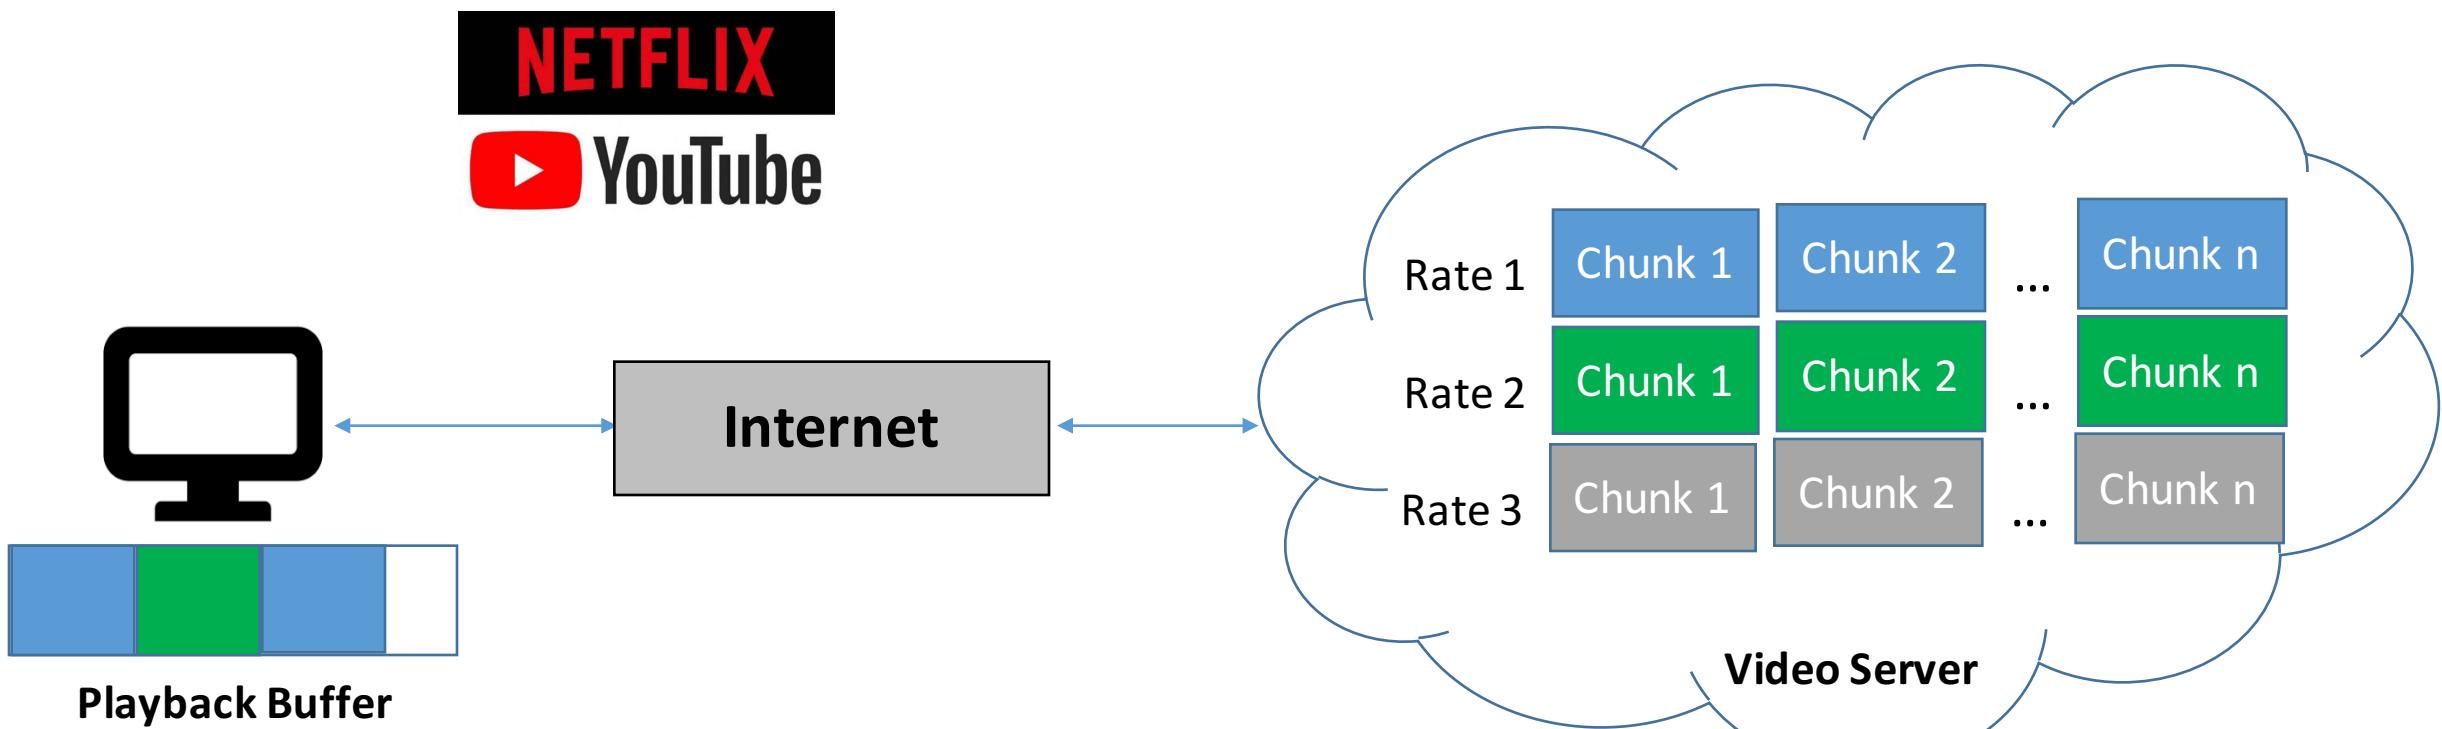
Sadjad Fouladi, John Emmons, Emre Orbay, Catherine Wu[†], Riad S. Wahby, Keith Winstein

Stanford University, [†]Saratoga High School

<https://snr.stanford.edu/salsify>

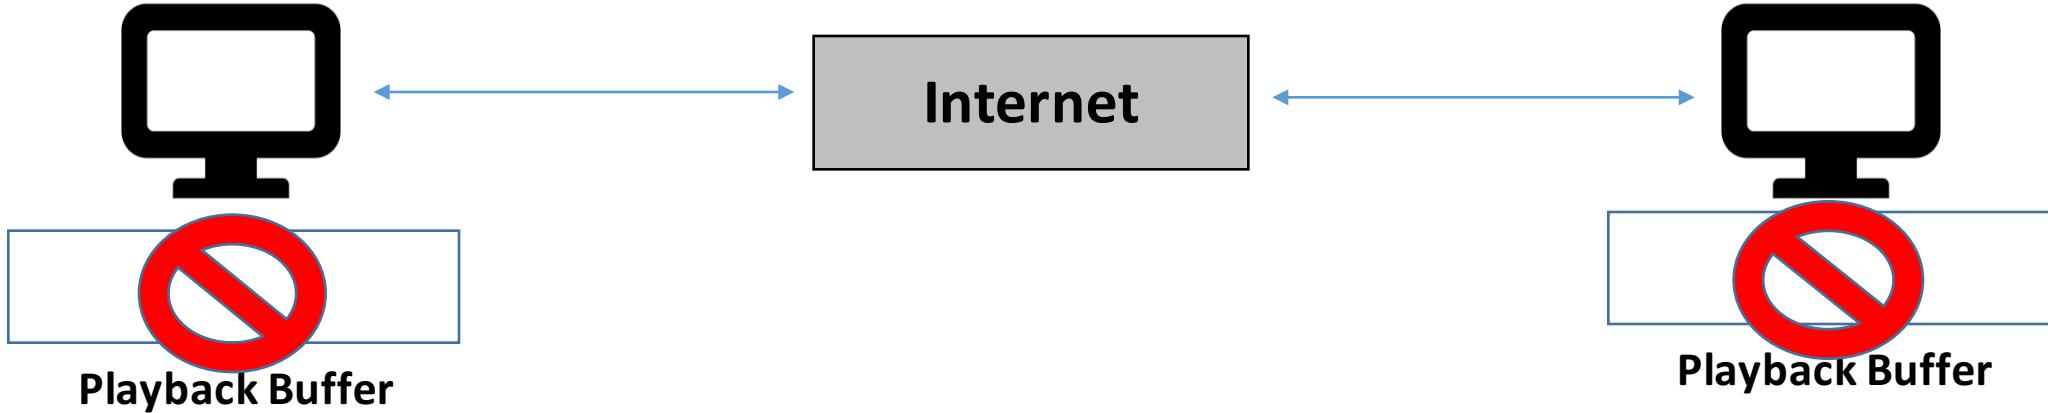
Stanford

Video Systems: Streaming



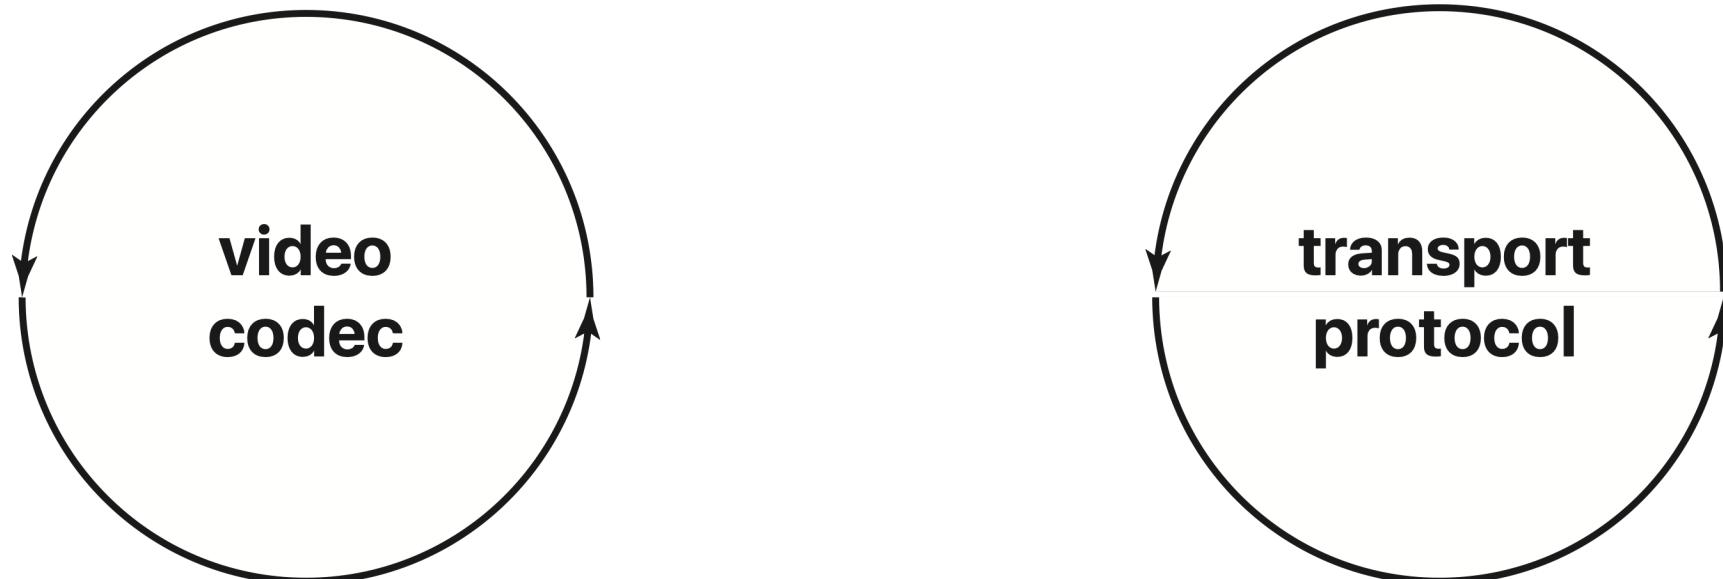
- Use **TCP**
 - Want reliability, not terribly latency-sensitive
- Video **encoded beforehand**
- Uses **playback buffer** (downloads ahead of time)
 - Can hide errors in network prediction

Video Systems: Conferencing

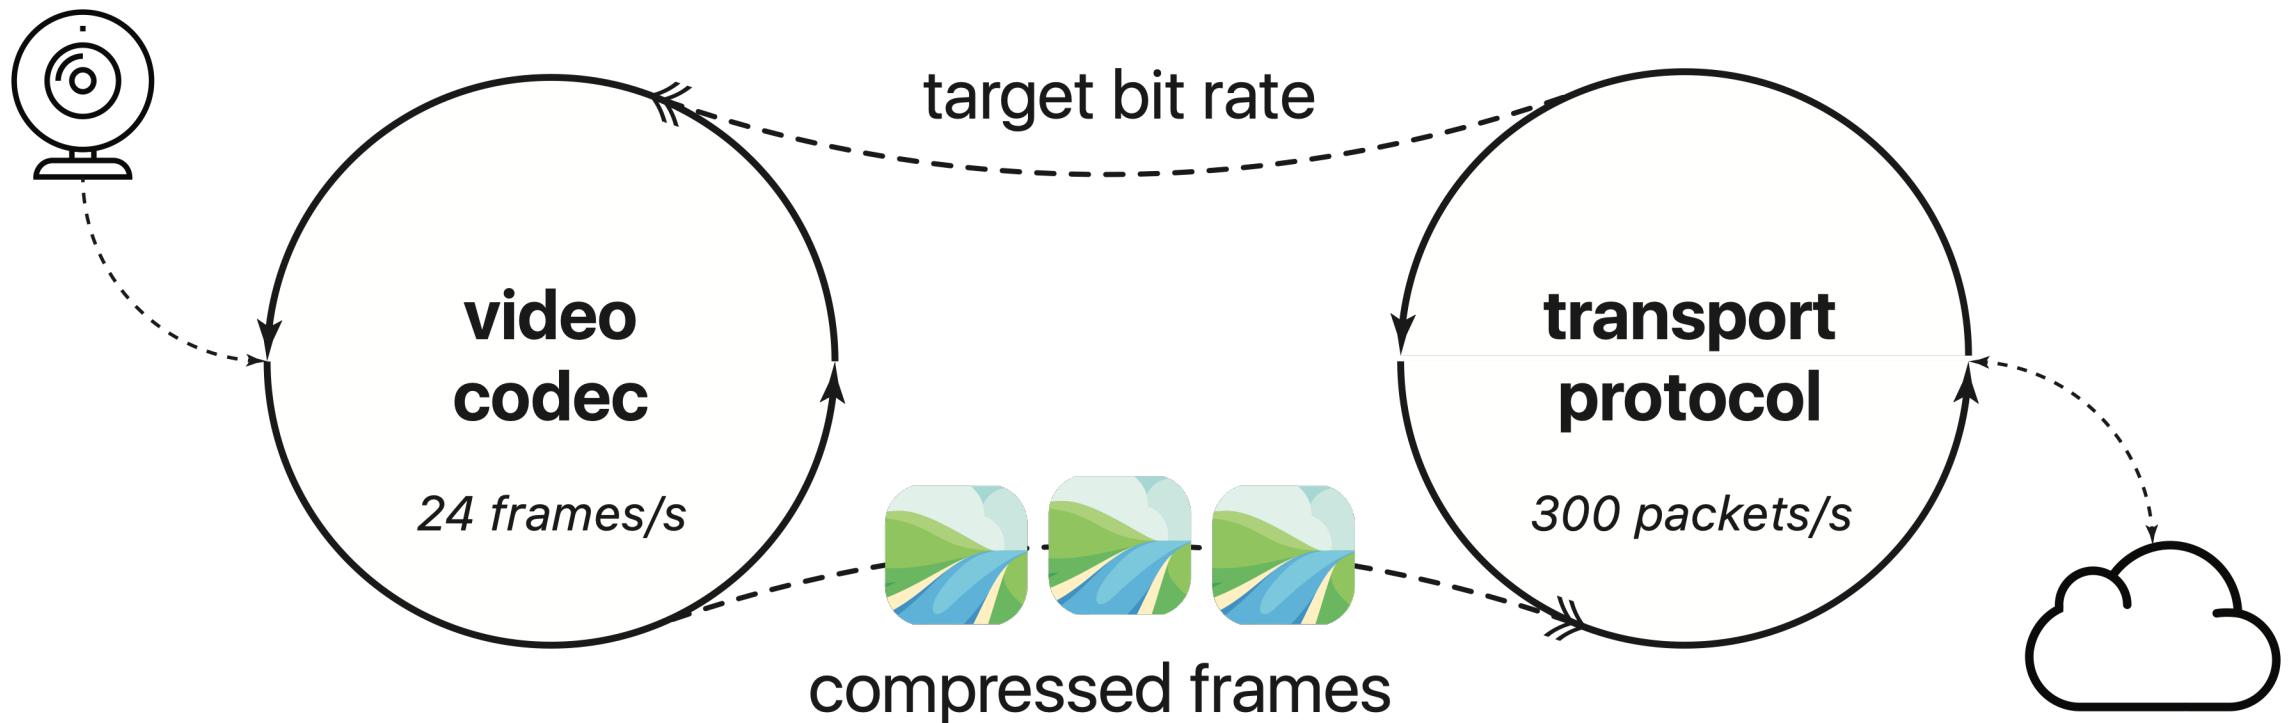


- Use **UDP**
 - Prioritize low latency over reliability
- Video **encoded on the fly**
- **No playback buffer**
 - Network prediction errors can be costly!

Today's systems combine two (*loosely-coupled*) components

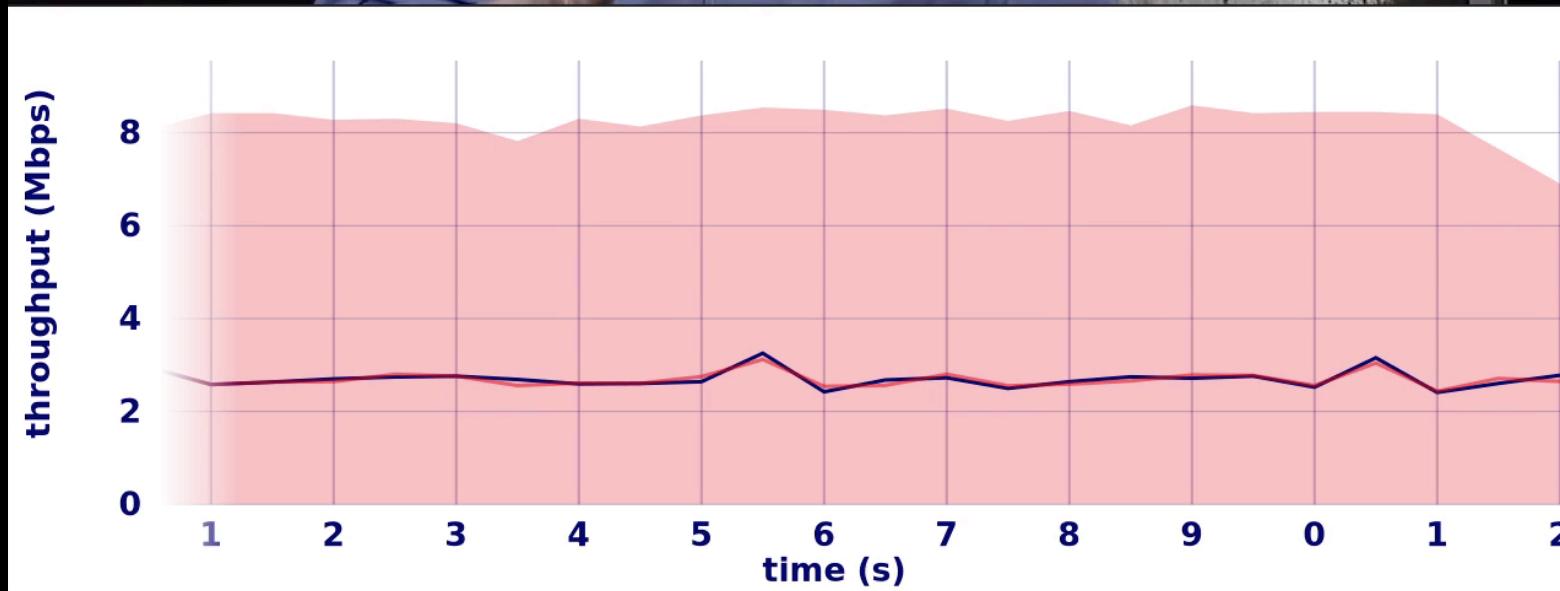


Two distinct modules, two separate control loops



Existing conferencing systems already break layering
we discussed! → expose network forecasts

WebRTC
(Chrome 65)



Current systems do not react fast enough to **network variations**, end up congesting the network, causing **stalls and glitches**.



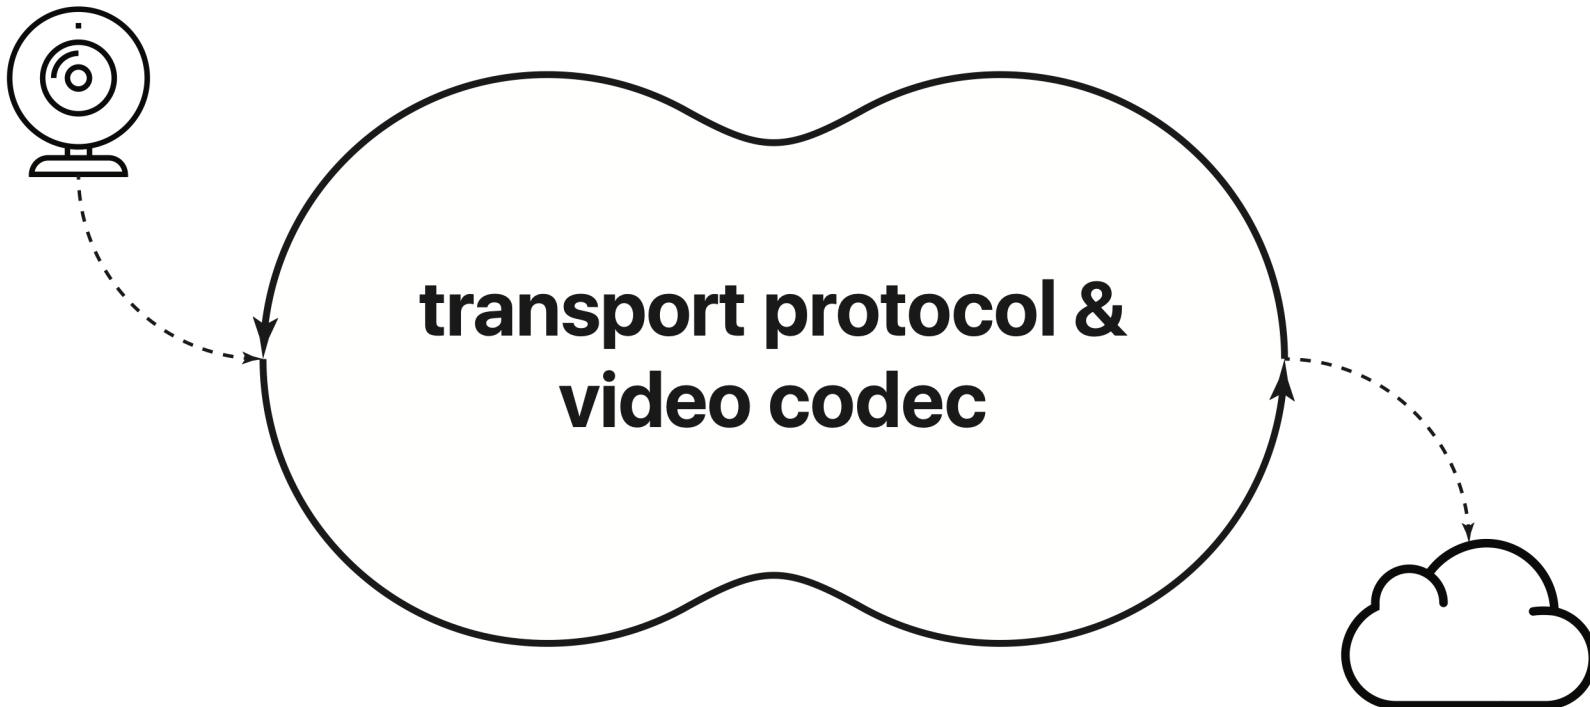
The problem: codec and transport are *too* decoupled

- The codec can only respond to changes in target bit rate over *coarse time intervals*.
 - Individual frames may cause packet loss/queueing.
 - The transport has little control over what codec produces.
- ⇒ The resulting system is slow to react to network variations.

Enter Salsify

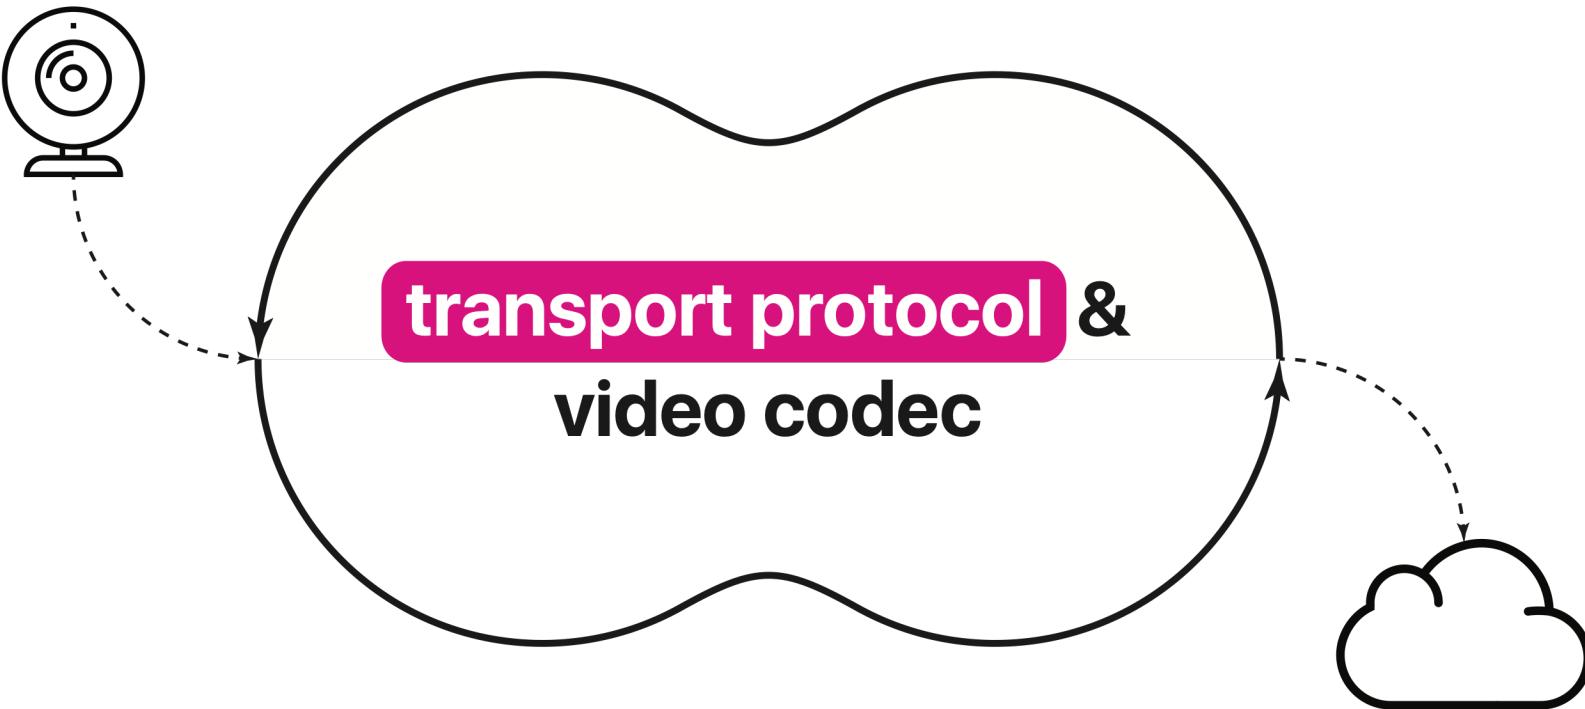
- Salsify is a new architecture for real-time Internet video.
 - Salsify tightly integrates a **video-aware transport protocol**, with a **functional video codec**, allowing it to **respond quickly to changing network conditions**.
- Salsify achieves **4.6× lower p95-delay** and **2.1 dB SSIM higher visual quality** on average when compared with FaceTime, Hangouts, Skype, and WebRTC.

Salsify explores a more tightly-integrated design



Salsify's architecture:

Video-aware transport protocol



Video-aware transport protocol

“What should be the size of the next frame?”

* without causing excessive delay

- There's no notion of bit rate, only the next frame size!
- Transport uses **packet inter-arrival time**, reported by the receiver.

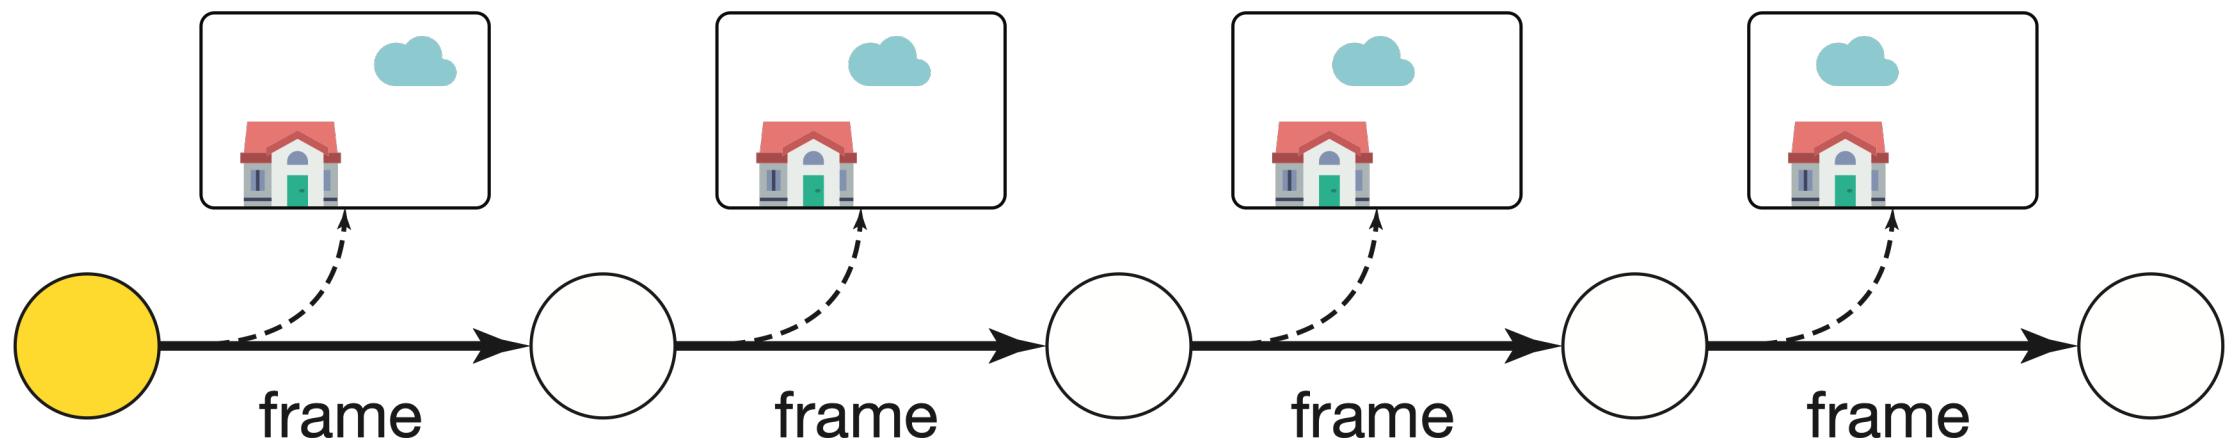
Transport tells us how big the next frame should be, but...

It's challenging for any codec to choose the appropriate quality settings upfront to meet a **target size**—they tend to over-/undershoot the target.

How to get an accurate frame out of an inaccurate codec

- Trial and error: Encode with different quality settings, pick the one that fits.
 - ***Not possible with existing codecs.***

After encoding a frame, the encoder goes through a state transition that is impossible to undo



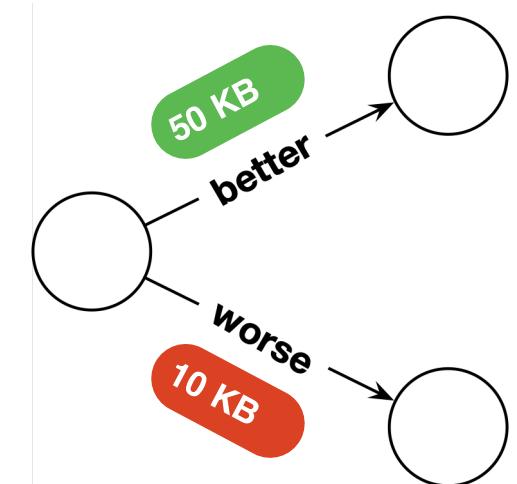
Functional video codec to the rescue

encode(state, \hat{a}) \rightarrow state', frame

Salsify's functional video codec exposes the state that can be saved/restored.

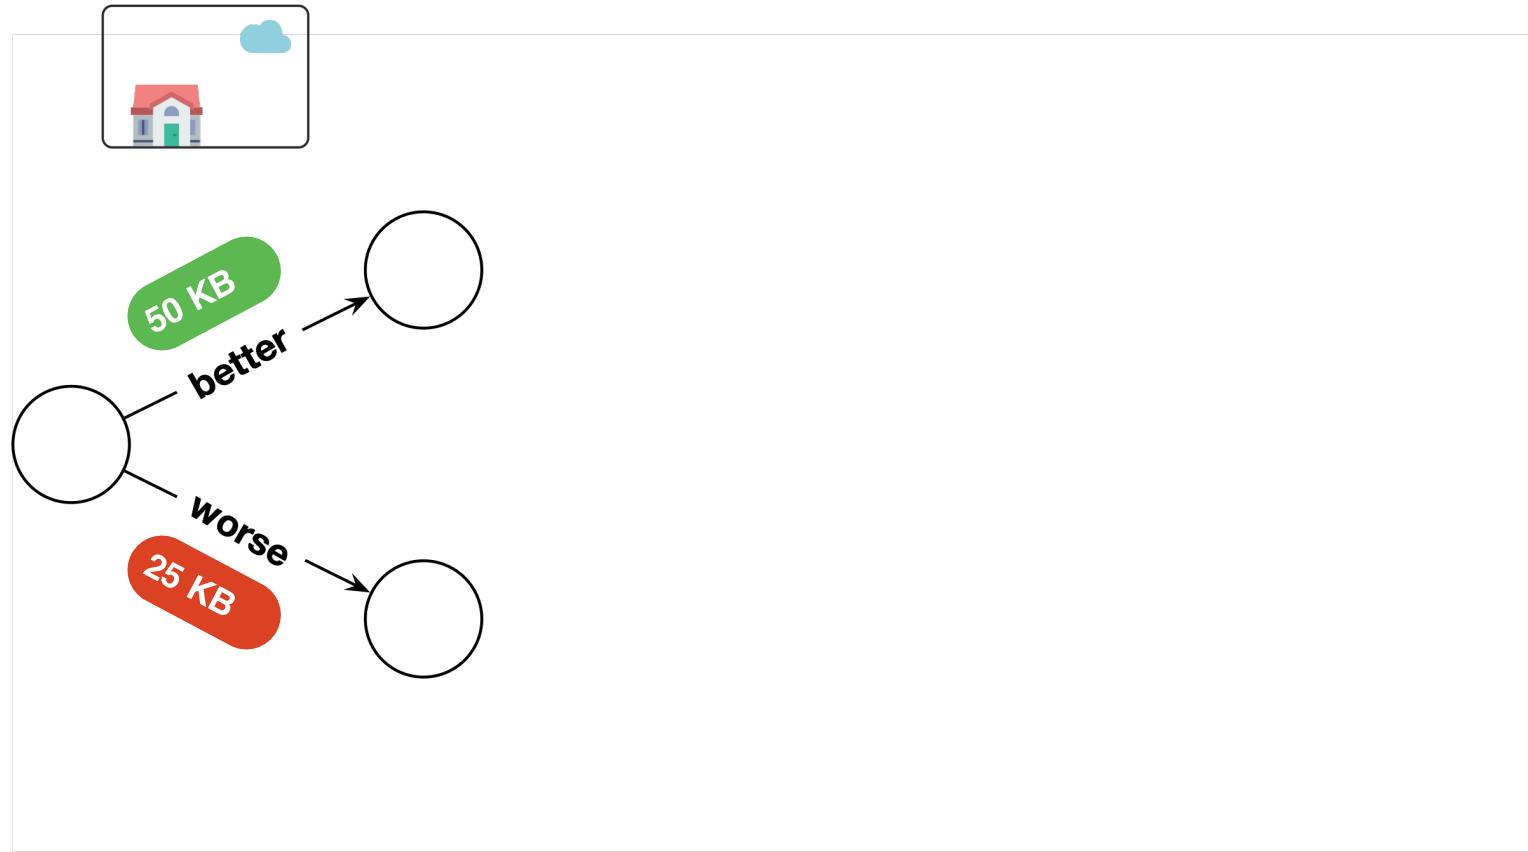
Order two, pick the one that fits!

- Salsify's functional video codec can **explore different execution paths** without committing to them.
- For each frame, codec presents the transport with *three* options:
 - ▲ A slightly-higher-quality version,
 - ▼ A slightly-lower-quality version,
 - ✖ Discarding the frame.



Codec → Transport

“Here’s two versions of the current frame.”

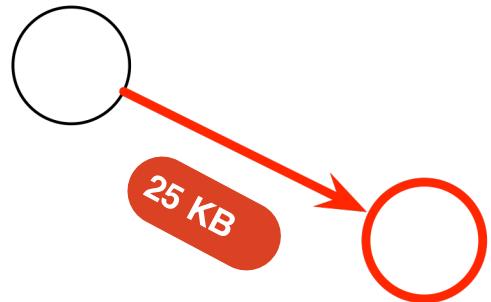
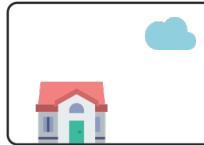


target frame size

30 KB

Transport → Codec

“I picked #2. Base the next frame on its state.”

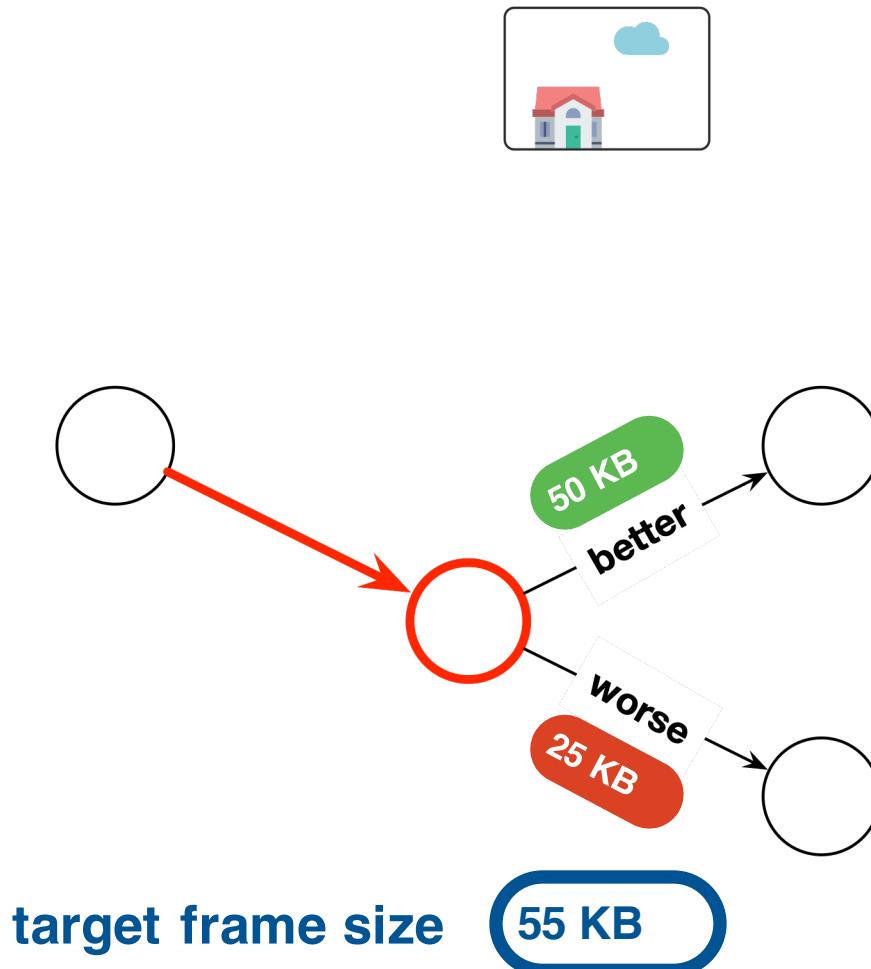


target frame size

30 KB

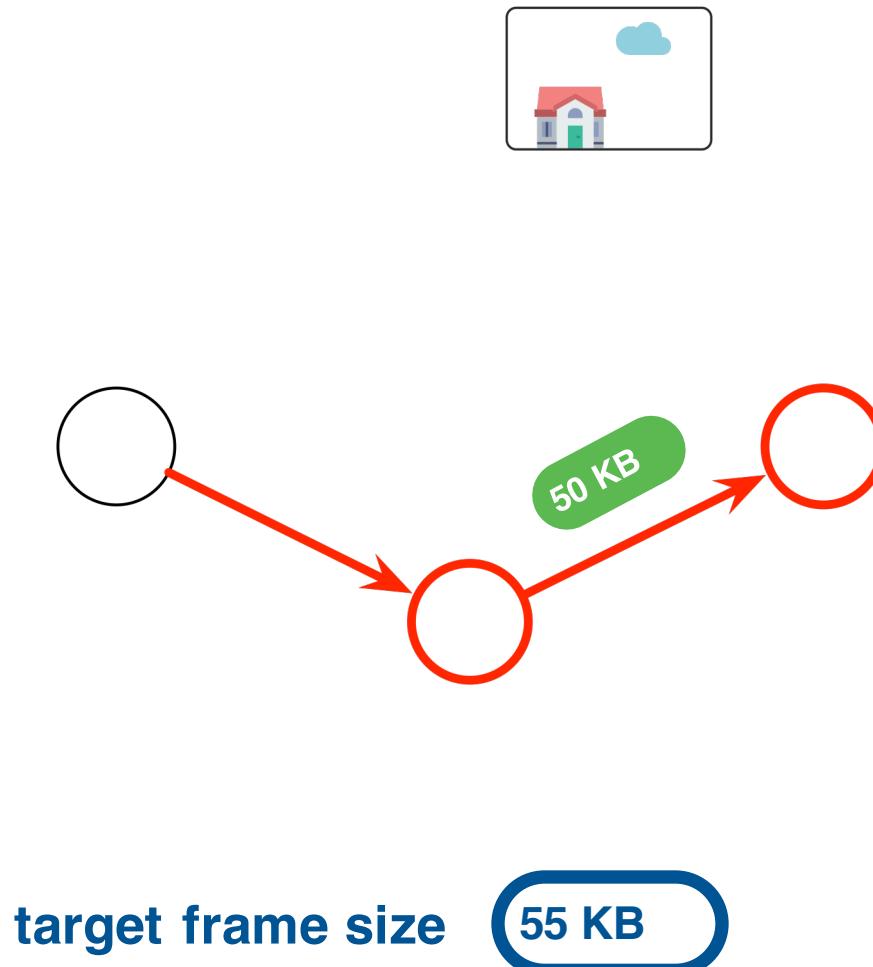
Codec → Transport

“Here’s two versions of the latest frame.”



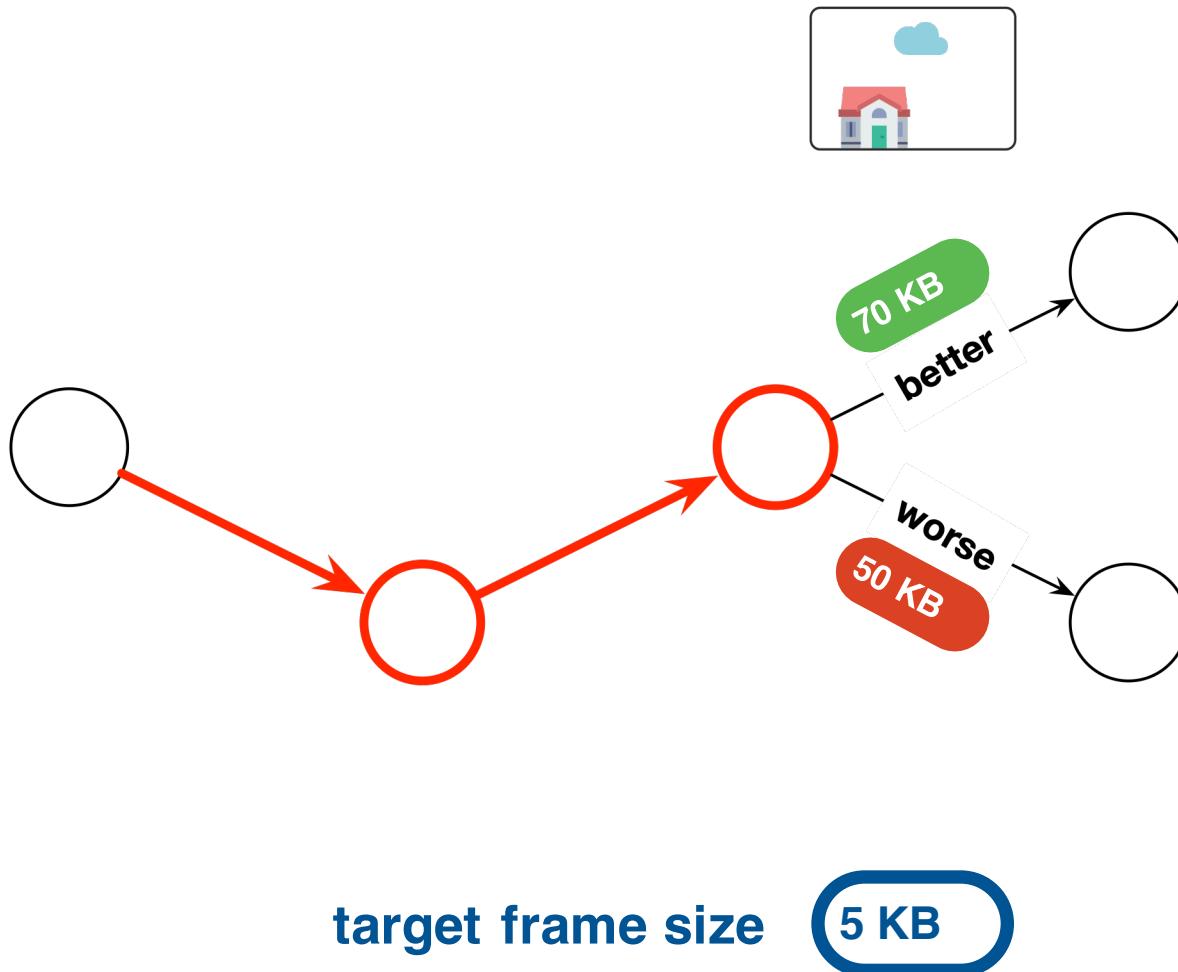
Transport → Codec

“I picked #1. Base next frame on its state.”



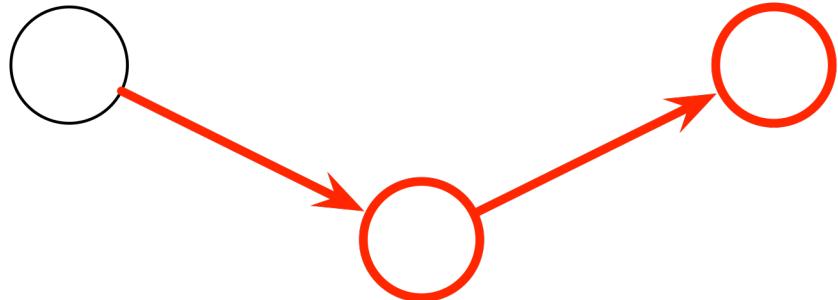
Codec → Transport

“Here’s two versions of the latest frame.”



Transport → Codec

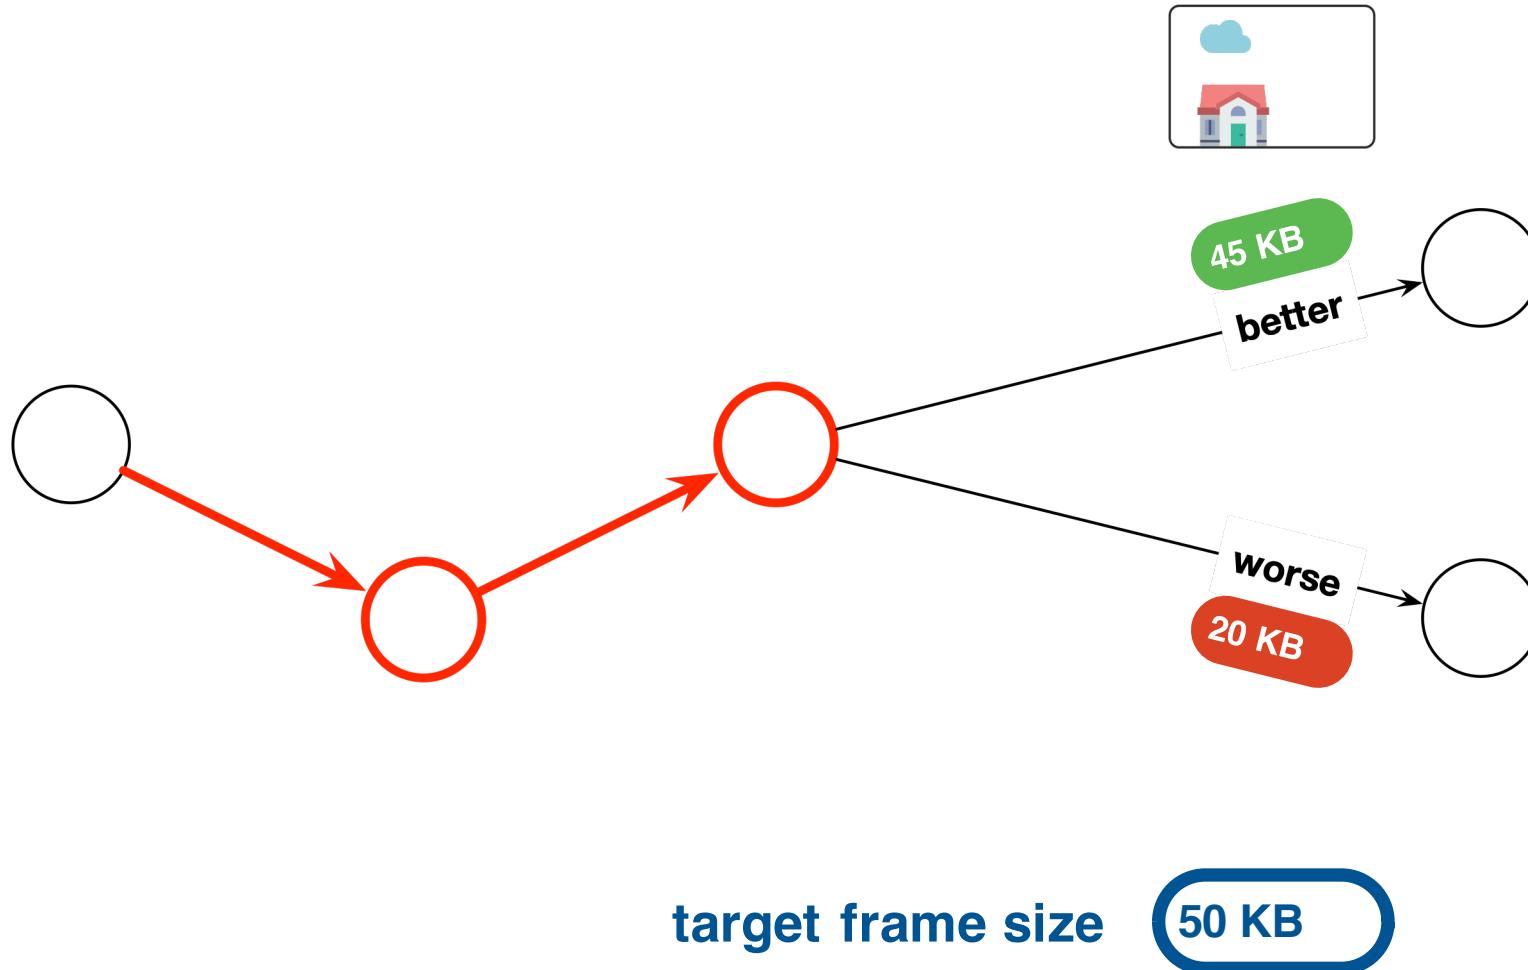
“Can’t send frames right now. Discard them.”



target frame size 5 KB

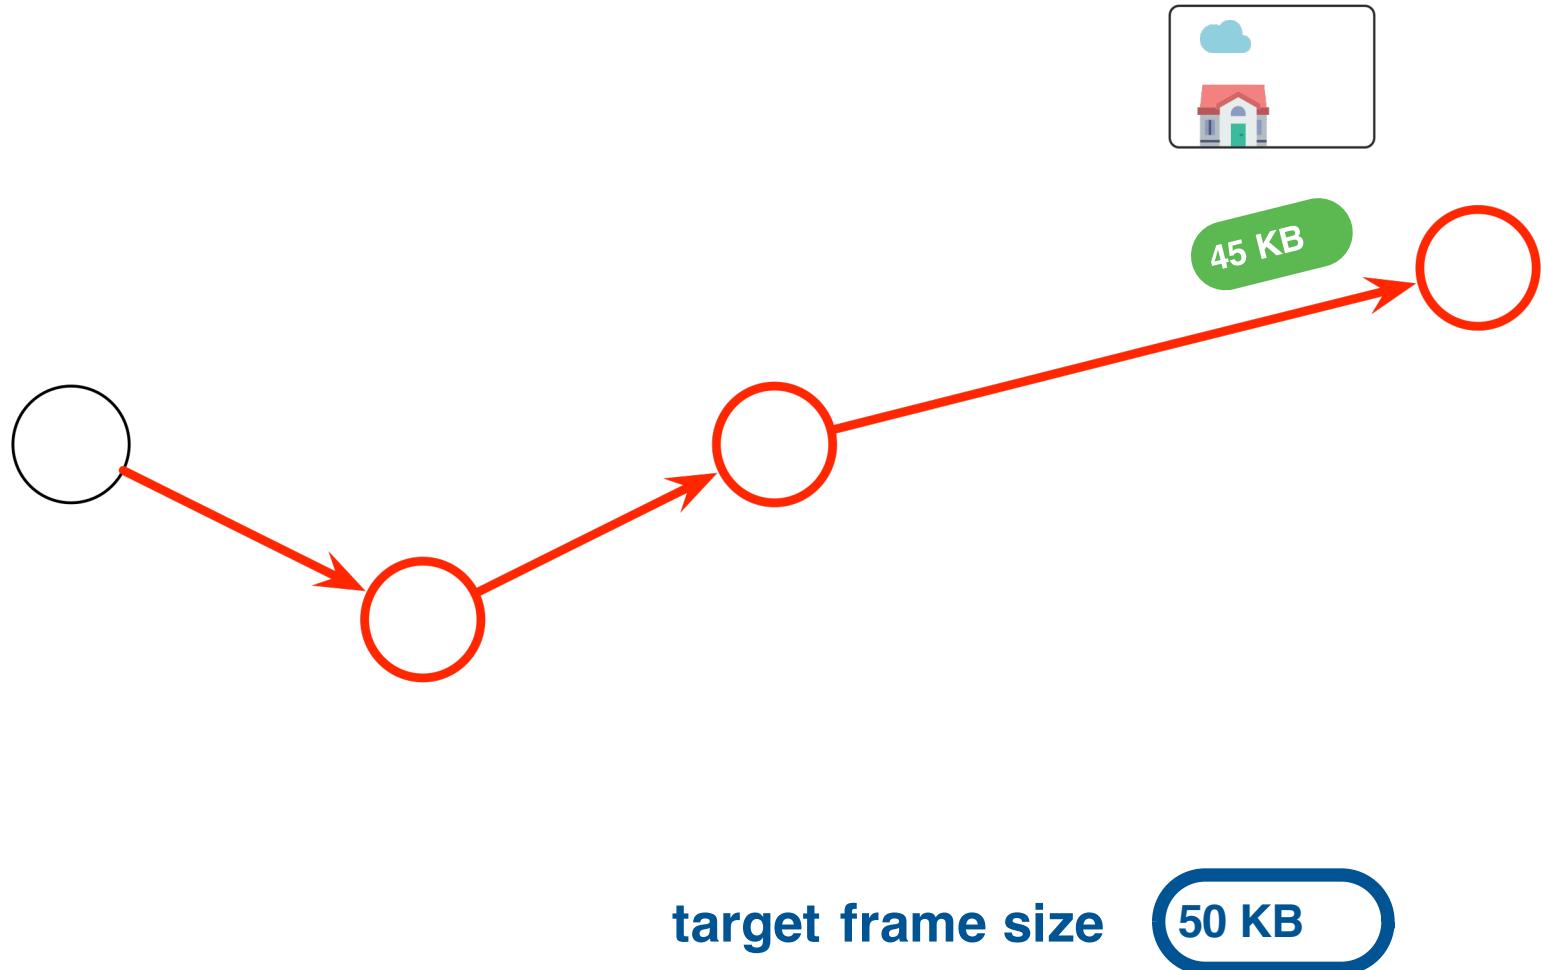
Codec → Transport

“Fine. Here’s 2 versions of the latest frame.”

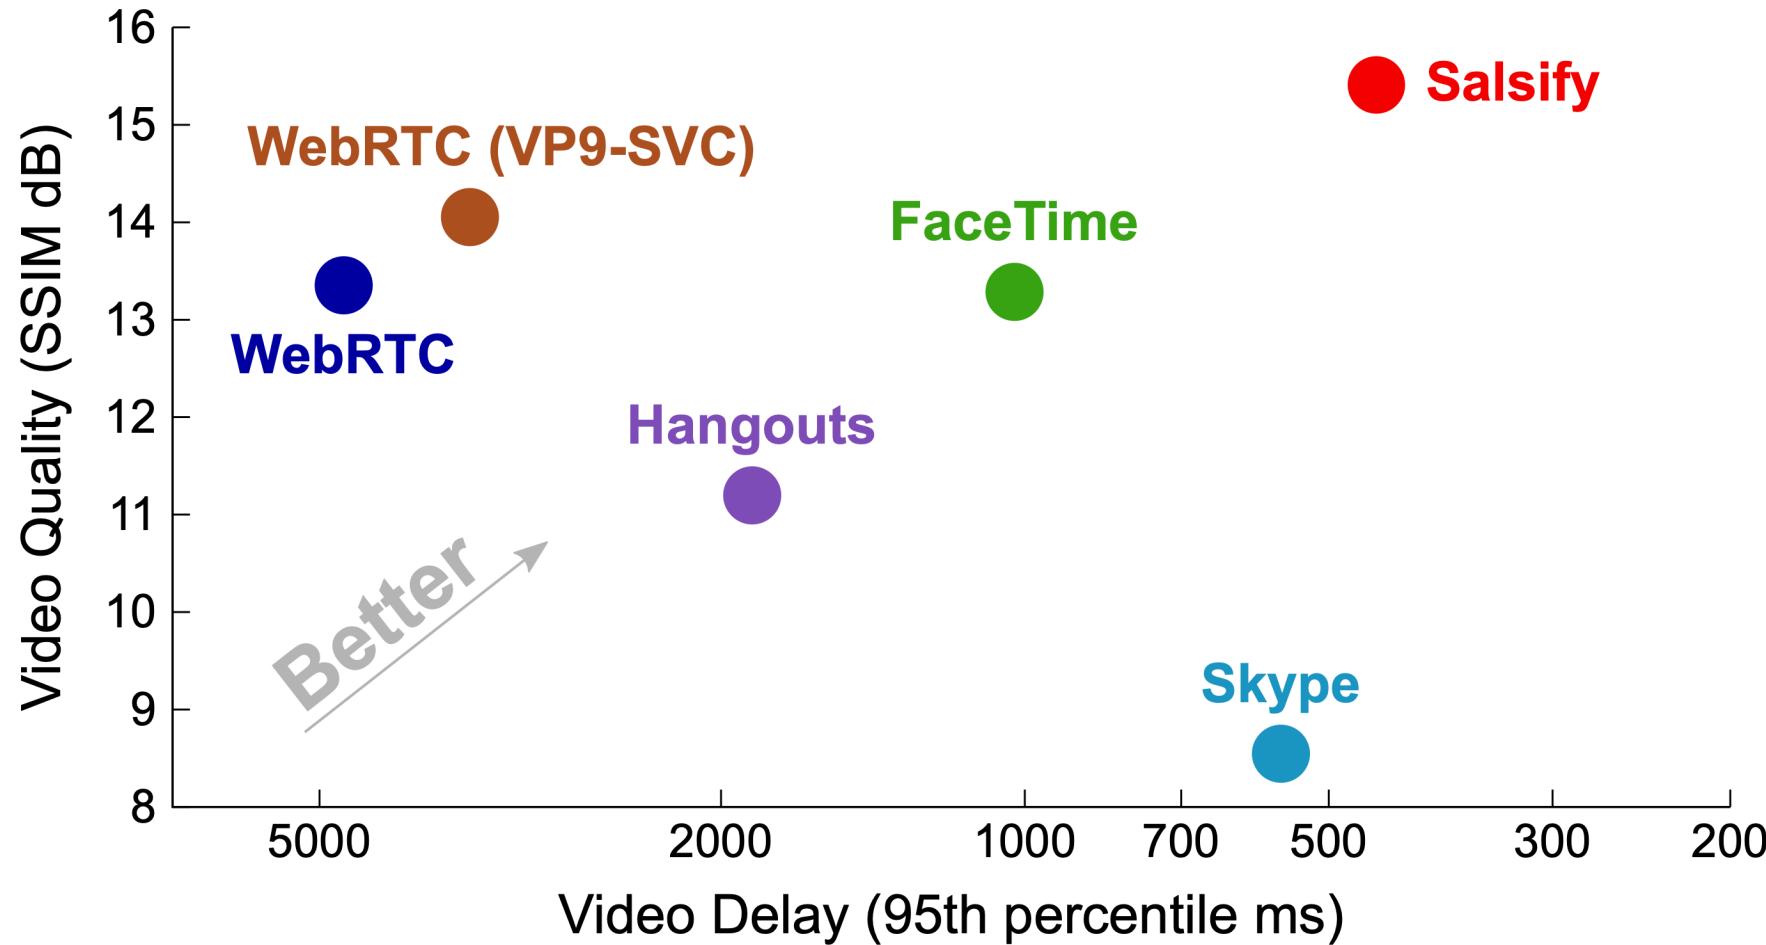


Transport → Codec

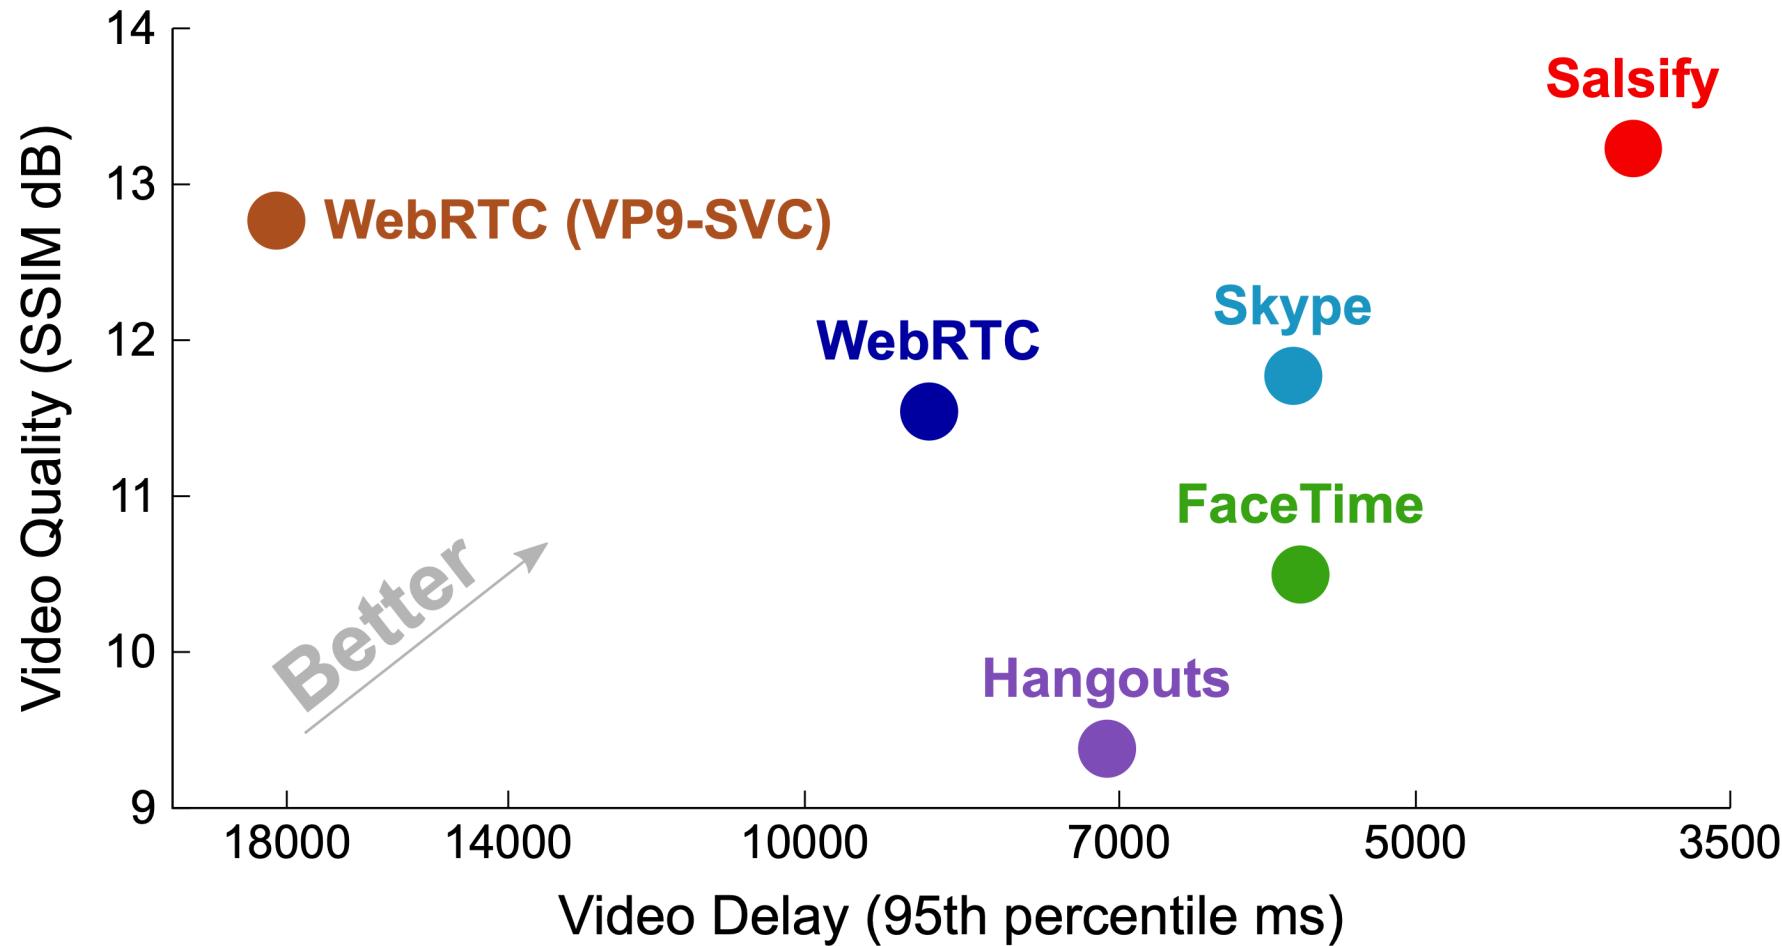
“I picked #1. Base next frame on its state.”



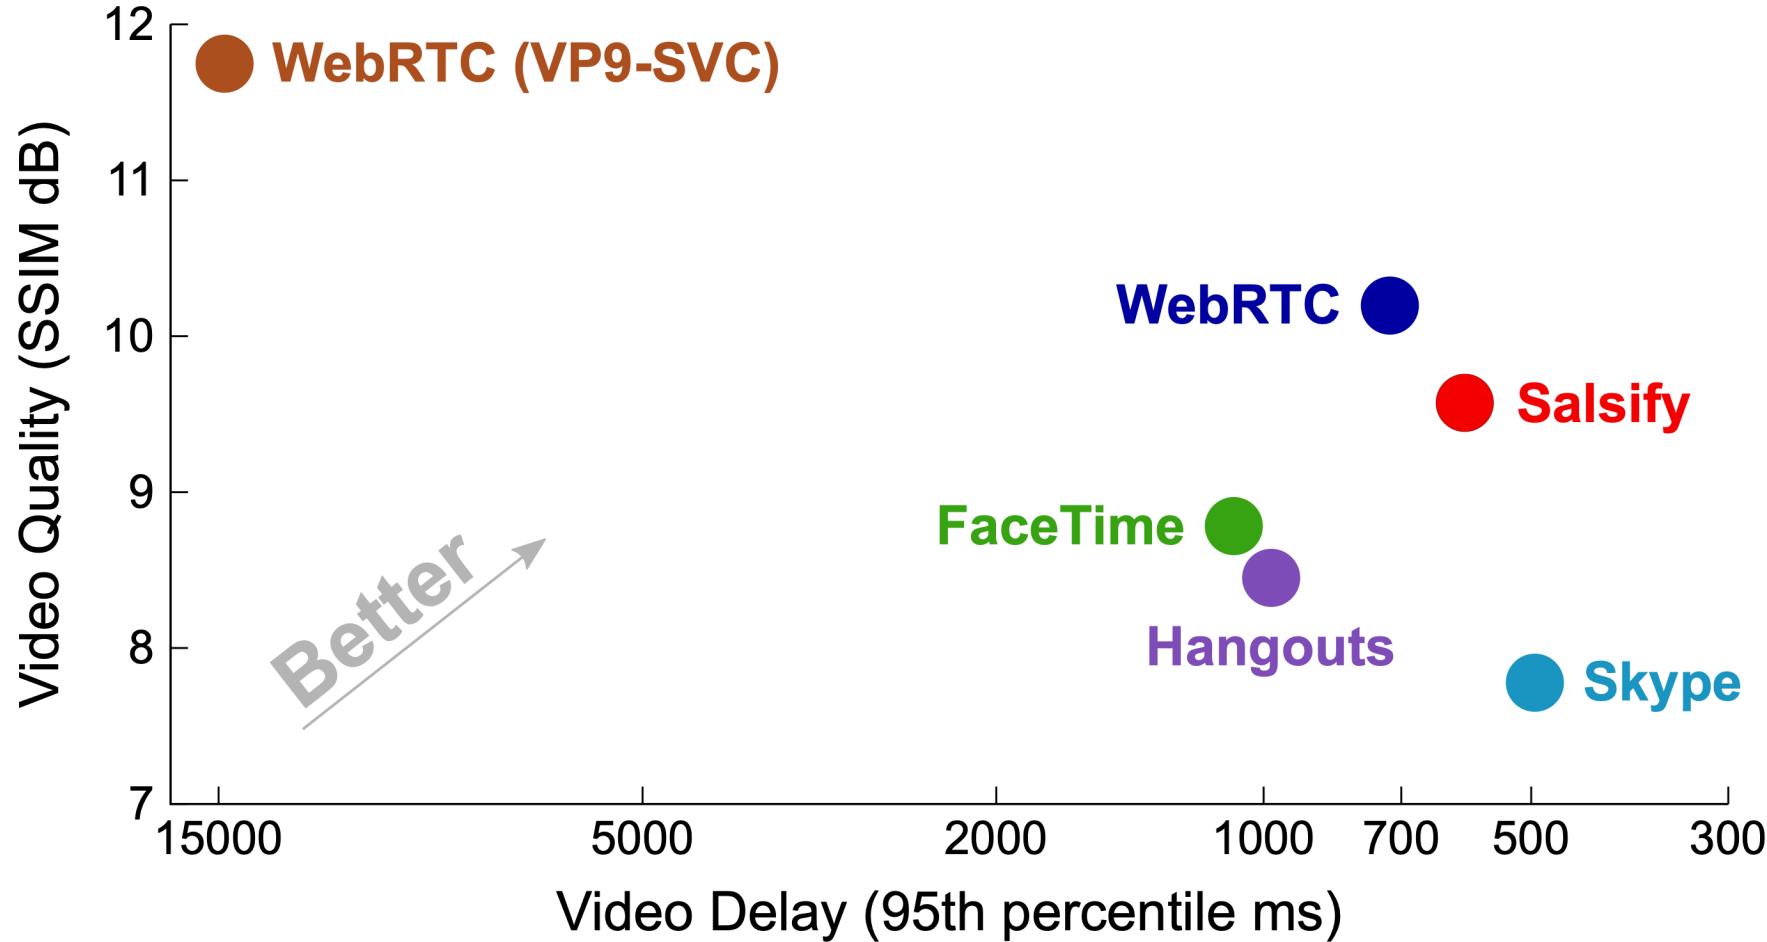
Evaluation results: AT&T LTE Trace



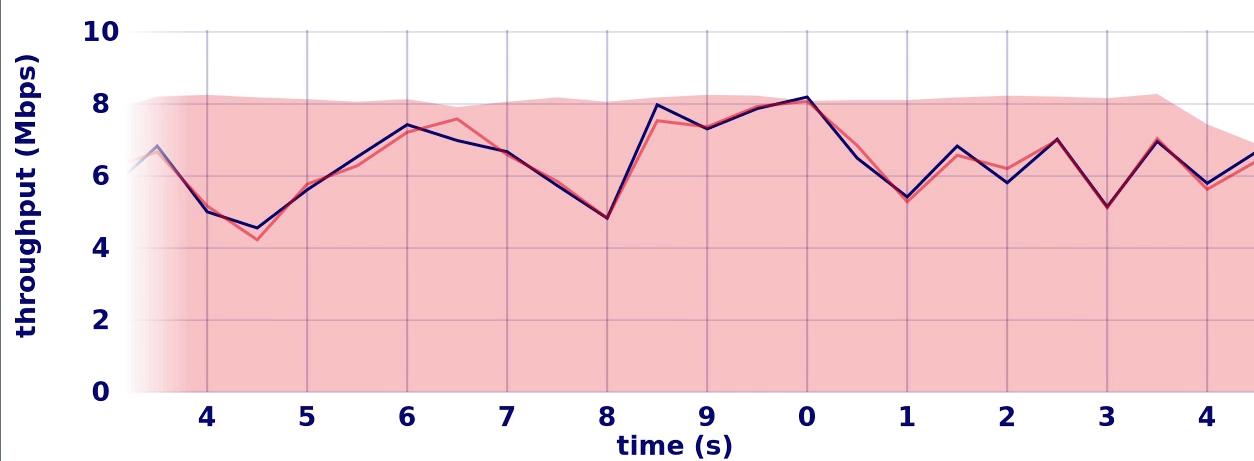
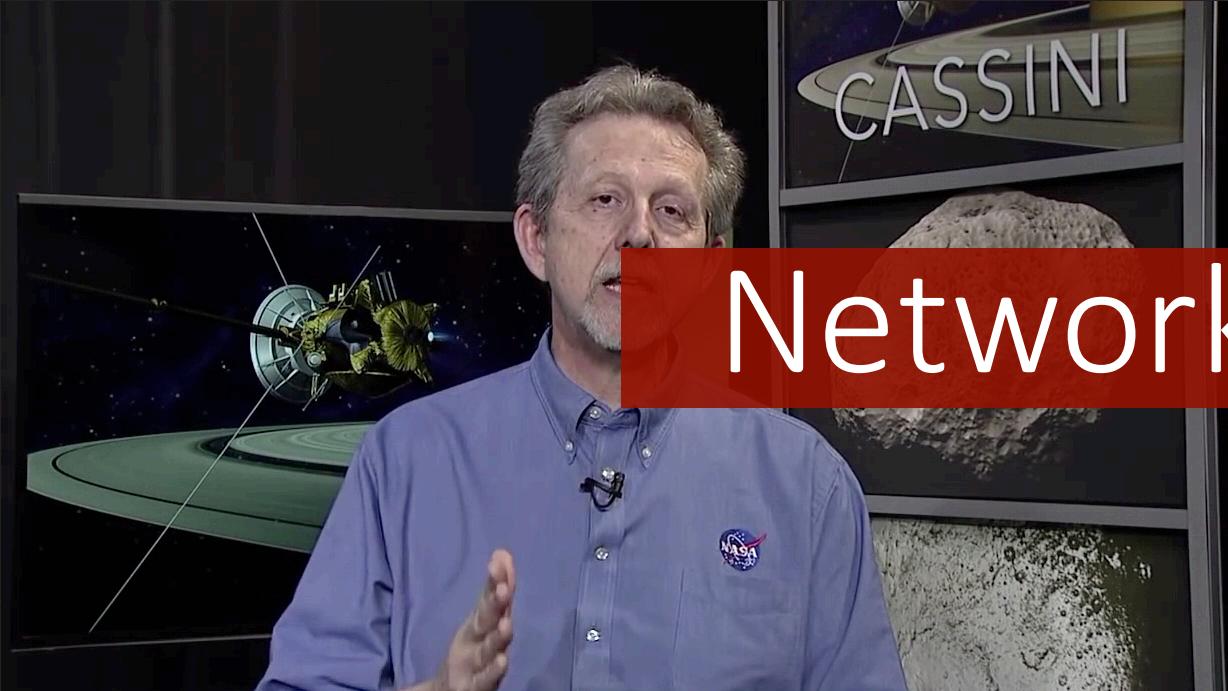
Evaluation results: T-Mobile UMTS Trace



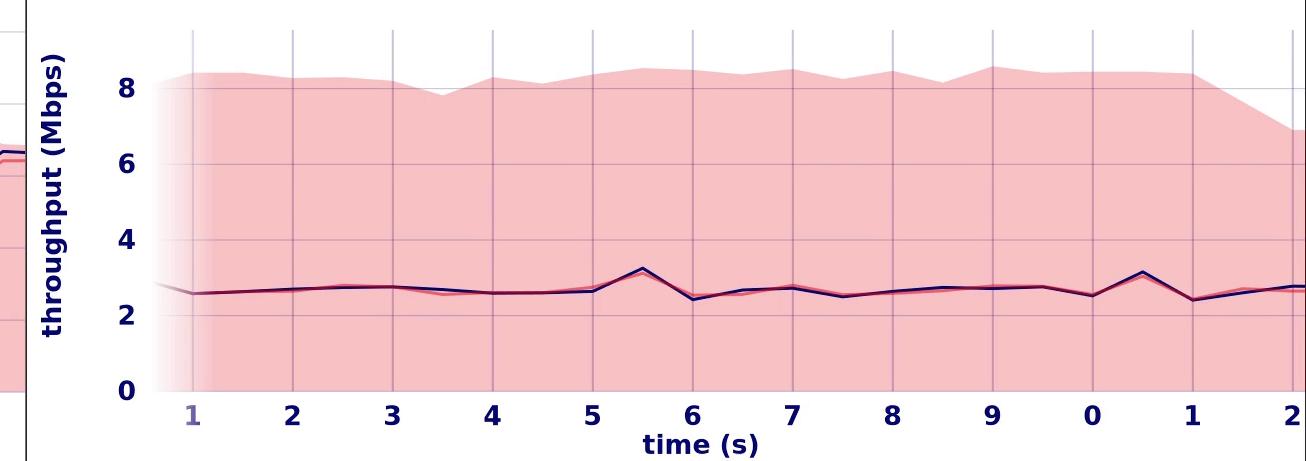
Evaluation results: Emulated Wi-Fi (no variations, only loss)



Network Variations



Salsify



WebRTC (Chrome 65)

Salsify, a new heart from old parts

- Individual component of Salsify are not exactly new:
 - The transport protocol is inspired by “packet pair” and “Sprout-EWMA”.
 - The video format, VP8, was finalized in 2008.
 - The functional video codec was described at NSDI’17.
- Salsify is a **new architecture** for real-time video that integrates these components in a way that responds quickly to network variations.

Improvements to *video codecs* may have reached the point of diminishing returns, but changes to the architecture of *video systems* can still yield significant benefits.

Takeaways

- Salsify is a new architecture for real-time Internet video.
 - Salsify tightly integrates a **video-aware transport protocol**, with a **functional video codec**, allowing it to **respond quickly to changing network conditions**.
- Salsify achieves **4.6× lower p95-delay** and **2.1 dB SSIM higher visual quality** on average when compared with FaceTime, Hangouts, Skype, and WebRTC.
- The code is open-source, and the paper and raw data are open-access:
<https://snr.stanford.edu/salsify>

Breaking Layers!

- **Expose more information to application**
 - WebRTC (and other transport protocols like Sprout) share forecasts of what they think network can support based on ACKs, etc.
- **Coordinate the control loops**
 - Transport and application layers each have their own designs and flaws; jointly operating them helps mask those issues
 - 2 big techniques: (1) late-binding decisions and (2) test alternatives