

Consistency



COS 316: Principles of Computer System Design


Amit Levy & Ravi Netravali

Why Do We Build Systems?


- ...
- Abstract away complexity

Distributed Systems are Highly Complex Internally

Sharding



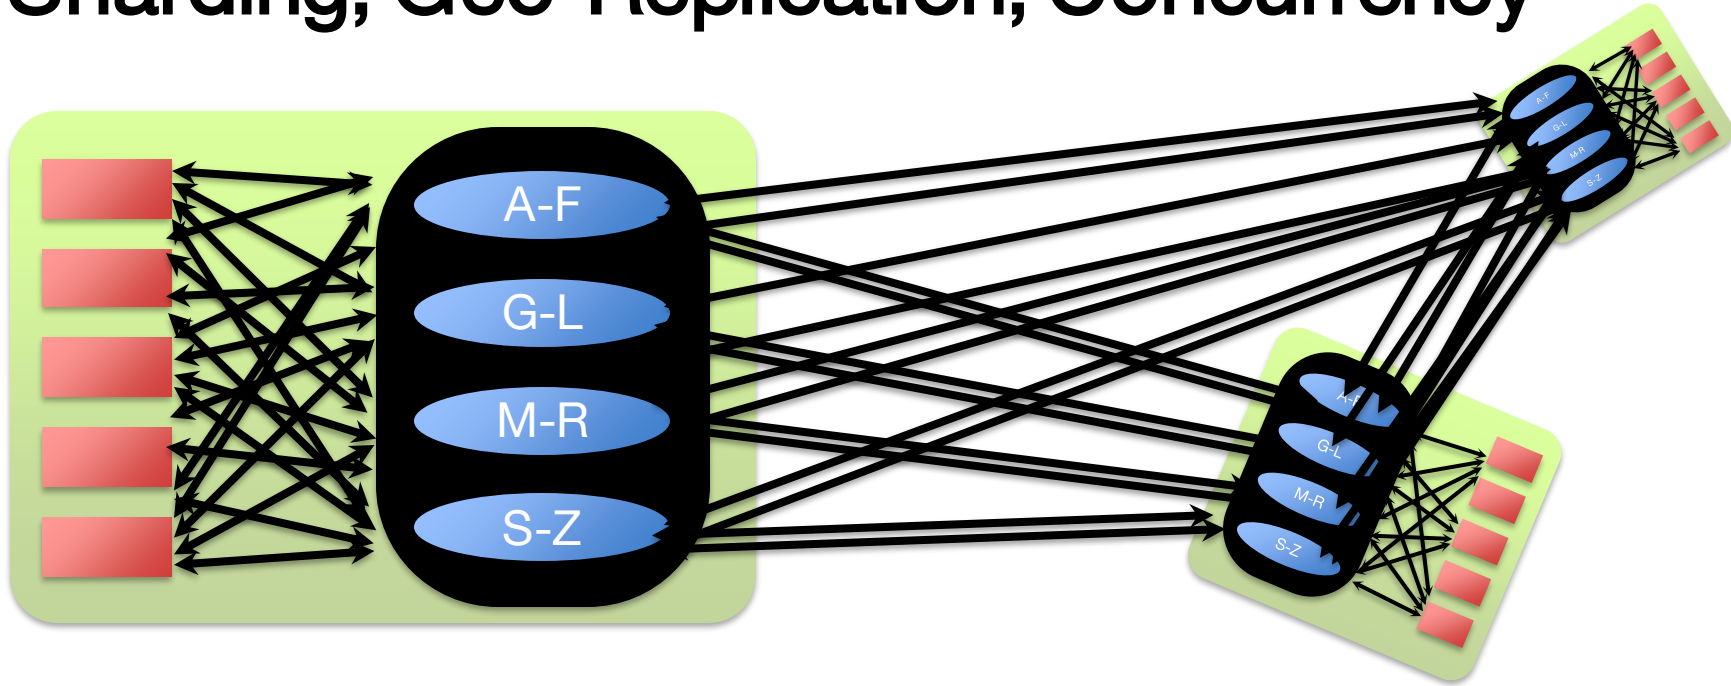
(Geo)-Replication



Concurrent access by many client

Distributed Systems are Highly Complex Internally

Sharding, Geo-Replication, Concurrency



Distributed Systems are Highly Complex Internally Sharding, Geo-Replication, Concurrency



Consistency Models:

Control how much of this
complexity is abstracted away

Consistency Models

- Contract between a (distributed) system and the applications that run on it
- A consistency model is a set of **guarantees** made by the distributed system

Stronger vs Weaker Consistency



Stronger vs Weaker Consistency

- Stronger consistency models
 - + Easier to write applications
 - System must hide many behaviors
 - Might be slow
- Fundamental tradeoffs between consistency & performance
 - (Discuss CAP, PRAM, SNOW in 418!)
- Weaker consistency models
 - Harder to write applications
 - Cannot (reasonably) write some applications
 - + System needs to hide few behaviors
 - + Can be faster!

Consistency Hierarchy

Linearizability



Causal+ Consistency



Eventual Consistency

Behaves like a single machine

Everyone sees related operations in the same order


Anything goes

Linearizability ==


“Appears to be a Single Processor”

- External client submitting requests and getting responses from the system can't tell this is not a single processor!
- Consistent with some **total order** over all operations
 - As though all requests processed one by one in some order
 - Such that...
- Order preserves the **real-time ordering** between operations
 - If operation A **completes** before operation B **begins**, then A is ordered before B in real-time
 - If neither A nor B completes before the other begins, then there is no real-time order
 - (But there must be *some* total order)

Real-Time Ordering Examples



Real-Time Ordering Examples



Linearizable?

P	$\vdash w(x=1) \dashv$	
P_B	$\vdash w(x=2) \dashv$	
P_C		$\vdash w(x=3) \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=3 \dashv$



w_1, w_2, r_2, w_3, r_3

Linearizable?

P	$\vdash w(x=1) \dashv$	
P_B	$\vdash w(x=2) \dashv$	
P_C		$\vdash w(x=3) \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=3 \dashv$
P_D	$\vdash r(x)=1 \dashv$	$\vdash r(x)=2 \dashv$



w_1, r_1, w_2, r_2, w_3

Linearizable?

P	$\vdash w(x=1) \dashv$	
P_B	$\vdash w(x=2) \dashv$	
P_C		$\vdash w(x=3) \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=3 \dashv$
P_D	$\vdash r(x)=1 \dashv$	$\vdash r(x)=2 \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=2 \dashv$



w_1, w_2, r_2, r_2, w_3

Linearizable?

P	$\vdash w(x=1) \dashv$	
P_B	$\vdash w(x=2) \dashv$	
P_C		$\vdash w(x=3) \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=3 \dashv$
P_D	$\vdash r(x)=1 \dashv$	$\vdash r(x)=2 \dashv$
P_D	$\vdash r(x)=2 \dashv$	$\vdash r(x)=2 \dashv$
P_D	$\vdash r(x)=1 \dashv$	$\vdash r(x)=3 \dashv$


 w_1, r_1, w_2, w_3, r_3

Linearizable?

$P \vdash w(x=1) \dashv$

$P_B \vdash w(x=2) \dashv$

$P_C \vdash w(x=3) \dashv$

$P_D \vdash r(x)=2 \dashv \vdash r(x)=3 \dashv$

$P_D \vdash r(x)=1 \dashv \vdash r(x)=2 \dashv$

$P_D \vdash r(x)=2 \dashv \vdash r(x)=2 \dashv$

$P_D \vdash r(x)=1 \dashv \vdash r(x)=3 \dashv$

$P_D \vdash r(x)=2 \dashv \vdash r(x)=1 \dashv$



Linearizable?

$P \vdash w(x=1) \dashv$

$P_B \vdash w(x=2) \dashv$

$P_C \vdash w(x=3) \dashv$

$P_D \vdash w(x=4) \dashv \vdash w(x=5) \dashv$

$P_E \vdash w(x=6) \dashv$

$P_F \vdash r(x)=2 \dashv \vdash r(x)=3 \dashv \vdash r(x)=6 \dashv \vdash r(x)=5 \dashv \quad \checkmark$

$w_1, w_2, r_2, w_4, w_3, r_3, w_6, r_6, w_5, r_5$

OR

$w_1, \mathbf{w}_4, w_2, r_2, w_3, r_3, w_6, r_6, w_5, r_5$

OR

$w_1, w_2, r_2, w_3, r_3, \mathbf{w}_4, w_6, r_6, w_5, r_5$

Linearizable?

$P \vdash w(x=1) \dashv$
 $P_B \vdash w(x=2) \dashv$
 $P_C \vdash w(x=3) \dashv$
 $P_D \vdash w(x=4) \dashv \vdash w(x=5) \dashv$
 $P_E \vdash w(x=6) \dashv$
 $P_G \vdash r(x)=2 \dashv \vdash r(x)=5 \dashv \vdash r(x)=6 \dashv \vdash r(x)=5 \dashv \quad \mathbf{X}$

Linearizable?

$P \vdash w(x=1) \dashv$

$P_B \vdash w(x=2) \dashv$

$P_C \vdash w(x=3) \dashv$

$P_D \vdash w(x=4) \dashv \vdash w(x=5) \dashv$

$P_E \vdash w(x=6) \dashv$

$P_H \vdash r(x)=4 \dashv \vdash r(x)=2 \dashv \vdash r(x)=3 \dashv \vdash r(x)=6 \dashv \checkmark$

$w_1, w_4, r_4, w_2, r_2, w_3, r_3, w_5, w_6, r_6$

Linearizable?

$P \vdash w(x=1) \dashv$
 $P_B \vdash w(x=2) \dashv$
 $P_C \vdash r(x)=1 \dashv$ **X**

Linearizability ==

“Appears to be a Single Processor”

- External client submitting requests and getting responses from the system can't tell this is not a single processor!
- Consistent with some **total order** over all operations
 - As though all requests processed one by one in some order
 - Such that...
- Order preserves the **real-time ordering** between operations
 - If operation A **completes** before operation B **begins**, then A is ordered before B in real-time
 - If neither A nor B completes before the other begins, then there is no real-time order
 - (But there must be *some* total order)

How to Provide Linearizability?

1. Use a single machine 😊

1. Use “state-machine replication” on top of a consensus protocol like Paxos
- Distributed system appears to be single machine that does not fail!!
 - Covered extensively in 418

2. ...

Consistency Hierarchy

Linearizability



Causal+ Consistency



Eventual Consistency

Behaves like a single machine

Everyone sees related operations in the same order

Anything goes

Causal+ Consistency Informally

1. • Potential causality: event a could have a causal effect on event b .
2. • Think: is there a path of information from a to b ?
 - a and b done by the same entity (e.g., me)
 - a is a write and b is a read of that write
 - + transitivity

Causal+ Sufficient



↓ Then ↓



Photo Added




↓ Then ↓



Purchase
retained



↓ Then ↓



Deletion
retained

Causal+ Not Sufficient

(Need Linearizability)

- Need a total order of operations
 - e.g., Alice's bank account ≥ 0
- Need a real-time ordering of operations
 - e.g., Alice changes her password, Eve cannot login with old password

Consistency Hierarchy

Linearizability



Causal+ Consistency



Eventual Consistency

Behaves like a single machine

Everyone sees related operations in the same order

Anything goes

Eventual Consistency

- Anything goes for now...
 - (If updates stop, eventually all copies of the data are the same)
- But, eventually consistent systems often try to provide consistency and often do
 - e.g., Facebook's TAO system provided linearizable results 99.9994% of the time [Lu et al. SOSP '15]
- “Good enough” sometimes
 - e.g., 99 vs 100 likes

Consistency Model Summary

- Consistency model specifies strength of abstraction
 - Linearizability ☐ Causal+ ☐ Eventual
 - Stronger hides more, but has worse performance
- When building an application, what do you need?
 - Select system(s) with necessary consistency
 - Always safe to pick stronger
- When building a system, what are your guarantees?
 - Must design system such that they always hold
 - Must confront fundamental tradeoffs with performance
 - What is more important?