


# Access Control Lists & Capabilities

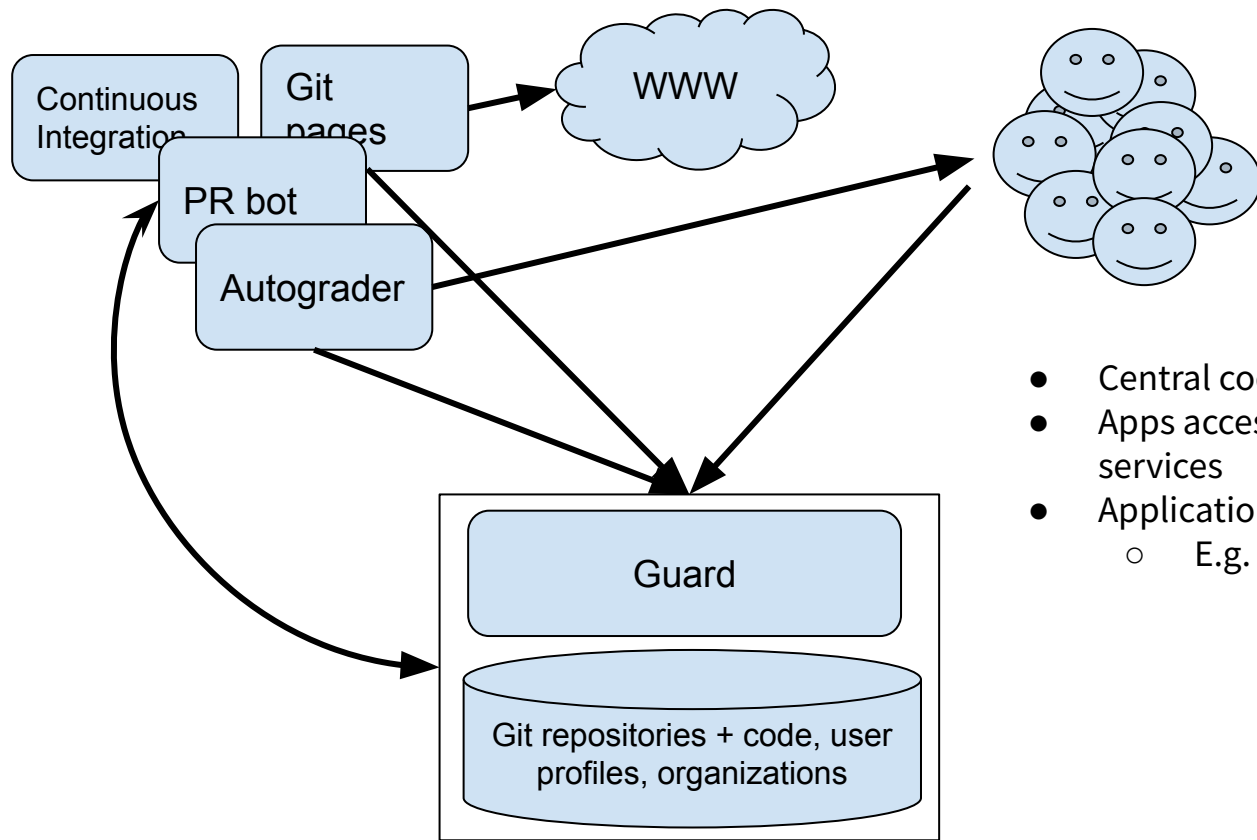
COS 316: Principles of Computer System Design

*Amit Levy & Jennifer Rexford*

# Last Time - The Guard Model



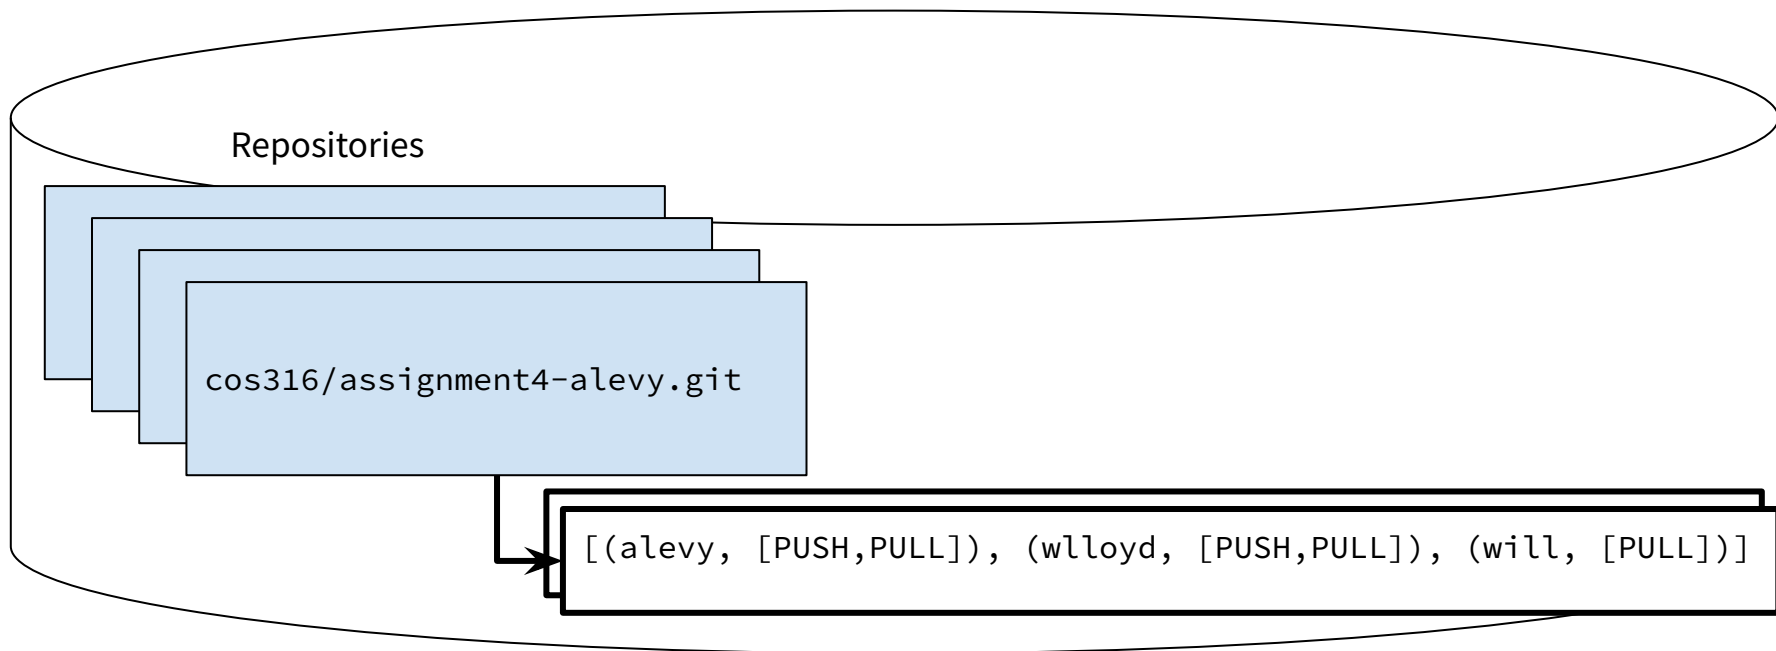
# Consider a GitHub-like Ecosystem



- Central code DB
- Apps access DB resources to provide extra services
- Application access must be restricted:
  - E.g. don't make private repos public

# Let's Start with User Permissions

Associate a list of (user, permissions) with each resource



# Implementing ACLs: Inline with Object

Repository Table

id	name	language	acl
1	cos316/assignment4-aalevy	Golang	"[(alevy, [PUSH,PULL]), (wlloyd, [PUSH,PULL]), ...]"
2	tock/tock	Rust	...
...	...	...	...

# Implementing ACLs: Normalize

ACL Table

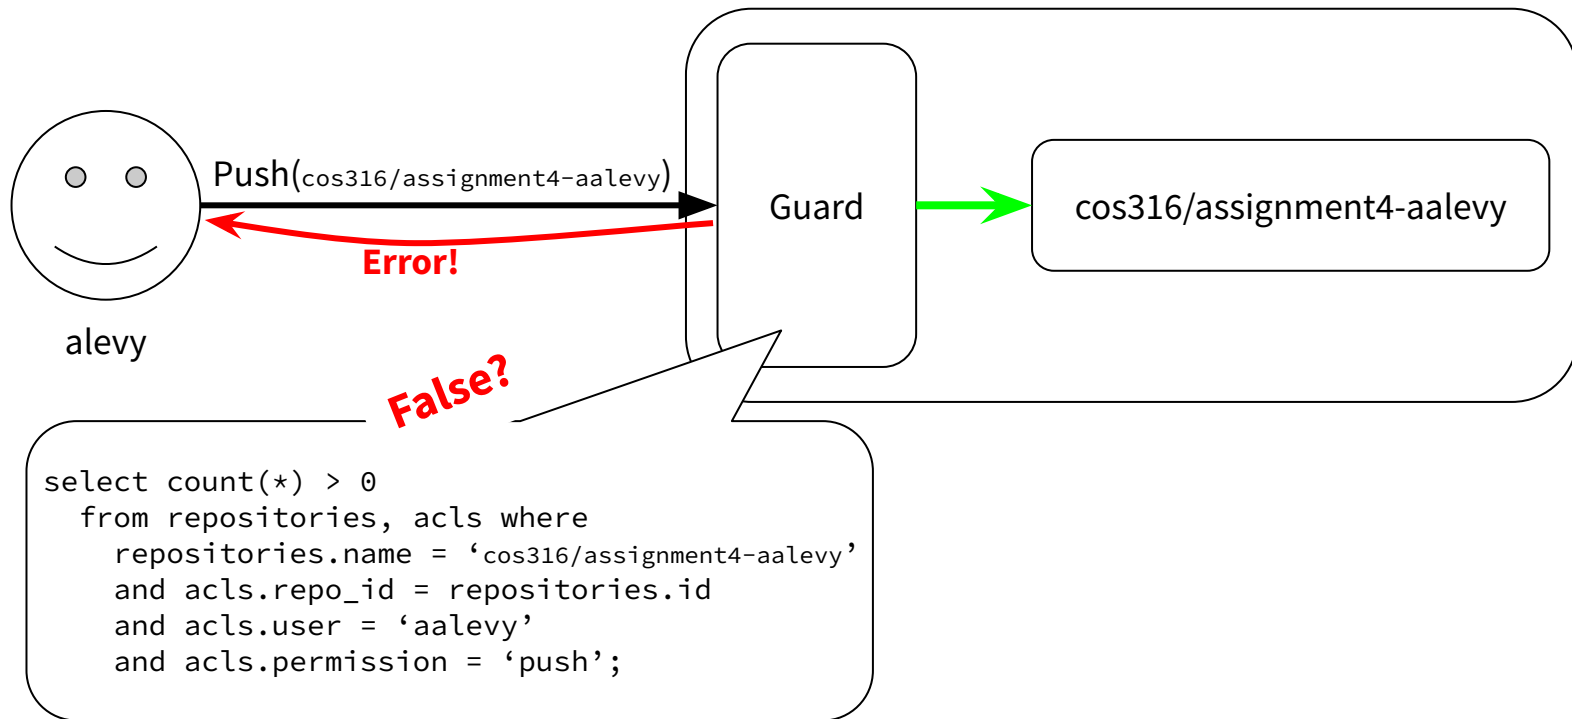
repo_id	user	permission
1	aalevy	push
1	kap	push
1	kap	pull
1	aalevy	pull
1	will	pull
2	aalevy	push
...	...	...

```
select (acls.user, acls.permission)
from repositories, acls where
  repositories.name = 'cos316/assignment4-aalevy'
and acls.repo_id = repositories.id;
```

Repository Table

id	name	language
1	cos316/assignment4-aalevy	Golang
2	tock/tock	Rust
...	...	...

# ACLs in Action



# Extending ACLs to Apps: a-la UNIX

- Applications act *on behalf of* users
- When an application makes a request, it uses a particular user's credentials
  - Either one user per application
  - Or different users for different requests
- Works great for:
  - Alternative UIs, e.g. the `git` client vs. the GitHub Web UI both act on behalf of users
- Why might this be suboptimal?



# Extending ACLs to Apps: Special Principles

- Create a unique principles for each app
  - E.g., the “autograder” principle
  - Acts just like a regular user
- When applications make request, they use their own, unique, credentials
- Add application principals to resource ACLs as desired
- Works when
  - Applications need to operate with more than one user’s access
    - E.g. the autograder needs to access private repositories owned by different students
  - and less than any one user’s access
    - E.g. the autograder shouldn’t be able to access non COS316 repositories

# Access Control Lists

## Advantages

- Simple to implement
- Simple to administer
- Easy to revoke access

## Drawbacks

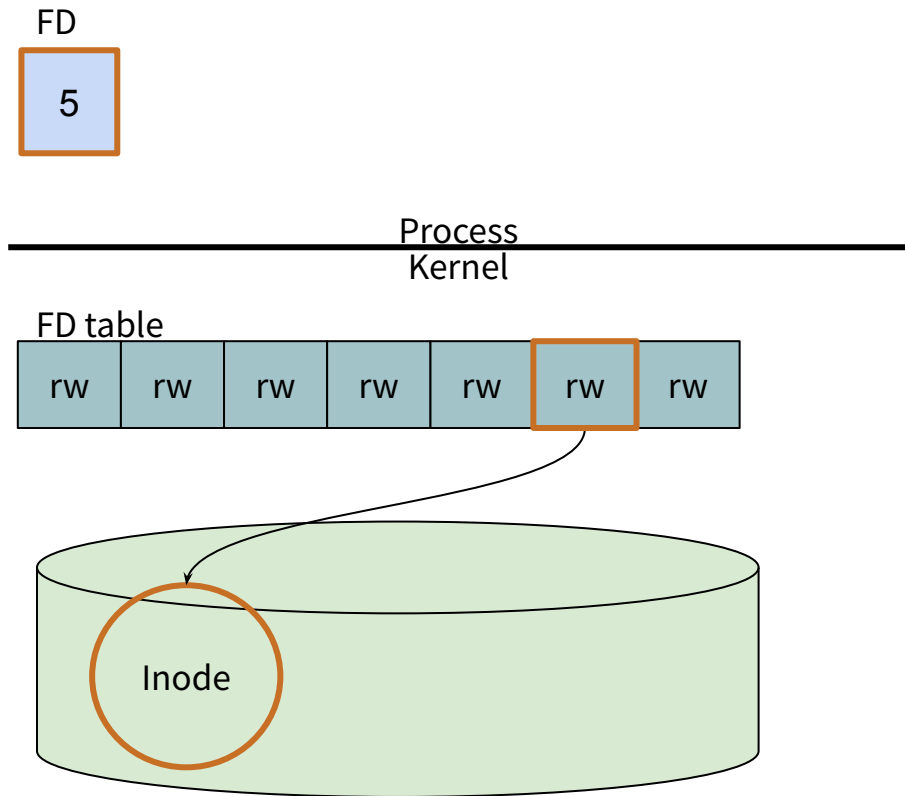
- Tradeoff granularity for simplicity
  - More granular permissions require more complex rules in the guard
- Doesn't scale well
  - E.g. need up to Users X Repos X Access Right entries in ACL table
- Centralized access control
  - Needs server's cooperation to delegate access

# An Alternative - Capabilities

“[A] token, ticket, or key that gives the possessor permission to access an entity or object in a computer system.” - *Capability-Based Computer Systems*


- Self-describing
  - Contains both object name and permitted operations
- Globally meaningful
  - Object and operation names are not subject-specific
- Transferrable
  - A subject can pass a capability to another (e.g. a sub-process, via IPC, a third-party app, etc)
  - Ideally can delegate subset of capabilities
- Unforgeable
  - Subjects cannot create capabilities with arbitrary permissions

# File Descriptors as Proto-Capabilities



- Unforgeable ✓
  - Process-level fd is just an index in a kernel structure
- Self-describing ✓
  - Kernel fd contains reference to inode + permissions
- Globally meaningful ✗
  - Fds are process-specific
- Transferrable ✓/✗
  - Via IPC sendmsg/recvmmsg

# Consider a GitHub-like Ecosystem



- Central code DB
- Apps access DB resources to provide extra services
- Application access must be restricted:
  - E.g. don't make private repos public

# User Permissions using Capabilities

Hand out communicable, unforgeable tokens encoding:

- Object
- Access right

Users store capabilities, not the database

E.g.

**“push(cos316/assignment4-aalevy)”**

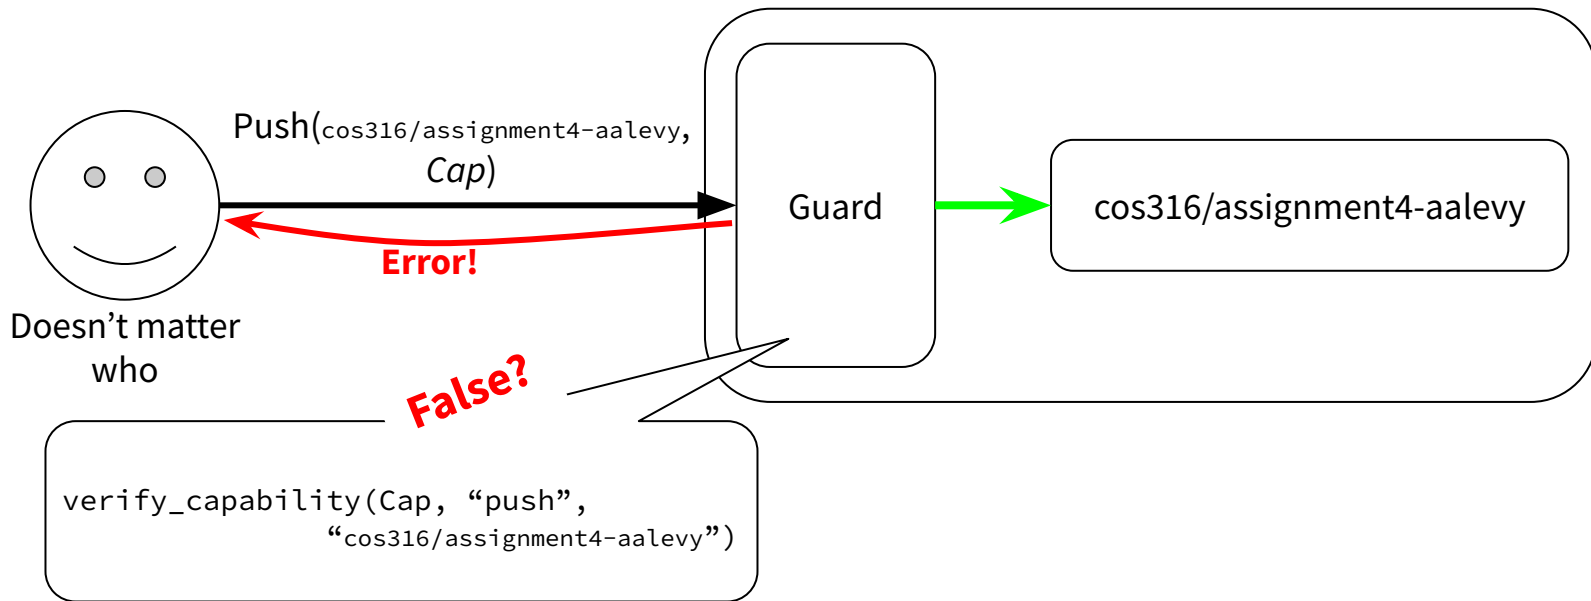
**“pull(cos316/assignment4-aalevy)”**

# Implementing Capabilities with HMAC

HMAC - a keyed-hash function: `hmac(secret_key, data)` hash of data

```
fn gen_capability(op, repo) {  
    hmac(db_secret, fmt.Sprintf("%s(%s)", op, repo))  
}  
  
fn verify_capability(cap, op, repo) {  
    cap == hmac(db_secret, fmt.Sprintf("%s(%s)", op, repo))  
}
```


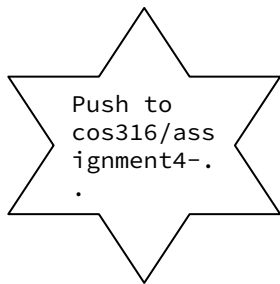

# Capabilities in Action






# Extending Capabilities to Applications

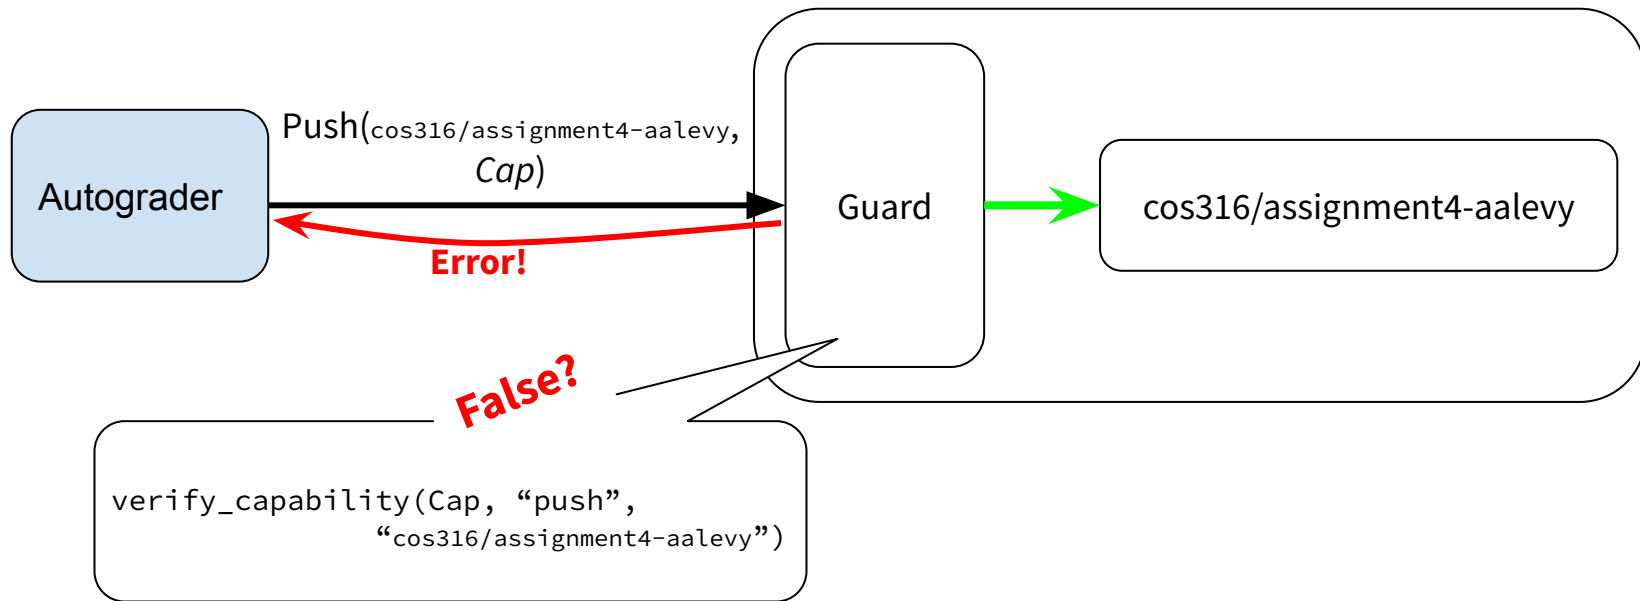
- Users can simply give applications a subset of their capabilities



aalevy



# Extending Capabilities to Applications



# Capabilities

## Advantages

- Decentralized access control
  - Anyone can “pass” anyone a capability
- Scales well
- Granular permissions are simple to check

## Drawbacks

- How do you revoke a capability?
- Moves complexity to users/clients
  - Users have to manage their capabilities now

# Capabilities In The Wild

- Operating Systems

- History of industry and research operating systems
- seL4
- FreeBSD's Capsicum
- Fuschia OS

- Web

- S3 Signed URLs
  - URL to private resources, contain signature, expiration, permitted HTTP methods, etc
- CDN-hosted images/videos (FB, Instagram, YouTube)
  - Browsing via Web page/app is protected by login+cookie, but media typically fetched unauthenticated

# Next time...

We still have a problem!

The autograder is allowed to:

- read all cos316/ repositories
- comment on all cos316/ repositories

Can code from a private repository end up in a comment on a public repository?