

Hails: Protecting Data Privacy in Untrusted Web Applications

COS 316: Principles of Computer System Design

Amit Levy & Wyatt Lloyd

Hails context

- Daniel B. Giffin, *Amit Levy*, Deian Stefan, David Terei, John Mitchell, David Mazières, & Alejandro Russo
- Developed 2010-~2015
 - First git commit from current version December 20th 2011
- Hails/Gitstar ('12) -> Gitstar Inc. ('14) (later renamed Intrinsic) -> VMWare Intrinsic
- Key ideas (aka how it makes MAC practical):
 - Complexity: Leverage Haskell language to build IFC as a library → easier to iterate
 - Performance: Leverage “purity” in Haskell to minimize security checks
 - Simplicity: Natural to *extract* policy from data, should be natural to use end-to-end policies on data



github:develop

Web platforms are **great**!
They allow third-party developers to
build apps that use our personal data.



Google
Developers

foursquare™
DEVELOPERS



Path Uploads Your Entire iPhone Contact List By Default

THE WALL STREET JOURNAL.

Home World Europe U.K. U.S. Business Markets Market Data

WHAT THEY KNOW | October 17, 2010, 8:33 p.m. ET

Facebook in Privacy Breach

Top-Ranked Applications Transmit Personal IDs, a Journal Investigation Finds

Web platforms are **scary**!
They allow third-party developers to
build apps that use our personal data.

 Symantec. | Connect

 COMMUNITY: Security  Blogs  Security Response

Facebook Applications Accidentally Leaking Access to Third Parties - Updated

The GitHub Blog

 March 4, 2012 mojombo

Public Key Security Vulnerability and Mitigation

Trust Concerns

- Don't know the developers
 - Cannot determine trustworthiness of apps
- They may be malicious or security-unaware
- Building secure web apps is hard
 - Even well-meaning authors cannot be trusted


Typical App Design

Use the MVC paradigm

Model: interface to data

View: renders pages

Controller: handles and responds to HTTP requests




Typical App Design

How is security policy specified and enforced?

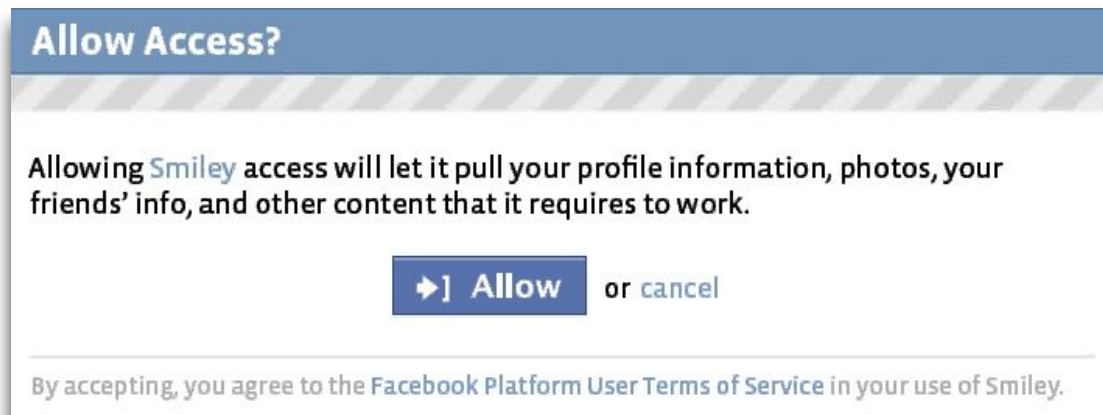
- E.g., only Jen's friends may see her email address

Intertwined throughout code

- Error prone and not scalable



Platform “solutions”



Users can decide to give an app access to data, but can't control how the app uses your data.

Path Uploads Your Entire iPhone Contact List By Default

Facebook in Privacy Breach

Top-Ranked Applications Transmit Personal IDs, a Journal Investigation Finds

Is there any hope for privacy on platforms?

Change the hosting model

- Current model
 - App developers host their own apps
 - Platform enforces security: terms of service
- New model
 - Platform provider hosts apps
 - Platform enforces security mandatorily: information flow control

Hails: A web platform framework

- Security policy is explicit and first-class
 - Specified as single concise module
- Users still trust core platform components
- Apps are untrusted
 - *Language-level* information flow control guarantees apps always obey policy

Hails vs Previous Systems

Aeolus, HiStar, Nexus, Jif, Ur/Web, ...

- No guide for structuring applications
- Policies are hard to write
- Not appropriate for dynamic systems, e.g., web
- Modify entire application stack

Goals

- Deployable
- Usable by Web developers
- Suitable for building extensible Web *platforms*
 - Enforcing policy across untrusted apps

Adding Policy to MVC


New programming paradigm: Model-**Policy**-View-Controller

- Policy specified alongside data model
 - Models are *partially* trusted to define the policy related to model data
- No policy code in View or Controller
 - Vast majority of bug-prone code
 - All* the code that third-party apps use to handle sensitive data

*Except the front-end code in the browser which, today, is *much* of the app's code

Two categories of code


Models-Policies (MPs)



Specify data model and policy on data

Users *trust* MPs they use to handle data

Views-Controllers (VCs)



Implement UI and other functionality

Users need not trust VCs with data


Policy enforced globally

Information flow control

- Policy specifies where data can flow
 - **Wrong:** app cannot read Jen's email address because it may leak it to Eve
 - **Right:** app can read Jen's email address, but only reveal it to Jen, Alice or Bob
- Policy follows data through system
- Runtime enforces policy end-to-end
 - E.g., when making an HTTP request

Case study: Gitstar

Gitstar [List Users](#) [List Projects](#)

 [deian](#)

hails


Fork

Haskell Web Platform Framework.

Repo: `git clone ssh://deian@gitstar.com/scs/hails.git`

[Wiki](#) [Code](#)




[Home](#) / [master](#) /

 update db conf file
v ^

d62d1c

1 file changed, with 1 addition and 1 deletion.
parent: 62b9a9

— Authored 2012-08-30T19:57:56-07:00 by Deian Stefan <deian@cs.stanford.edu>

name	size
 Hails	--
 examples	--
 tests	--


Case study: Gitstar

GitStar provides

- MPs that specify projects and users
- VC for managing projects and users

Third-party authors provide


- Code viewer
- Wiki
- Follower app
- etc.



Models-Policies (MPs)



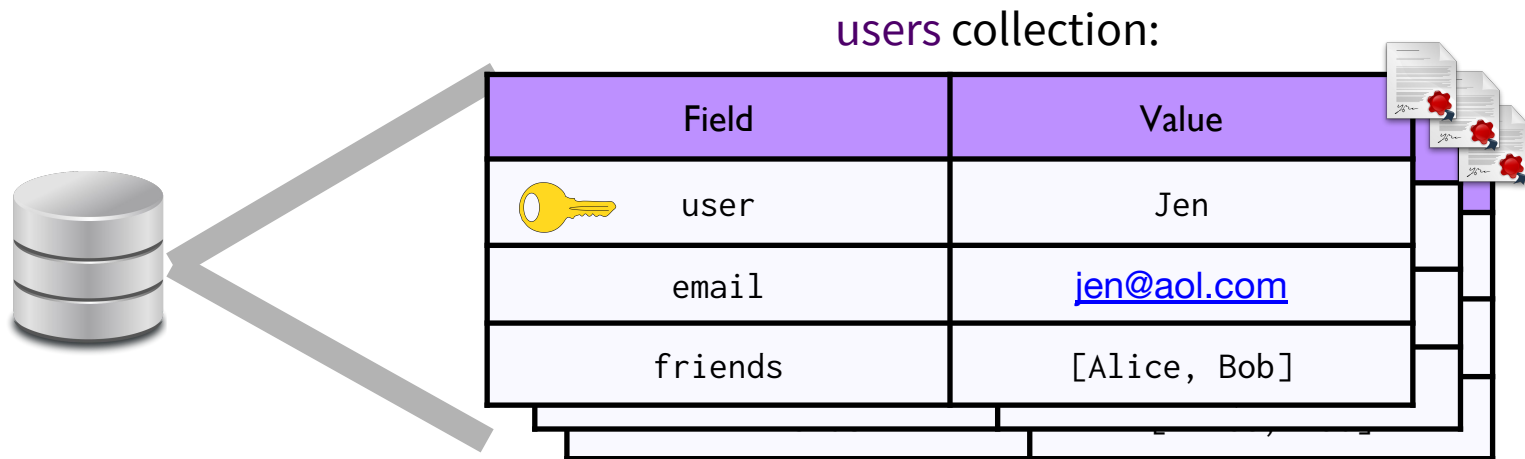
Views-Controllers (VCs)



Model-Policy (MP)

Data model: document-oriented

- Collection: set of documents
- Document: set of field-value pairs




Model-Policy (MP)

- Policy specifies restrictions on:
 - Collections, documents, fields
 - E.g., only Jen may modify her profile
 - E.g., only Jen and her friends may read her email address
- Policy composes
 - E.g., to read document you must be able to read the collection

Example: Enforcing policy

- MP:

Field	Value
 user	Jen
email	jen@aol.com
friends	[Alice, Bob]

Policy: Only Jen, Alice and Bob can read

GitStar User MP

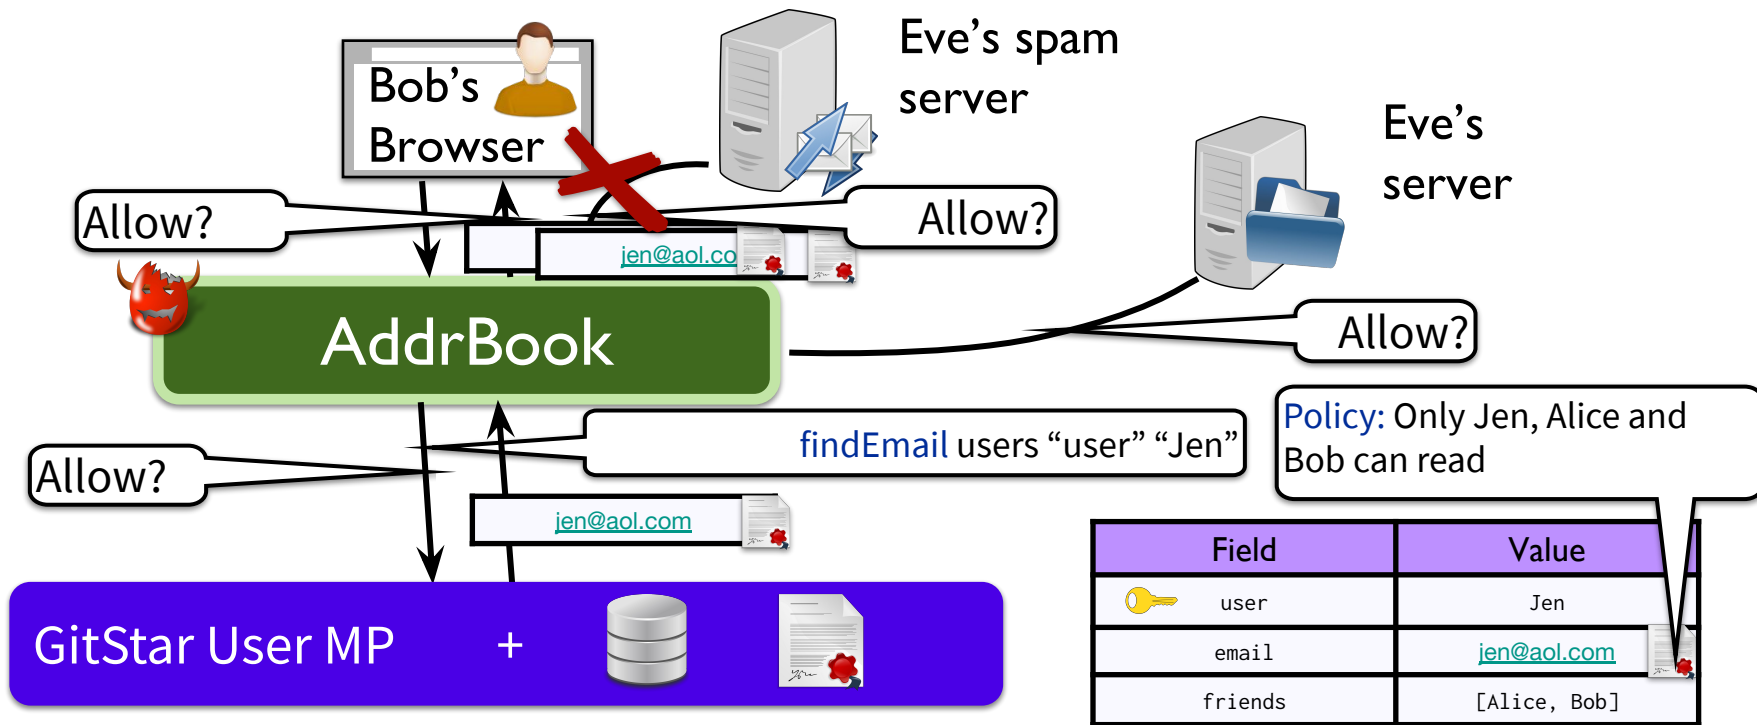


- Eve's untrusted address book VC:



AddrBook


Example: Enforcing policy



Policy specified in terms of data

Web app data models already encode policy

- Ownership
- Relationships between users
- ...

Field		Value
 user		Jen
email		jen@aol.com
friends		[Alice, Bob]

Policy: Only user can modify

Policy: Only user and friends can read

Example: Policy specification

```
collection "users" $ do
```

```
  access $ do
```

```
    readers ==> anybody
```

```
    writers ==> anybody
```

```
  field "user" key
```

```
  document $ λdoc -> do
```

```
    readers ==> anybody
```

```
    writers ==> ("user" `from` doc)
```

```
  field "email" $ labeled $ λdoc -> do
```

```
    readers ==> ("user" `from` doc)
```

```
      V fromList ("friends" `from` doc)
```

```
    writers ==> anybody
```

Collection is public
Only Jen can modify


Index document types
document types Alice and Bob

user names can read Jen's email
address

Models-Policies (MPs)



Views-Controllers (VCs)




View-Controller (VC)

- A VC is a request handler
- Provide application functionality
 - E.g., source code browser, blog editor, ...
- Invoke MPs to store/fetch user data
- Bugs in VCs are never vulnerabilities
 - Runtime enforces security policy

Models-Policies (MPs)



Views-Controllers (VCs)



Implications of MPVC

- Users: choose VCs based on functionality
- Developers: build apps on top of existing user-data
 - Models and policies are reusable

Gitstar List Users List Projects

hails

Haskell Web Platform Framework.

Repo: `git clone ssh://deian@gitstar.com/scs/hails.git`

Wiki Code

🏠 / master / examples / SimpleParams.hs

```
1. {-# LANGUAGE OverloadedStrings #-}
2. module SimpleParams (server) where
3.
4. import qualified Data.ByteString.Lazy.Char8 as L8
5.
6. import      LIO
7. import      LIO.DCLLabel
8. import      Hails.HttpServer
9. import      Hails.Data.Hson
10.
11. server :: Application
12. server _ lreq = do
13.   req <- unlabel lreq
14.   let ldoc = labeledRequestToLabeledDocument lreq
15.   doc <- unlabel ldoc
16.   return $ case pathInfo req of
17.     ("login":_) => Response temporaryRedirect307
18.       [ ("x-hails-login","", (hLocation,"/")) ""
19.     => Response ok200 [ $ toHttpl (label10f lreq, req) (label10f ldoc, doc)
```

Gitstar List Users List Projects

hails

Haskell Web Platform Framework.

Repo: `git clone ssh://deian@gitstar.com/scs/hails.git`

Wiki Code

Home Pages

Hails: Protecting Data Privacy in Untrusted Web Apps

Hails is a platform and web framework that leverages Information Flow Control (IFC) to support untrusted and mutually distrustful web applications interacting and processing private data.

Resources

- [Brief motivation and architecture overview](#)
- [Tutorial](#) (slightly outdated)

Installation

You can compile and install Hails as usual with `cabal-dev`:

```
$ cabal-dev install-deps
$ cabal-dev install
```

Launching an app

Every define your main application (named `main`) in `src/Main.hs`. You can launch it with:

GitStar Project MP

+




Implementation

- Hails is a Haskell library
 - Quick turnaround on API design
 - Developers can use existing tools and libraries
- Hails runtime system
 - Provides HTTP server that invokes VC
 - Enforces information flow at the language-level

Evaluation: Usability

- ✓ MPVC simplifies reasoning about security when building a platform
- ✓ Hails renders common security bugs futile
E.g., mass assignment vulnerability
- Need scaffolding tools
- ~~● Writing raw policy is hard~~
 - ✓ Writing policy with DSL is simpler

Performance evaluation



Conclusions

Current platforms: functionality vs. privacy

Hails platforms guarantee security across apps

- Hosts apps on platform
- Make policy explicit
- Enforce policy with information flow control