

# COS 316

## Precept:

# Socket Programming

# Naming

- Why is naming important? What would happen if we didn't name things?
- Naming is important because it gives us a way to find and access things
  - how Amazon knows how to deliver packages to you
  - how to access stored objects in memory
- Socket
  - their names enable the system (and others) to know how to find/contact it

# *Abstraction* Clarification

- “A way of modeling things”
- Don’t worry about the exact implementation
- Focus on the paradigm
- Socket abstraction

# What are Sockets/Connections?

- **Connection**

- A process on one host (host A) communicates with a process on another host (host B) via a connection
- A communication channel
- Another abstraction

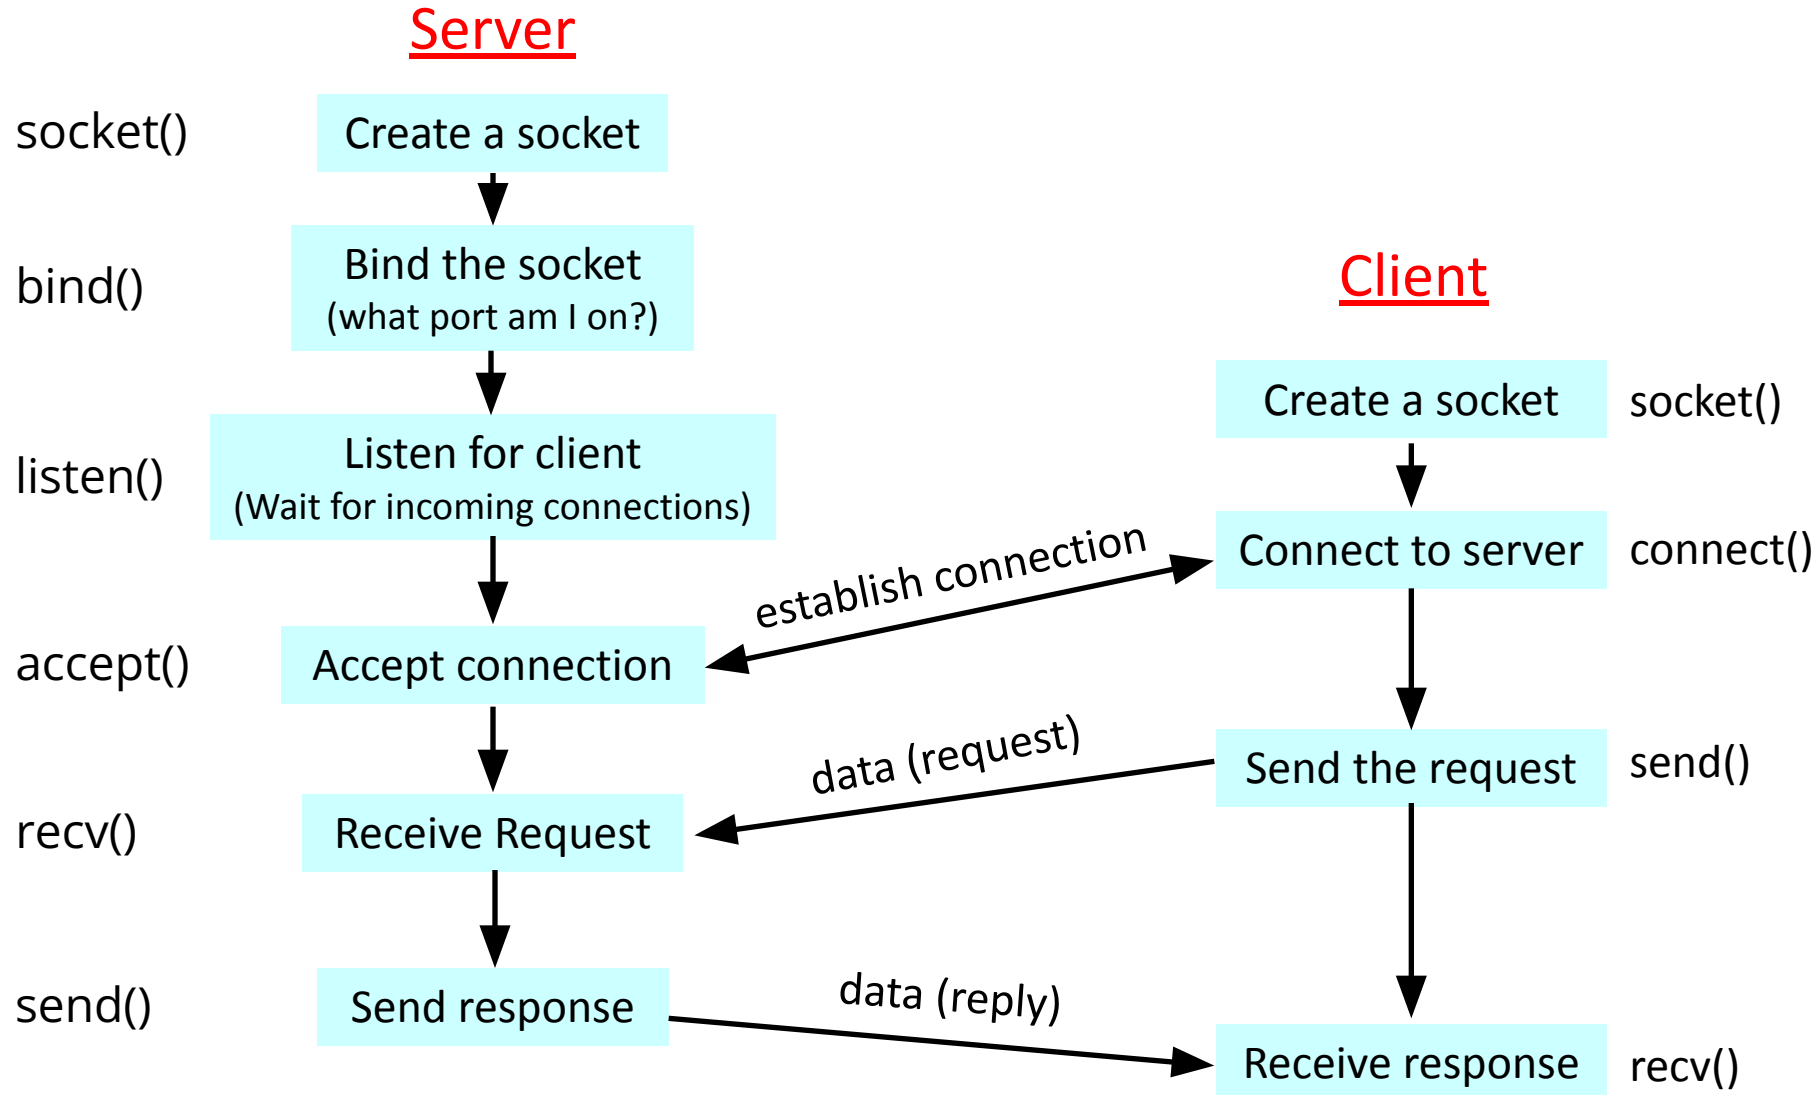
- **Socket**

- In order for host A to start a connection with host B, host A needs to know where and how to contact host B
- This endpoint on host B is what we call a socket

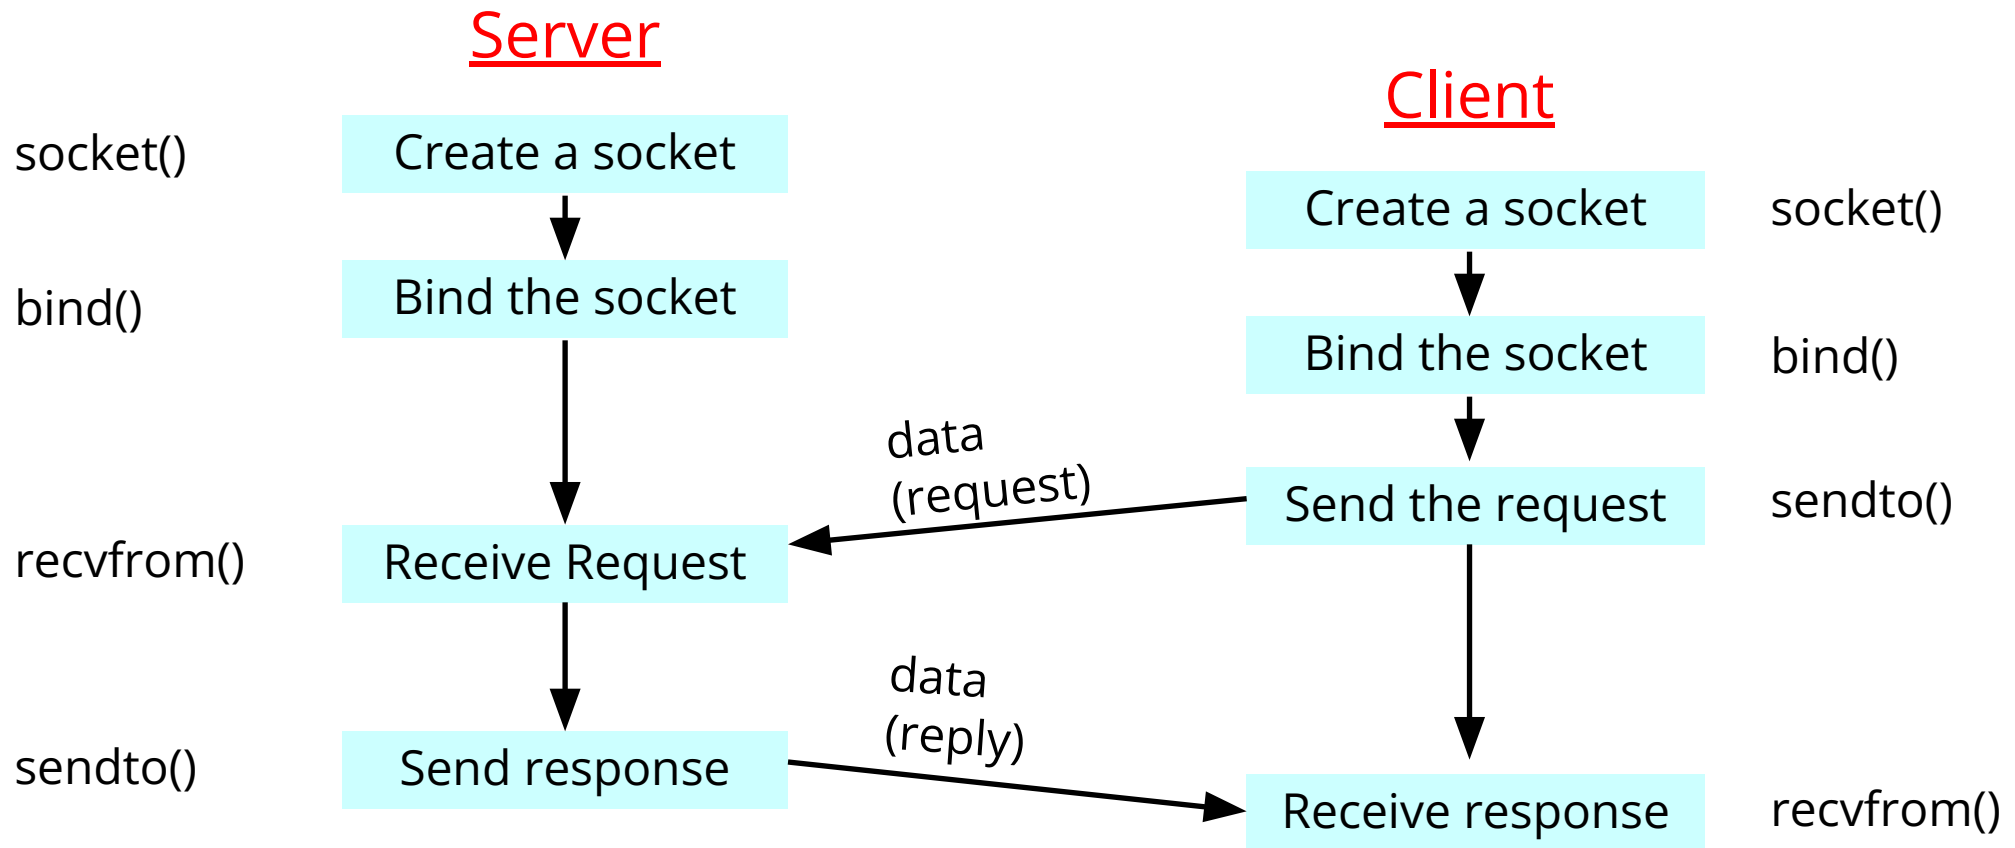
# Client - Server Communication

- Client “sometimes on”
  - Initiates a request to the server when interested
  - E.g., Web browser on your laptop or cell phone
  - Doesn’t communicate directly with other clients
  - Needs to know server’s address
- Server is “always on”
  - Handles services requests from many client hosts
  - E.g., Web server for the [www.cnn.com](http://www.cnn.com) Web site
  - Doesn’t initiate contact with the clients
  - Needs fixed, known address

# Stream Sockets (TCP): Connection-oriented



# Datagram Sockets (UDP): Connectionless



# Assignment 1

- Employ the client - server architecture
- Two files you'll modify: **client.go** and **server.go**
- Having a client send bytes to a server
- Implement the Stream Sockets (TCP): Connection-oriented



# The net package

- net.Listen receives the ip, port, and protocol, and returns a net.Listener
- net.Listener#Accept waits for connections from clients
  - Once a client connects, net.Accept returns a net.Conn to be used for communication
- net.Dial connects to the given ip and port, with the specified protocol.
  - Once it is connected, net.Dial returns a net.Conn to be used for communication

# Socket Server/Client: Go

## SERVER

- `socket, err := net.Listen("tcp4", "127.0.0.1:8080")`
  - `net.Listen` performs the C `socket`, `bind` and `listen` system calls
  - `socket` is of type `net.Listener`
- `connection, err := server.Accept()`
  - `net.Accept` accepts an incoming client request
  - `connection` is of type `net.Conn`

## CLIENT

- `connection, err := net.Dial("tcp4", "127.0.0.1:8080")`
  - Creates a TCP socket, establish connection
  - `connection` is of type `net.Conn`

# net.Conn

- **net.Conn.Read** reads from the connection
  - Wrap the connection in [bufio.Reader](#)
- **net.Conn.Write** writes to the connection
- **net.Conn.Close** closes the connection

## `net/http` (Useful in Future)

- A collection of useful functions for handling and processing http requests

# Tips and Common gotcha

- `fmt.Sprintf` could be handy
- Don't print the entire buffer
- Convert bytes to string when print
- Client needs to `close()` at end of connection
- EOF is not a character, it's a type of error

# Resources

- <https://beej.us/guide/bgipc/html/multi/unixsock.html>

# Echo Demo Code

- The one shown in Precept

```

1  package main
2
3  import (
4      "fmt"
5      "log"
6      "net"
7  )
8
9  func main() {
10     ln, err := net.Listen("tcp", "localhost:8080")
11     if err != nil {
12         log.Fatalf("Failed to setup a listener - %v\n", err)
13     }
14     defer ln.Close()
15     conn, err := ln.Accept()
16     if err != nil {
17         log.Fatalf("Failed to accept connection - %v\n", err)
18     }
19     defer conn.Close()
20     buf := make([]byte, 1024)
21     _, err = conn.Read(buf)
22     if err != nil {
23         log.Fatalf("Failed to read from connection %v\n", err)
24     }
25     fmt.Println(string(buf))
26 }

```

server.go



```
1 package main
2
3 import (
4     "fmt"
5     "log"
6     "net"
7 )
8
9 func main() {
10     conn, err := net.Dial("tcp", "localhost:8080")
11     if err != nil {
12         log.Fatalf("Failed to connect to server - %v\n", err)
13     }
14     fmt.Fprintf(conn, "Hello world!")
15 }
```

client.go

# Backup Slides

# High-level Architecture

- **Application**

- Read data from and write data to the socket
- Interpret the data (e.g., render a Web page)

- **Transport**

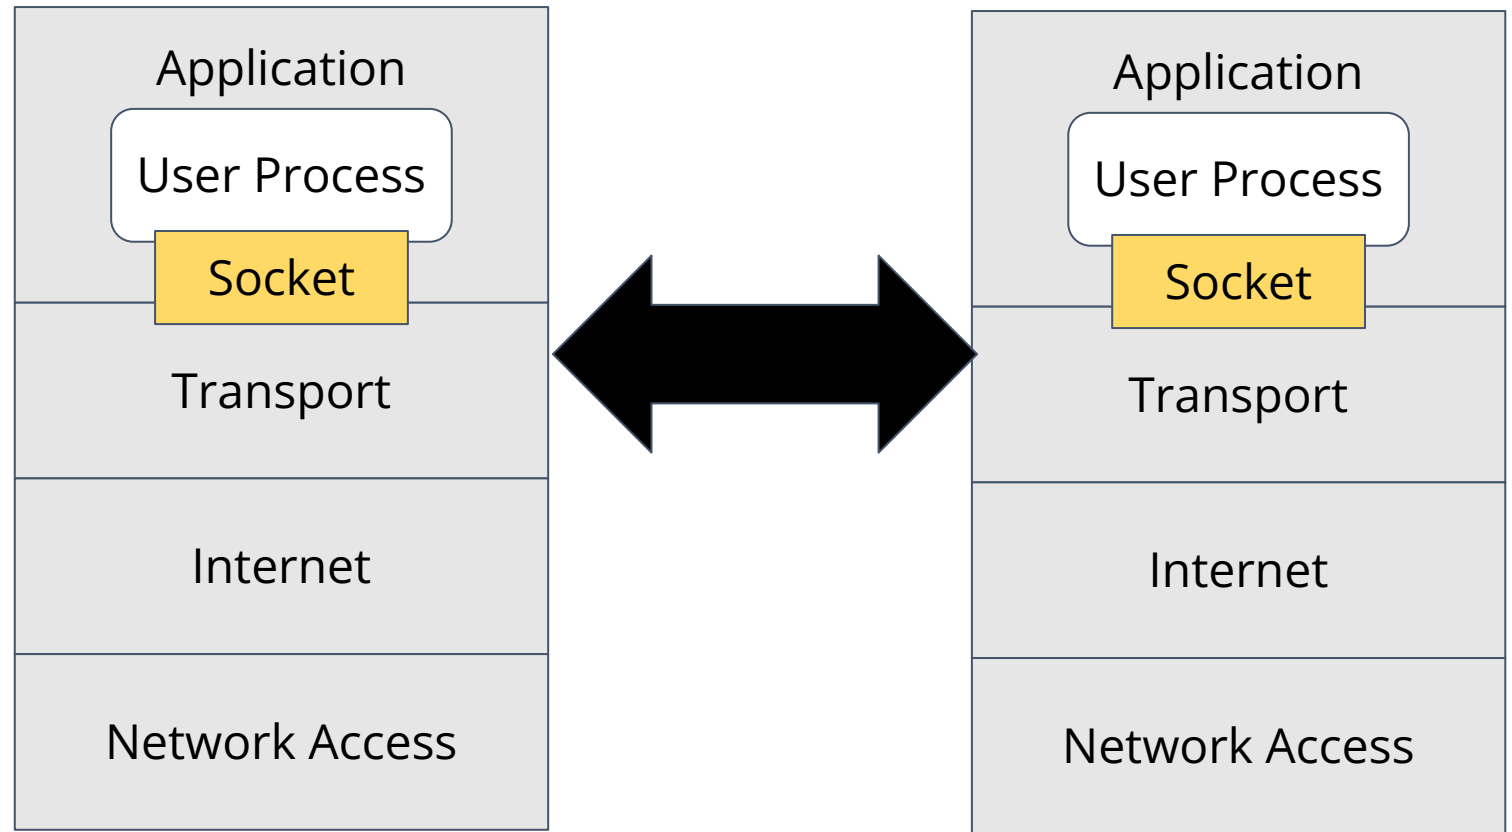
- Deliver data to the destination socket
- Based on the destination port number (e.g., 80)

- **Internet**

- Deliver data packet to the destination host
- Based on the destination IP address

- **Network Access**

- Transmit data between devices
- Encapsulate IP packet into frames transmitted by the network
- Map IP addresses into physical addresses




# Terminology

- IP (IPv4) Addresses
  - Hosts mapped to 32 bit IP addresses:  
aaaaaaaa.bbbbbbbb.cccccccc.dddddd
  - E.g., 128.112.136.51
  - Various special IP addresses, e.g., 127.0.0.1
- Domain names
  - IP addresses are mapped to an identification string
  - E.g., www.cs.princeton.edu
  - E.g., localhost
- Port - a unique communication end point on a host, named by a 16-bit integer, and associated with a process
- Connections
  - A process on one host communicates with another process on another host over a connection
  - Clients and servers communicate by sending streams of bytes over connections
  - E.g., using TCP or UDP
- Socket - end-point of a connection
  - Sending message from one process to another
    - Message must traverse the underlying network
  - Process sends and receives through a “socket”
    - In essence, the doorway leading in/out of the house
  - Socket as an Application Programming Interface
    - Supports the creation of network applications
- Stream Socket (TCP - Transmission Control Protocol)
  - Stream of bytes
  - Reliable
  - Connection-oriented
- Datagram Socket (UDP - User Datagram Protocol)
  - Collection of messages
  - Best effort
  - Connectionless

# Socket Identification

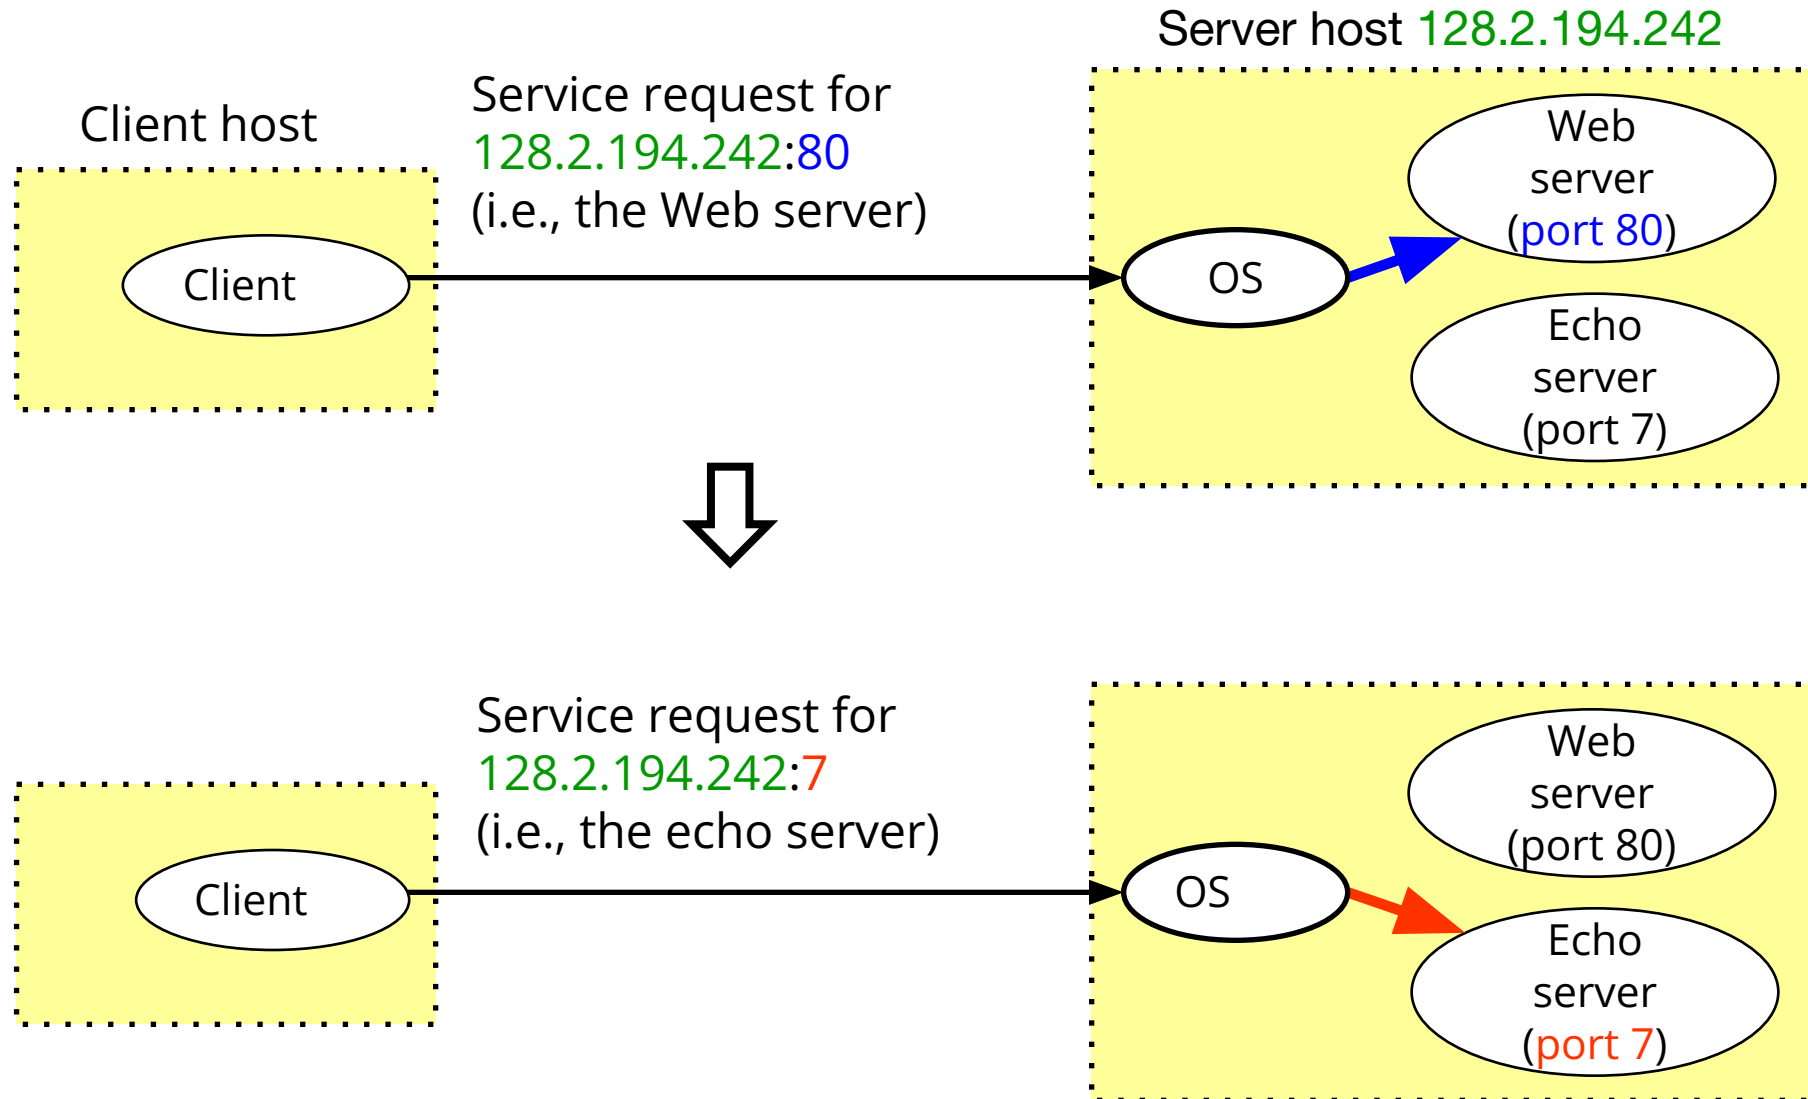
- Receiving host
  - Destination **address** that uniquely identifies host
  - **IP address**: 32-bit quantity
- Receiving socket
  - Host may be running many different processes
  - Destination **port** that uniquely identifies socket
  - **Port number**: 16-bits



# Knowing What Port Number To Use

- Popular applications have well-known ports
  - E.g., port 80 for Web and port 25 for e-mail
  - See <http://www.iana.org/assignments/port-numbers>
- Well-known vs. ephemeral ports
  - Server has a well-known port (e.g., port 80)
    - Between 0 and 1023 (requires root to use)
  - Client picks an unused ephemeral (i.e., temporary) port
    - Between 1024 and 65535
- “5 tuple” uniquely identifies traffic between hosts
  - Two IP addresses and two port numbers
  - + underlying transport protocol (e.g., TCP or UDP)

# Using Ports to Identify Services



# Worksheet



# Example C Server and Client

# Byte Order

- Network byte order
  - Big Endian
- Host byte order
  - Big Endian *or* Little Endian
- Functions to deal with this
  - `htons()` & `htonl()` (host to network short and long)
  - `ntohs()` & `ntohl()` (network to host short and long)
- When to worry?
  - putting data onto the wire
  - pulling data off the wire

# Server: Server Preparing its Socket

- Create a socket

- `int socket(int domain, int type, int protocol)`

- Bind socket to the local address and port number

- `int bind(int sock_fd, struct sockaddr *server_address, socklen_t addrlen )`

# Server: Allowing Clients to Wait

- Many client requests may arrive
  - Server cannot handle them all at the same time
  - Server could reject the requests, or let them wait
- Define how many connections can be pending
  - `int listen(int socket_fd, int backlog)`
    - Arguments: socket descriptor and acceptable backlog
    - Returns a 0 on success, and -1 on error
    - Listen is **non-blocking**: returns immediately
- What if too many clients arrive?
  - Some requests don't get through
  - The Internet makes no promises...
  - And the client can always try again

# Server: Accepting Client Connection

- Now all the server can do is wait...
  - Waits for connection request to arrive
  - **Blocking** until the request arrives
  - And then accepting the new request
- Accept a new connection from a client
  - `int accept(int sockfd, struct sockaddr *addr, socketlen_t *addrlen)`
    - Arguments: sockfd, structure that will provide client address and port, and length of the structure
    - Returns descriptor of socket for this new connection

# Client and Server: Closing Connection

- Once the connection is open
  - Both sides can read and write
  - Two unidirectional streams of data
  - In practice, client writes first, and server reads
  - ... then server writes, and client reads, and so on
- Closing down the connection
  - Either side can close the connection
  - ... using the `int close(int sockfd)`
- What about the data still “in flight”
  - Data in flight still reaches the other end
  - So, server can `close()` before client finishes reading

# Server: One Request at a Time?

- Serializing requests is inefficient
  - Server can process just one request at a time
  - All other clients must wait until previous one is done
  - What makes this inefficient?
- May need to time share the server machine
  - Alternate between servicing different requests
    - Do a little work on one request, then switch when you are waiting for some other resource (e.g., reading file from disk)
    - “Nonblocking I/O”
  - Or, use a different process/thread for each request
    - Allow OS to share the CPU(s) across processes
  - Or, some hybrid of these two approaches

# Handle Multiple Clients using `fork()`

- Steps to handle multiple clients
  - Go to a loop and accept connections using `accept()`
  - After a connection is established, call `fork()` to create a new child process to handle it
  - Go back to listen for another socket in the parent process
  - `close()` when you are done.
- Want to know more?
  - Checkout out *[Beej's guide to network programming](#)*



# Sockets in Go