


COS 316

Precept:

Cache Eviction (Replacement)

Overview of Web Caching

- Basic idea:
 - Bring objects “closer” to clients
- Three primary features:
 - Reduce network bandwidth
 - Reduce client-perceived delays
 - Reduce load on server
- Cache Replacement Strategy
 - When a cache becomes full, which object should be **evicted/replaced**?



Cache Eviction Algorithms

- High level
 - Client requests a new object
 - If object is in cache, return the object
 - If object is not in cache:
 - Get object from server/provider and return the object
 - Cache full:
 - Identify an object in cache to evict
 - Evict the object in the cache
 - Replace with new object
 - Cache not full:
 - Admit the new object to the cache

Cache Eviction Algorithms

- Least recently used (LRU): Evict the object from the cache whose last request is the oldest
- First-in, First-out (FIFO): Evict the object from the cache that has been in the cache the longest
- Many others...

LRU



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 8
size: 10
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 47

LRU



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 8
size: 10
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 47

LRU



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 8
size: 10
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 50
Cache size = 47

FIFO



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 8
size: 10
request: 14:00
admit: 14:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 53

FIFO



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 8
size: 10
request: 14:00
admit: 14:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 53

FIFO



id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 8
size: 10
request: __: __
admit: __: __

Current
time:
14:00

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 6
size: 2
request: 13:00
admit: 11:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 11:30
admit: 11:30

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 45

id: 8
size: 10
request: 14:00
admit: 14:00

id: 3
size: 10
request: 13:45
admit: 13:45

id: 1
size: 3
request: 15:01
admit: 12:01

id: 4
size: 5
request: 11:53
admit: 11:33

id: 11
size: 8
request: 14:00
admit: 14:00

id: 7
size: 17
request: 13:30
admit: 13:30

Cache capacity = 55
Cache size = 53

Experiments

- > `cd <COS316-Public repo> # directory containing Vagrantfile`
- > `git pull # update with precept5 code and data`
- > `vagrant up ## if instance is not already started`
- > `vagrant ssh`
- > `cd precepts/precept5/webcachesim-master`
- > `make`

Trace File Form

- Request traces must be given in a space-separated format with three columns
- time - long long int
- id - long long int, used to uniquely identify objects
- size should be a long long int, object's size in bytes

- Example

time	id	size
1	1	120
2	2	64
3	1	120
4	3	14
4	1	120

- See test.tr

Using the Simulator^{*}

```
> ./webcachesim test.tr LRU 1000
```

```
LRU:1000 bytes, 10492 reqs, 8495 hits, 81 hits/reqs(%)
```

```
> ./webcachesim test.tr FIFO 1000
```

```
FIFO:1000 bytes, 10492 reqs, 8206 hits, 78 hits/reqs(%)
```

^{*} Derived from <https://github.com/dasebe/webcachesim>

Experiments

- Trace data from a production CDN
 - cd1-10M.tr *
 - 10 million requests / Object sizes from 10 byte to .7GB
- LIFO and FIFO
- Vary cache sizes
- 16000000
- 32000000
- 64000000
- 128000000
- 256000000
- 512000000
- 1024000000
- 2048000000
- 4096000000
- Create a Google Sheet
- Three columns
- SIZE LRU FIFO
- Copy results accordingly
- Select three columns to create line chart

* Practical Bounds on Optimal Caching with Variable Object Sizes Daniel S. Berger, Nathan Beckmann, Mor Harchol-Balter. ACM SIGMETRICS, June 2018