

COS 316

Precept:

Concurrency

Precept Objectives

- Learn about basics of concurrent programs and programming
- Understand how to use a few Golang mechanisms for developing concurrent programs

Overview of Concurrency

Sequential programs:

- Single thread of control
- Subprograms / tasks - don't overlap in time - executed one after another

Concurrent programs

- Multiple threads of control
- Subprograms / tasks - may (conceptually) overlap in time - (appear to be) executed at the same time

- Recall from lecture
 - Computer with a single processor can have multiple processes at once
 - OS schedules different processes - giving illusion that multiple processes are running simultaneously
- Note - parallel architectures can have N processes running simultaneously on N processors

Operating System - Review

- Allows many processes to execute concurrently
- Ensures each process's physical address space does not overlap
- Ensures all processes get fair share of processor time and resources
- Processes can run concurrently and (context) switch
- User's perspective: appears that processes run in parallel although they don't

Context Switch - Review

- Control flow changes from one process to another
 - E.g., switching from processA to processB
- Overhead:
 - Before each switch OS has to save the state (context) of currently running process and restore it when next time its execution gets resumed

Scheduling Policies - Discussion

- What are the three policies outlined in lecture?
- What are some other possible policies?

Threads vs Processes

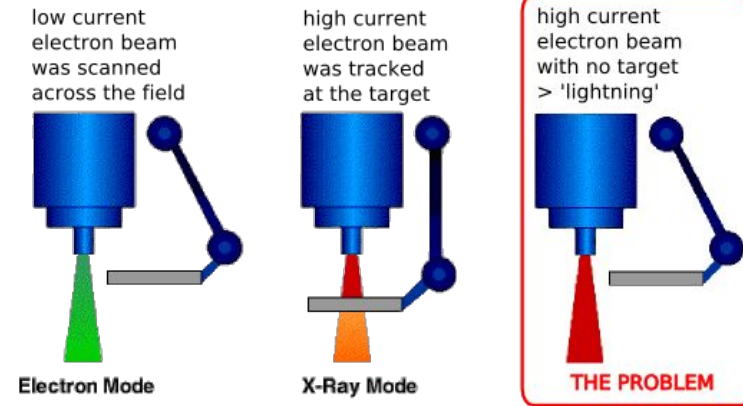
- Processes
 - Process context switching time is long (switching between processes is slow because memory access)
- Threads
 - thread is a “lightweight” process
 - thread shares some of the context with other threads in a process
- Shared process context among process threads:
 - Virtual memory
 - File descriptors
- Private context for each thread:
 - Stack
 - Data registers
 - Code (PC)
- Switching between threads is faster because there is less context - less data that has to be read/written from/to memory

Why Concurrency?

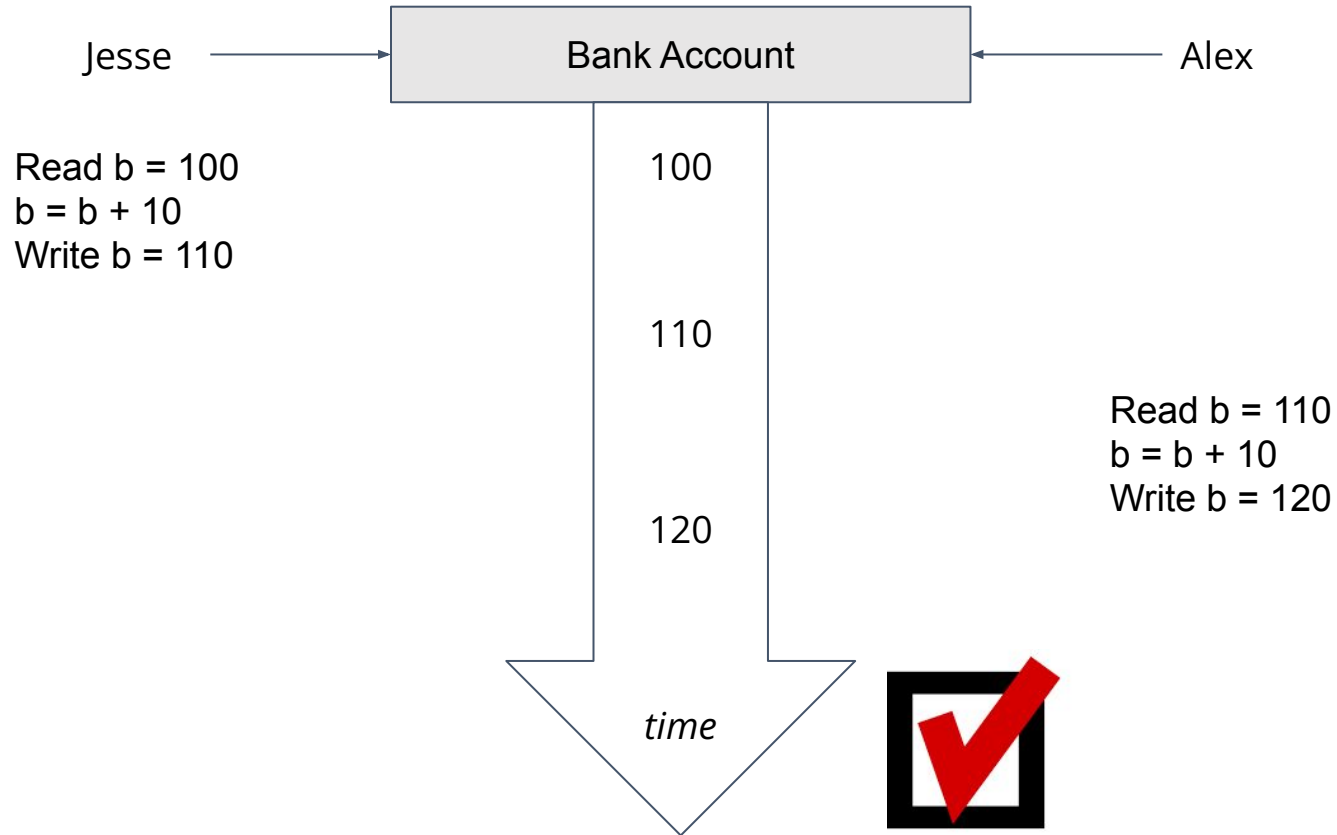
- Performance gain
 - Google search queries
- Application throughput
 - Throughput = amount of work that a computer can do in a given time period
 - When one task is waiting (blocking) for I/O another task can continue its execution
- Model real-world structures
 - Multiple sensors
 - Multiple events
 - Multiple activities

Tradeoffs - Concurrent Programming

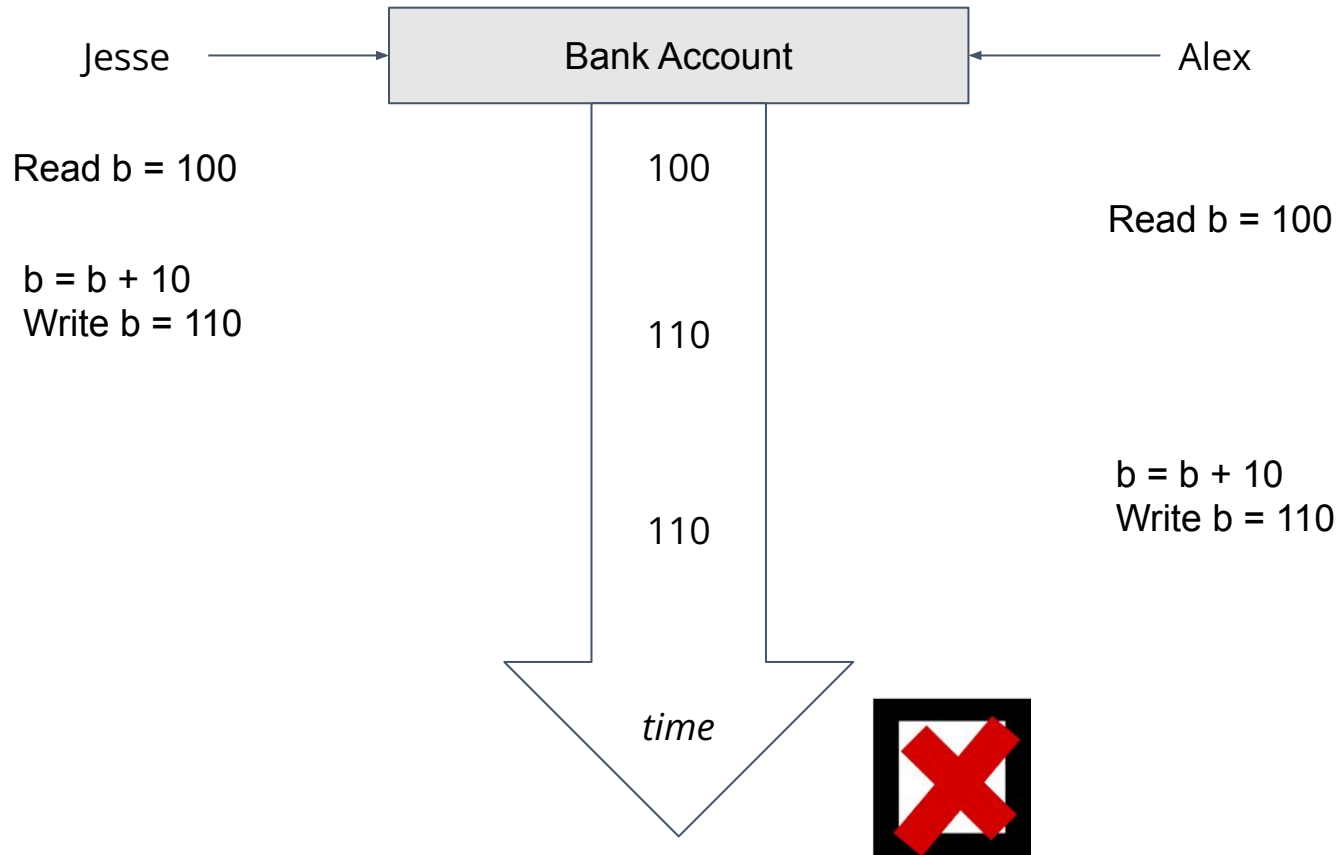
- Complex
- Error-prone
- Hard to debug
- Famous bugs related to concurrency:
 - Therac-25
 - Computerized radiation therapy machine
 - Two “safe modes” of operations, but a race condition caused a severely unsafe mode
 - Patients died from overdoses
 - See [N.G. Leveson and C.S. Turner. An investigation of the Therac-25 accidents. Computer. July 1993](http://www.nrc.gov/readingrm/doc/40101main.pdf)



Example



Example



Go and Concurrency

- Goroutines
- The sync package - <https://golang.org/pkg/sync>
 - `sync.WaitGroup`
 - `sync.Mutex`

Goroutines

- A lightweight thread managed by the Go runtime
- Many goroutines execute within a single OS thread
 - One goroutine is created automatically to execute the `main()`
 - Other goroutines are created using the `go` keyword
 - Order of execution depends on the Go scheduler
 - Go takes a process with main thread and schedules / switches goroutines within that thread
- Compare
 - Sequential Program
 - <https://play.golang.org/p/PLeCGtRp2QB>
 - Concurrent program
 - https://play.golang.org/p/sDitCEr_3vX

Go Runtime Scheduler

- Schedules goroutines inside an OS thread (main thread)
 - Like a little OS inside a single OS thread
- Go runs on a main thread and switches goroutines on one thread
- Go runtime scheduler uses Logical Processor which is mapped to a thread
- Typically there is one Logical Processor which is mapped to a main thread
- Since all these goroutines are running on one thread, we don't have parallelism but concurrency
- Increase number of Logical Processors - mapped to different threads and OS can map those threads to different cores
- Program can determine how many Logical Processors will be there; default is 1 (so we'll have concurrent execution of routines) but can be increased (so we might have parallel goroutines execution - if OS schedules running different threads on different cores)

Goroutines - Exiting

- goroutine exits when code associated with its function returns
- When the main goroutine is complete, **all** other goroutines exit, even if they are not finished
 - goroutines are forced to exit when main goroutine exits
 - goroutine may not complete its execution because main completes early
- Execution order of goroutines is non-deterministic

Exercises

- Recall the exercise:
- https://play.golang.org/p/sDitCEr_3vX
- Switch the order of the calls from

```
go say("world")  
say("hello")
```

```
say("hello")  
go say("world")
```

- What happens?
- How to fix?

Synchronization

- Synchronization is when multiple threads agree on a timing of an event
- Global *events* whose execution is viewed by all threads, simultaneously
- One goroutine does not know the timing of other goroutines
- Synchronization introduces some global events that every thread sees at the same time

Synchronization and Go

- type WaitGroup
 - func (wg *WaitGroup) Add(delta int)
 - func (wg *WaitGroup) Done()
 - func (wg *WaitGroup) Wait()
- type Mutex
 - func (m *Mutex) Lock()
 - func (m *Mutex) Unlock()
- Channels
 - See COS 418

WaitGroup

- Forces a goroutine to wait for other goroutines
- WaitGroup - a group of goroutines that a goroutine has to wait for
- A goroutine will not continue until all goroutines from WaitGroup finish
- Can wait on one or more other goroutines

- Create a WaitGroup
var wg sync.WaitGroup
- Set the size of the WaitGroup
wg.Add(num_goroutines)
- Pass a pointer to the WaitGroup to each go routine
func f(wg *sync.WaitGroup)
- When goroutine completes, invoke Done
wg.Done()
- Invoke Wait - blocks until all goroutines complete
wg.Wait()

WaitGroup Exercises

Consider this program:

```
func doWork(id int, sec int) {  
    fmt.Printf("goroutine %d - entered. ", id)  
    fmt.Printf("Sleep for %d seconds.\n", sec)  
    time.Sleep(time.Duration(sec) * time.Second)  
    fmt.Printf("goroutine %d - exits. ", id)  
    fmt.Printf("Slept for %d seconds\n", sec)  
}  
  
func main() {  
    rand.Seed(time.Now().UnixNano())  
    for i := 1; i <= 5; i++ {  
        go doWork(i, rand.Intn(5) + 1)  
    }  
    fmt.Println("Main goroutine exit")  
}
```

- Run the program

<https://play.golang.org/p/nb8lJC3lylt>

- Modify the program so that each worker prints its:
 - Enter statement
 - Exit statement

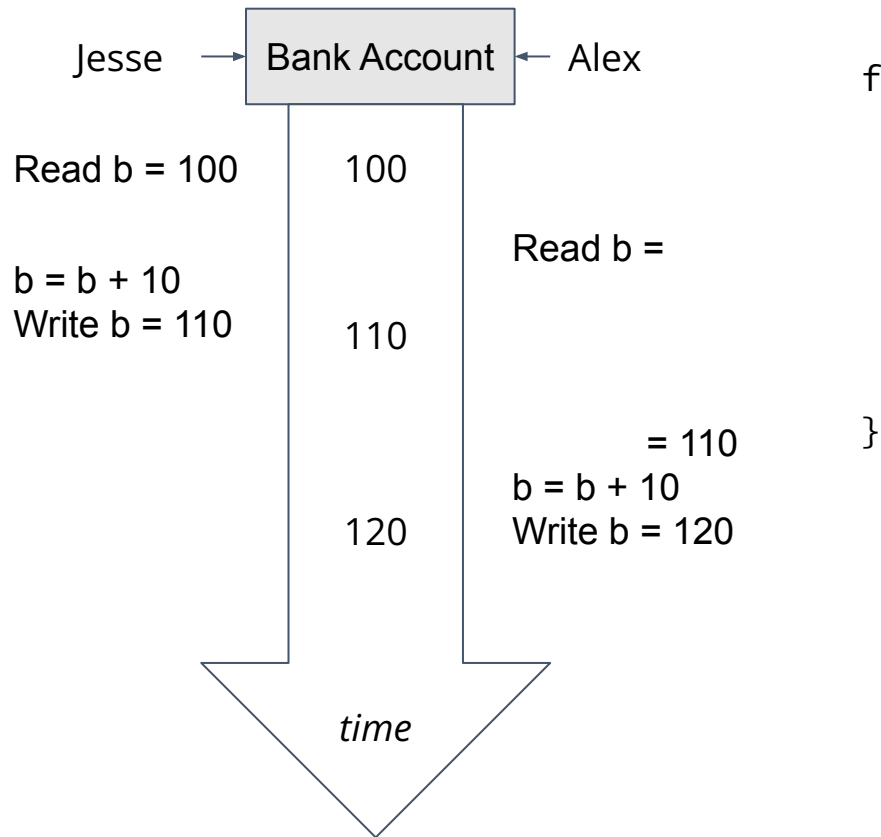
Mutex (Mutual Exclusion)

- Sharing variables between goroutines (concurrently) can cause problems
- Two goroutines writing to the same shared variable can interfere with each other
- Function/goroutine is said to be concurrency-safe if can be executed concurrently with other goroutines without interfering improperly with them
 - e.g., it will not alter variables in other goroutines in some unexpected/unintended/unsafe way

Sync.Mutex

- A mutex ensures *mutual exclusion*
- Uses a binary semaphore
 - If flag is up → shared variable is in use by somebody
- Only one goroutine can write into variable at a time
- Once goroutine is done with using shared variable it has to put the flag down
 - if flag is down → shared variable is available
- If another goroutine see that flag is down it knows it can use the shared variable but first it has to put the flag up

Back to our example



```
func Deposit(amount) {
```

```
    lock balanceLock
```

```
    read balance
```

```
    balance = balance + amount
```

```
    write balance
```

```
    unlock balanceLock
```

```
}
```

} CRITICAL
SECTION

Sync.Mutex

- **Lock()**
 - Puts the flag up (if none of other goroutines has already put the flag up)
 - Notifies others that shared variable is in use
 - If second goroutine also calls **Lock()** it will be blocked, it has to wait until first goroutine releases the lock
 - Note - any number of goroutines (not just two) competing to **Lock()**
 - **Unlock()**
 - Puts the flag down
 - Notifies others that it is done with using shared variable
 - When **Unlock()** is called, a blocked **Lock()** can proceed
 - In general: put **Lock()** at the beginning of the critical section and call **Unlock()** at the end of it; ensures that only one goroutine will be in critical section region
- Create a Mutex
var mut sync.Mutex
 - To lock a critical section
mut.Lock()
 - To unlock a critical section
mut.Unlock()

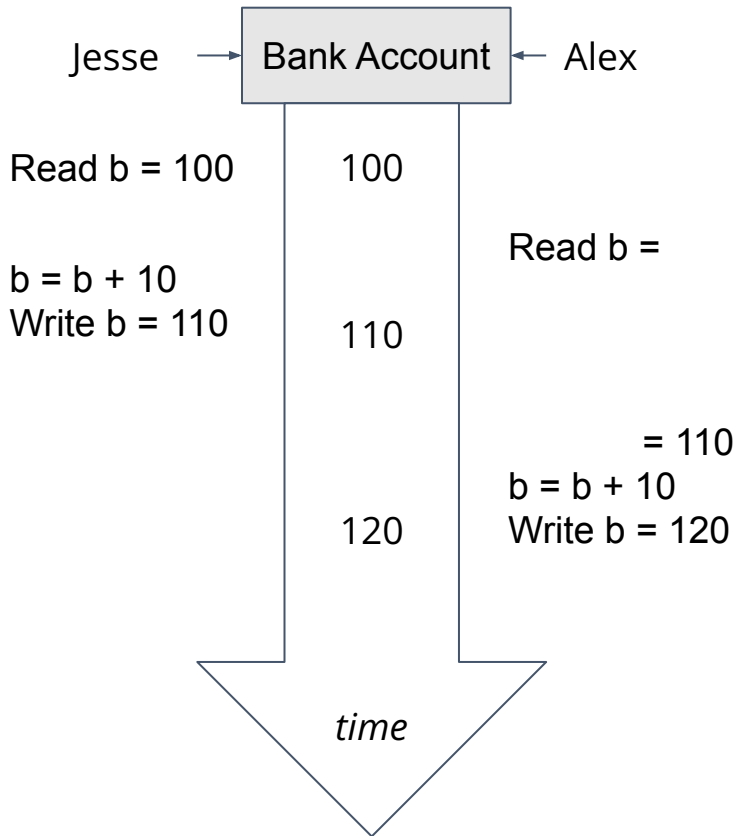
Mutex Exercise

Consider:

```
var i int = 0
var wg sync.WaitGroup
func inc() {
    i = i + 1
    wg.Done()
}
func main() {
    wg.Add(2)
    go inc()
    go inc()
    wg.Wait()
    fmt.Println(i)
}
```

- Run the program
<https://play.golang.org/p/hNevYkKDp30>
- Is it concurrency-safe? Discuss.
- Consider this program
<https://play.golang.org/p/c-D5UiTmgnX>
- Copy this program to your local machine - build and then execute multiple times
 - Not different behavior than Go playground
- Use `Lock()` and `Unlock()` to make these programs concurrency-safe

Mutex Exercise - Bank Account



- Make this code concurrency-safe

<https://play.golang.org/p/sbc5lly8cq>

Interesting Example

Consider:

```
var mu sync.Mutex
func funcA() {
    mu.Lock()
    funcB()
    mu.Unlock()
}
func funcB() {
    mu.Lock()
    fmt.Println("Hello, World")
    mu.Unlock()
}
func main() {
    funcA()
}
```

- Run the program

https://play.golang.org/p/c2Qgo-W_4mP

- Discuss.

Next Week - Dining Philosophers

References

Derived from:

<http://www.bojankomazec.com/2019/02/concurrency-in-go-notes-on-coursera.html>